

# Quick introduction into SAT/SMT solvers and symbolic execution

Dennis Yurichev <dennis(a)yurichev.com>

October 3, 2017

## Contents

<b>1 This is a draft!</b>	<b>4</b>
<b>2 Thanks</b>	<b>4</b>
<b>3 Praise</b>	<b>4</b>
<b>4 Introduction</b>	<b>4</b>
<b>5 Is it a hype? Yet another fad?</b>	<b>4</b>
<b>6 SMT<sup>1</sup>-solvers</b>	<b>4</b>
6.1 School-level system of equations . . . . .	4
6.2 Another school-level system of equations . . . . .	6
6.3 Connection between SAT <sup>2</sup> and SMT solvers . . . . .	6
<b>7 Alphametics</b>	<b>6</b>
7.1 Generating de Bruijn sequences using Z3 . . . . .	8
7.2 Zebra puzzle (AKA <sup>3</sup> Einstein puzzle) . . . . .	10
7.3 Sudoku puzzle . . . . .	13
7.3.1 The first idea . . . . .	13
7.3.2 The second idea . . . . .	17
7.3.3 Conclusion . . . . .	18
7.3.4 Homework . . . . .	19
7.3.5 Further reading . . . . .	19
7.3.6 Sudoku as a SAT problem . . . . .	19
7.4 Solving Problem Euler 31: “Coin sums” . . . . .	19
7.5 Using Z3 theorem prover to prove equivalence of some weird alternative to XOR operation . . . . .	20
7.5.1 In SMT-LIB form . . . . .	21
7.5.2 Using universal quantifier . . . . .	21
7.5.3 How the expression works . . . . .	22
7.5.4 In SAT . . . . .	22
7.6 Dietz’s formula . . . . .	22
7.7 Cracking LCG <sup>4</sup> with Z3 . . . . .	23
7.8 Solving pipe puzzle using Z3 SMT-solver . . . . .	25
7.8.1 Generation . . . . .	26
7.8.2 Solving . . . . .	27
7.9 Cracking Minesweeper with Z3 SMT solver . . . . .	30
7.9.1 The method . . . . .	30
7.9.2 The code . . . . .	31
7.10 Recalculating micro-spreadsheet using Z3Py . . . . .	34
7.10.1 Unsat core . . . . .	35

---

<sup>1</sup>Satisfiability modulo theories

<sup>2</sup>Boolean satisfiability problem

<sup>3</sup>Also Known As

<sup>4</sup>Linear congruential generator

7.10.2	Stress test	35
7.10.3	The files	37
7.11	Discrete tomography	37
7.12	Simplifying long and messy expressions using Mathematica and Z3	40
7.13	Solving XKCD 287 using Z3	41
7.14	Making smallest possible test suite using Z3	42
7.15	Package manager and Z3	44
7.16	Cracking simple XOR cipher with Z3	46
7.17	Balanced Gray code and Z3 SMT solver	49
7.17.1	Duke Nukem 3D from 1990s	54
7.18	Integer factorization using Z3 SMT solver	55
7.19	Tiling puzzle and Z3 SMT solver	56
<b>8</b>	<b>Program synthesis</b>	<b>59</b>
8.1	Synthesis of simple program using Z3 SMT-solver	60
8.1.1	Few notes	62
8.1.2	The code	62
8.2	Rockey dongle: finding unknown algorithm using only input/output pairs	62
8.2.1	Conclusion	67
8.2.2	The files	67
8.2.3	Further work	67
8.2.4	Exercise	67
<b>9</b>	<b>Toy decompiler</b>	<b>67</b>
9.1	Introduction	67
9.2	Data structure	67
9.3	Simple examples	69
9.4	Dealing with compiler optimizations	73
9.4.1	Division using multiplication	78
9.5	Obfuscation/deobfuscation	79
9.6	Tests	84
9.6.1	Evaluating expressions	84
9.6.2	Using Z3 SMT-solver for testing	85
9.7	My other implementations of toy decompiler	86
9.7.1	Even simpler toy decompiler	87
9.8	Difference between toy decompiler and commercial-grade one	87
9.9	Further reading	88
9.10	The files	88
<b>10</b>	<b>Symbolic execution</b>	<b>88</b>
10.1	Symbolic computation	88
10.1.1	Rational data type	89
10.2	Symbolic execution	90
10.2.1	Swapping two values using XOR	90
10.2.2	Change endianness	90
10.2.3	Fast Fourier transform	92
10.2.4	Cyclic redundancy check	94
10.2.5	Linear congruential generator	97
10.2.6	Path constraint	98
10.2.7	Division by zero	100
10.2.8	Merge sort	100
10.2.9	Extending Expr class	102
10.2.10	Conclusion	102
10.3	Further reading	102

<b>11 KLEE</b>	<b>103</b>
11.1 Installation	103
11.2 School-level equation	103
11.3 Zebra puzzle	104
11.4 Sudoku	108
11.5 Unit test: HTML/CSS color	112
11.6 Unit test: strcmp() function	114
11.7 UNIX date/time	116
11.8 Inverse function for base64 decoder	120
11.9 CRC (Cyclic redundancy check)	123
11.9.1 Buffer alteration case #1	123
11.9.2 Buffer alteration case #2	124
11.9.3 Recovering input data for given CRC32 value of it	125
11.9.4 In comparison with other hashing algorithms	126
11.10 ZSS decompressor	126
11.11 strtodx() from RetroBSD	129
11.12 Unit testing: simple expression evaluator (calculator)	132
11.13 Regular expressions	137
11.14 Exercise	138
<b>12 (Amateur) cryptography</b>	<b>138</b>
12.1 Serious cryptography	138
12.1.1 Attempts to break “serious” crypto	141
12.2 Amateur cryptography	141
12.2.1 Bugs	143
12.2.2 XOR ciphers	143
12.2.3 Other features	143
12.2.4 Examples	143
12.3 Case study: simple hash function	144
12.3.1 Manual decompiling	144
12.3.2 Now let’s use the Z3	147
<b>13 SAT-solvers</b>	<b>150</b>
13.1 CNF form	150
13.2 Example: 2-bit adder	151
13.2.1 MiniSat	153
13.2.2 CryptoMiniSat	154
13.3 Picosat	155
13.4 Eight queens puzzle	155
13.4.1 POPCNT1	155
13.4.2 Eight queens	156
13.4.3 Counting all solutions	158
13.4.4 Skipping symmetrical solutions	158
13.5 Sudoku in SAT	158
13.5.1 Getting rid of one POPCNT1 function call	162
13.6 Zebra puzzle as a SAT problem	162
13.7 Cracking Minesweeper with SAT solver	166
13.7.1 Simple <i>population count</i> function	166
13.7.2 Minesweeper	169
13.8 Conway’s “Game of Life”	172
13.8.1 Reversing back state of “Game of Life”	172
13.8.2 Finding “still lives”	178
13.8.3 The source code	185
13.9 Simplest SAT solver in ~120 lines	186
13.10 Integer factorization using SAT solver	188
13.10.1 Binary adder in SAT	188
13.10.2 Binary multiplier in SAT	190
13.10.3 Glueing all together	192
13.10.4 Division using multiplier	193

13.10. Breaking <b>RSA!</b> <sup>5</sup>	194
13.10. Further reading	194
13.11. Proving bizarre XOR alternative using SAT solver	194
<b>14 MaxSAT</b>	<b>196</b>
14.1 Gray code in MaxSAT	196
<b>15 Acronyms used</b>	<b>198</b>

## 1 This is a draft!

This is very early draft, but still can be interesting for someone.

Latest version is always available at [http://yurichev.com/writings/SAT\\_SMT\\_draft-EN.pdf](http://yurichev.com/writings/SAT_SMT_draft-EN.pdf). Russian version is at [http://yurichev.com/writings/SAT\\_SMT\\_draft-RU.pdf](http://yurichev.com/writings/SAT_SMT_draft-RU.pdf).

New parts are appearing here from time to time, see: [https://github.com/dennis714/SAT\\_SMT\\_article/blob/master/ChangeLog](https://github.com/dennis714/SAT_SMT_article/blob/master/ChangeLog).

For news about updates, you may subscribe my twitter<sup>6</sup>, facebook<sup>7</sup>, or github repo<sup>8</sup>.

## 2 Thanks

Leonardo Mendonça de Moura<sup>9</sup>, Nikolaj Bjørner<sup>10</sup>, Armin Biere<sup>11</sup> and Mate Soos<sup>12</sup>, for help.

## 3 Praise

“An excellent source of well-worked through and motivating examples of using Z3’s python interface.”<sup>13</sup> (Nikolaj Bjorner, one of Z3’s author).

## 4 Introduction

**SAT/SMT** solvers can be viewed as solvers of huge systems of equations. The difference is that **SMT** solvers takes systems in arbitrary format, while **SAT** solvers are limited to boolean equations in **CNF**<sup>14</sup> form.

A lot of real world problems can be represented as problems of solving system of equations.

## 5 Is it a hype? Yet another fad?

Some people say, this is just another hype. No, **SAT** is old enough and fundamental to **CS**<sup>15</sup>. The reason of increased interest to it is that computers gets faster over the last couple decades, so there are attempts to solve old problems using **SAT/SMT**, which were inaccessible in past.

<sup>5</sup>**RSA!**

<sup>6</sup><https://twitter.com/yurichev>

<sup>7</sup><https://www.facebook.com/dennis.yurichev.5>

<sup>8</sup>[https://github.com/dennis714/SAT\\_SMT\\_article](https://github.com/dennis714/SAT_SMT_article)

<sup>9</sup><https://www.microsoft.com/en-us/research/people/leonardo/>

<sup>10</sup><https://www.microsoft.com/en-us/research/people/nbjorner/>

<sup>11</sup><http://fmv.jku.at/biere/>

<sup>12</sup><https://www.msoos.org/>

<sup>13</sup><https://github.com/Z3Prover/z3/wiki>

<sup>14</sup>Conjunctive normal form

<sup>15</sup>Computer science

## 6 SMT-solvers

### 6.1 School-level system of equations

I've got this school-level system of equations copypasted from Wikipedia <sup>16</sup>:

$$\begin{aligned}3x + 2y - z &= 1 \\ 2x - 2y + 4z &= -2 \\ -x + \frac{1}{2}y - z &= 0\end{aligned}$$

Will it be possible to solve it using Z3? Here it is:

```
#!/usr/bin/python
from z3 import *

x = Real('x')
y = Real('y')
z = Real('z')
s = Solver()
s.add(3*x + 2*y - z == 1)
s.add(2*x - 2*y + 4*z == -2)
s.add(-x + 0.5*y - z == 0)
print s.check()
print s.model()
```

We see this after run:

```
sat
[z = -2, y = -2, x = 1]
```

If we change any equation in some way so it will have no solution, `s.check()` will return “unsat”.

I've used “Real” *sort* (some kind of data type in SMT-solvers) because the last expression equals to  $\frac{1}{2}$ , which is, of course, a real number. For the integer system of equations, “Int” *sort* would work fine.

Python (and other high-level PL<sup>17</sup>s like C#) interface is highly popular, because it's practical, but in fact, there is a standard language for SMT-solvers called SMT-LIB <sup>18</sup>.

Our example rewritten to it looks like this:

```
(declare-const x Real)
(declare-const y Real)
(declare-const z Real)
(assert (= (- (+ (* 3 x) (* 2 y)) z) 1))
(assert (= (+ (- (* 2 x) (* 2 y)) (* 4 z)) -2))
(assert (= (- (+ (- 0 x) (* 0.5 y)) z) 0))
(check-sat)
(get-model)
```

This language is very close to LISP, but is somewhat hard to read for untrained eyes.

Now we run it:

```
% z3 -smt2 example.smt
sat
(model
  (define-fun z () Real
    (- 2.0))
  (define-fun y () Real
    (- 2.0))
  (define-fun x () Real
    1.0)
)
```

So when you look back to my Python code, you may feel that these 3 expressions could be executed. This is not true: Z3Py API offers overloaded operators, so expressions are constructed and passed into the guts of Z3 without any execution <sup>19</sup>. I would call it “embedded DSL<sup>20</sup>”.

<sup>16</sup>[https://en.wikipedia.org/wiki/System\\_of\\_linear\\_equations](https://en.wikipedia.org/wiki/System_of_linear_equations)

<sup>17</sup>Programming Language

<sup>18</sup><http://smtlib.cs.uiowa.edu/papers/smt-lib-reference-v2.5-r2015-06-28.pdf>

<sup>19</sup><https://github.com/Z3Prover/z3/blob/6e852762baf568af2aad1e35019fdf41189e4e12/src/api/python/z3.py>

<sup>20</sup>Domain-specific language

Same thing for Z3 C++ API, you may find there “operator+” declarations and many more <sup>21</sup>. Z3 API<sup>22</sup>s for Java, ML and .NET are also exist <sup>23</sup>.

Z3Py tutorial: <https://github.com/ericpony/z3py-tutorial>.

Z3 tutorial which uses SMT-LIB language: <http://rise4fun.com/Z3/tutorial/guide>.

## 6.2 Another school-level system of equations

I’ve found this somewhere at Facebook:

$$\begin{aligned} \text{circle} + \text{circle} &= 10 \\ \text{circle} \times \text{square} + \text{square} &= 12 \\ \text{circle} \times \text{square} - \text{triangle} \times \text{circle} &= \text{circle} \\ \text{triangle} &= ? \end{aligned}$$

Figure 1: System of equations

It’s that easy to solve it in Z3:

```
#!/usr/bin/python
from z3 import *

circle, square, triangle = Ints('circle square triangle')
s = Solver()
s.add(circle+circle==10)
s.add(circle*square+square==12)
s.add(circle*square-triangle*circle==circle)
print s.check()
print s.model()
```

```
sat
[triangle = 1, square = 2, circle = 5]
```

## 6.3 Connection between SAT and SMT solvers

SMT-solvers are frontends to SAT solvers, i.e., they translating input SMT expressions into CNF and feed SAT-solver with it. Translation process is sometimes called “bit blasting”. Some SMT-solvers uses external SAT-solver: STP uses MiniSAT or CryptoMiniSAT as backend. Some other SMT-solvers (like Z3) has their own SAT solver.

## 7 Alphametics

According to Donald Knuth, the term “Alphametics” was coined by J. A. H. Hunter. This is a puzzle: what decimal digits in 0..9 range must be assigned to each letter, so the following equation will be true?

```
  SEND
+ MORE
-----
 MONEY
```

<sup>21</sup><https://github.com/Z3Prover/z3/blob/6e852762baf568af2aad1e35019fdf41189e4e12/src/api/c%2B%2B/z3%2B%2B.h>

<sup>22</sup>Application programming interface

<sup>23</sup><https://github.com/Z3Prover/z3/tree/6e852762baf568af2aad1e35019fdf41189e4e12/src/api>

This is easy for Z3:

```
from z3 import *

# SEND+MORE=MONEY

D, E, M, N, O, R, S, Y = Ints('D, E, M, N, O, R, S, Y')

s = Solver()

s.add(Distinct(D, E, M, N, O, R, S, Y))
s.add(And(D >= 0, D <= 9))
s.add(And(E >= 0, E <= 9))
s.add(And(M >= 0, M <= 9))
s.add(And(N >= 0, N <= 9))
s.add(And(O >= 0, O <= 9))
s.add(And(R >= 0, R <= 9))
s.add(And(S >= 0, S <= 9))
s.add(And(Y >= 0, Y <= 9))

s.add(1000*S+100*E+10*N+D + 1000*M+100*O+10*R+E == 10000*M+1000*O+100*N+10*E+Y)

print s.check()
print s.model()
```

Output:

```
sat
[E, = 5,
 S, = 9,
 M, = 1,
 N, = 6,
 D, = 7,
 R, = 8,
 O, = 0,
 Y = 2]
```

Another one, also from **TAOCP!**<sup>24</sup> volume IV (<http://www-cs-faculty.stanford.edu/~uno/fasc2b.ps.gz>):

```
from z3 import *

# VIOLIN+VIOLIN+VIOLA = TRIO+SONATA

A, I, L, N, O, R, S, T, V = Ints('A, I, L, N, O, R, S, T, V')

s = Solver()

s.add(Distinct(A, I, L, N, O, R, S, T, V))
s.add(And(A >= 0, A <= 9))
s.add(And(I >= 0, I <= 9))
s.add(And(L >= 0, L <= 9))
s.add(And(N >= 0, N <= 9))
s.add(And(O >= 0, O <= 9))
s.add(And(R >= 0, R <= 9))
s.add(And(S >= 0, S <= 9))
s.add(And(T >= 0, T <= 9))
s.add(And(V >= 0, V <= 9))

VIOLIN, VIOLA, SONATA, TRIO = Ints('VIOLIN, VIOLA, SONATA, TRIO')

s.add(VIOLIN == 100000*V+10000*I+1000*O+100*L+10*I+N)
s.add(VIOLA == 10000*V+1000*I+100*O+10*L+A)
s.add(SONATA == 100000*S+10000*O+1000*N+100*A+10*T+A)
s.add(TRIO == 1000*T+100*R+10*I+O)

s.add(VIOLIN+VIOLIN+VIOLA == TRIO+SONATA)

print s.check()
print s.model()
```

```
sat
[L, = 6,
```

---

<sup>24</sup>**TAOCP!**

```

S, = 7,
N, = 2,
T, = 1,
I, = 5,
V = 3,
A, = 8,
R, = 9,
O, = 4,
TRIO = 1954,
SONATA, = 742818,
VIOLA, = 35468,
VIOLIN, = 354652]

```

This puzzle I've found in pySMT examples:

```

from z3 import *

# H+E+L+L+O = W+O+R+L+D = 25

H, E, L, O, W, R, D = Ints ('H, E, L, O, W, R, D')

s=Solver()

s.add(Distinct(H, E, L, O, W, R, D))
s.add(And(H>=1, H<=9))
s.add(And(E>=1, E<=9))
s.add(And(L>=1, L<=9))
s.add(And(O>=1, O<=9))
s.add(And(W>=1, W<=9))
s.add(And(R>=1, R<=9))
s.add(And(D>=1, D<=9))

s.add(H+E+L+L+O == W+O+R+L+D == 25)

print s.check()
print s.model()

```

```

sat
[E, = 4, D = 6, O, = 2, W, = 3, R, = 5, L, = 1, H, = 9]

```

## 7.1 Generating de Bruijn sequences using Z3

The following piece of quite esoteric code calculates number of leading zero bits <sup>25</sup>:

```

int v[64]=
{ -1,31, 8,30, -1, 7,-1,-1, 29,-1,26, 6, -1,-1, 2,-1,
  -1,28,-1,-1, -1,19,25,-1, 5,-1,17,-1, 23,14, 1,-1,
   9,-1,-1,-1, 27,-1, 3,-1, -1,-1,20,-1, 18,24,15,10,
  -1,-1, 4,-1, 21,-1,16,11, -1,22,-1,12, 13,-1, 0,-1 };

int LZCNT(uint32_t x)
{
    x |= x >> 1;
    x |= x >> 2;
    x |= x >> 4;
    x |= x >> 8;
    x |= x >> 16;
    x *= 0x4badf0d;
    return v[x >> 26];
}

```

(This is usually done using simpler algorithm, but it will contain conditional jumps, which is bad for CPUs starting at RISC. There are no conditional jumps in this algorithm.)

Read more about it: [https://yurichev.com/blog/de\\_bruijn/](https://yurichev.com/blog/de_bruijn/). The magic number used here is called *de Bruijn sequence*, and I once used bruteforce to find it (one of the results was `0x4badf0d`, which is used here). But what if we need magic number for 64-bit values? Bruteforce is not an option here.

If you already read about these sequences in my blog or in other sources, you can see that the 32-bit magic number is a number consisting of 5-bit overlapping chunks, and all chunks must be unique, i.e., must not be repeating.

<sup>25</sup>[https://en.wikipedia.org/wiki/Find\\_first\\_set](https://en.wikipedia.org/wiki/Find_first_set)



For 64-bit magic number, these are 6-bit overlapping chunks.

To find the magic number, one can find a Hamiltonian path of a de Bruijn graph. But I've found that Z3 is also can do this, though, overkill, but this is more illustrative.

```
#!/usr/bin/python
from z3 import *

out = BitVec('out', 64)

tmp=[]
for i in range(64):
    tmp.append((out>>i)&0x3F)

s=Solver()

# all overlapping 6-bit chunks must be distinct:
s.add(Distinct(*tmp))
# MSB must be zero:
s.add((out&0x8000000000000000)==0)

print s.check()

result=s.model()[out].as_long()
print "0x%x" % result

# print overlapping 6-bit chunks:
for i in range(64):
    t=(result>>i)&0x3F
    print " "*(63-i) + format(t, 'b').zfill(6)
```

We just enumerate all overlapping 6-bit chunks and tell Z3 that they must be unique (see `Distinct`).  
Output:

```
sat
0x79c52dd0991abf60

                                100000
                                110000
                                011000
                                101100
                                110110
                                111011
                                111101
                                111110
                                111111
                                011111
                                101111
                                010111
                                101011
                                010101
                                101010
                                110101
                                011010
                                001101
                                000110
                                100011
                                010001
                                001000
                                100100
                                110010
                                011001
                                001100
                                100110
                                010011
                                001001
                                000100
                                000010
                                100001
                                010000
                                101000
                                110100
                                111010
                                011101
                                101110
                                110111
                                011011
```

```

        101101
        010110
        001011
        100101
        010010
        101001
        010100
        001010
        000101
        100010
        110001
        111000
        011100
        001110
        100111
        110011
        111001
        111100
        011110
        001111
        000111
        000011
        000001
        000000

```

Overlapping chunks are clearly visible. So the magic number is `0x79c52dd0991abf60`. Let's check:

```

#include <stdint.h>
#include <stdio.h>
#include <assert.h>

#define MAGIC 0x79c52dd0991abf60

int magic_tbl[64];

// returns single bit position counting from LSB
// not works for i==0
int bitpos (uint64_t i)
{
    return magic_tbl[(MAGIC/i) & 0x3F];
};

// count trailing zeroes
// not works for i==0
int tzcnt (uint64_t i)
{
    uint64_t a=i & (-i);
    return magic_tbl[(MAGIC/a) & 0x3F];
};

int main()
{
    // construct magic table
    // may be omitted in production code
    for (int i=0; i<64; i++)
        magic_tbl[(MAGIC/(1ULL<<i)) & 0x3F]=i;

    // test
    for (int i=0; i<64; i++)
    {
        printf ("input=0x%llx, result=%d\n", 1ULL<<i, bitpos (1ULL<<i));
        assert(bitpos(1ULL<<i)==i);
    };
    assert(tzcnt (0xFFFF0000)==16);
    assert(tzcnt (0xFFFF0010)==4);
};

```

That works!

More about de Bruijn sequences: [https://yurichev.com/blog/de\\_bruijn/](https://yurichev.com/blog/de_bruijn/), <https://chessprogramming.wikispaces.com/De+Bruijn+sequence>, <https://chessprogramming.wikispaces.com/De+Bruijn+Sequence+Generator>.

## 7.2 Zebra puzzle (AKA Einstein puzzle)

Zebra puzzle is a popular puzzle, defined as follows:

1. There are five houses.
2. The Englishman lives in the red house.
3. The Spaniard owns the dog.
4. Coffee is drunk in the green house.
5. The Ukrainian drinks tea.
6. The green house is immediately to the right of the ivory house.
7. The Old Gold smoker owns snails.
8. Kools are smoked in the yellow house.
9. Milk is drunk in the middle house.
10. The Norwegian lives in the first house.
11. The man who smokes Chesterfields lives in the house next to the man with the fox.
12. Kools are smoked in the house next to the house where the horse is kept.
13. The Lucky Strike smoker drinks orange juice.
14. The Japanese smokes Parliaments.
15. The Norwegian lives next to the blue house.

Now, who drinks water? Who owns the zebra?

In the interest of clarity, it must be added that each of the five houses is painted a different color, and their inhabitants are of different national extractions, own different pets, drink different beverages and smoke different brands of American cigarettes [sic]. One other thing: in statement 6, right means your right.

( [https://en.wikipedia.org/wiki/Zebra\\_Puzzle](https://en.wikipedia.org/wiki/Zebra_Puzzle) )

It's a very good example of CSP<sup>26</sup>.

We would encode each entity as integer variable, representing number of house.

Then, to define that Englishman lives in red house, we will add this constraint: `Englishman == Red`, meaning that number of a house where Englishmen resides and which is painted in red is the same.

To define that Norwegian lives next to the blue house, we don't really know, if it is at left side of blue house or at right side, but we know that house numbers are different by just 1. So we will define this constraint:

`Norwegian==Blue-1 OR Norwegian==Blue+1`.

We will also need to limit all house numbers, so they will be in range of 1..5.

We will also use `Distinct` to show that all various entities of the same type are all has different house numbers.

```
#!/usr/bin/env python
from z3 import *

Yellow, Blue, Red, Ivory, Green=Ints('Yellow Blue Red Ivory Green')
Norwegian, Ukrainian, Englishman, Spaniard, Japanese=Ints('Norwegian Ukrainian Englishman Spaniard Japanese')
Water, Tea, Milk, OrangeJuice, Coffee=Ints('Water Tea Milk OrangeJuice Coffee')
Kools, Chesterfield, OldGold, LuckyStrike, Parliament=Ints('Kools Chesterfield OldGold LuckyStrike Parliament')
Fox, Horse, Snails, Dog, Zebra=Ints('Fox Horse Snails Dog Zebra')

s = Solver()

# colors are distinct for all 5 houses:
s.add(Distinct(Yellow, Blue, Red, Ivory, Green))

# all nationalities are living in different houses:
s.add(Distinct(Norwegian, Ukrainian, Englishman, Spaniard, Japanese))

# so are beverages:
s.add(Distinct(Water, Tea, Milk, OrangeJuice, Coffee))
```

<sup>26</sup>Constraint satisfaction problem

```

# so are cigarettes:
s.add(Distinct(Kools, Chesterfield, OldGold, LuckyStrike, Parliament))

# so are pets:
s.add(Distinct(Fox, Horse, Snails, Dog, Zebra))

# limits.
# adding two constraints at once (separated by comma) is the same
# as adding one And() constraint with two subconstraints
s.add(Yellow>=1, Yellow<=5)
s.add(Blue>=1, Blue<=5)
s.add(Red>=1, Red<=5)
s.add(Ivory>=1, Ivory<=5)
s.add(Green>=1, Green<=5)

s.add(Norwegian>=1, Norwegian<=5)
s.add(Ukrainian>=1, Ukrainian<=5)
s.add(Englishman>=1, Englishman<=5)
s.add(Spaniard>=1, Spaniard<=5)
s.add(Japanese>=1, Japanese<=5)

s.add(Water>=1, Water<=5)
s.add(Tea>=1, Tea<=5)
s.add(Milk>=1, Milk<=5)
s.add(OrangeJuice>=1, OrangeJuice<=5)
s.add(Coffee>=1, Coffee<=5)

s.add(Kools>=1, Kools<=5)
s.add(Chesterfield>=1, Chesterfield<=5)
s.add(OldGold>=1, OldGold<=5)
s.add(LuckyStrike>=1, LuckyStrike<=5)
s.add(Parliament>=1, Parliament<=5)

s.add(Fox>=1, Fox<=5)
s.add(Horse>=1, Horse<=5)
s.add(Snails>=1, Snails<=5)
s.add(Dog>=1, Dog<=5)
s.add(Zebra>=1, Zebra<=5)

# main constraints of the puzzle:

# 2.The Englishman lives in the red house.
s.add(Englishman==Red)

# 3.The Spaniard owns the dog.
s.add(Spaniard==Dog)

# 4.Coffee is drunk in the green house.
s.add(Coffee==Green)

# 5.The Ukrainian drinks tea.
s.add(Ukrainian==Tea)

# 6.The green house is immediately to the right of the ivory house.
s.add(Green==Ivory+1)

# 7.The Old Gold smoker owns snails.
s.add(OldGold==Snails)

# 8.Kools are smoked in the yellow house.
s.add(Kools==Yellow)

# 9.Milk is drunk in the middle house.
s.add(Milk==3) # i.e., 3rd house

# 10.The Norwegian lives in the first house.
s.add(Norwegian==1)

# 11.The man who smokes Chesterfields lives in the house next to the man with the fox.
s.add(Or(Chesterfield==Fox+1, Chesterfield==Fox-1)) # left or right

# 12.Kools are smoked in the house next to the house where the horse is kept.
s.add(Or(Kools==Horse+1, Kools==Horse-1)) # left or right

# 13.The Lucky Strike smoker drinks orange juice.
s.add(LuckyStrike==OrangeJuice)

```

```
# 14.The Japanese smokes Parliaments.
s.add(Japanese==Parliament)

# 15.The Norwegian lives next to the blue house.
s.add(Or(Norwegian==Blue+1, Norwegian==Blue-1)) # left or right

r=s.check()
print r
if r==unsat:
    exit(0)
m=s.model()
print(m)
```

When we run it, we got correct result:

```
sat
[Snails = 3,
Blue = 2,
Ivory = 4,
OrangeJuice = 4,
Parliament = 5,
Yellow = 1,
Fox = 1,
Zebra = 5,
Horse = 2,
Dog = 4,
Tea = 2,
Water = 1,
Chesterfield = 2,
Red = 3,
Japanese = 5,
LuckyStrike = 4,
Norwegian = 1,
Milk = 3,
Kools = 1,
OldGold = 3,
Ukrainian = 2,
Coffee = 5,
Green = 5,
Spaniard = 4,
Englishman = 3]
```

## 7.3 Sudoku puzzle

Sudoku puzzle is a 9\*9 grid with some cells filled with values, some are empty:

		5	3					
8							2	
	7			1		5		
4					5	3		
	1			7				6
		3	2				8	
	6		5					9
		4					3	
					9	7		

Unsolved Sudoku

Numbers of each row must be unique, i.e., it must contain all 9 numbers in range of 1..9 without repetition. Same story for each column and also for each 3\*3 square.

This puzzle is good candidate to try [SMT](#) solver on, because it's essentially an unsolved system of equations.

### 7.3.1 The first idea

The only thing we must decide is that how to determine in one expression, if the input 9 variables has all 9 unique numbers? They are not ordered or sorted, after all.

From the school-level arithmetics, we can devise this idea:

$$\underbrace{10^{i_1} + 10^{i_2} + \dots + 10^{i_9}}_9 = 1111111110 \quad (1)$$

Take each input variable, calculate  $10^{i_i}$  and sum them all. If all input values are unique, each will be settled at its own place. Even more than that: there will be no holes, i.e., no skipped values. So, in case of Sudoku, 1111111110 number will be final result, indicating that all 9 input values are unique, in range of 1..9.

Exponentiation is heavy operation, can we use binary operations? Yes, just replace 10 with 2:

$$\underbrace{2^{i_1} + 2^{i_2} + \dots + 2^{i_9}}_9 = 1111111110_2 \quad (2)$$

The effect is just the same, but the final value is in base 2 instead of 10.

Now a working example:

```
import sys
from z3 import *

"""
coordinates:
-----
00 01 02 | 03 04 05 | 06 07 08
10 11 12 | 13 14 15 | 16 17 18
20 21 22 | 23 24 25 | 26 27 28
-----
30 31 32 | 33 34 35 | 36 37 38
40 41 42 | 43 44 45 | 46 47 48
50 51 52 | 53 54 55 | 56 57 58
-----
60 61 62 | 63 64 65 | 66 67 68
70 71 72 | 73 74 75 | 76 77 78
80 81 82 | 83 84 85 | 86 87 88
-----
"""

s=Solver()

# using Python list comprehension, construct array of arrays of BitVec instances:
cells=[[BitVec('cell%d%d' % (r, c), 16) for c in range(9)] for r in range(9)]

# http://www.norvig.com/sudoku.html
# http://www.mirror.co.uk/news/weird-news/worlds-hardest-sudoku-can-you-242294
puzzle="..53....8.....2..7..1.5..4....53...1..7...6..32...8..6.5....9..4....3.....97.."

# process text line:
current_column=0
current_row=0
for i in puzzle:
    if i!='.':
        s.add(cells[current_row][current_column]==BitVecVal(int(i),16))
        current_column=current_column+1
    if current_column==9:
        current_column=0
        current_row=current_row+1

one=BitVecVal(1,16)
mask=BitVecVal(0b1111111110,16)

# for all 9 rows
for r in range(9):
    s.add(((one<<cells[r][0]) +
            (one<<cells[r][1]) +
            (one<<cells[r][2]) +
            (one<<cells[r][3]) +
            (one<<cells[r][4]) +
            (one<<cells[r][5]) +
            (one<<cells[r][6]) +
            (one<<cells[r][7]) +
```

```

        (one<<cells[r][8]))==mask)

# for all 9 columns
for c in range(9):
    s.add(((one<<cells[0][c]) +
            (one<<cells[1][c]) +
            (one<<cells[2][c]) +
            (one<<cells[3][c]) +
            (one<<cells[4][c]) +
            (one<<cells[5][c]) +
            (one<<cells[6][c]) +
            (one<<cells[7][c]) +
            (one<<cells[8][c]))==mask)

# enumerate all 9 squares
for r in range(0, 9, 3):
    for c in range(0, 9, 3):
        # add constraints for each 3*3 square:
        s.add((one<<cells[r+0][c+0]) +
                (one<<cells[r+0][c+1]) +
                (one<<cells[r+0][c+2]) +
                (one<<cells[r+1][c+0]) +
                (one<<cells[r+1][c+1]) +
                (one<<cells[r+1][c+2]) +
                (one<<cells[r+2][c+0]) +
                (one<<cells[r+2][c+1]) +
                (one<<cells[r+2][c+2]))==mask)

#print s.check()
s.check()
#print s.model()
m=s.model()

for r in range(9):
    for c in range(9):
        sys.stdout.write (str(m[cells[r][c]])+" ")
    print ""

```

( [https://github.com/dennis714/SAT\\_SMT\\_article/blob/master/SMT/sudoku\\_plus.py](https://github.com/dennis714/SAT_SMT_article/blob/master/SMT/sudoku_plus.py) )

```

% time python sudoku_plus.py
1 4 5 3 2 7 6 9 8
8 3 9 6 5 4 1 2 7
6 7 2 9 1 8 5 4 3
4 9 6 1 8 5 3 7 2
2 1 8 4 7 3 9 5 6
7 5 3 2 9 6 4 8 1
3 6 7 5 4 2 8 1 9
9 8 4 7 6 1 2 3 5
5 2 1 8 3 9 7 6 4

real    0m11.717s
user    0m10.896s
sys      0m0.068s

```

Even more, we can replace summing operation to logical OR:

$$\underbrace{2^{i_1} \vee 2^{i_2} \vee \dots \vee 2^{i_9}}_9 = 111111110_2 \quad (3)$$

```

import sys
from z3 import *

"""
coordinates:
-----
00 01 02 | 03 04 05 | 06 07 08
10 11 12 | 13 14 15 | 16 17 18
20 21 22 | 23 24 25 | 26 27 28
-----
30 31 32 | 33 34 35 | 36 37 38
40 41 42 | 43 44 45 | 46 47 48
50 51 52 | 53 54 55 | 56 57 58
-----
60 61 62 | 63 64 65 | 66 67 68

```

```

70 71 72 | 73 74 75 | 76 77 78
80 81 82 | 83 84 85 | 86 87 88
-----
"""

s=Solver()

# using Python list comprehension, construct array of arrays of BitVec instances:
cells=[[BitVec('cell%d%d' % (r, c), 16) for c in range(9)] for r in range(9)]

# http://www.norvig.com/sudoku.html
# http://www.mirror.co.uk/news/weird-news/worlds-hardest-sudoku-can-you-242294
puzzle="..53....8.....2..7..1.5..4....53...1..7...6..32...8..6.5....9..4....3.....97.."

# process text line:
current_column=0
current_row=0
for i in puzzle:
    if i!='.':
        s.add(cells[current_row][current_column]==BitVecVal(int(i),16))
        current_column=current_column+1
    if current_column==9:
        current_column=0
        current_row=current_row+1

one=BitVecVal(1,16)
mask=BitVecVal(0b111111110,16)

# for all 9 rows
for r in range(9):
    s.add(((one<<cells[r][0]) |
            (one<<cells[r][1]) |
            (one<<cells[r][2]) |
            (one<<cells[r][3]) |
            (one<<cells[r][4]) |
            (one<<cells[r][5]) |
            (one<<cells[r][6]) |
            (one<<cells[r][7]) |
            (one<<cells[r][8]))==mask)

# for all 9 columns
for c in range(9):
    s.add(((one<<cells[0][c]) |
            (one<<cells[1][c]) |
            (one<<cells[2][c]) |
            (one<<cells[3][c]) |
            (one<<cells[4][c]) |
            (one<<cells[5][c]) |
            (one<<cells[6][c]) |
            (one<<cells[7][c]) |
            (one<<cells[8][c]))==mask)

# enumerate all 9 squares
for r in range(0, 9, 3):
    for c in range(0, 9, 3):
        # add constraints for each 3*3 square:
        s.add(one<<cells[r+0][c+0] |
              one<<cells[r+0][c+1] |
              one<<cells[r+0][c+2] |
              one<<cells[r+1][c+0] |
              one<<cells[r+1][c+1] |
              one<<cells[r+1][c+2] |
              one<<cells[r+2][c+0] |
              one<<cells[r+2][c+1] |
              one<<cells[r+2][c+2]==mask)

#print s.check()
s.check()
#print s.model()
m=s.model()

for r in range(9):
    for c in range(9):
        sys.stdout.write (str(m[cells[r][c]])+" ")
    print ""

```



( [https://github.com/dennis714/SAT\\_SMT\\_article/blob/master/SMT/sudoku\\_or.py](https://github.com/dennis714/SAT_SMT_article/blob/master/SMT/sudoku_or.py) )  
Now it works much faster. Z3 handles OR operation over bit vectors better than addition?

```
% time python sudoku_or.py
1 4 5 3 2 7 6 9 8
8 3 9 6 5 4 1 2 7
6 7 2 9 1 8 5 4 3
4 9 6 1 8 5 3 7 2
2 1 8 4 7 3 9 5 6
7 5 3 2 9 6 4 8 1
3 6 7 5 4 2 8 1 9
9 8 4 7 6 1 2 3 5
5 2 1 8 3 9 7 6 4

real    0m1.429s
user    0m1.393s
sys      0m0.036s
```

The puzzle I used as example is dubbed as one of the hardest known <sup>27</sup> (well, for humans). It took  $\approx 1.4$  seconds on my Intel Core i3-3110M 2.4GHz notebook to solve it.

### 7.3.2 The second idea

My first approach is far from effective, I did what first came to my mind and worked. Another approach is to use `distinct` command from SMT-LIB, which tells Z3 that some variables must be distinct (or unique). This command is also available in Z3 Python interface.

I've rewritten my first Sudoku solver, now it operates over *Int sort*, it has `distinct` commands instead of bit operations, and now also other constraint added: each cell value must be in 1..9 range, because, otherwise, Z3 will offer (although correct) solution with too big and/or negative numbers.

```
import sys
from z3 import *

"""
-----
00 01 02 | 03 04 05 | 06 07 08
10 11 12 | 13 14 15 | 16 17 18
20 21 22 | 23 24 25 | 26 27 28
-----
30 31 32 | 33 34 35 | 36 37 38
40 41 42 | 43 44 45 | 46 47 48
50 51 52 | 53 54 55 | 56 57 58
-----
60 61 62 | 63 64 65 | 66 67 68
70 71 72 | 73 74 75 | 76 77 78
80 81 82 | 83 84 85 | 86 87 88
-----
"""

s=Solver()

# using Python list comprehension, construct array of arrays of BitVec instances:
cells=[[Int('cell%d%d' % (r, c)) for c in range(9)] for r in range(9)]

# http://www.norvig.com/sudoku.html
# http://www.mirror.co.uk/news/weird-news/worlds-hardest-sudoku-can-you-242294
puzzle="..53.....8.....2..7..1.5..4....53...1..7...6..32...8..6.5....9..4....3.....97.."

# process text line:
current_column=0
current_row=0
for i in puzzle:
    if i!='.':
        s.add(cells[current_row][current_column]==int(i))
        current_column=current_column+1
        if current_column==9:
            current_column=0
            current_row=current_row+1

# this is important, because otherwise, Z3 will report correct solutions with too big and/or negative numbers in
cells
```

<sup>27</sup><http://www.mirror.co.uk/news/weird-news/worlds-hardest-sudoku-can-you-242294>

```

for r in range(9):
    for c in range(9):
        s.add(cells[r][c]>=1)
        s.add(cells[r][c]<=9)

# for all 9 rows
for r in range(9):
    s.add(Distinct(cells[r][0],
                    cells[r][1],
                    cells[r][2],
                    cells[r][3],
                    cells[r][4],
                    cells[r][5],
                    cells[r][6],
                    cells[r][7],
                    cells[r][8]))

# for all 9 columns
for c in range(9):
    s.add(Distinct(cells[0][c],
                    cells[1][c],
                    cells[2][c],
                    cells[3][c],
                    cells[4][c],
                    cells[5][c],
                    cells[6][c],
                    cells[7][c],
                    cells[8][c]))

# enumerate all 9 squares
for r in range(0, 9, 3):
    for c in range(0, 9, 3):
        # add constraints for each 3*3 square:
        s.add(Distinct(cells[r+0][c+0],
                        cells[r+0][c+1],
                        cells[r+0][c+2],
                        cells[r+1][c+0],
                        cells[r+1][c+1],
                        cells[r+1][c+2],
                        cells[r+2][c+0],
                        cells[r+2][c+1],
                        cells[r+2][c+2]))

#print s.check()
s.check()
#print s.model()
m=s.model()

for r in range(9):
    for c in range(9):
        sys.stdout.write (str(m[cells[r][c]])+" ")
    print ""

```

( [https://github.com/dennis714/SAT\\_SMT\\_article/blob/master/SMT/sudoku2.py](https://github.com/dennis714/SAT_SMT_article/blob/master/SMT/sudoku2.py) )

```

% time python sudoku2.py
1 4 5 3 2 7 6 9 8
8 3 9 6 5 4 1 2 7
6 7 2 9 1 8 5 4 3
4 9 6 1 8 5 3 7 2
2 1 8 4 7 3 9 5 6
7 5 3 2 9 6 4 8 1
3 6 7 5 4 2 8 1 9
9 8 4 7 6 1 2 3 5
5 2 1 8 3 9 7 6 4

real    0m0.382s
user    0m0.346s
sys     0m0.036s

```

That's much faster.

### 7.3.3 Conclusion

SMT-solvers are so helpful, is that our Sudoku solver has nothing else, we have just defined relationships between variables (cells).

### 7.3.4 Homework

As it seems, true Sudoku puzzle is the one which has only one solution. The piece of code I've included here shows only the first one. Using the method described earlier (7.4, also called “model counting”), try to find more solutions, or prove that the solution you have just found is the only one possible.

### 7.3.5 Further reading

<http://www.norvig.com/sudoku.html>

### 7.3.6 Sudoku as a SAT problem

It's also possible to represent Sudoku puzzle as a huge CNF equation and use SAT-solver to find solution, but it's just trickier.

Some articles about it: *Building a Sudoku Solver with SAT*<sup>28</sup>, Tjark Weber, *A SAT-based Sudoku Solver*<sup>29</sup>, Ines Lynce, Joel Ouaknine, *Sudoku as a SAT Problem*<sup>30</sup>, Gihwon Kwon, Himanshu Jain, *Optimized CNF Encoding for Sudoku Puzzles*<sup>31</sup>.

SMT-solver can also use SAT-solver in its core, so it does all mundane translating work. As a “compiler”, it may not do this in the most efficient way, though.

## 7.4 Solving Problem Euler 31: “Coin sums”

(This text was first published in my blog<sup>32</sup> at 10-May-2013.)

In England the currency is made up of pound, £, and pence, p, and there are eight coins in general circulation:

1p, 2p, 5p, 10p, 20p, 50p, £1 (100p) and £2 (200p). It is possible to make £2 in the following way:

1£1 + 150p + 220p + 15p + 12p + 31p How many different ways can £2 be made using any number of coins?

( [Problem Euler 31 — Coin sums](#) )

Using Z3 for solving this is overkill, and also slow, but nevertheless, it works, showing all possible solutions as well. The piece of code for blocking already found solution and search for next, and thus, counting all solutions, was taken from Stack Overflow answer<sup>33</sup>. This is also called “model counting”. Constraints like “a>=0” must be present, because Z3 solver will find solutions with negative numbers.

```
#!/usr/bin/python

from z3 import *

a,b,c,d,e,f,g,h = Ints('a b c d e f g h')
s = Solver()
s.add(1*a + 2*b + 5*c + 10*d + 20*e + 50*f + 100*g + 200*h == 200,
      a>=0, b>=0, c>=0, d>=0, e>=0, f>=0, g>=0, h>=0)
result=[]
```

<sup>28</sup>[http://ocw.mit.edu/courses/electrical-engineering-and-computer-science/6-005-elements-of-software-construction-fall-assignments/MIT6\\_005F11\\_ps4.pdf](http://ocw.mit.edu/courses/electrical-engineering-and-computer-science/6-005-elements-of-software-construction-fall-assignments/MIT6_005F11_ps4.pdf)

<sup>29</sup><https://www.lri.fr/~conchon/mpri/weber.pdf>

<sup>30</sup><http://sat.inesc-id.pt/~ines/publications/aimath06.pdf>

<sup>31</sup><http://www.cs.cmu.edu/~hjain/papers/sudoku-as-SAT.pdf>

<sup>32</sup><http://dennisyurichev.blogspot.de/2013/05/in-england-currency-is-made-up-of-pound.html>

<sup>33</sup><http://stackoverflow.com/questions/11867611/z3py-checking-all-solutions-for-equation>, another question: <http://stackoverflow.com/questions/13395391/z3-finding-all-satisfying-models>

```

while True:
    if s.check() == sat:
        m = s.model()
        print m
        result.append(m)
        # Create a new constraint the blocks the current model
        block = []
        for d in m:
            # d is a declaration
            if d.arity() > 0:
                raise Z3Exception("uninterpreted functions are not supported")
            # create a constant from declaration
            c=d()
            #print c, m[d]
            if is_array(c) or c.sort().kind() == Z3_UNINTERPRETED_SORT:
                raise Z3Exception("arrays and uninterpreted sorts are not supported")
            block.append(c != m[d])
        #print "new constraint:",block
        s.add(Or(block))
    else:
        print len(result)
        break

```

Works very slow, and this is what it produces:

```

[h = 0, g = 0, f = 0, e = 0, d = 0, c = 0, b = 0, a = 200]
[f = 1, b = 5, a = 0, d = 1, g = 1, h = 0, c = 2, e = 1]
[f = 0, b = 1, a = 153, d = 0, g = 0, h = 0, c = 1, e = 2]
...
[f = 0, b = 31, a = 33, d = 2, g = 0, h = 0, c = 17, e = 0]
[f = 0, b = 30, a = 35, d = 2, g = 0, h = 0, c = 17, e = 0]
[f = 0, b = 5, a = 50, d = 2, g = 0, h = 0, c = 24, e = 0]

```

73682 results in total.

## 7.5 Using Z3 theorem prover to prove equivalence of some weird alternative to XOR operation

(The test was first published in my blog at April 2015: <http://blog.yurichev.com/node/86>).

There is a “A Hacker’s Assistant” program<sup>34</sup> (*Aha!*) written by Henry Warren, who is also the author of the great “Hacker’s Delight” book.

The *Aha!* program is essentially *superoptimizer*<sup>35</sup>, which blindly brute-force a list of some generic RISC CPU instructions to achieve shortest possible (and jumpless or branch-free) CPU code sequence for desired operation. For example, *Aha!* can find jumpless version of `abs()` function easily.

Compiler developers use superoptimization to find shortest possible (and/or jumpless) code, but I tried to do otherwise—to find longest code for some primitive operation. I tried *Aha!* to find equivalent of basic XOR operation without usage of the actual XOR instruction, and the most bizarre example *Aha!* gave is:

```

Found a 4-operation program:
add    r1,ry,rx
and     r2,ry,rx
mul     r3,r2,-2
add     r4,r3,r1
Expr: (((y & x)*-2) + (y + x))

```

And it’s hard to say, why/where we can use it, maybe for obfuscation, I’m not sure. I would call this *suboptimization* (as opposed to *superoptimization*). Or maybe *superdeoptimization*.

But my another question was also, is it possible to prove that this is correct formula at all? The *Aha!* checking some input/output values against XOR operation, but of course, not all the possible values. It is 32-bit code, so it may take very long time to try all possible 32-bit inputs to test it.

We can try Z3 theorem prover for the job. It’s called *prover*, after all.

So I wrote this:

```

#!/usr/bin/python
from z3 import *

x = BitVec('x', 32)

```

<sup>34</sup><http://www.hackersdelight.org/>

<sup>35</sup><http://en.wikipedia.org/wiki/Superoptimization>

```

y = BitVec('y', 32)
output = BitVec('output', 32)
s = Solver()
s.add(x^y==output)
s.add(((y & x)*0xFFFFFFE) + (y + x)!=output)
print s.check()

```

In plain English language, this means “are there any case for  $x$  and  $y$  where  $x \oplus y$  doesn't equals to  $((y \& x) * -2) + (y + x)$ ?” ...and Z3 prints “unsat”, meaning, it can't find any counterexample to the equation. So this *Aha!* result is proved to be working just like XOR operation.

Oh, I also tried to extend the formula to 64 bit:

```

#!/usr/bin/python
from z3 import *

x = BitVec('x', 64)
y = BitVec('y', 64)
output = BitVec('output', 64)
s = Solver()
s.add(x^y==output)
s.add(((y & x)*0xFFFFFFE) + (y + x)!=output)
print s.check()

```

Nope, now it says “sat”, meaning, Z3 found at least one counterexample. Oops, it's because I forgot to extend -2 number to 64-bit value:

```

#!/usr/bin/python
from z3 import *

x = BitVec('x', 64)
y = BitVec('y', 64)
output = BitVec('output', 64)
s = Solver()
s.add(x^y==output)
s.add(((y & x)*0xFFFFFFFFFFE) + (y + x)!=output)
print s.check()

```

Now it says “unsat”, so the formula given by *Aha!* works for 64-bit code as well.

### 7.5.1 In SMT-LIB form

Now we can rephrase our expression to more suitable form:  $(x + y - ((x \& y) << 1))$ . It also works well in Z3Py:

```

#!/usr/bin/python
from z3 import *

x = BitVec('x', 64)
y = BitVec('y', 64)
output = BitVec('output', 64)
s = Solver()
s.add(x^y==output)
s.add((x + y - ((x & y)<<1)) != output)
print s.check()

```

Here is how to define it in SMT-LIB way:

```

(declare-const x (_ BitVec 64))
(declare-const y (_ BitVec 64))
(assert
  (not
    (=
      (bvsb
        (bvadd x y)
        (bvshl (bvand x y) (_ bv1 64)))
      (bvxor x y)
    )
  )
)
(check-sat)

```

### 7.5.2 Using universal quantifier

Z3 supports universal quantifier `exists`, which is true if at least one set of variables satisfied underlying condition:

```
(declare-const x (_ BitVec 64))
(declare-const y (_ BitVec 64))
(assert
  (exists ((x (_ BitVec 64)) (y (_ BitVec 64)))
    (not (=
      (bvsub
        (bvadd x y)
        (bvshl (bvand x y) (_ bv1 64))
      )
      (bvxor x y)
    ))
  )
)
(check-sat)
```

It returns “unsat”, meaning, Z3 couldn’t find any counterexample of the equation, i.e., it’s not exist.

This is also known as  $\exists$  in mathematical logic lingo.

Z3 also supports universal quantifier `forall`, which is true if the equation is true for all possible values. So we can rewrite our SMT-LIB example as:

```
(declare-const x (_ BitVec 64))
(declare-const y (_ BitVec 64))
(assert
  (forall ((x (_ BitVec 64)) (y (_ BitVec 64)))
    (=
      (bvsub
        (bvadd x y)
        (bvshl (bvand x y) (_ bv1 64))
      )
      (bvxor x y)
    )
  )
)
(check-sat)
```

It returns “sat”, meaning, the equation is correct for all possible 64-bit `x` and `y` values, like them all were checked.

Mathematically speaking:  $\forall n \in \mathbb{N} (x \oplus y = (x + y - ((x \& y) << 1)))$  <sup>36</sup>

### 7.5.3 How the expression works

First of all, binary addition can be viewed as binary XORing with carrying (13.2). Here is an example: let’s add 2 (10b) and 2 (10b). XORing these two values resulting 0, but there is a carry generated during addition of two second bits. That carry bit is propagated further and settles at the place of the 3rd bit: 100b. 4 (100b) is hence a final result of addition.

If the carry bits are not generated during addition, the addition operation is merely XORing. For example, let’s add 1 (1b) and 2 (10b).  $1 + 2$  equals to 3, but  $1 \oplus 2$  is also 3.

If the addition is XORing plus carry generation and application, we should eliminate effect of carrying somehow here. The first part of the expression  $(x + y)$  is addition, the second  $((x \& y) << 1)$  is just calculation of every carry bit which was used during addition. If to subtract carry bits from the result of addition, the only XOR effect is left then.

It’s hard to say how Z3 proves this: maybe it just simplifies the equation down to single XOR using simple boolean algebra rewriting rules?

### 7.5.4 In SAT

See also: 13.11.

<sup>36</sup>  $\forall$  means *equation must be true for all possible values*, which are choosen from natural numbers ( $\mathbb{N}$ ).

## 7.6 Dietz's formula

One of the impressive examples of *Aha!* work is finding of Dietz's formula<sup>37</sup>, which is the code of computing average number of two numbers without overflow (which is important if you want to find average number of numbers like 0xFFFFF00 and so on, using 32-bit registers).

Taking this in input:

```
int userfun(int x, int y) {    // To find Dietz's formula for
                               // the floor-average of two
                               // unsigned integers.
    return ((unsigned long long)x + (unsigned long long)y) >> 1;
}
```

...the *Aha!* gives this:

```
Found a 4-operation program:
and  r1,ry,rx
xor  r2,ry,rx
shrs r3,r2,1
add  r4,r3,r1
Expr: (((y ^ x) >>s 1) + (y & x))
```

And it works correctly<sup>38</sup>. But how to prove it?

We will place Dietz's formula on the left side of equation and  $x + y/2$  (or  $x + y >> 1$ ) on the right side:

$$\forall n \in 0..2^{64} - 1. (x \& y) + (x \oplus y) >> 1 = x + y >> 1$$

One important thing is that we can't operate on 64-bit values on right side, because result will overflow. So we will zero extend inputs on right side by 1 bit (in other words, we will just 1 zero bit before each value). The result of Dietz's formula will also be extended by 1 bit. Hence, both sides of the equation will have a width of 65 bits:

```
(declare-const x (_ BitVec 64))
(declare-const y (_ BitVec 64))
(assert
  (forall ((x (_ BitVec 64)) (y (_ BitVec 64)))
    (=
      ((_ zero_extend 1)
        (bvadd
          (bvand x y)
          (bvlshr (bvxor x y) (_ bv1 64))
        ))
      (bvlshr
        (bvadd ((_ zero_extend 1) x) ((_ zero_extend 1) y))
        (_ bv1 65)
      )
    )
  )
)
(check-sat)
```

Z3 says "sat".

65 bits are enough, because the result of addition of two biggest 64-bit values has width of 65 bits:

0xFF...FF + 0xFF...FF = 0x1FF...FE .

As in previous example about XOR equivalent, `(not (= ... ))` and `exists` can also be used here instead of `forall` .

## 7.7 Cracking LCG with Z3

(This text is first appeared in my blog in June 2015 at <http://yurichev.com/blog/modulo/>.)

<sup>37</sup><http://aggregate.org/MAGIC/#Average%20of%20Integers>

<sup>38</sup>For those who interesting how it works, its mechanics is closely related to the weird XOR alternative we just saw. That's why I placed these two pieces of text one after another.

There are well-known weaknesses of LCG<sup>39</sup>, but let's see, if it would be possible to crack it straightforwardly, without any special knowledge. We will define all relations between LCG states in terms of Z3. Here is a test program:

```
#include <stdlib.h>
#include <stdio.h>
#include <time.h>

int main()
{
    int i;

    srand(time(NULL));

    for (i=0; i<10; i++)
        printf ("%d\n", rand()%100);
};
```

It is printing 10 pseudorandom numbers in 0..99 range:

```
37
29
74
95
98
40
23
58
61
17
```

Let's say we are observing only 8 of these numbers (from 29 to 61) and we need to predict next one (17) and/or previous one (37).

The program is compiled using MSVC 2013 (I choose it because its LCG is simpler than that in Glib):

```
.text:0040112E rand      proc near
.text:0040112E          call    __getptd
.text:00401133          imul    ecx, [eax+0x14], 214013
.text:0040113A          add     ecx, 2531011
.text:00401140          mov     [eax+14h], ecx
.text:00401143          shr     ecx, 16
.text:00401146          and     ecx, 7FFFh
.text:0040114C          mov     eax, ecx
.text:0040114E          retn
.text:0040114E rand      endp
```

Let's define LCG in Z3Py:

```
#!/usr/bin/python
from z3 import *

output_prev = BitVec('output_prev', 32)
state1 = BitVec('state1', 32)
state2 = BitVec('state2', 32)
state3 = BitVec('state3', 32)
state4 = BitVec('state4', 32)
state5 = BitVec('state5', 32)
state6 = BitVec('state6', 32)
state7 = BitVec('state7', 32)
state8 = BitVec('state8', 32)
state9 = BitVec('state9', 32)
state10 = BitVec('state10', 32)
output_next = BitVec('output_next', 32)

s = Solver()

s.add(state2 == state1*214013+2531011)
s.add(state3 == state2*214013+2531011)
s.add(state4 == state3*214013+2531011)
s.add(state5 == state4*214013+2531011)
s.add(state6 == state5*214013+2531011)
```

<sup>39</sup>[http://en.wikipedia.org/wiki/Linear\\_congruential\\_generator#Advantages\\_and\\_disadvantages\\_of\\_LCGs](http://en.wikipedia.org/wiki/Linear_congruential_generator#Advantages_and_disadvantages_of_LCGs), <http://www.reteam.org/papers/e59.pdf>, <http://stackoverflow.com/questions/8569113/why-1103515245-is-used-in-rand-8574774#8574774>



```

s.add(state7 == state6*214013+2531011)
s.add(state8 == state7*214013+2531011)
s.add(state9 == state8*214013+2531011)
s.add(state10 == state9*214013+2531011)

s.add(output_prev==URem((state1>>16)&0x7FFF,100))
s.add(URem((state2>>16)&0x7FFF,100)==29)
s.add(URem((state3>>16)&0x7FFF,100)==74)
s.add(URem((state4>>16)&0x7FFF,100)==95)
s.add(URem((state5>>16)&0x7FFF,100)==98)
s.add(URem((state6>>16)&0x7FFF,100)==40)
s.add(URem((state7>>16)&0x7FFF,100)==23)
s.add(URem((state8>>16)&0x7FFF,100)==58)
s.add(URem((state9>>16)&0x7FFF,100)==61)
s.add(output_next==URem((state10>>16)&0x7FFF,100))

print(s.check())
print(s.model())

```

*URem* states for *unsigned remainder*. It works for some time and gave us correct result!

```

sat
[state3 = 2276903645,
 state4 = 1467740716,
 state5 = 3163191359,
 state7 = 4108542129,
 state8 = 2839445680,
 state2 = 998088354,
 state6 = 4214551046,
 state1 = 1791599627,
 state9 = 548002995,
 output_next = 17,
 output_prev = 37,
 state10 = 1390515370]

```

I added  $\approx 10$  states to be sure result will be correct. It may be not in case of smaller set of information.

That is the reason why [LCG](#) is not suitable for any security-related task. This is why cryptographically secure pseudorandom number generators exist: they are designed to be protected against such simple attack. Well, at least if [NSA](#)<sup>40</sup> don't get involved <sup>41</sup>.

Security tokens like “RSA SecurID” can be viewed just as [CPRNG](#)<sup>42</sup> with a secret seed. It shows new pseudorandom number each minute, and the server can predict it, because it knows the seed. Imagine if such token would implement [LCG](#)—it would be much easier to break!

## 7.8 Solving pipe puzzle using Z3 SMT-solver

“Pipe puzzle” is a popular puzzle (just google “pipe puzzle” and look at images).

This is how shuffled puzzle looks like:

<sup>40</sup>National Security Agency

<sup>41</sup>[https://en.wikipedia.org/wiki/Dual\\_EC\\_DRBG](https://en.wikipedia.org/wiki/Dual_EC_DRBG)

<sup>42</sup>Cryptographically Secure Pseudorandom Number Generator




Figure 2: Shuffled puzzle

...and solved:





Figure 3: Solved puzzle

Let's try to find a way to solve it.

### 7.8.1 Generation

First, we need to generate it. Here is my quick idea on it. Take 8\*16 array of cells. Each cell may contain some type of block. There are joints between cells:



Blue lines are horizontal joints, red lines are vertical joints. We just set each joint to 0/false (absent) or 1/true (present), randomly.

Once set, it's now easy to find type for each cell. There are:

joints	our internal name	angle	symbol
0	type 0	0°	(space)
2	type 2a	0°	
2	type 2a	90°	-
2	type 2b	0°	┌
2	type 2b	90°	┐
2	type 2b	180°	└
2	type 2b	270°	┘
3	type 3	0°	┌┐
3	type 3	90°	┐┌
3	type 3	180°	└└
3	type 3	270°	┘┘
4	type 4	0°	┌┐┌┐

*Dangling* joints can be preset at a first stage (i.e., cell with only one joint), but they are removed recursively, these cells are transforming into empty cells. Hence, at the end, all cells has at least two joints, and the whole plumbing system has no connections with outer world—I hope this would make things clearer.

The C source code of generator is here: [https://github.com/dennis714/SAT\\_SMT\\_article/tree/master/SMT/pipe/generator](https://github.com/dennis714/SAT_SMT_article/tree/master/SMT/pipe/generator). All horizontal joints are stored in the global array *hjoints[]* and vertical in *vjoints[]*.

The C program generates ANSI-colored output like it has been showed above (7.8, 7.8) plus an array of types, with no angle information about each cell:

```
[
["0", "0", "2b", "3", "2a", "2a", "2a", "3", "3", "2a", "3", "2b", "2b", "2b", "0", "0"],
["2b", "2b", "3", "2b", "0", "0", "2b", "3", "3", "3", "3", "4", "2b", "0", "0"],
["3", "4", "2b", "0", "0", "0", "3", "2b", "2b", "4", "2b", "3", "4", "2b", "2b"],
["2b", "4", "3", "2a", "3", "3", "3", "2b", "2b", "3", "3", "3", "2a", "2b", "4", "3"],
["0", "2b", "3", "2b", "3", "4", "2b", "3", "3", "2b", "3", "3", "3", "0", "2a", "2a"],
["0", "0", "2b", "2b", "0", "3", "3", "4", "3", "4", "3", "3", "3", "2b", "3", "3"],
["0", "2b", "3", "2b", "0", "3", "3", "4", "3", "4", "3", "0", "3", "4", "3"],
["0", "2b", "3", "3", "2a", "3", "2b", "2b", "3", "3", "3", "3", "2a", "3", "3", "2b"],
]
```

## 7.8.2 Solving

First of all, we would think about 8\*16 array of cells, where each has four bits: "T" (top), "B" (bottom), "L" (left), "R" (right). Each bit represents half of joint.

	0]	1]	2]	3]	4]	5]	6]	7]	8]	9]	10]	11]	12]	13]	14]	15]
	[...]	[...]	[...]	[...]	[...]	[...]	[...]	[...]	[...]	[...]	[...]	[...]	[...]	[...]	[...]	[...]
[0, ...]	L	T	R	L	T	R	L	T	R	L	T	R	L	T	R	L
[1, ...]	L	T	R	L	T	R	L	T	R	L	T	R	L	T	R	L
[2, ...]	L	T	R	L	T	R	L	T	R	L	T	R	L	T	R	L
[3, ...]	L	T	R	L	T	R	L	T	R	L	T	R	L	T	R	L
[4, ...]	L	T	R	L	T	R	L	T	R	L	T	R	L	T	R	L
[5, ...]	L	T	R	L	T	R	L	T	R	L	T	R	L	T	R	L
[6, ...]	L	T	R	L	T	R	L	T	R	L	T	R	L	T	R	L
[7, ...]	L	T	R	L	T	R	L	T	R	L	T	R	L	T	R	L

Now we define arrays of each of four half-joints + angle information:

```
HEIGHT=8
WIDTH=16

# if T/B/R/L is Bool instead of Int, Z3 solver will work faster
T=[[Bool('cell_%d_%d_top' % (r, c)) for c in range(WIDTH)] for r in range(HEIGHT)]
B=[[Bool('cell_%d_%d_bottom' % (r, c)) for c in range(WIDTH)] for r in range(HEIGHT)]
R=[[Bool('cell_%d_%d_right' % (r, c)) for c in range(WIDTH)] for r in range(HEIGHT)]
L=[[Bool('cell_%d_%d_left' % (r, c)) for c in range(WIDTH)] for r in range(HEIGHT)]
A=[[Int('cell_%d_%d_angle' % (r, c)) for c in range(WIDTH)] for r in range(HEIGHT)]
```

We know that if each of half-joints is present, corresponding half-joint must be also present, and vice versa. We define this using these constraints:

```
# shorthand variables for True and False:
t=True
f=False

# "top" of each cell must be equal to "bottom" of the cell above
# "bottom" of each cell must be equal to "top" of the cell below
# "left" of each cell must be equal to "right" of the cell at left
# "right" of each cell must be equal to "left" of the cell at right
for r in range(HEIGHT):
    for c in range(WIDTH):
        if r!=0:
            s.add(T[r][c]==B[r-1][c])
        if r!=HEIGHT-1:
            s.add(B[r][c]==T[r+1][c])
        if c!=0:
            s.add(L[r][c]==R[r][c-1])
        if c!=WIDTH-1:
            s.add(R[r][c]==L[r][c+1])

# "left" of each cell of first column shouldn't have any connection
# so is "right" of each cell of the last column
for r in range(HEIGHT):
    s.add(L[r][0]==f)
    s.add(R[r][WIDTH-1]==f)

# "top" of each cell of the first row shouldn't have any connection
# so is "bottom" of each cell of the last row
for c in range(WIDTH):
    s.add(T[0][c]==f)
    s.add(B[HEIGHT-1][c]==f)
```

Now we'll enumerate all cells in the initial array (7.8.1). First two cells are empty there. And the third one has type "2b". This is "┐" and it can be oriented in 4 possible ways. And if it has angle 0°, bottom and right half-joints are present, others are absent. If it has angle 90°, it looks like "┌", and bottom and left half-joints are present, others are absent.

In plain English: "if cell of this type has angle 0°, these half-joints must be present **OR** if it has angle 90°, these half-joints must be present, **OR**, etc, etc."

Likewise, we define all these rules for all types and all possible angles:

```
for r in range(HEIGHT):
    for c in range(WIDTH):
        ty=cells_type[r][c]

        if ty=="0":
            s.add(A[r][c]==f)
            s.add(T[r][c]==f, B[r][c]==f, L[r][c]==f, R[r][c]==f)

        if ty=="2a":
            s.add(Or(And(A[r][c]==0, L[r][c]==f, R[r][c]==f, T[r][c]==t, B[r][c]==t), # ┌
                    And(A[r][c]==90, L[r][c]==t, R[r][c]==t, T[r][c]==f, B[r][c]==f))) # ┐

        if ty=="2b":
            s.add(Or(And(A[r][c]==0, L[r][c]==f, R[r][c]==t, T[r][c]==f, B[r][c]==t), # ┐
                    And(A[r][c]==90, L[r][c]==t, R[r][c]==f, T[r][c]==f, B[r][c]==t), # ┌
                    And(A[r][c]==180, L[r][c]==t, R[r][c]==f, T[r][c]==t, B[r][c]==f), # └
                    And(A[r][c]==270, L[r][c]==f, R[r][c]==t, T[r][c]==t, B[r][c]==f))) # ┘

        if ty=="3":
            s.add(Or(And(A[r][c]==0, L[r][c]==f, R[r][c]==t, T[r][c]==t, B[r][c]==t), # ┌
                    And(A[r][c]==90, L[r][c]==t, R[r][c]==t, T[r][c]==f, B[r][c]==t), # ┐
                    And(A[r][c]==180, L[r][c]==t, R[r][c]==f, T[r][c]==t, B[r][c]==t), # └
                    And(A[r][c]==270, L[r][c]==t, R[r][c]==t, T[r][c]==t, B[r][c]==f))) # ┘

        if ty=="4":
            s.add(A[r][c]==0)
            s.add(T[r][c]==t, B[r][c]==t, L[r][c]==t, R[r][c]==t) # ┣
```

Full source code is here: [https://github.com/dennis714/SAT\\_SMT\\_article/blob/master/SMT/pipe/solver/solve\\_pipe\\_puzzle1.py](https://github.com/dennis714/SAT_SMT_article/blob/master/SMT/pipe/solver/solve_pipe_puzzle1.py).

It produces this result (prints angle for each cell and (pseudo)graphical representation):

```
sat
 0  0  0  90  90  90  90  90  90  90  90  90  0  90  0  0
 0  90  0 180  0  0  0 180  0  90 180  0  0 180  0  0
 0  0 180  0  0  0  0 180 270  0 180  0  0  90  0  90
270  0  90  90  90  90 270  90  0 270  90 180  0 270  0 180
 0 270 180  0 270  0  90  0 180  0 270 270 180  0  0  0
 0  0 270 180  0  0 270  0 270  0  90  90 270  90  0 180
 0  0  90  90  0  0  90  0  90  0  0 180  0  0  0 180
 0 270 270 270  90 270 180 270 270 270 270 270  90 270 270 180

  r  T  -  -  -  T  T  -  T  r  r  r
r  r  ┌  ┐      r  ┌  ┌  T  ┌  ┌  ┌  ┐
┌  ┌  ┐      ┌  ┐  ┌  ┌  ┐  ┌  ┌  r  r
┌  ┌  T  -  T  T  ┌  r  ┌  T  ┌  ┌  ┌  ┌
  ┌  r  ┌  ┌  r  ┌  ┌  r  ┌  ┌  ┌  ┌  ┌
  ┌  ┐  ┌  ┌  ┌  ┌  T  T  ┌  r  ┌  ┌
r  T  r  ┌  T  ┌  T  ┌  ┌  ┌  ┌  ┌  ┌
┌  ┌  ┌  -  ┌  ┌  ┌  ┌  ┌  -  ┌  ┌  ┌
```

Figure 4: Solver script output

It worked  $\approx 4$  seconds on my old and slow Intel Atom N455 1.66GHz. Is it fast? I don't know, but again, what is really cool, we do not know about any mathematical background of all this, we just defined cells, (half-)joints and defined relations between them.

Now the next question is, how many solutions are possible? Using method described earlier (7.4), I've altered solver script <sup>43</sup> and solver said two solutions are possible.

Let's compare these two solutions using gvimdiff:

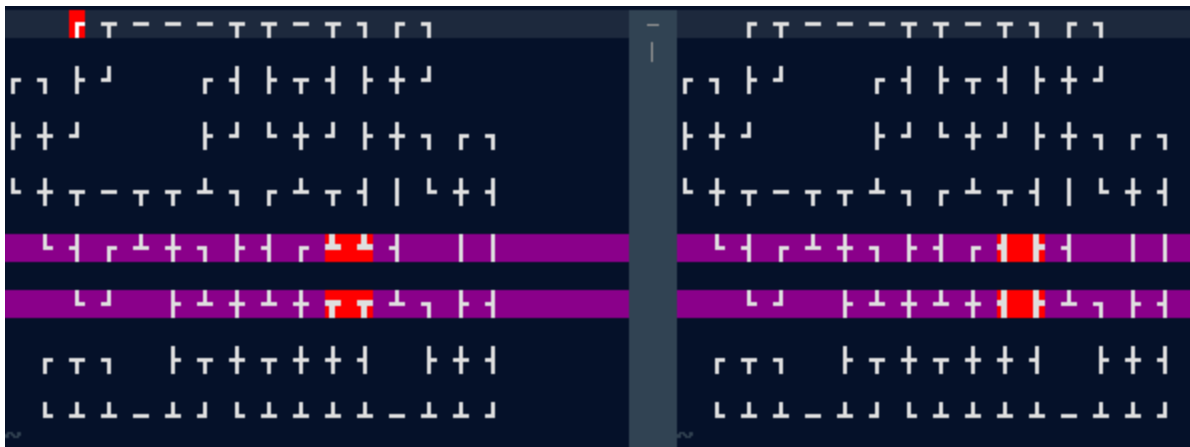


Figure 5: gvimdiff output (pardon my red cursor at left pane at left-top corner)

4 cells in the middle can be orientated differently. Perhaps, other puzzles may produce different results.

P.S. *Half-joint* is defined as boolean type. But in fact, the first version of the solver has been written using integer type for half-joints, and 0 was used for False and 1 for True. I did it so because I wanted to make source code tidier and narrower without using long words like “False” and “True”. And it worked, but slower. Perhaps, Z3 handles boolean data types faster? Better? Anyway, I writing this to note that integer type can also be used instead of boolean, if needed.

## 7.9 Cracking Minesweeper with Z3 SMT solver

For those who are not very good at playing Minesweeper (like me), it's possible to predict bombs' placement without touching debugger.

Here is a clicked somewhere and I see revealed empty cells and cells with known number of “neighbours”:



What we have here, actually? Hidden cells, empty cells (where bombs are not present), and empty cells with numbers, which shows how many bombs are placed nearby.

<sup>43</sup>[https://github.com/dennis714/SAT\\_SMT\\_article/blob/master/SMT/pipe/solver/solve\\_pipe\\_puzzle2.py](https://github.com/dennis714/SAT_SMT_article/blob/master/SMT/pipe/solver/solve_pipe_puzzle2.py)

### 7.9.1 The method

Here is what we can do: we will try to place a bomb to all possible hidden cells and ask Z3 SMT solver, if it can disprove the very fact that the bomb can be placed there.

Take a look at this fragment. "?" mark is for hidden cell, "." is for empty cell, number is a number of neighbours.

	C1	C2	C3
R1	?	?	?
R2	?	3	.
R3	?	1	.

So there are 5 hidden cells. We will check each hidden cell by placing a bomb there. Let's first pick top/left cell:

	C1	C2	C3
R1	*	?	?
R2	?	3	.
R3	?	1	.

Then we will try to solve the following system of equations ( $RrCc$  is cell of row  $r$  and column  $c$ ):

- $R1C2 + R2C1 + R2C2 = 1$  (because we placed bomb at  $R1C1$ )
- $R2C1 + R2C2 + R3C1 = 1$  (because we have "1" at  $R3C2$ )
- $R1C1 + R1C2 + R1C3 + R2C1 + R2C2 + R2C3 + R3C1 + R3C2 + R3C3 = 3$  (because we have "3" at  $R2C2$ )
- $R1C2 + R1C3 + R2C2 + R2C3 + R3C2 + R3C3 = 0$  (because we have "." at  $R2C3$ )
- $R2C2 + R2C3 + R3C2 + R3C3 = 0$  (because we have "." at  $R3C3$ )

As it turns out, this system of equations is satisfiable, so there could be a bomb at this cell. But this information is not interesting to us, since we want to find cells we can freely click on. And we will try another one. And if the equation will be unsatisfiable, that would imply that a bomb cannot be there and we can click on it.

### 7.9.2 The code

```
#!/usr/bin/python

known=[
"01?10001?",
"01?100011",
"011100000",
"000000000",
"111110011",
"?????1001?",
"?????3101?",
"?????211?",
"?????????"]

from z3 import *
import sys

WIDTH=len(known[0])
HEIGHT=len(known)

print "WIDTH=", WIDTH, "HEIGHT=", HEIGHT

def chk_bomb(row, col):

    s=Solver()

    cells=[[Int('cell_r=%d_c=%d' % (r,c)) for c in range(WIDTH+2)] for r in range(HEIGHT+2)]

    # make border
    for c in range(WIDTH+2):
```

```

s.add(cells[0][c]==0)
s.add(cells[HEIGHT+1][c]==0)
for r in range(HEIGHT+2):
    s.add(cells[r][0]==0)
    s.add(cells[r][WIDTH+1]==0)

for r in range(1,HEIGHT+1):
    for c in range(1,WIDTH+1):

        t=known[r-1][c-1]
        if t in "012345678":
            s.add(cells[r][c]==0)
            # we need empty border so the following expression would be able to work for all possible cases:
            s.add(cells[r-1][c-1] + cells[r-1][c] + cells[r-1][c+1] + cells[r][c-1] + cells[r][c+1] + cells[
                r+1][c-1] + cells[r+1][c] + cells[r+1][c+1]==int(t))

# place bomb:
s.add(cells[row][col]==1)

result=str(s.check())
if result=="unsat":
    print "row=%d col=%d, unsat!" % (row, col)

# enumerate all hidden cells:
for r in range(1,HEIGHT+1):
    for c in range(1,WIDTH+1):
        if known[r-1][c-1]=="?":
            chk_bomb(r, c)

```

The code is almost self-explanatory. We need border for the same reason, why Conway's "Game of Life" implementations also has border (to make calculation function simpler). Whenever we know that the cell is free of bomb, we put zero there. Whenever we know number of neighbours, we add a constraint, again, just like in "Game of Life": number of neighbours must be equal to the number we have seen in the Minesweeper. Then we place bomb somewhere and check.

Let's run:

```

row=1 col=3, unsat!
row=6 col=2, unsat!
row=6 col=3, unsat!
row=7 col=4, unsat!
row=7 col=9, unsat!
row=8 col=9, unsat!

```

These are cells where I can click safely, so I did:



Now we have more information, so we update input:

```

known=[
"01110001?",
"01?100011",
"011100000",
"000000000",
"111110011",

```



```
"?11?1001?",
"???331011",
"?????2110",
"???????10"]
```

I run it again:

```
row=7 col=1, unsat!
row=7 col=2, unsat!
row=7 col=3, unsat!
row=8 col=3, unsat!
row=9 col=5, unsat!
row=9 col=6, unsat!
```

I click on these cells again:



I update it again:

```
known=[
"01110001?",
"01?100011",
"011100000",
"000000000",
"111110011",
"?11?1001?",
"222331011",
"??2??2110",
"????22?10"]
```

```
row=8 col=2, unsat!
row=9 col=4, unsat!
```



This is last update:

```
known=[
"01110001?",
"01?100011",
"011100000",
"000000000",
"111110011",
"?11?1001?",
"222331011",
"?22??2110",
"???322?10"]
```

...last result:

```
row=9 col=1, unsat!
row=9 col=2, unsat!
```

Voila!



The source code: [https://github.com/dennis714/SAT\\_SMT\\_article/blob/master/SMT/minesweeper/minesweeper\\_solver.py](https://github.com/dennis714/SAT_SMT_article/blob/master/SMT/minesweeper/minesweeper_solver.py).

Some discussion on HN: <https://news.ycombinator.com/item?id=13797375>.

See also: cracking Minesweeper using SAT solver: [13.7](#).

## 7.10 Recalculating micro-spreadsheet using Z3Py

There is a nice exercise<sup>44</sup>: write a program to recalculate micro-spreadsheet, like this one:

1	0	B0+B2	A0*B0*C0
123	10	12	11
667	A0+B1	(C1*A0)*122	A3+C2

As it turns out, though overkill, this can be solved using Z3 with little effort:

```
#!/usr/bin/python

from z3 import *
import sys, re

# MS Excel or LibreOffice style.
# first top-left cell is A0, not A1
def coord_to_name(R, C):
    return "ABCDEFGHIJKLMNOPQRSTUVWXYZ"[R]+str(C)

# open file and parse it as list of lists:
f=open(sys.argv[1],"r")
# filter(None, ...) to remove empty sublists:
ar=filter(None, [item.rstrip().split() for item in f.readlines()])
f.close()
```

<sup>44</sup>Blog post in Russian: <http://thesz.livejournal.com/280784.html>

```

WIDTH=len(ar[0])
HEIGHT=len(ar)

# cells{} is a dictionary with keys like "A0", "B9", etc:
cells={}
for R in range(HEIGHT):
    for C in range(WIDTH):
        name=coord_to_name(R, C)
        cells[name]=Int(name)

s=Solver()

cur_R=0
cur_C=0

for row in ar:
    for c in row:
        # string like "A0+B2" becomes "cells["A0"]+cells["B2"]":
        c=re.sub(r'([A-Z]{1}[0-9]+)', r'cells["\1"]', c)
        st="cells[\\"%s\\"]==%s" % (coord_to_name(cur_R, cur_C), c)
        # evaluate string. Z3Py expression is constructed at this step:
        e=eval(st)
        # add constraint:
        s.add (e)
        cur_C=cur_C+1
    cur_R=cur_R+1
    cur_C=0

result=str(s.check())
print result
if result=="sat":
    m=s.model()
    for r in range(HEIGHT):
        for c in range(WIDTH):
            sys.stdout.write (str(m[cells[coord_to_name(r, c)]])+"\t")
            sys.stdout.write ("\n")

```

( <https://github.com/dennis714/yurichev.com/blob/master/blog/spreadsheet/1.py> )

All we do is just creating pack of variables for each cell, named A0, B1, etc, of integer type. All of them are stored in *cells[]* dictionary. Key is a string. Then we parse all the strings from cells, and add to list of constraints  $A0=123$  (in case of number in cell) or  $A0=B1+C2$  (in case of expression in cell). There is a slight preparation: string like  $A0+B2$  becomes *cells["A0"]+cells["B2"]*.

Then the string is evaluated using Python *eval()* method, which is highly dangerous<sup>45</sup>: imagine if end-user could add a string to cell other than expression? Nevertheless, it serves our purposes well, because this is a simplest way to pass a string with expression into Z3.

Z3 do the job with little effort:

```

% python 1.py test1
sat
1      0      135      82041
123    10      12       11
667    11     1342     83383

```

### 7.10.1 Unsat core

Now the problem: what if there is circular dependency? Like:

1	0	$B0+B2$	$A0*B0$
123	10	12	11
$C1+123$	$C0*123$	$A0*122$	$A3+C2$

Two first cells of the last row (C0 and C1) are linked to each other. Our program will just tells “unsat”, meaning, it couldn’t satisfy all constraints together. We can’t use this as error message reported to end-user, because it’s highly unfriendly.

However, we can fetch *unsat core*, i.e., list of variables which Z3 finds conflicting.

```

...
s=Solver()

```

<sup>45</sup><http://stackoverflow.com/questions/1832940/is-using-eval-in-python-a-bad-practice>

```
s.set(unsat_core=True)
...
    # add constraint:
    s.assert_and_track(e, coord_to_name(cur_R, cur_C))
...
if result=="sat":
...
else:
    print s.unsat_core()
```

( <https://github.com/dennis714/yurichev.com/blob/master/blog/spreadsheet/2.py> )

We should explicitly turn on unsat core support and use *assert\_and\_track()* instead of *add()* method, because this feature slows down the whole process, and is turned off by default. That works:

```
% python 2.py test_circular
unsat
[C0, C1]
```

Perhaps, these variables could be removed from the 2D array, marked as *unresolved* and the whole spreadsheet could be recalculated again.

### 7.10.2 Stress test

How to generate large random spreadsheet? What we can do. First, create random DAG<sup>46</sup>, like this one:

---

<sup>46</sup>Directed acyclic graph




Figure 6: Random DAG

Arrows will represent information flow. So a vertex (node) which has no incoming arrows to it (indegree=0), can be set to a random number. Then we use topological sort to find dependencies between vertices. Then we assign spreadsheet cell names to each vertex. Then we generate random expression with random operations/numbers/cells to each cell, with the use of information from topological sorted graph.

Wolfram Mathematica:

```
(* Utility functions *)
In[1]:= findSublistBeforeElementByValue[lst_,element_]:=lst[[ 1;;Position[lst, element][[1]][[1]]-1]]

(* Input in ∞1.. range. 1->A0, 2->A1, etc *)
In[2]:= vertexToName[x_,width_]:=StringJoin[FromCharacterCode[ToCharacterCode["A"]][[1]]+Floor[(x-1)/width]],
ToString[Mod[(x-1),width]]]

In[3]:= randomNumberAsString[]:=ToString[RandomInteger[{1,1000}]]

In[4]:= interleaveListWithRandomNumbersAsStrings[lst_]:=Riffle[lst,Table[randomNumberAsString[],Length[lst]-1]]

(* We omit division operation because micro-spreadsheet evaluator can't handle division by zero *)
In[5]:= interleaveListWithRandomOperationsAsStrings[lst_]:=Riffle[lst,Table[RandomChoice[{"+", "-", "*"}],Length[
lst]-1]]
```

```

In[6]:= randomNonNumberExpression[g_,vertex_]:=StringJoin[interleaveListWithRandomOperationsAsStrings[
interleaveListWithRandomNumbersAsStrings[Map[vertexToName[#,WIDTH]&,pickRandomNonDependentVertices[g,vertex
]]]]]

In[7]:= pickRandomNonDependentVertices[g_,vertex_]:=DeleteDuplicates[RandomChoice[
findSublistBeforeElementByValue[TopologicalSort[g],vertex],RandomInteger[{1,5}]]]

In[8]:= assignNumberOrExpr[g_,vertex_]:=If[VertexInDegree[g,vertex]==0,randomNumberAsString[],
randomNonNumberExpression[g,vertex]]

(* Main part *)
(* Create random graph *)
In[21]:= WIDTH=7;HEIGHT=8;TOTAL=WIDTH*HEIGHT
Out[21]= 56

In[24]:= g=DirectedGraph[RandomGraph[BernoulliGraphDistribution[TOTAL,0.05]],"Acyclic"];

...

(* Generate random expressions and numbers *)
In[26]:= expressions=Map[assignNumberOrExpr[g,#]&,VertexList[g]];

(* Make 2D table of it *)
In[27]:= t=Partition[expressions,WIDTH];

(* Export as tab-separated values *)
In[28]:= Export["/home/dennis/1.txt",t,"TSV"]
Out[28]= /home/dennis/1.txt

In[29]:= Grid[t,Frame->All,Alignment->Left]

```

Here is an output from *Grid[]*:

846	499	A3*913-H4	808	278	303	D
B4*860+D2	999	59	442	425	A5*163+B2+127*C2*927*D3*213+C1	5
G6*379-C3-436-C4-289+H6	972	804	D2	G5+108-F1*413-D3	B5	G
F2	E0	B6-731-D3+791+B4*92+C1	551	F4*922*C2+760*A6-992+B4-184-A4	B1-624-E3	F
519	G1*402+D1*107*G3-458*A1	D3	B4	B3*811-D3*345+E0	B5	H
F5-531+B5-222*E4	9	B5+106*B6+600-B1	E3	A5+866*F6+695-A3*226+C6	F4*102*E4*998-H0	B
C3-956*A5	G4*408-D3*290*B6-899*G5+400+F1	B2-701+H6	A3+782*A5+46-B3-731+C1	42	287	H
B4-792*H4*407+F6-425-E1	D2	D3	F2-327*G4*35*E1	E1+376*A6-606*F6*554+C5	E3	F

Using this script, I can generate random spreadsheet of  $26 \cdot 500 = 13000$  cells, which seems to be processed in couple of seconds.

### 7.10.3 The files

The files, including Mathematica notebook: <https://github.com/dennis714/yurichev.com/tree/master/blog/spreadsheet>.

## 7.11 Discrete tomography

How computed tomography (CT scan) actually works? A human body is bombarded by X-rays in various angles by X-ray tube in rotating torus. X-ray detectors are also located in torus, and all the information is recorded.

Here is we can simulate simple tomograph. An “i” character is rotating and will be “enlighten” at 4 angles. Let’s imagine, character is bombarded by X-ray tube at left. All asterisks in each row is then summed and sum is “received” by X-ray detector at the right.

```

WIDTH= 11 HEIGHT= 11
angleπ=(/4)*0
**      2
**      2
      0
***     3
**      2
**      2
**      2
**      2
**      2
**      2
****    4
      0
[2, 2, 0, 3, 2, 2, 2, 2, 2, 4, 0] ,

```

```

angleπ=(/4)*1
0
0
* 1
** 2
* 1
** 2
** 2
**** 4
* 1
* 1
0
[0, 0, 1, 2, 1, 2, 2, 4, 1, 1, 0] ,
angleπ=(/4)*2
0
0
0
0
* 1
** ***** 9
** ***** 9
* * 2
0
0
0
[0, 0, 0, 0, 1, 9, 9, 2, 0, 0, 0] ,
angleπ=(/4)*3
0
0
* 1
** 2
** * 3
*** 3
** 2
0
** 2
* 1
0
[0, 0, 1, 2, 3, 3, 2, 0, 2, 1, 0] ,

```

( The source code: [https://github.com/dennis714/SAT\\_SMT\\_article/blob/master/SMT/tomo/gen.py](https://github.com/dennis714/SAT_SMT_article/blob/master/SMT/tomo/gen.py) )

All we got from our toy-level tomograph is 4 vectors, these are sums of all asterisks in rows for 4 angles:

```

[2, 2, 0, 3, 2, 2, 2, 2, 2, 4, 0] ,
[0, 0, 1, 2, 1, 2, 2, 4, 1, 1, 0] ,
[0, 0, 0, 0, 1, 9, 9, 2, 0, 0, 0] ,
[0, 0, 1, 2, 3, 3, 2, 0, 2, 1, 0] ,

```

How do we recover initial image? We are going to represent 11\*11 matrix, where sum of each row must be equal to some value we already know. Then we rotate matrix, and do this again.

The “rotate” function has been taken from the generation program, because, due to Python’s dynamic typization nature, it’s not important for the function to what operate on: strings, characters, or Z3 variable instances, so it works very well for all of them.

```

#-*- coding: utf-8 -*-

import math, sys
from z3 import *

# https://en.wikipedia.org/wiki/Rotation_matrix
def rotate(pic, angle):
    WIDTH=len(pic[0])
    HEIGHT=len(pic)
    #print WIDTH, HEIGHT
    assert WIDTH==HEIGHT
    ofs=WIDTH/2

    out = [[0 for x in range(WIDTH)] for y in range(HEIGHT)]

    for x in range(-ofs,ofs):
        for y in range(-ofs,ofs):
            newX = int(round(math.cos(angle)*x - math.sin(angle)*y,3))+ofs
            newY = int(round(math.sin(angle)*x + math.cos(angle)*y,3))+ofs

```

```

        # clip at boundaries, hence min(..., HEIGHT-1)
        out[min(newX,HEIGHT-1)][min(newY,WIDTH-1)]=pic[x+ofs][y+ofs]
    return out

vectors=[
[2, 2, 0, 3, 2, 2, 2, 2, 2, 4, 0] ,
[0, 0, 1, 2, 1, 2, 2, 4, 1, 1, 0] ,
[0, 0, 0, 0, 1, 9, 9, 2, 0, 0, 0] ,
[0, 0, 1, 2, 3, 3, 2, 0, 2, 1, 0]]

WIDTH = HEIGHT = len(vectors[0])

s=Solver()
cells=[[Int('cell_r=%d_c=%d' % (r,c)) for c in range(WIDTH)] for r in range(HEIGHT)]

# monochrome picture, only 0's or 1's:
for c in range(WIDTH):
    for r in range(HEIGHT):
        s.add(Or(cells[r][c]==0, cells[r][c]==1))

def all_zeroes_in_vector(vec):
    for v in vec:
        if v!=0:
            return False
    return True

ANGLES=len(vectors)
for a in range(ANGLES):
    angle=a*(math.pi/ANGLES)
    rows=rotate(cells, angle)
    r=0
    for row in rows:
        # skip empty rows:
        if all_zeroes_in_vector(row)==False:
            # sum of row must be equal to the corresponding element of vector:
            s.add(Sum(*row)==vectors[a][r])
            r=r+1

print s.check()
m=s.model()
for r in range(HEIGHT):
    for c in range(WIDTH):
        if str(m[cells[r][c]])=="1":
            sys.stdout.write("*")
        else:
            sys.stdout.write(" ")
    print ""

```

( The source code: [https://github.com/dennis714/SAT\\_SMT\\_article/blob/master/SMT/tomo/solve.py](https://github.com/dennis714/SAT_SMT_article/blob/master/SMT/tomo/solve.py) )

That works:

```

% python solve.py
sat
**
**

***
**
**
**
**
**
**
****

```

In other words, all SMT-solver does here is solving a system of equations.  
So, 4 angles are enough. What if we could use only 3 angles?

```

WIDTH= 11 HEIGHT= 11
anglen=(/3)*0
**      2
**      2
      0
***     3
**      2

```



```

**      2
**      2
**      2
**      2
****    4
        0
[2, 2, 0, 3, 2, 2, 2, 2, 2, 4, 0] ,
angle $\pi$ =(/3)*1
        0
        0
        0
**      2
**      2
***     3
****    4
      **  2
      *   1
        0
        0
[0, 0, 0, 2, 2, 3, 4, 2, 1, 0, 0] ,
angle $\pi$ =(/3)*2
        0
        0
        0
      **  2
      **  2
*****  5
**      2
**      2
*       1
        0
        0
[0, 0, 0, 2, 2, 5, 2, 2, 1, 0, 0] ,

```

No, it's not enough:

```

% time python solve3.py
sat
*  *
  **
    * **
  **
*  *
  **
   *  *
*  *
****

```

However, the result is correct, but only 3 vectors allows too many possible “initial images”, and Z3 SMT-solver finds first.

Further reading: [https://en.wikipedia.org/wiki/Discrete\\_tomography](https://en.wikipedia.org/wiki/Discrete_tomography), [https://en.wikipedia.org/wiki/2-satisfiability#Discrete\\_tomography](https://en.wikipedia.org/wiki/2-satisfiability#Discrete_tomography).

## 7.12 Simplifying long and messy expressions using Mathematica and Z3

...which can be results of Hex-Rays and/or manual rewriting.

I've added to my RE4B book about Wolfram Mathematica capabilities to minimize expressions <sup>47</sup>.

Today I stumbled upon this Hex-Rays output:

```

if ( ( x != 7 || y!=0 ) && ( x < 6 || x > 7) )
{
    ...
};

```

Both Mathematica and Z3 (using “simplify” command) can't make it shorter, but I've got gut feeling, that there is something redundant.

Let's take a look at the right part of the expression. If  $x$  must be less than 6 OR greater than 7, then it can hold any value except 6 AND 7, right? So I can rewrite this manually:

<sup>47</sup>[https://github.com/dennis714/RE-for-beginners/blob/cd85356051937e87f90967cc272248084808223b/other/hexrays\\_EN.tex#L412](https://github.com/dennis714/RE-for-beginners/blob/cd85356051937e87f90967cc272248084808223b/other/hexrays_EN.tex#L412), <https://beginners.re/>

```
if ( ( x != 7 || y!=0 ) && x != 6 && x != 7 )
{
    ...
};
```

And this is what Mathematica can simplify:

```
In[]:= BooleanMinimize[(x != 7 || y != 0) && (x != 6 && x != 7)]
Out[]:= x != 6 && x != 7
```

$y$  gets reduced.

But am I really right? And why Mathematica and Z3 didn't simplify this at first place?

I can use Z3 to prove that these expressions are equal to each other:

```
#!/usr/bin/env python

from z3 import *

x=Int('x')
y=Int('y')

s=Solver()

exp1=And(Or(x!=7, y!=0), Or(x<6, x>7))
exp2=And(x!=6, x!=7)

s.add(exp1!=exp2)

print simplify(exp1) # no luck

print s.check()
print s.model()
```

Z3 can't find counterexample, so it says "unsat", meaning, these expressions are equivalent to each other. So I've rewritten this expression in my code, tests has been passed, etc.

Yes, using both Mathematica and Z3 is overkill, and this is basic boolean algebra, but after ~10 hours of sitting at a computer you can make really dumb mistakes, and additional proof your piece of code is correct is never unwanted.

## 7.13 Solving XKCD 287 using Z3

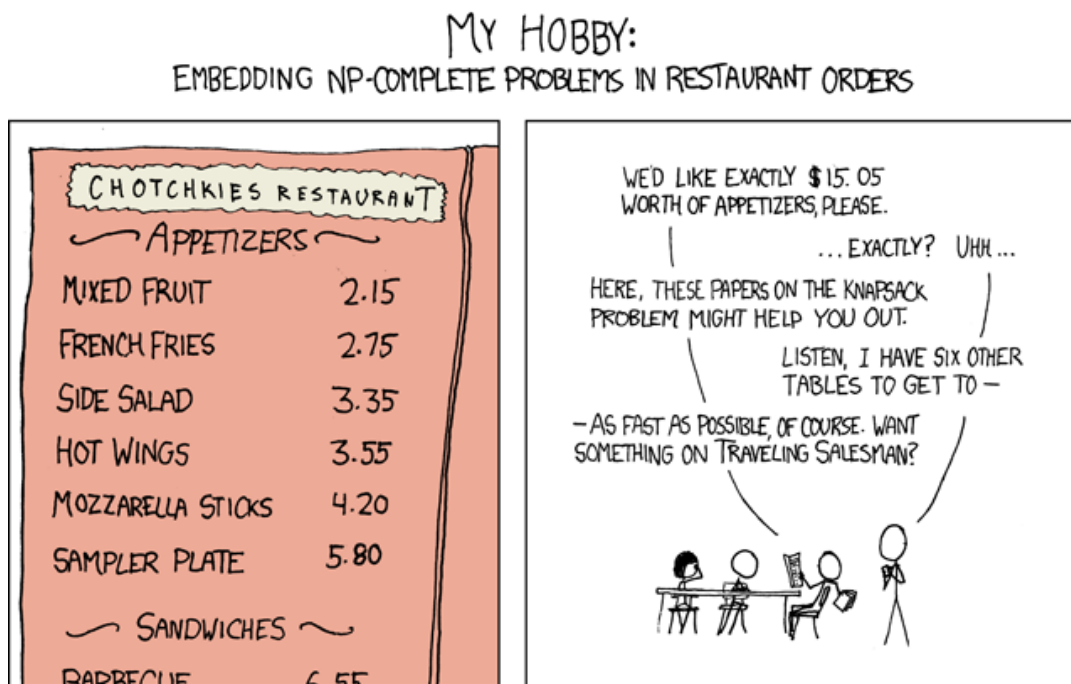


Figure 7: xkcd #287

( <https://www.xkcd.com/287/> )

The problem is to solve the following equation:  $2.15a + 2.75b + 3.35c + 3.55d + 4.20e + 5.80f == 15.05$ , where  $a..f$  are integers. So this is a linear diophantine equation.

```
#!/usr/bin/python

from z3 import *

a,b,c,d,e,f = Ints('a b c d e f')
s = Solver()
s.add(215*a + 275*b + 335*c + 355*d + 420*e + 580*f == 1505, a>=0, b>=0, c>=0, d>=0, e>=0, f>=0)

results=[]

# enumerate all possible solutions:
while True:
    if s.check() == sat:
        m = s.model()
        print m
        results.append(m)
        block = []
        for d in m:
            c=d()
            block.append(c != m[d])
        s.add(Or(block))
    else:
        print "results total=", len(results)
        break
```

( The source code: [https://github.com/dennis714/SAT\\_SMT\\_article/blob/master/SMT/xkcd287/xkcd287.py](https://github.com/dennis714/SAT_SMT_article/blob/master/SMT/xkcd287/xkcd287.py) )

There are just 2 solutions:

```
[f = 0, b = 0, a = 7, c = 0, d = 0, e = 0]
[f = 1, b = 0, a = 1, c = 0, d = 2, e = 0]
results total= 2
```

Wolfram Mathematica can solve the equation as well:

```
In[]:= FindInstance[2.15 a + 2.75 b + 3.35 c + 3.55 d + 4.20 e + 5.80 f == 15.05 &&
  a >= 0 && b >= 0 && c >= 0 && d >= 0 && e >= 0 && f >= 0,
  {a, b, c, d, e, f}, Integers, 1000]

Out[]= {{a -> 1, b -> 0, c -> 0, d -> 2, e -> 0, f -> 1},
  {a -> 7, b -> 0, c -> 0, d -> 0, e -> 0, f -> 0}}
```

1000 means “find at most 1000 solutions”, but only 2 are found. See also: <http://reference.wolfram.com/language/ref/FindInstance.html>.

Other ways to solve it: <https://stackoverflow.com/questions/141779/solving-the-np-complete-problem-in>  
[http://www.explainxkcd.com/wiki/index.php/287:\\_NP-Complete](http://www.explainxkcd.com/wiki/index.php/287:_NP-Complete).

## 7.14 Making smallest possible test suite using Z3

I once worked on rewriting large piece of code into pure C, and there were a tests, several thousands. Testing process was painfully slow, so I thought if the test suite can be minimized somehow.

What we can do is to run each test and get code coverage (information about which lines of code was executed and which are not). Then the task is to make such test suite, where coverage is maximum, and number of tests is minimal.

In fact, this is *set cover problem* (also known as *hitting set problem*). While simpler algorithms exist (see Wikipedia<sup>48</sup>), it is also possible to solve with SMT-solver.

First, I took LZSS<sup>49</sup> compression/decompression code<sup>50</sup> for the example, from Apple sources. Such routines are not easy to test. Here is my version of it: [https://github.com/dennis714/SAT\\_SMT\\_article/blob/master/SMT/set\\_cover/compression.c](https://github.com/dennis714/SAT_SMT_article/blob/master/SMT/set_cover/compression.c). I've added random generation of input data to be compressed.

<sup>48</sup>[https://en.wikipedia.org/wiki/Set\\_cover\\_problem](https://en.wikipedia.org/wiki/Set_cover_problem)

<sup>49</sup>Lempel-Ziv-Storer-Szymanski

<sup>50</sup>[https://github.com/opensource-apple/kext\\_tools/blob/master/compression.c](https://github.com/opensource-apple/kext_tools/blob/master/compression.c)

Random generation is dependent of some kind of input seed. Standard `srand()` / `rand()` are not recommended to be used, but for such simple task as ours, it's OK. I'll generate<sup>51</sup> 1000 tests with 0..999 seeds, that would produce random data to be compressed/decompressed/checked.

After the compression/decompression routine has finished its work, GNU gcov utility is executed, which produces result like this:

```
...
3395: 189:      for (i = 1; i < F; i++) {
3395: 190:          if ((cmp = key[i] - sp->text_buf[p + i]) != 0)
2565: 191:              break;
-: 192:      }
2565: 193:      if (i > sp->match_length) {
1291: 194:          sp->match_position = p;
1291: 195:          if ((sp->match_length = i) >= F)
#####: 196:              break;
-: 197:      }
2565: 198:  }
#####: 199:      sp->parent[r] = sp->parent[p];
#####: 200:      sp->lchild[r] = sp->lchild[p];
#####: 201:      sp->rchild[r] = sp->rchild[p];
#####: 202:      sp->parent[sp->lchild[p]] = r;
#####: 203:      sp->parent[sp->rchild[p]] = r;
#####: 204:      if (sp->rchild[sp->parent[p]] == p)
#####: 205:          sp->rchild[sp->parent[p]] = r;
...
```

A leftmost number is an execution count for each line. ##### means the line of code hasn't been executed at all. The second column is a line number.

Now the Z3Py script, which will parse all these 1000 gcov results and produce minimal *hitting set*:

```
#!/usr/bin/env python

import re, sys
from z3 import *

TOTAL_TESTS=1000

# read gcov result and return list of lines executed:
def process_file (fname):
    lines=[]
    f=open(fname,"r")

    while True:
        l=f.readline().rstrip()
        m = re.search('^[^*]([0-9]+): ([0-9]+):.*$', l)
        if m!=None:
            lines.append(int(m.group(2)))
        if len(l)==0:
            break

    f.close()
    return lines

# k=test number; v=list of lines executed
stat={}
for test in range(TOTAL_TESTS):
    stat[test]=process_file("compression.c.gcov."+str(test))

# that will be a list of all lines in all tests:
all_lines=set()
# k=line, v=list of tests, which trigger that line:
tests_for_line={}

for test in stat:
    all_lines|=set(stat[test])
    for line in stat[test]:
        tests_for_line[line]=tests_for_line.get(line, []) + [test]

# int variable for each test:
tests=[Int('test_%d' % (t)) for t in range(TOTAL_TESTS)]

# this is optimization problem, so Optimize() instead of Solver():
```

<sup>51</sup>[https://github.com/dennis714/yurichev.com/blob/master/blog/set\\_cover/gen\\_gcov\\_tests.sh](https://github.com/dennis714/yurichev.com/blob/master/blog/set_cover/gen_gcov_tests.sh)

```

opt = Optimize()

# each test variable is either 0 (absent) or 1 (present):
for t in tests:
    opt.add(Or(t==0, t==1))

# we know which tests can trigger each line
# so we enumerate all tests when preparing expression for each line
# we form expression like "test_1==1 OR test_2==1 OR ..." for each line:
for line in list(all_lines):
    expressions=[tests[s]==1 for s in tests_for_line[line]]
    # expression is a list which unfolds as list of arguments into Z3's Or() function (see asterisk)
    # that results in big expression like "Or(test_1==1, test_2==1, ...)"
    # the expression is then added as a constraint:
    opt.add(Or(*expressions))

# we need to find a such solution, where minimal number of all "test_X" variables will have 1
# "*tests" unfolds to a list of arguments: [test_1, test_2, test_3,...]
# "Sum(*tests)" unfolds to the following expression: "Sum(test_1, test_2, ...)"
# the sum of all "test_X" variables should be as minimal as possible:
h=opt.minimize(Sum(*tests))

print (opt.check())
m=opt.model()

# print all variables set to 1:
for t in tests:
    if m[t].as_long()==1:
        print (t)

```

And what it produces (~19s on my old Intel Quad-Core Xeon E3-1220 3.10GHz):

```

% time python set_cover.py
sat
test_7
test_48
test_134
python set_cover.py  18.95s user 0.03s system 99% cpu 18.988 total

```

We need just these 3 tests to execute (almost) all lines in the code: looks impressive, given the fact, that it would be notoriously hard to pick these tests by hand! The result can be checked easily, again, using gcov utility.

This is sometimes also called MaxSAT/MaxSMT — the problem is to find solution, but the solution where some variable/expression is maximal as possible, or minimal as possible.

Also, the code gives incorrect results on Z3 4.4.1, but working correctly on Z3 4.5.0 (so please upgrade). This is relatively fresh feature in Z3, so probably it was not stable in previous versions?

The files: [https://github.com/dennis714/SAT\\_SMT\\_article/tree/master/SMT/set\\_cover](https://github.com/dennis714/SAT_SMT_article/tree/master/SMT/set_cover).

Further reading: [https://en.wikipedia.org/wiki/Set\\_cover\\_problem](https://en.wikipedia.org/wiki/Set_cover_problem), [https://en.wikipedia.org/wiki/Maximum\\_satisfiability\\_problem](https://en.wikipedia.org/wiki/Maximum_satisfiability_problem), [https://en.wikipedia.org/wiki/Optimization\\_problem](https://en.wikipedia.org/wiki/Optimization_problem).

## 7.15 Package manager and Z3

Here is simplified example. We have libA, libB, libC and libD, available in various versions (and flavors). We're going to install programA and programB, which use these libraries.

```

#!/usr/bin/env python

from z3 import *

s=Optimize()

libA=Int('libA')
# libA's version is 1..5 or 999 (which means library will not be installed):
s.add(Or(And(libA>=1, libA<=5),libA==999))

libB=Int('libB')
# libB's version is 1, 4, 5 or 999:
s.add(Or(libB==1, libB==4, libB==5, libB==999))

libC=Int('libC')
# libC's version is 10, 11, 14 or 999:
s.add(Or(libC==10, libC==11, libC==14, libC==999))

```

```

# libC is dependent on libA
# libC v10 is dependent on libA v1..3, but not newer
# libC v11 requires at least libA v3
# libC v14 requires at least libA v5
s.add(If(libC==10, And(libA>=1, libA<=3), True))
s.add(If(libC==11, libA>=3, True))
s.add(If(libC==14, libA>=5, True))

libD=Int('libD')
# libD's version is 1..10
s.add(Or(And(libD>=1, libD<=10),libD==999))

programA=Int('programA')
# programA came as v1 or v2:
s.add(Or(programA==1, programA==2))

# programA is dependent on libA, libB and libC
# programA v1 requires libA v2 (only this version), libB v4 or v5, libC v10:
s.add(If(programA==1, And(libA==2, Or(libB==4, libB==5), libC==10), True))
# programA v2 requires these libraries: libA v3, libB v5, libC v11
s.add(If(programA==2, And(libA==3, libB==5, libC==11), True))

programB=Int('programB')
# programB came as v7 or v8:
s.add(Or(programB==7, programB==8))

# programB v7 requires libA at least v2 and libC at least v10:
s.add(If(programB==7, And(libA>=2, libC>=10), True))
# programB v8 requires libA at least v6 and libC at least v11:
s.add(If(programB==8, And(libA>=6, libC>=11), True))

s.add(programA==1)
s.add(programB==7) # change this to 8 to make it unsat

# we want latest libraries' versions.
# if the library is not required, its version is "pulled up" to 999,
# and 999 means the library is not needed to be installed
s.maximize(Sum(libA,libB,libC,libD))

print s.check()
print s.model()

```

( The source code: [https://github.com/dennis714/SAT\\_SMT\\_article/blob/master/SMT/dep/dependency.py](https://github.com/dennis714/SAT_SMT_article/blob/master/SMT/dep/dependency.py) )

The output:

```

sat
[libB = 5,
 libD = 999,
 libC = 10,
 programB = 7,
 programA = 1,
 libA = 2]

```

999 means that there is no need to install libD, it's not required by other packages.

Change version of ProgramB to v8 and it will says "unsat", meaning, there is a conflict: ProgramA requires libA v2, but ProgramB v8 eventually requires newer libA.

Still, there is a work to do: "unsat" message is somewhat useless to end user, some information about conflicting items should be printed.

Here is my another optimization problem example: [7.14](#).

More about using SAT/SMT solvers in package managers: <https://research.swtch.com/version-sat>, <https://cseweb.ucsd.edu/~lerner/papers/opium.pdf>.

Now in the opposite direction: forcing aptitude package manager to solve Sudoku:

<http://web.archive.org/web/20160326062818/http://algebraicthunk.net/~dburrows/blog/entry/package>

Some readers may ask, how to order libraries/programs/packages to be installed? This is simpler problem, which is often solved by topological sorting. The algorithm reorders graph in such a way so that vertices not depended on anything will be on the top of queue. Next, there will be vertices dependend on vertices from the previous layer. And so on.

*make* UNIX utility does this while constructing order of items to be processed. Even more: older *make* utilities offloaded the job to the external utility (*tsort*). Some older UNIX has it, at least some versions of

## 7.16 Cracking simple XOR cipher with Z3

Here is a problem: a text encrypted with simple XOR cipher. Trying all possible keys is not an option.

Relationships between plain text, cipher text and key can be described using simple system of equations. But we can do more: we can ask Z3 to find such a key (array of bytes), so the plain text will have as many lowercase letters (a...z) as possible, but a solution will still satisfy all conditions.

```
import sys, hexdump
from z3 import *

def xor_strings(s,t):
    # https://en.wikipedia.org/wiki/XOR_cipher#Example_implementation
    """xor two strings together"""
    return "".join(chr(ord(a)^ord(b)) for a,b in zip(s,t))

def read_file(fname):
    file=open(fname, mode='rb')
    content=file.read()
    file.close()
    return content

def chunks(l, n):
    """divide input buffer by n-len chunks"""
    n = max(1, n)
    return [l[i:i + n] for i in range(0, len(l), n)]

def print_model(m, KEY_LEN, key):
    # fetch key from model:
    test_key="".join(chr(int(obj_to_string(m[key[i]]))) for i in range(KEY_LEN))
    print "key="
    hexdump.hexdump(test_key)

    # decrypt using the key:
    tmp=chunks(cipher_file, KEY_LEN)
    plain_attempt="".join(map(lambda x: xor_strings(x, test_key), tmp))
    print "plain="
    hexdump.hexdump(plain_attempt)

def try_len(KEY_LEN, cipher_file):
    cipher_len=len(cipher_file)
    print "len=", KEY_LEN
    s=Optimize()

    # variables for each byte of key:
    key=[BitVec('key_%d' % i, 8) for i in range (KEY_LEN)]
    # variables for each byte of input cipher text:
    cipher=[BitVec('cipher_%d' % i, 8) for i in range (cipher_len)]
    # variables for each byte of input plain text:
    plain=[BitVec('plain_%d' % i, 8) for i in range (cipher_len)]
    # variable for each byte of plain text: 1 if the byte in 'a'...'z' range:
    az_in_plain=[Int('az_in_plain_%d' % i) for i in range (cipher_len)]

    for i in range(cipher_len):
        # assign each byte of cipher text from the input file:
        s.add(cipher[i]==ord(cipher_file[i]))
        # plain text is cipher text XOR-ed with key:
        s.add(plain[i]==cipher[i]^key[i % KEY_LEN])
        # each byte must be in printable range, or CR or LF:
        s.add(Or(And(plain[i]>=0x20, plain[i]<=0x7E),plain[i]==0xA,plain[i]==0xD))
        # 1 if in 'a'...'z' range, 0 otherwise:
        s.add(az_in_plain[i]==If(And(plain[i]>=ord('a'),plain[i]<=ord('z')), 1, 0))

    # find solution, where the sum of all az_in_plain variables is maximum:
    s.maximize(Sum(*az_in_plain))

    if s.check()==unsat:
        return
    m=s.model()
    print_model(m, KEY_LEN, key)
```

<sup>52</sup><http://netbsd.gw.com/cgi-bin/man-cgi/man?tsort+1+NetBSD-current>

```
cipher_file=read_file (sys.argv[1])
```

```
for i in range(1,20):  
    try_len(i, cipher_file)
```

```
#try_len(17, cipher_file)
```

( The source code: [https://github.com/dennis714/SAT\\_SMT\\_article/blob/master/SMT/XOR/1.py](https://github.com/dennis714/SAT_SMT_article/blob/master/SMT/XOR/1.py) )  
Let's try it on a small 350-bytes file<sup>53</sup>:

```
% python 1.py cipher1.txt  
len= 1  
len= 2  
len= 3  
len= 4  
len= 5  
  
...  
  
len= 16  
len= 17  
key=  
00000000: 90 A0 21 52 48 84 FB FF 86 83 CF 50 46 12 7A F9 ...!RH.....PF.z.  
00000010: 36                                     6  
plain=  
00000000: 4D 72 2E 22 54 63 65 72 6F 6F 63 6D 27 48 6F 6C Mr."Tceroo cm'Hol  
00000010: 6A 65 73 2C 22 70 63 6F 20 74 61 73 26 72 73 75 jes,"pco tas&rsu  
00000020: 61 6B 6C 79 20 74 62 79 79 20 6F 61 74 63 27 69 akly tbyy oatc'i  
00000030: 6E 20 73 68 65 20 6F 68 79 6E 69 6D 67 73 2A 27 n she ohynings*'  
00000040: 73 61 76 62 0D 0A 75 72 68 65 20 74 6B 6F 73 63 savb..urhe tkosc  
00000050: 27 6E 6F 74 27 69 6E 66 70 62 7A 75 65 6D 74 20 'not'infpbzuent  
00000060: 69 64 63 61 73 6E 6F 6E 73 22 70 63 65 6E 23 68 idcasnons"pcen#h  
00000070: 65 26 70 61 73 20 72 70 20 61 6E 6B 2B 6E 69 64 e&pas rp ank+nid  
00000080: 68 74 2A 27 77 61 73 27 73 65 61 76 62 6F 0D 0A ht*'was'seavbo..  
00000090: 62 74 20 72 6F 65 20 62 75 65 61 6B 64 66 78 74 bt roe bueakdfxt  
000000A0: 20 77 61 62 6A 62 2E 20 49 27 73 74 6F 6D 63 2B wabjb. I'stomc+  
000000B0: 75 70 6C 6E 20 72 6F 65 20 68 62 61 72 74 6A 2A upln roe hbartj*  
000000C0: 79 75 67 23 61 6E 62 27 70 69 63 6C 65 64 20 77 yug#anb'picled w  
000000D0: 77 2B 74 68 66 0D 0A 75 73 69 63 6B 27 77 68 69 w+thf..usick'whi  
000000E0: 61 6F 2B 6F 75 71 20 76 6F 74 69 74 6F 75 20 68 ao+ouq votitou h  
000000F0: 61 66 27 67 65 66 77 20 62 63 6F 69 6E 64 27 68 af'gef w bcoind'h  
00000100: 69 6D 22 73 63 65 20 6D 69 67 6E 73 20 62 65 61 im"sce migns bea  
00000110: 6F 72 65 2C 27 42 74 20 74 61 73 26 66 0D 0A 66 ore,'Bt tas&f..f  
00000120: 6E 6E 65 2C 22 73 63 69 63 68 20 70 6F 62 63 65 nne,"scich pobce  
00000130: 20 68 66 20 77 6D 68 6F 2C 20 61 75 6C 64 68 75 hf wmho, auldh  
00000140: 73 2D 6F 65 61 64 67 63 27 20 6F 65 20 74 6E 62 s-oadgc' oe tnb  
00000150: 20 73 6F 75 74 20 77 6A 6E 68 68 20 6A 73 sout wjnhh js  
len= 18  
len= 19
```

This is not readable. But what is interesting, the solution exist only for 17-byte key.

What do we know about English language texts? Digits are rare there, so we can *minimize* them in plain text.

There are so called *digraphs*—a very popular combinations of two letters. The most popular in English are: *th*, *he*, *in* and *er*. We can count them in plain text and *maximize* them:

```
...  
  
# ... for each byte of plain text: 1 if the byte is digit:  
digits_in_plain=[Int('digits_in_plain_%d' % i) for i in range (cipher_len)]  
# ... for each byte of plain text: 1 if the byte + next byte is "th" characters:  
th_in_plain=[Int('th_in_plain_%d' % i) for i in range (cipher_len-1)]  
# ... etc:  
he_in_plain=[Int('he_in_plain_%d' % i) for i in range (cipher_len-1)]  
in_in_plain=[Int('in_in_plain_%d' % i) for i in range (cipher_len-1)]  
er_in_plain=[Int('er_in_plain_%d' % i) for i in range (cipher_len-1)]  
  
...  
  
for i in range(cipher_len-1):  
    # ... for each byte of plain text: 1 if the byte + next byte is "th" characters:  
    s.add(th_in_plain[i]==(If(And(plain[i]==ord('t'),plain[i+1]==ord('h')), 1, 0)))  
    # ... etc:
```

<sup>53</sup>[https://github.com/dennis714/SAT\\_SMT\\_article/blob/master/SMT/XOR/cipher1.txt](https://github.com/dennis714/SAT_SMT_article/blob/master/SMT/XOR/cipher1.txt)



```

s.add(he_in_plain[i]==(If(And(plain[i]==ord('h'),plain[i+1]==ord('e')), 1, 0)))
s.add(in_in_plain[i]==(If(And(plain[i]==ord('i'),plain[i+1]==ord('n')), 1, 0)))
s.add(er_in_plain[i]==(If(And(plain[i]==ord('e'),plain[i+1]==ord('r')), 1, 0)))

# find solution, where the sum of all az_in_plain variables is maximum:
s.maximize(Sum(*az_in_plain))
# ... and sum of digits_in_plain is minimum:
s.minimize(Sum(*digits_in_plain))

# "maximize" presence of "th", "he", "in" and "er" digraphs:
s.maximize(Sum(*th_in_plain))
s.maximize(Sum(*he_in_plain))
s.maximize(Sum(*in_in_plain))
s.maximize(Sum(*er_in_plain))

```

...

( The source code: [https://github.com/dennis714/SAT\\_SMT\\_article/blob/master/SMT/XOR/2.py](https://github.com/dennis714/SAT_SMT_article/blob/master/SMT/XOR/2.py) )  
Now this is something familiar:

```

len= 17
key=
00000000: 90 A0 22 50 4F 8F FB FF 85 83 CF 56 41 12 7A FE ...P0.....VA.z.
00000010: 31                                     1
plain=
00000000: 4D 72 2D 20 53 68 65 72 6C 6F 63 6B 20 48 6F 6B Mr- Sherlock Hok
00000010: 6D 65 73 2F 20 77 68 6F 20 77 61 73 20 75 73 75 mes/ who was usu
00000020: 66 6C 6C 79 23 76 65 72 79 20 6C 61 74 65 20 69 ffly#very late i
00000030: 6E 27 74 68 65 23 6D 6F 72 6E 69 6E 67 73 2C 20 n'the#mornings,
00000040: 73 61 71 65 0D 0A 76 70 6F 6E 20 74 68 6F 73 65 sage..vpon those
00000050: 20 6E 6F 73 20 69 6E 65 72 65 71 75 65 6E 74 20 nos inerequent
00000060: 6F 63 63 61 74 69 6F 6E 70 20 77 68 65 6E 20 68 occationp when h
00000070: 65 20 77 61 73 27 75 70 20 62 6C 6C 20 6E 69 67 e was'up bll nig
00000080: 68 74 2C 20 77 61 74 20 73 65 62 74 65 64 0D 0A ht, wat sebted..
00000090: 61 74 20 74 68 65 20 65 72 65 61 68 66 61 73 74 at the ereahfast
000000A0: 20 74 61 62 6C 65 2E 20 4E 20 73 74 6C 6F 64 20 table. N stlod
000000B0: 75 70 6F 6E 20 74 68 65 20 6F 65 61 72 77 68 2D upon the oearwh-
000000C0: 72 75 67 20 61 6E 64 20 70 69 64 6B 65 64 23 75 rug and picked#u
000000D0: 70 20 74 68 65 0D 0A 73 74 69 63 6C 20 77 68 6A p the..sticl whj
000000E0: 63 68 20 6F 75 72 20 76 69 73 69 74 68 72 20 68 ch our visithr h
000000F0: 62 64 20 6C 65 66 74 20 62 65 68 69 6E 63 20 68 bd left behinc h
00000100: 69 6E 20 74 68 65 20 6E 69 67 68 74 20 62 62 66 in the night bbf
00000110: 6F 72 66 2E 20 49 74 20 77 61 73 20 61 0D 0A 61 orf. It was a..a
00000120: 69 6E 65 2F 20 74 68 69 63 6B 20 70 69 65 63 65 ine/ thick piece
00000130: 27 6F 66 20 74 6F 6F 64 2C 20 62 75 6C 62 6F 75 'of tood, bulbou
00000140: 73 2A 68 65 61 67 65 64 2C 20 6F 66 20 74 68 65 s*heaged, of the
00000150: 20 73 68 72 74 20 74 68 69 63 68 20 69 73 shrt thich is

```

Several characters are wrong. But we can fix them, adding these conditions:

```

...
# 3 known characters of plain text:
s.add(plain[0xf]==ord('l'))
s.add(plain[0x20]==ord('a'))
s.add(plain[0x57]==ord('f'))
...

```

( The source code: [https://github.com/dennis714/SAT\\_SMT\\_article/blob/master/SMT/XOR/3.py](https://github.com/dennis714/SAT_SMT_article/blob/master/SMT/XOR/3.py) )  
This key is seems correct:

```

len= 17
key=
00000000: 90 A0 21 50 4F 8F FB FF 85 83 CF 56 41 12 7A F9 ...!P0.....VA.z.
00000010: 31                                     1
plain=
00000000: 4D 72 2E 20 53 68 65 72 6C 6F 63 6B 20 48 6F 6C Mr. Sherlock Hol
00000010: 6D 65 73 2C 20 77 68 6F 20 77 61 73 20 75 73 75 mes, who was usu
00000020: 61 6C 6C 79 20 76 65 72 79 20 6C 61 74 65 20 69 ally very late i
00000030: 6E 20 74 68 65 20 6D 6F 72 6E 69 6E 67 73 2C 20 n the mornings,
00000040: 73 61 76 65 0D 0A 75 70 6F 6E 20 74 68 6F 73 65 save..upon those
00000050: 20 6E 6F 74 20 69 6E 66 72 65 71 75 65 6E 74 20 not infrequent
00000060: 6F 63 63 61 73 69 6F 6E 73 20 77 68 65 6E 20 68 occasions when h
00000070: 65 20 77 61 73 20 75 70 20 61 6C 6C 20 6E 69 67 e was up all nig
00000080: 68 74 2C 20 77 61 73 20 73 65 61 74 65 64 0D 0A ht, was seated..
00000090: 61 74 20 74 68 65 20 62 72 65 61 6B 66 61 73 74 at the breakfast

```

000000A0:	20 74 61 62 6C 65 2E 20	49 20 73 74 6F 6F 64 20	table. I stood
000000B0:	75 70 6F 6E 20 74 68 65	20 68 65 61 72 74 68 2D	upon the hearth-
000000C0:	72 75 67 20 61 6E 64 20	70 69 63 6B 65 64 20 75	rug and picked u
000000D0:	70 20 74 68 65 0D 0A 73	74 69 63 6B 20 77 68 69	p the..stick whi
000000E0:	63 68 20 6F 75 72 20 76	69 73 69 74 6F 72 20 68	ch our visitor h
000000F0:	61 64 20 6C 65 66 74 20	62 65 68 69 6E 64 20 68	ad left behind h
00000100:	69 6D 20 74 68 65 20 6E	69 67 68 74 20 62 65 66	im the night bef
00000110:	6F 72 65 2E 20 49 74 20	77 61 73 20 61 0D 0A 66	ore. It was a..f
00000120:	69 6E 65 2C 20 74 68 69	63 6B 20 70 69 65 63 65	ine, thick piece
00000130:	20 6F 66 20 77 6F 6F 64	2C 20 62 75 6C 62 6F 75	of wood, bulbou
00000140:	73 2D 68 65 61 64 65 64	2C 20 6F 66 20 74 68 65	s-headed, of the
00000150:	20 73 6F 72 74 20 77 68	69 63 68 20 69 73	sort which is

So this is correct 17-byte XOR-key.

Needless to say, that the bigger ciphertext for analysis we have, the better. That 350-byte file is in fact the beginning of bigger file I prepared ([https://github.com/dennis714/SAT\\_SMT\\_article/blob/master/SMT/XOR/cipher2.txt](https://github.com/dennis714/SAT_SMT_article/blob/master/SMT/XOR/cipher2.txt), 12903 bytes). And a correct key for it can be found for it without additional *heuristics* we used here.

SMT solver is overkill for this. I once solved this problem naively, and it was much faster: [https://yurichev.com/blog/XOR\\_mask\\_2/](https://yurichev.com/blog/XOR_mask_2/). Nevertheless, this is yet another demonstration of yet another optimization problem.

The files: [https://github.com/dennis714/SAT\\_SMT\\_article/tree/master/SMT/XOR](https://github.com/dennis714/SAT_SMT_article/tree/master/SMT/XOR).

## 7.17 Balanced Gray code and Z3 SMT solver

Suppose, you are making a rotary encoder. This is a device that can signal its angle in some form, like:




Figure 8: Rotary encoder

( The image has been taken from Wikipedia: [https://en.wikipedia.org/wiki/Gray\\_code](https://en.wikipedia.org/wiki/Gray_code) )  
Click on [bigger image](#).

This is a rotary (shaft) encoder: [https://en.wikipedia.org/wiki/Rotary\\_encoder](https://en.wikipedia.org/wiki/Rotary_encoder).



Figure 9: Cropped and photoshopped version

( Source: [http://homepages.dordt.edu/~ddeboer//S10/304/c\\_at\\_d/304S10\\_RC\\_TRK.HTM](http://homepages.dordt.edu/~ddeboer//S10/304/c_at_d/304S10_RC_TRK.HTM) )  
Click on [bigger one](#).

There are pins and tracks on rotating wheel. How would you do this? Easiest way is to use binary code. But it has a problem: when a wheel is rotating, in a moment of transition from one state to another, several bits may be changed, hence, undesirable state may be present for a short period of time. This is bad. To deal with it, Gray code was invented: only 1 bit is changed during rotation. Like:

Decimal	Binary	Gray
0	0000	0000
1	0001	0001
2	0010	0011
3	0011	0010
4	0100	0110
5	0101	0111
6	0110	0101
7	0111	0100
8	1000	1100
9	1001	1101
10	1010	1111
11	1011	1110
12	1100	1010
13	1101	1011
14	1110	1001
15	1111	1000

Now the second problem. Look at the picture again. It has a lot of bit changes on the outer circles. And this is electromechanical device. Surely, you may want to make tracks as long as possible, to reduce wearing of both tracks and pins. This is a first problem. The second: wearing should be even across all tracks (this is balanced Gray code).

How we can find a table for all states using Z3:

```

#!/usr/bin/env python

from z3 import *

BITS=5

# how many times a run of bits for each bit can be changed (max).
# it can be 4 for 4-bit Gray code or 8 for 5-bit code.
# 12 for 6-bit code (maybe even less)
CHANGES_MAX=8

ROWS=2**BITS
MASK=ROWS-1 # 0x1f for 5 bits, 0xf for 4 bits, etc

def bool_to_int (b):
    if b==True:
        return 1
    return 0

s=Solver()

# add a constraint: Hamming distance between two bitvectors must be 1
# i.e., two bitvectors can differ in only one bit.
# for 4 bits it works like that:
#     s.add(Or(
#         And(a3!=b3,a2==b2,a1==b1,a0==b0),
#         And(a3==b3,a2!=b2,a1==b1,a0==b0),
#         And(a3==b3,a2==b2,a1!=b1,a0==b0),
#         And(a3==b3,a2==b2,a1==b1,a0!=b0)))
def hamming1(l1, l2):
    assert len(l1)==len(l2)
    r=[]
    for i in range(len(l1)):
        t=[]
        for j in range(len(l1)):
            if i==j:
                t.append(l1[j]!=l2[j])
            else:
                t.append(l1[j]==l2[j])
        r.append(And(t))
    s.add(Or(r))

# add a constraint: bitvectors must be different.
# for 4 bits works like this:
#     s.add(Or(a3!=b3, a2!=b2, a1!=b1, a0!=b0))
def not_eq(l1, l2):
    assert len(l1)==len(l2)
    t=[l1[i]!=l2[i] for i in range(len(l1))]
    s.add(Or(t))

code=[[Bool('code_%d_%d' % (r,c)) for c in range(BITS)] for r in range(ROWS)]
ch=[[Bool('ch_%d_%d' % (r,c)) for c in range(BITS)] for r in range(ROWS)]

# each rows must be different from a previous one and a next one by 1 bit:
for i in range(ROWS):
    # get bits of the current row:
    lst1=[code[i][bit] for bit in range(BITS)]
    # get bits of the next row.
    # important: if the current row is the last one, (last+1)&MASK==0, so we overlap here:
    lst2=[code[(i+1)&MASK][bit] for bit in range(BITS)]
    hamming1(lst1, lst2)

# no row must be equal to any another row:
for i in range(ROWS):
    for j in range(ROWS):
        if i==j:
            continue
        lst1=[code[i][bit] for bit in range(BITS)]
        lst2=[code[j][bit] for bit in range(BITS)]
        not_eq(lst1, lst2)

# 1 in ch[] table means that run of 1's has been changed to run of 0's, or back.
# "run" change detected using simple XOR:
for i in range(ROWS):
    for bit in range(BITS):
        # row overlapping works here as well:

```



```

        s.add(ch[i][bit]==Xor(code[i][bit],code[(i+1)&MASK][bit]))
# only CHANGES_MAX of 1 bits is allowed in ch[] table for each bit:
for bit in range(BITS):
    t=[ch[i][bit] for i in range(ROWS)]
    # this is a dirty hack.
    # AtMost() takes arguments like:
    # AtMost(v1, v2, v3, v4, 2) <- this means, only 2 booleans (or less) from the list can be True.
    # but we need to pass a list here.
    # so a CHANGES_MAX number is appended to a list and a new list is then passed as arguments list:
    s.add(AtMost(*(t+[CHANGES_MAX])))

result=s.check()
if result==unsat:
    exit(0)
m=s.model()

# get the model.
print "code table:"

for i in range(ROWS):
    for bit in range(BITS):
        # comma at the end means "no newline":
        print bool_to_int(is_true(m[code[i][BITS-1-bit]])),
    print ""

print "ch table:"

stat={}

for i in range(ROWS):
    for bit in range(BITS):
        x=is_true(m[ch[i][BITS-1-bit]])
        if x:
            # increment if bit is present in dict, set 1 if not present
            stat[bit]=stat.get(bit, 0)+1
        # comma at the end means "no newline":
        print bool_to_int(x),
    print ""

print "stat (bit number: number of changes): ", stat

```

( The source code: <https://github.com/dennis714/yurichev.com/blob/master/blog/gray/gray.py> )

For 4 bits, 4 changes is enough:

```

code table:
0 1 0 1
0 0 0 1
0 0 1 1
0 0 1 0
1 0 1 0
1 0 1 1
1 1 1 1
1 1 0 1
1 0 0 1
1 0 0 0
0 0 0 0
0 1 0 0
1 1 0 0
1 1 1 0
0 1 1 0
0 1 1 1
ch table:
0 1 0 0
0 0 1 0
0 0 0 1
1 0 0 0
0 0 0 1
0 1 0 0
0 0 1 0
0 1 0 0
0 0 0 1
1 0 0 0

```

```
0 1 0 0
1 0 0 0
0 0 1 0
1 0 0 0
0 0 0 1
0 0 1 0
stat (bit number: count of changes): {0: 4, 1: 4, 2: 4, 3: 4}
```

8 changes for 5 bits: <https://github.com/dennis714/yurichev.com/blob/master/blog/gray/5.txt>.  
12 for 6 bits (or maybe even less): <https://github.com/dennis714/yurichev.com/blob/master/blog/gray/6.txt>.

### 7.17.1 Duke Nukem 3D from 1990s




Figure 10: Duke Nukem 3D

Another application of Gray code:

with.inspiring@flair.and.erudition (Mike Naylor) wrote:

>In Duke Nukem, you often come upon panels that have four buttons in a row,  
>all in their "off" position. Each time you "push" a button, it toggles from  
>one state to the other. The object is to find the unique combination that  
>unlocks something in the game.

>My question is: What is the most efficient order in which to push the  
>buttons so that every combination is tested with no wasted effort?

A Gray Code. :-)

(Oh, you wanted to know what one would be? How about:

```
0000
0001
0011
```

```
0010
0110
0111
0101
0100
1100
1101
1111
1110
1010
1000
1001
1011
```

Or, if you prefer, with buttons A,B,C,D: D,C,D,B,D,C,D,A,D,C,D,B,C,D,C  
It isn't the "canonical" Gray code (or if it is, it is by Divine  
Providence), but it works.

Douglas Limmer -- lim...@math.orst.edu  
"No wonder these mathematical wizards were nuts - went off the beam -  
he'd be pure squirrel-food if he had half that stuff in \_his\_ skull!"  
E. E. "Doc" Smith, \_Second Stage Lensmen\_

( <https://groups.google.com/forum/#!topic/rec.puzzles/Dh2H-pGJcbI> )

Obviously, using our solution, you can minimize all movements in this ancient videogame, for 4 switches, that would be  $4*4=16$  switches. With our solution (balanced Gray code), wearing would be even across all 4 switches.

The same problem for MaxSAT: [14.1](#).

## 7.18 Integer factorization using Z3 SMT solver

Integer factorization is method of breaking a composite (non-prime number) into prime factors. Like  $12345 = 3*4*823$ .

Though for small numbers, this task can be accomplished by Z3:

```
#!/usr/bin/env python

import random
from z3 import *
from operator import mul

def factor(n):
    print "factoring",n

    in1,in2,out=Ints('in1 in2 out')

    s=Solver()
    s.add(out==n)
    s.add(in1*in2==out)
    # inputs cannot be negative and must be non-1:
    s.add(in1>1)
    s.add(in2>1)

    if s.check()==unsat:
        print n,"is prime (unsat)"
        return [n]
    if s.check()==unknown:
        print n,"is probably prime (unknown)"
        return [n]

    m=s.model()
    # get inputs of multiplier:
    in1_n=m[in1].as_long()
    in2_n=m[in2].as_long()

    print "factors of", n, "are", in1_n, "and", in2_n
    # factor factors recursively:
    rt=sorted(factor(in1_n) + factor(in2_n))
    # self-test:
    assert reduce(mul, rt, 1)==n
    return rt

# infinite test:
```



```
def test():
    while True:
        print factor (random.randrange(1000000000))

#test()

print factor(1234567890)
```

( The source code: [https://github.com/dennis714/yurichev.com/blob/master/blog/factor/factor\\_z3.py](https://github.com/dennis714/yurichev.com/blob/master/blog/factor/factor_z3.py) )

When factoring 1234567890 recursively:

```
% time python z.py
factoring 1234567890
factors of 1234567890 are 342270 and 3607
factoring 342270
factors of 342270 are 2 and 171135
factoring 2
2 is prime (unsat)
factoring 171135
factors of 171135 are 3803 and 45
factoring 3803
3803 is prime (unsat)
factoring 45
factors of 45 are 3 and 15
factoring 3
3 is prime (unsat)
factoring 15
factors of 15 are 5 and 3
factoring 5
5 is prime (unsat)
factoring 3
3 is prime (unsat)
factoring 3607
3607 is prime (unsat)
[2, 3, 3, 5, 3607, 3803]
python z.py 19.30s user 0.02s system 99% cpu 19.443 total
```

So,  $1234567890 = 2 \cdot 3 \cdot 3 \cdot 5 \cdot 3607 \cdot 3803$ .

One important note: there is no primality test, no lookup tables, etc. Prime number is a number for which " $x \cdot y = \text{prime}$ " (where  $x > 1$  and  $y > 1$ ) diophantine equation (which allows only integers in solution) has no solutions. It can be solved for real numbers, though.

Z3 is [not yet good enough for non-linear integer arithmetic](#) and sometimes returns "unknown" instead of "unsat", but, as Leonardo de Moura (one of Z3's author) commented about this:

```
...Z3 will solve the problem as a real problem. If no real solution is found, we know there is no integer solution.
If a solution is found, Z3 will check if the solution is really assigning integer values to integer variables.
If that is not the case, it will return unknown to indicate it failed to solve the problem.
```

( <https://stackoverflow.com/questions/13898175/how-does-z3-handle-non-linear-integer-arithmetic> )

Probably, this is the case: we getting "unknown" in the case when a number cannot be factored, i.e., it's prime.

It's also very slow. Wolfram Mathematica can factor number around  $2^{80}$  in a matter of seconds. Still, I've written this for demonstration.

The problem of breaking **RSA** is a problem of factorization of very large numbers, up to  $2^{4096}$ . It's currently not possible to do this in practice.

See also: integer factorization using SAT solver ([13.10](#)).

## 7.19 Tiling puzzle and Z3 SMT solver

This is classic problem: given 12 polyomino tiles, cover mutilated chessboard with them (it has 60 squares with no central 4 squares).

The problem is covered at least in [Donald E. Knuth - Dancing Links](#), and this Z3 solution has been inspired by it.

Another thing I've added: graph coloring. You see, my script gives correct solutions, but somewhat unpleasant visually. So I used colored pseudographics. There are 12 tiles, it's not a problem to assign 12 colors to them. But there is another heavily used SAT problem: graph coloring.

Given a graph, assign a color to each vertex/node, so that colors wouldn't be equal in adjacent nodes. The problem can be solved easily in SMT: assign variable to each vertex. If two vertices are connected, add a constraint:  $vertex1\_color \neq vertex2\_color$ . As simple as that. In my case, each polyomino is vertex and if polyomino is adjacent to another polyomino, an edge/link is added between vertices. So I did, and output is now colored.

But this is planar graph (i.e., a graph which is, if represented in two-dimensional space has no intersected edges/links). And here is a famous four color theorem can be used. The solution of tiled polyomios is in fact like planar graph, or, a map, like a world map. Theorem states that any planar graph (or map) can be colored only 4 colors.

This is true, even more, several tilings can be colors with only 3 colors:



Figure 11:

Now the classic: 12 pentominos and "mutilated" chess board, several solutions:



Figure 12:




Figure 13:

The source code: [https://github.com/dennis714/yurichev.com/blob/master/blog/tiling\\_Z3/tiling.py](https://github.com/dennis714/yurichev.com/blob/master/blog/tiling_Z3/tiling.py).

Further reading: [https://en.wikipedia.org/wiki/Exact\\_cover#Pentomino\\_tiling](https://en.wikipedia.org/wiki/Exact_cover#Pentomino_tiling).

Four-color theorem has an interesting story, it has been finally proved in 2005 by Coq proof assistant: [https://en.wikipedia.org/wiki/Four\\_color\\_theorem](https://en.wikipedia.org/wiki/Four_color_theorem).

## 8 Program synthesis

Program synthesis is a process of automatic program generation, in accordance with some specific goals.

## 8.1 Synthesis of simple program using Z3 SMT-solver

Sometimes, multiplication operation can be replaced with a several operations of shifting/addition/subtraction. Compilers do so, because pack of instructions can be executed faster.

For example, multiplication by 19 is replaced by GCC 5.4 with pair of instructions: `lea edx, [eax+eax*8]` and


`lea eax, [eax+edx*2]`. This is sometimes also called “superoptimization”.

Let’s see if we can find a shortest possible instructions pack for some specified multiplier.

As I’ve already wrote once, SMT-solver can be seen as a solver of huge systems of equations. The task is to construct such system of equations, which, when solved, could produce a short program. I will use electronics analogy here, it can make things a little simpler.

First of all, what our program will be consting of? There will be 3 operations allowed: ADD/SUB/SHL. Only registers allowed as operands, except for the second operand of SHL (which could be in 1..31 range). Each register will be assigned only once (as in [SSA<sup>54</sup>](#)).

And there will be some “magic block”, which takes all previous register states, it also takes operation type, operands and produces a value of next register’s state.



Now let’s take a look on our schematics on top level:

```

0 -> blk -> blk -> blk .. -> blk -> 0
1 -> blk -> blk -> blk .. -> blk -> multiplier


```

Each block takes previous state of registers and produces new states. There are two chains. First chain takes 0 as state of R0 at the very beginning, and the chain is supposed to produce 0 at the end (since zero multiplied by any value is still zero). The second chain takes 1 and must produce multiplier as the state of very last register (since 1 multiplied by multiplier must equal to multiplier).

Each block is “controlled” by operation type, operands, etc. For each column, there is each own set.

Now you can view these two chains as two equations. The ultimate goal is to find such state of all operation types and operands, so the first chain will equal to 0, and the second to multiplier.

Let’s also take a look into “magic block” inside:



Each selector can be viewed as a simple multipositional switch. If operation is SHL, a value in range of 1..31 is used as second operand.

So you can imagine this electric circuit and your goal is to turn all switches in such a state, so two chains will have 0 and multiplier on output. This sounds like logic puzzle in some way. Now we will try to use Z3 to solve this puzzle.

First, we define all variables:

```

R=[BitVec('S_{}_{}_c'.format(s, c), 32) for s in range(MAX_STEPS)] for c in range(CHAINS)]
op=[Int('op_{}_s'.format(s)) for s in range(MAX_STEPS)]
op1_reg=[Int('op1_reg_{}_s'.format(s)) for s in range(MAX_STEPS)]
op2_reg=[Int('op2_reg_{}_s'.format(s)) for s in range(MAX_STEPS)]
op2_imm=[BitVec('op2_imm_{}_s'.format(s), 32) for s in range(MAX_STEPS)]

```

<sup>54</sup>Static single assignment form

R[][] is registers state for each chain and each step.

On contrary, `op / op1_reg / op2_reg / op2_imm` variables are defined for each step, but for both chains, since both chains at each column has the same operation/operands.

Now we must limit count of operations, and also, register's number for each step must not be bigger than step number, in other words, instruction at each step is allowed to access only registers which were already set before:

```
for s in range(1, STEPS):
    # for each step
    sl.add(And(op[s]>=0, op[s]<=2))
    sl.add(And(op1_reg[s]>=0, op1_reg[s]<s))
    sl.add(And(op2_reg[s]>=0, op2_reg[s]<s))
    sl.add(And(op2_imm[s]>=1, op2_imm[s]<=31))
```

Fix register of first step for both chains:

```
for c in range(CHAINS):
    # for each chain:
    sl.add(R[c][0]==chain_inputs[c])
    sl.add(R[c][STEPS-1]==chain_inputs[c]*multiplier)
```

Now let's add "magic blocks":

```
for s in range(1, STEPS):
    sl.add(R[c][s]==simulate_op(R,c, op[s], op1_reg[s], op2_reg[s], op2_imm[s]))
```

Now how "magic block" is defined?

```
def selector(R, c, s):
    # for all MAX STEPS:
    return If(s==0, R[c][0],
            If(s==1, R[c][1],
            If(s==2, R[c][2],
            If(s==3, R[c][3],
            If(s==4, R[c][4],
            If(s==5, R[c][5],
            If(s==6, R[c][6],
            If(s==7, R[c][7],
            If(s==8, R[c][8],
            If(s==9, R[c][9],
            0)))))) # default

def simulate_op(R, c, op, op1_reg, op2_reg, op2_imm):
    op1_val=selector(R,c,op1_reg)
    return If(op==0, op1_val + selector(R, c, op2_reg),
            If(op==1, op1_val - selector(R, c, op2_reg),
            If(op==2, op1_val << op2_imm,
            0))) # default
```

This is very important to understand: if the operation is ADD/SUB, `op2_imm`'s value is just ignored. Otherwise, if operation is SHL, value of `op2_reg` is ignored. Just like in case of digital circuit.

The code: [https://github.com/dennis714/SAT\\_SMT\\_article/blob/master/pgm\\_synth/mult.py](https://github.com/dennis714/SAT_SMT_article/blob/master/pgm_synth/mult.py).

Now let's see how it works:

```
% ./mult.py 12
multiplier= 12
attempt, STEPS= 2
unsat
attempt, STEPS= 3
unsat
attempt, STEPS= 4
sat!
r1=SHL r0, 2
r2=SHL r1, 1
r3=ADD r1, r2
tests are OK
```

The first step is always a step containing 0/1, or, r0. So when our solver reporting about 4 steps, this means 3 instructions.

Something harder:

```
% ./mult.py 123
multiplier= 123
attempt, STEPS= 2
unsat
attempt, STEPS= 3
unsat
attempt, STEPS= 4
unsat
attempt, STEPS= 5
sat!
r1=SHL r0, 2
r2=SHL r1, 5
r3=SUB r2, r1
r4=SUB r3, r0
tests are OK
```

Now the code multiplying by 1234:

```
r1=SHL r0, 6
r2=ADD r0, r1
r3=ADD r2, r1
r4=SHL r2, 4
r5=ADD r2, r3
r6=ADD r5, r4
```

Looks great, but it took  $\approx 23$  seconds to find it on my Intel Xeon CPU E31220 @ 3.10GHz. I agree, this is far from practical usage. Also, I'm not quite sure that this piece of code will work faster than a single multiplication instruction. But anyway, it's a good demonstration of SMT solvers capabilities.

The code multiplying by 12345 ( $\approx 150$  seconds):

```
r1=SHL r0, 5
r2=SHL r0, 3
r3=SUB r2, r1
r4=SUB r1, r3
r5=SHL r3, 9
r6=SUB r4, r5
r7=ADD r0, r6
```

Multiplication by 123456 ( $\approx 8$  minutes!):

```
r1=SHL r0, 9
r2=SHL r0, 13
r3=SHL r0, 2
r4=SUB r1, r2
r5=SUB r3, r4
r6=SHL r5, 4
r7=ADD r1, r6
```

### 8.1.1 Few notes

I've removed SHR instruction support, simply because the code multiplying by a constant makes no use of it. Even more: it's not a problem to add support of constants as second operand for all instructions, but again, you wouldn't find a piece of code which does this job and uses some additional constants. Or maybe I wrong?

Of course, for another job you'll need to add support of constants and other operations. But at the same time, it will work slower and slower. So I had to keep [ISA<sup>55</sup>](#) of this toy [CPU<sup>56</sup>](#) as compact as possible.

### 8.1.2 The code

[https://github.com/dennis714/SAT\\_SMT\\_article/blob/master/pgm\\_synth/mult.py](https://github.com/dennis714/SAT_SMT_article/blob/master/pgm_synth/mult.py).

## 8.2 Rockey dongle: finding unknown algorithm using only input/output pairs

(This text was first published in August 2012 in my blog: <http://blog.yurichev.com/node/71>.)

Some smartcards can execute Java or .NET code - that's the way to hide your sensitive algorithm into chip that is very hard to break (decapsulate). For example, one may encrypt/decrypt data files by hidden crypto

<sup>55</sup>Instruction Set Architecture

<sup>56</sup>Central processing unit

algorithm rendering software piracy of such software close to impossible — an encrypted data file created on software with connected smartcard would be impossible to decrypt on cracked version of the same software. (This leads to many nuisances, though.)

That's what called *black box*.

Some software protection dongles offers this functionality too. One example is Rockey 4<sup>57</sup>.




Figure 14: Rockey 4 dongle

This is a small dongle connected via USB. It contains some user-defined memory but also memory for user algorithms.

The virtual (toy) CPU for these algorithms is very simple: it offers only 8 16-bit registers (however, only 4 can be set and read) and 8 operations (addition, subtraction, cyclic left shifting, multiplication, OR, XOR, AND, negation).

Second instruction argument can be a constant (from 0 to 63) instead of register.

Each algorithm is described by string like

`A=A+B, B=C*13, D=D^A, C=B*55, C=C&A, D=D|A, A=A*9, A=A&B.`

There are no memory, stack, conditional/unconditional jumps, etc.

Each algorithm, obviously, can't have side effects, so they are actually *pure functions* and their results can be *memoized*.

By the way, as it has been mentioned in Rockey 4 manual, first and last instruction cannot have constants. Maybe that's because these fields used for some internal data: each algorithm start and end should be marked somehow internally anyway.

Would it be possible to reveal hidden impossible-to-read algorithm only by recording input/output dongle traffic? Common sense tells us "no". But we can try anyway.

Since, my goal wasn't to break into some Rockey-protected software, I was interested only in limits (which algorithms could we find), so I make some things simpler: we will work with only 4 16-bit registers, and there will be only 6 operations (add, subtract, multiply, OR, XOR, AND).

Let's first calculate, how much information will be used in brute-force case.

There are 384 of all possible instructions in `reg=reg,op,reg` format for 4 registers and 6 operations, and also 6144 instructions in `reg=reg,op,constant` format. Remember that constant limited to 63 as maximal value? That helps us for a little.

So, there are 6528 of all possible instructions. This means, there are  $\approx 1.1 \cdot 10^{19}$  5-instruction algorithms. That's too much.

How can we express each instruction as system of equations? While remembering some school mathematics, I wrote this:

```
Function one_step()=
# Each Bx is integer, but may be only 0 or 1.

# only one of B1..B4 and B5..B9 can be set
reg1=B1*A + B2*B + B3*C + B4*D
reg_or_constant2=B5*A + B6*B + B7*C + B8*D + B9*constant
reg1 should not be equal to reg_or_constant2

# Only one of B10..B15 can be set
result=result+B10*(reg1*reg2)
result=result+B11*(reg1^reg2)
result=result+B12*(reg1+reg2)
result=result+B13*(reg1-reg2)
```

<sup>57</sup><http://www.rockey.nl/en/rockey.html>



```
result=result+B14*(reg1|reg2)
result=result+B15*(reg1&reg2)
```

B16 - true if register isn't updated in this part  
 B17 - true if register is updated in this part  
 (B16 cannot be equal to B17)  
 A=B16\*A + B17\*result  
 B=B18\*A + B19\*result  
 C=B20\*A + B21\*result  
 D=B22\*A + B23\*result

That's how we can express each instruction in algorithm.

5-instructions algorithm can be expressed like this:

```
one_step (one_step (one_step (one_step (one_step (input_registers))))).
```

Let's also add five known input/output pairs and we'll get system of equations like this:

```
one_step (one_step (one_step (one_step (one_step (input_1))))==output_1
one_step (one_step (one_step (one_step (one_step (input_2))))==output_2
one_step (one_step (one_step (one_step (one_step (input_3))))==output_3
one_step (one_step (one_step (one_step (one_step (input_4))))==output_4
.. etc
```

So the question now is to find  $5 \cdot 2^3$  boolean values satisfying known input/output pairs.

I wrote small utility to probe Rocky 4 algorithm with random numbers, it produce results in form:

```
RY_CALCULATE1: (input) p1=30760 p2=18484 p3=41200 p4=61741 (output) p1=49244 p2=11312 p3=27587 p4=12657
RY_CALCULATE1: (input) p1=51139 p2=7852 p3=53038 p4=49378 (output) p1=58991 p2=34134 p3=40662 p4=9869
RY_CALCULATE1: (input) p1=60086 p2=52001 p3=13352 p4=45313 (output) p1=46551 p2=42504 p3=61472 p4=1238
RY_CALCULATE1: (input) p1=48318 p2=6531 p3=51997 p4=30907 (output) p1=54849 p2=20601 p3=31271 p4=44794
```

p1/p2/p3/p4 are just another names for A/B/C/D registers.

Now let's start with Z3. We will need to express Rocky 4 toy CPU in Z3Py (Z3 Python [API](#)) terms.

It can be said, my Python script is divided into two parts:

- constraint definitions (like, *output\_1 should be n for input\_1=m, constant cannot be greater than 63, etc*);
- functions constructing system of equations.

This piece of code define some kind of *structure* consisting of 4 named 16-bit variables, each represent register in our toy CPU.

```
Registers_State=Datatype ('Registers_State')
Registers_State.declare('cons', ('A', BitVecSort(16)), ('B', BitVecSort(16)), ('C', BitVecSort(16)), ('D',
    BitVecSort(16)))
Registers_State=Registers_State.create()
```

These enumerations define two new types (or *sorts* in Z3's terminology):

```
Operation, (OP_MULT, OP_MINUS, OP_PLUS, OP_XOR, OP_OR, OP_AND) = EnumSort('Operation', ('OP_MULT', 'OP_MINUS', '
    OP_PLUS', 'OP_XOR', 'OP_OR', 'OP_AND'))

Register, (A, B, C, D) = EnumSort('Register', ('A', 'B', 'C', 'D'))
```

This part is very important, it defines all variables in our system of equations. `op_step` is type of operation in instruction. `reg_or_constant` is selector between register and constant in second argument — *False* if it's a register and *True* if it's a constant. `reg_step` is a destination register of this instruction. `reg1_step` and `reg2_step` are just registers at arg1 and arg2. `constant_step` is constant (in case it's used in instruction instead of arg2).

```
op_step=[Const('op_step%s' % i, Operation) for i in range(STEPS)]
reg_or_constant_step=[Bool('reg_or_constant_step%s' % i) for i in range(STEPS)]
reg_step=[Const('reg_step%s' % i, Register) for i in range(STEPS)]
reg1_step=[Const('reg1_step%s' % i, Register) for i in range(STEPS)]
reg2_step=[Const('reg2_step%s' % i, Register) for i in range(STEPS)]
constant_step = [BitVec('constant_step%s' % i, 16) for i in range(STEPS)]
```

Adding constraints is very simple. Remember, I wrote that each constant cannot be larger than 63?

```
# according to Rockey 4 dongle manual, arg2 in first and last instructions cannot be a constant
s.add (reg_or_constant_step[0]==False)
s.add (reg_or_constant_step[STEPS-1]==False)

...

for x in range(STEPS):
    s.add (constant_step[x]>=0, constant_step[x]<=63)
```

Known input/output values are added as constraints too.  
Now let's see how to construct our system of equations:

```
# Register, Registers_State -> int
def register_selector(register, input_registers):
    return If(register==A, Registers_State.A(input_registers),
              If(register==B, Registers_State.B(input_registers),
                If(register==C, Registers_State.C(input_registers),
                  If(register==D, Registers_State.D(input_registers),
                    0)))) # default
```

This function returning corresponding register value from *structure*. Needless to say, the code above is not executed. `If()` is Z3Py function. The code only declares the function, which will be used in another. Expression declaration resembling LISP [PL](#) in some way.

Here is another function where `register_selector()` is used:

```
# Bool, Register, Registers_State, int -> int
def register_or_constant_selector(register_or_constant, register, input_registers, constant):
    return If(register_or_constant==False, register_selector(register, input_registers), constant)
```

The code here is never executed too. It only constructs one small piece of very big expression. But for the sake of simplicity, one can think all these functions will be called during bruteforce search, many times, at fastest possible speed.

```
# Operation, Bool, Register, Register, Int, Registers_State -> int
def one_op(op, register_or_constant, reg1, reg2, constant, input_registers):
    arg1=register_selector(reg1, input_registers)
    arg2=register_or_constant_selector(register_or_constant, reg2, input_registers, constant)
    return If(op==OP_MULT, arg1*arg2,
              If(op==OP_MINUS, arg1-arg2,
                If(op==OP_PLUS, arg1+arg2,
                  If(op==OP_XOR, arg1^arg2,
                    If(op==OP_OR, arg1|arg2,
                      If(op==OP_AND, arg1&arg2,
                        0)))))) # default
```

Here is the expression describing each instruction. `new_val` will be assigned to destination register, while all other registers' values are copied from input registers' state:

```
# Bool, Register, Operation, Register, Register, Int, Registers_State -> Registers_State
def one_step(register_or_constant, register_assigned_in_this_step, op, reg1, reg2, constant, input_registers):
    new_val=one_op(op, register_or_constant, reg1, reg2, constant, input_registers)
    return If (register_assigned_in_this_step==A, Registers_State.cons (new_val,
                                                                        Registers_State.B(input_registers),
                                                                        Registers_State.C(input_registers),
                                                                        Registers_State.D(input_registers)),
              If (register_assigned_in_this_step==B, Registers_State.cons (Registers_State.A(input_registers),
                                                                        new_val,
                                                                        Registers_State.C(input_registers),
                                                                        Registers_State.D(input_registers)),
                If (register_assigned_in_this_step==C, Registers_State.cons (Registers_State.A(input_registers),
                                                                        Registers_State.B(input_registers),
                                                                        new_val,
                                                                        Registers_State.D(input_registers)),
                  If (register_assigned_in_this_step==D, Registers_State.cons (Registers_State.A(input_registers),
                                                                        Registers_State.B(input_registers),
                                                                        Registers_State.C(input_registers),
                                                                        new_val),
                    Registers_State.cons(0,0,0,0)))) # default
```

This is the last function describing a whole *n*-step program:

```
def program(input_registers, STEPS):
    cur_input=input_registers
```

```

for x in range(STEPS):
    cur_input=one_step (reg_or_constant_step[x], reg_step[x], op_step[x], reg1_step[x], reg2_step[x],
        constant_step[x], cur_input)
    return cur_input

```

Again, for the sake of simplicity, it can be said, now Z3 will try each possible registers/operations/constants against this expression to find such combination which satisfy all input/output pairs. Sounds absurdic, but this is close to reality. SAT/SMT-solvers indeed tries them all. But the trick is to prune search tree as early as possible, so it will work for some reasonable time. And this is hardest problem for solvers.

Now let's start with very simple 3-step algorithm:  $B=A^D$ ,  $C=D*D$ ,  $D=A*C$ . Please note: register **A** left unchanged. I programmed Rocky 4 dogle with the algorithm, and recorded algorithm outputs are:

```

RY_CALCULATE1: (input) p1=8803 p2=59946 p3=36002 p4=44743 (output) p1=8803 p2=36004 p3=7857 p4=24691
RY_CALCULATE1: (input) p1=5814 p2=55512 p3=52155 p4=55813 (output) p1=5814 p2=52403 p3=33817 p4=4038
RY_CALCULATE1: (input) p1=25206 p2=2097 p3=55906 p4=22705 (output) p1=25206 p2=15047 p3=10849 p4=43702
RY_CALCULATE1: (input) p1=10044 p2=14647 p3=27923 p4=7325 (output) p1=10044 p2=15265 p3=47177 p4=20508
RY_CALCULATE1: (input) p1=15267 p2=2690 p3=47355 p4=56073 (output) p1=15267 p2=57514 p3=26193 p4=53395

```

It took about one second and only 5 pairs above to find algorithm (on my quad-core Xeon E3-1220 3.1GHz, however, Z3 solver working in single-thread mode):

```

B = A ^ D
C = D * D
D = C * A

```

Note the last instruction: **C** and **A** registers are swapped comparing to version I wrote by hand. But of course, this instruction is working in the same way, because multiplication is commutative operation.

Now if I try to find 4-step program satisfying to these values, my script will offer this:

```

B = A ^ D
C = D * D
D = A * C
A = A | A

```

...and that's really fun, because the last instruction do nothing with value in register **A**, it's like **NOP**<sup>58</sup>—but still, algorithm is correct for all values given.

Here is another 5-step algorithm:  $B=B^D$ ,  $C=A*22$ ,  $A=B*19$ ,  $A=A\&42$ ,  $D=B\&C$  and values:

```

RY_CALCULATE1: (input) p1=61876 p2=28737 p3=28636 p4=50362 (output) p1=32 p2=46331 p3=50552 p4=33912
RY_CALCULATE1: (input) p1=46843 p2=43355 p3=39078 p4=24552 (output) p1=8 p2=63155 p3=47506 p4=45202
RY_CALCULATE1: (input) p1=22425 p2=51432 p3=40836 p4=14260 (output) p1=0 p2=65372 p3=34598 p4=34564
RY_CALCULATE1: (input) p1=44214 p2=45766 p3=19778 p4=59924 (output) p1=2 p2=22738 p3=55204 p4=20608
RY_CALCULATE1: (input) p1=27348 p2=49060 p3=31736 p4=59576 (output) p1=0 p2=22300 p3=11832 p4=1560

```

It took 37 seconds and we've got:

```

B = D ^ B
C = A * 22
A = B * 19
A = A & 42
D = C & B

```

$A=A\&42$  was correctly deduced (look at these five p1's at output (assigned to output **A** register): 32,8,0,2,0)

6-step algorithm  $A=A+B$ ,  $B=C*13$ ,  $D=D^A$ ,  $C=C\&A$ ,  $D=D|B$ ,  $A=A\&B$  and values:

```

RY_CALCULATE1: (input) p1=4110 p2=35411 p3=54308 p4=47077 (output) p1=32832 p2=50644 p3=36896 p4=60884
RY_CALCULATE1: (input) p1=12038 p2=7312 p3=39626 p4=47017 (output) p1=18434 p2=56386 p3=2690 p4=64639
RY_CALCULATE1: (input) p1=48763 p2=27663 p3=12485 p4=20563 (output) p1=10752 p2=31233 p3=8320 p4=31449
RY_CALCULATE1: (input) p1=33174 p2=38937 p3=54005 p4=38871 (output) p1=4129 p2=46705 p3=4261 p4=48761
RY_CALCULATE1: (input) p1=46587 p2=36275 p3=6090 p4=63976 (output) p1=258 p2=13634 p3=906 p4=48966

```

90 seconds and we've got:

```

A = A + B
B = C * 13
D = D ^ A
D = B | D
C = C & A
A = B & A

```

<sup>58</sup>No Operation

But that was simple, however. Some 6-step algorithms are not possible to find, for example:  
`A=A^B, A=A*9, A=A^C, A=A*19, A=A^D, A=A&B`. Solver was working too long (up to several hours), so I didn't even know is it possible to find it anyway.

### 8.2.1 Conclusion

This is in fact an exercise in program synthesis.

Some short algorithms for tiny CPUs are really possible to find using so small set set of data. Of course it's still not possible to reveal some complex algorithm, but this method definitely should not be ignored.

### 8.2.2 The files

Rockey 4 dongle programmer and reader, Rocky 4 manual, Z3Py script for finding algorithms, input/output pairs: [https://github.com/dennis714/SAT\\_SMT\\_article/tree/master/pgm\\_synth/rockey\\_files](https://github.com/dennis714/SAT_SMT_article/tree/master/pgm_synth/rockey_files).

### 8.2.3 Further work

Perhaps, constructing LISP-like S-expression can be better than a program for toy-level CPU.

It's also possible to start with smaller constants and then proceed to bigger. This is somewhat similar to increasing password length in password brute-force cracking.

### 8.2.4 Exercise

<https://challenges.re/25/>.

## 9 Toy decompiler

### 9.1 Introduction

A modern-day compiler is a product of hundreds of developer/year. At the same time, toy compiler can be an exercise for a student for a week (or even weekend).

Likewise, commercial decompiler like Hex-Rays can be extremely complex, while toy decompiler like this one, can be easy to understand and remake.

The following decompiler written in Python, supports only short basic blocks, with no jumps. Memory is also not supported.


### 9.2 Data structure

Our toy decompiler will use just one single data structure, representing expression tree.

Many programming textbooks has an example of conversion from Fahrenheit temperature to Celsius, using the following formula:

$$celsius = (fahrenheit - 32) \cdot \frac{5}{9}$$

This expression can be represented as a tree:



How to store it in memory? We see here 3 types of nodes: 1) numbers (or values); 2) arithmetical operations; 3) symbols (like "INPUT").

Many developers with OOP<sup>59</sup> in their mind will create some kind of class. Other developer maybe will use "variant type".

I'll use simplest possible way of representing this structure: a Python tuple. First element of tuple can be a string: either "EXPR\_OP" for operation, "EXPR\_SYMBOL" for symbol or "EXPR\_VALUE" for value. In case of symbol or value, it follows the string. In case of operation, the string followed by another tuples.

Node type and operation type are stored as plain strings—to make debugging output easier to read.

There are *constructors* in our code, in OOP sense:

```
def create_val_expr (val):
    return ("EXPR_VALUE", val)

def create_symbol_expr (val):
    return ("EXPR_SYMBOL", val)

def create_binary_expr (op, op1, op2):
    return ("EXPR_OP", op, op1, op2)
```

There are also *accessors*:

```
def get_expr_type(e):
    return e[0]

def get_symbol (e):
    assert get_expr_type(e)=="EXPR_SYMBOL"
    return e[1]

def get_val (e):
    assert get_expr_type(e)=="EXPR_VALUE"
    return e[1]

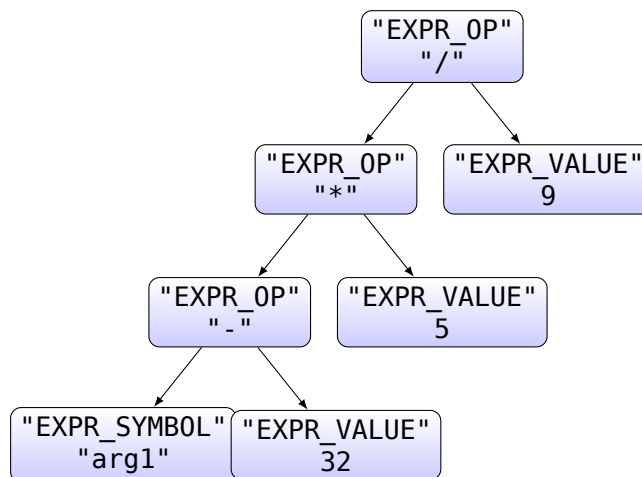
def is_expr_op(e):
    return get_expr_type(e)=="EXPR_OP"

def get_op (e):
    assert is_expr_op(e)
    return e[1]

def get_op1 (e):
    assert is_expr_op(e)
    return e[2]

def get_op2 (e):
    assert is_expr_op(e)
    return e[3]
```

The temperature conversion expression we just saw will be represented as:



...or as Python expression:

---

<sup>59</sup>Object-oriented programming

```
( 'EXPR_OP', '/',
  ( 'EXPR_OP', '*',
    ( 'EXPR_OP', '-', ( 'EXPR_SYMBOL', 'arg1' ), ( 'EXPR_VALUE', 32 ) ),
    ( 'EXPR_VALUE', 5 ) ),
  ( 'EXPR_VALUE', 9 ) )
```

In fact, this is [AST<sup>60</sup>](#) in its simplest form. [ASTs](#) are used heavily in compilers.

## 9.3 Simple examples

Let's start with simplest example:

```
mov    rax, rdi
imul   rax, rsi
```

At start, these symbols are assigned to registers: RAX=initial\_RAX, RBX=initial\_RBX, RDI=arg1, RSI=arg2, RDX=arg3, RCX=arg4.

When we handle MOV instruction, we just copy expression from RDI to RAX. When we handle IMUL instruction, we create a new expression, adding together expressions from RAX and RSI and putting result into RAX again.

I can feed this to decompiler and we will see how register's state is changed through processing:

```
python td.py --show-registers --python-expr tests/mul.s

...

line=[mov    rax, rdi]
rcx=('EXPR_SYMBOL', 'arg4')
rsi=('EXPR_SYMBOL', 'arg2')
rbx=('EXPR_SYMBOL', 'initial_RBX')
rdx=('EXPR_SYMBOL', 'arg3')
rdi=('EXPR_SYMBOL', 'arg1')
rax=('EXPR_SYMBOL', 'arg1')

line=[imul   rax, rsi]
rcx=('EXPR_SYMBOL', 'arg4')
rsi=('EXPR_SYMBOL', 'arg2')
rbx=('EXPR_SYMBOL', 'initial_RBX')
rdx=('EXPR_SYMBOL', 'arg3')
rdi=('EXPR_SYMBOL', 'arg1')
rax=('EXPR_OP', '*', ( 'EXPR_SYMBOL', 'arg1' ), ( 'EXPR_SYMBOL', 'arg2' ))

...

result=('EXPR_OP', '*', ( 'EXPR_SYMBOL', 'arg1' ), ( 'EXPR_SYMBOL', 'arg2' ))
```

IMUL instruction is mapped to "\*" string, and then new expression is constructed in `handle_binary_op()`, which puts result into RAX.

In this output, the data structures are dumped using Python `str()` function, which does mostly the same, as `print()`.

Output is bulky, and we can turn off Python expressions output, and see how this internal data structure can be rendered neatly using our internal `expr_to_string()` function:

```
python td.py --show-registers tests/mul.s

...

line=[mov    rax, rdi]
rcx=arg4
rsi=arg2
rbx=initial_RBX
rdx=arg3
rdi=arg1
rax=arg1

line=[imul   rax, rsi]
rcx=arg4
rsi=arg2
```

<sup>60</sup>Abstract syntax tree

```
rbx=initial_RBX
rdx=arg3
rdi=arg1
rax=(arg1 * arg2)

...

result=(arg1 * arg2)
```

Slightly advanced example:

```
imul    rdi, rsi
lea     rax, [rdi+rdx]
```

LEA instruction is treated just as ADD.

```
python td.py --show-registers --python-expr tests/mul_add.s
```

```
...

line=[imul    rdi, rsi]
rcx=('EXPR_SYMBOL', 'arg4')
rsi=('EXPR_SYMBOL', 'arg2')
rbx=('EXPR_SYMBOL', 'initial_RBX')
rdx=('EXPR_SYMBOL', 'arg3')
rdi=('EXPR_OP', '*', ('EXPR_SYMBOL', 'arg1'), ('EXPR_SYMBOL', 'arg2'))
rax=('EXPR_SYMBOL', 'initial_RAX')

line=[lea     rax, [rdi+rdx]]
rcx=('EXPR_SYMBOL', 'arg4')
rsi=('EXPR_SYMBOL', 'arg2')
rbx=('EXPR_SYMBOL', 'initial_RBX')
rdx=('EXPR_SYMBOL', 'arg3')
rdi=('EXPR_OP', '*', ('EXPR_SYMBOL', 'arg1'), ('EXPR_SYMBOL', 'arg2'))
rax=('EXPR_OP', '+', ('EXPR_OP', '*', ('EXPR_SYMBOL', 'arg1'), ('EXPR_SYMBOL', 'arg2')), ('EXPR_SYMBOL', 'arg3'))

...

result=('EXPR_OP', '+', ('EXPR_OP', '*', ('EXPR_SYMBOL', 'arg1'), ('EXPR_SYMBOL', 'arg2')), ('EXPR_SYMBOL', 'arg3'))
```

And again, let's see this expression dumped neatly:

```
python td.py --show-registers tests/mul_add.s
```

```
...

result=((arg1 * arg2) + arg3)
```

Now another example, where we use 2 input arguments:

```
imul    rdi, rdi, 1234
imul    rsi, rsi, 5678
lea     rax, [rdi+rsi]
```

```
python td.py --show-registers --python-expr tests/mul_add3.s
```

```
...

line=[imul    rdi, rdi, 1234]
rcx=('EXPR_SYMBOL', 'arg4')
rsi=('EXPR_SYMBOL', 'arg2')
rbx=('EXPR_SYMBOL', 'initial_RBX')
rdx=('EXPR_SYMBOL', 'arg3')
rdi=('EXPR_OP', '*', ('EXPR_SYMBOL', 'arg1'), ('EXPR_VALUE', 1234))
rax=('EXPR_SYMBOL', 'initial_RAX')

line=[imul    rsi, rsi, 5678]
rcx=('EXPR_SYMBOL', 'arg4')
rsi=('EXPR_OP', '*', ('EXPR_SYMBOL', 'arg2'), ('EXPR_VALUE', 5678))
rbx=('EXPR_SYMBOL', 'initial_RBX')
rdx=('EXPR_SYMBOL', 'arg3')
rdi=('EXPR_OP', '*', ('EXPR_SYMBOL', 'arg1'), ('EXPR_VALUE', 1234))
rax=('EXPR_SYMBOL', 'initial_RAX')
```

```

line=[lea      rax, [rdi+rsi]]
rcx=('EXPR_SYMBOL', 'arg4')
rsi=('EXPR_OP', '*', ('EXPR_SYMBOL', 'arg2'), ('EXPR_VALUE', 5678))
rbx=('EXPR_SYMBOL', 'initial_RBX')
rdx=('EXPR_SYMBOL', 'arg3')
rdi=('EXPR_OP', '*', ('EXPR_SYMBOL', 'arg1'), ('EXPR_VALUE', 1234))
rax=('EXPR_OP', '+', ('EXPR_OP', '*', ('EXPR_SYMBOL', 'arg1'), ('EXPR_VALUE', 1234)), ('EXPR_OP', '*', ('
    EXPR_SYMBOL', 'arg2'), ('EXPR_VALUE', 5678)))

...

result=('EXPR_OP', '+', ('EXPR_OP', '*', ('EXPR_SYMBOL', 'arg1'), ('EXPR_VALUE', 1234)), ('EXPR_OP', '*', ('
    EXPR_SYMBOL', 'arg2'), ('EXPR_VALUE', 5678)))

```

...and now neat output:

```

python td.py --show-registers tests/mul_add3.s

...

result=((arg1 * 1234) + (arg2 * 5678))

```

Now conversion program:

```

mov    rax, rdi
sub    rax, 32
imul   rax, 5
mov    rbx, 9
idiv   rbx

```

You can see, how register's state is changed over execution (or parsing).

Raw:

```

python td.py --show-registers --python-expr tests/fahr_to_celsius.s

...

line=[mov      rax, rdi]
rcx=('EXPR_SYMBOL', 'arg4')
rsi=('EXPR_SYMBOL', 'arg2')
rbx=('EXPR_SYMBOL', 'initial_RBX')
rdx=('EXPR_SYMBOL', 'arg3')
rdi=('EXPR_SYMBOL', 'arg1')
rax=('EXPR_SYMBOL', 'arg1')

line=[sub      rax, 32]
rcx=('EXPR_SYMBOL', 'arg4')
rsi=('EXPR_SYMBOL', 'arg2')
rbx=('EXPR_SYMBOL', 'initial_RBX')
rdx=('EXPR_SYMBOL', 'arg3')
rdi=('EXPR_SYMBOL', 'arg1')
rax=('EXPR_OP', '-', ('EXPR_SYMBOL', 'arg1'), ('EXPR_VALUE', 32))

line=[imul     rax, 5]
rcx=('EXPR_SYMBOL', 'arg4')
rsi=('EXPR_SYMBOL', 'arg2')
rbx=('EXPR_SYMBOL', 'initial_RBX')
rdx=('EXPR_SYMBOL', 'arg3')
rdi=('EXPR_SYMBOL', 'arg1')
rax=('EXPR_OP', '*', ('EXPR_OP', '-', ('EXPR_SYMBOL', 'arg1'), ('EXPR_VALUE', 32)), ('EXPR_VALUE', 5))

line=[mov      rbx, 9]
rcx=('EXPR_SYMBOL', 'arg4')
rsi=('EXPR_SYMBOL', 'arg2')
rbx=('EXPR_VALUE', 9)
rdx=('EXPR_SYMBOL', 'arg3')
rdi=('EXPR_SYMBOL', 'arg1')
rax=('EXPR_OP', '*', ('EXPR_OP', '-', ('EXPR_SYMBOL', 'arg1'), ('EXPR_VALUE', 32)), ('EXPR_VALUE', 5))

line=[idiv     rbx]
rcx=('EXPR_SYMBOL', 'arg4')
rsi=('EXPR_SYMBOL', 'arg2')
rbx=('EXPR_VALUE', 9)
rdx=('EXPR_OP', '%', ('EXPR_OP', '*', ('EXPR_OP', '-', ('EXPR_SYMBOL', 'arg1'), ('EXPR_VALUE', 32)), ('
    EXPR_VALUE', 5)), ('EXPR_VALUE', 9))

```



```

rdi=('EXPR_SYMBOL', 'arg1')
rax=('EXPR_OP', '/', ('EXPR_OP', '*', ('EXPR_OP', '-', ('EXPR_SYMBOL', 'arg1'), ('EXPR_VALUE', 32)), ('EXPR_VALUE', 5)), ('EXPR_VALUE', 9))

...

result=('EXPR_OP', '/', ('EXPR_OP', '*', ('EXPR_OP', '-', ('EXPR_SYMBOL', 'arg1'), ('EXPR_VALUE', 32)), ('EXPR_VALUE', 5)), ('EXPR_VALUE', 9))

```

Neat:

```
python td.py --show-registers tests/fahr_to_celsius.s
```

```

...

line=[mov      rax, rdi]
rcx=arg4
rsi=arg2
rbx=initial_RBX
rdx=arg3
rdi=arg1
rax=arg1

line=[sub      rax, 32]
rcx=arg4
rsi=arg2
rbx=initial_RBX
rdx=arg3
rdi=arg1
rax=(arg1 - 32)

line=[imul     rax, 5]
rcx=arg4
rsi=arg2
rbx=initial_RBX
rdx=arg3
rdi=arg1
rax=((arg1 - 32) * 5)

line=[mov      rbx, 9]
rcx=arg4
rsi=arg2
rbx=9
rdx=arg3
rdi=arg1
rax=((arg1 - 32) * 5)

line=[idiv     rbx]
rcx=arg4
rsi=arg2
rbx=9
rdx=((arg1 - 32) * 5) % 9
rdi=arg1
rax=((arg1 - 32) * 5) / 9

...

result=((arg1 - 32) * 5) / 9)

```

It is interesting to note that IDIV instruction also calculates remainder of division, and it is placed into RDX register. It's not used, but is available for use.

This is how quotient and remainder are stored in registers:

```

def handle_unary_DIV_IDIV (registers, op1):
    op1_expr=register_or_number_in_string_to_expr (registers, op1)
    current_RAX=registers["rax"]
    registers["rax"]=create_binary_expr ("/", current_RAX, op1_expr)
    registers["rdx"]=create_binary_expr ("%", current_RAX, op1_expr)

```

Now this is `align2grain()` function<sup>61</sup>:

```

; uint64_t align2grain (uint64_t i, uint64_t grain)
;     return ((i + grain-1) & ~(grain-1));

```

<sup>61</sup>Taken from <https://docs.oracle.com/javase/specs/jvms/se6/html/Compiling.doc.html>

```

; rdi=i
; rsi=grain

sub     rsi, 1
add     rdi, rsi
not     rsi
and     rdi, rsi
mov     rax, rdi

```

```

...

line=[sub     rsi, 1]
rcx=arg4
rsi=(arg2 - 1)
rbx=initial_RBX
rdx=arg3
rdi=arg1
rax=initial_RAX

line=[add     rdi, rsi]
rcx=arg4
rsi=(arg2 - 1)
rbx=initial_RBX
rdx=arg3
rdi=(arg1 + (arg2 - 1))
rax=initial_RAX

line=[not     rsi]
rcx=arg4
rsi=(~(arg2 - 1))
rbx=initial_RBX
rdx=arg3
rdi=(arg1 + (arg2 - 1))
rax=initial_RAX

line=[and     rdi, rsi]
rcx=arg4
rsi=(~(arg2 - 1))
rbx=initial_RBX
rdx=arg3
rdi=((arg1 + (arg2 - 1)) & ~(arg2 - 1))
rax=initial_RAX

line=[mov     rax, rdi]
rcx=arg4
rsi=(~(arg2 - 1))
rbx=initial_RBX
rdx=arg3
rdi=((arg1 + (arg2 - 1)) & ~(arg2 - 1))
rax=((arg1 + (arg2 - 1)) & ~(arg2 - 1))

...

result=((arg1 + (arg2 - 1)) & ~(arg2 - 1))

```

## 9.4 Dealing with compiler optimizations

The following piece of code ...

```

mov     rax, rdi
add     rax, rax

```

...will be transformed into  $(arg1 + arg1)$  expression. It can be reduced to  $(arg1 * 2)$ . Our toy decompiler can identify patterns like such and rewrite them.

```

# X+X -> X*2
def reduce_ADD1 (expr):
    if is_expr_op(expr) and get_op (expr)=="+" and get_op1 (expr)==get_op2 (expr):
        return dbg_print_reduced_expr ("reduce_ADD1", expr, create_binary_expr ("*", get_op1 (expr),
            create_val_expr (2)))

    return expr # no match

```

This function will just test, if the current node has *EXPR\_OP* type, operation is “+” and both children are equal to each other. By the way, since our data structure is just tuple of tuples, Python can compare them using plain “==” operation. If the testing is finished successfully, current node is then replaced with a new expression: we take one of children, we construct a node of *EXPR\_VALUE* type with “2” number in it, and then we construct a node of *EXPR\_OP* type with “\*”.

`dbg_print_reduced_expr()` serving solely debugging purposes—it just prints the old and the new (reduced) expressions.

Decompiler is then traverse expression tree recursively in *deep-first search* fashion.

```
def reduce_step (e):
    if is_expr_op (e)==False:
        return e # expr isn't EXPR_OP, nothing to reduce (we don't reduce EXPR_SYMBOL and EXPR_VAL)

    if is_unary_op(get_op(e)):
        # recreate expr with reduced operand:
        return reducers(create_unary_expr (get_op(e), reduce_step (get_op1 (e))))
    else:
        # recreate expr with both reduced operands:
        return reducers(create_binary_expr (get_op(e), reduce_step (get_op1 (e)), reduce_step (get_op2 (e))))

...

# same as "return ...(reduce_MUL1 (reduce_ADD1 (reduce_ADD2 (... expr))))"
reducers=compose([
    ...
    reduce_ADD1, ...
    ...])

def reduce (e):
    print "going to reduce " + expr_to_string (e)
    new_expr=reduce_step(e)
    if new_expr==e:
        return new_expr # we are done here, expression can't be reduced further
    else:
        return reduce(new_expr) # reduced expr has been changed, so try to reduce it again
```

Reduction functions called again and again, as long, as expression changes.

Now we run it:

```
python td.py tests/add1.s

...

going to reduce (arg1 + arg1)
reduction in reduce_ADD1() (arg1 + arg1) -> (arg1 * 2)
going to reduce (arg1 * 2)
result=(arg1 * 2)
```

So far so good, now what if we would try this piece of code?

```
mov    rax, rdi
add    rax, rax
add    rax, rax
add    rax, rax
```

```
python td.py tests/add2.s

...

working out tests/add2.s
going to reduce (((arg1 + arg1) + (arg1 + arg1)) + ((arg1 + arg1) + (arg1 + arg1)))
reduction in reduce_ADD1() (arg1 + arg1) -> (arg1 * 2)
reduction in reduce_ADD1() (arg1 + arg1) -> (arg1 * 2)
reduction in reduce_ADD1() ((arg1 * 2) + (arg1 * 2)) -> ((arg1 * 2) * 2)
reduction in reduce_ADD1() (arg1 + arg1) -> (arg1 * 2)
reduction in reduce_ADD1() (arg1 + arg1) -> (arg1 * 2)
reduction in reduce_ADD1() ((arg1 * 2) + (arg1 * 2)) -> ((arg1 * 2) * 2)
reduction in reduce_ADD1() (((arg1 * 2) * 2) + ((arg1 * 2) * 2)) -> (((arg1 * 2) * 2) * 2)
going to reduce (((arg1 * 2) * 2) * 2)
result=(((arg1 * 2) * 2) * 2)
```

This is correct, but too verbose.

We would like to rewrite  $(X*n)*m$  expression to  $X*(n*m)$ , where  $n$  and  $m$  are numbers. We can do this by adding another function like `reduce_ADD1()`, but there is much better option: we can make matcher for tree. You can think about it as regular expression matcher, but over trees.

```
def bind_expr (key):
    return ("EXPR_WILDCARD", key)

def bind_value (key):
    return ("EXPR_WILDCARD_VALUE", key)

def match_EXPR_WILDCARD (expr, pattern):
    return {pattern[1] : expr} # return {key : expr}

def match_EXPR_WILDCARD_VALUE (expr, pattern):
    if get_expr_type (expr)!="EXPR_VALUE":
        return None
    return {pattern[1] : get_val(expr)} # return {key : expr}

def is_commutative (op):
    return op in ["+", "*", "&", "|", "^"]

def match_two_ops (op1_expr, op1_pattern, op2_expr, op2_pattern):
    m1=match (op1_expr, op1_pattern)
    m2=match (op2_expr, op2_pattern)
    if m1==None or m2==None:
        return None # one of match for operands returned False, so we do the same
    # join two dicts from both operands:
    rt={}
    rt.update(m1)
    rt.update(m2)
    return rt

def match_EXPR_OP (expr, pattern):
    if get_expr_type(expr)!=get_expr_type(pattern): # be sure, both EXPR_OP.
        return None
    if get_op (expr)!=get_op (pattern): # be sure, ops type are the same.
        return None

    if (is_unary_op(get_op(expr))):
        # match unary expression.
        return match (get_op1 (expr), get_op1 (pattern))
    else:
        # match binary expression.

        # first try match operands as is.
        m=match_two_ops (get_op1 (expr), get_op1 (pattern), get_op2 (expr), get_op2 (pattern))
        if m!=None:
            return m
        # if matching unsuccessful, AND operation is commutative, try also swapped operands.
        if is_commutative (get_op (expr))==False:
            return None
        return match_two_ops (get_op1 (expr), get_op2 (pattern), get_op2 (expr), get_op1 (pattern))

# returns dict in case of success, or None
def match (expr, pattern):
    t=get_expr_type(pattern)
    if t=="EXPR_WILDCARD":
        return match_EXPR_WILDCARD (expr, pattern)
    elif t=="EXPR_WILDCARD_VALUE":
        return match_EXPR_WILDCARD_VALUE (expr, pattern)
    elif t=="EXPR_SYMBOL":
        if expr==pattern:
            return {}
        else:
            return None
    elif t=="EXPR_VALUE":
        if expr==pattern:
            return {}
        else:
            return None
    elif t=="EXPR_OP":
        return match_EXPR_OP (expr, pattern)
    else:
        raise AssertionError
```

Now how we will use it:

```
# (X*A)*B -> X*(A*B)
def reduce_MUL1 (expr):
    m=match (expr, create_binary_expr ("*", (create_binary_expr ("*", bind_expr ("X"), bind_value ("A"))),
        bind_value ("B")))
    if m==None:
        return expr # no match

    return dbg_print_reduced_expr ("reduce_MUL1", expr, create_binary_expr ("*",
        m["X"], # new op1
        create_val_expr (m["A"] * m["B"]))) # new op2
```

We take input expression, and we also construct pattern to be matched. Matcher works recursively over both expressions synchronously. Pattern is also expression, but can use two additional node types: *EXPR\_WILDCARD* and *EXPR\_WILDCARD\_VALUE*. These nodes are supplied with keys (stored as strings). When matcher encounters *EXPR\_WILDCARD* in pattern, it just stashes current expression and will return it. If matcher encounters *EXPR\_WILDCARD\_VALUE*, it does the same, but only in case the current node has *EXPR\_VALUE* type.

`bind_expr()` and `bind_value()` are functions which create nodes with the types we have seen.

All this means, `reduce_MUL1()` function will search for the expression in form  $(X*A)*B$ , where  $A$  and  $B$  are numbers. In other cases, matcher will return input expression untouched, so these reducing function can be chained.

Now when `reduce_MUL1()` encounters (sub)expression we are interesting in, it will return dictionary with keys and expressions. Let's add `print m` call somewhere before return and rerun:

```
python td.py tests/add2.s

...
going to reduce (((arg1 + arg1) + (arg1 + arg1)) + ((arg1 + arg1) + (arg1 + arg1)))
reduction in reduce_ADD1() (arg1 + arg1) -> (arg1 * 2)
reduction in reduce_ADD1() (arg1 + arg1) -> (arg1 * 2)
reduction in reduce_ADD1() ((arg1 * 2) + (arg1 * 2)) -> ((arg1 * 2) * 2)
{'A': 2, 'X': ('EXPR_SYMBOL', 'arg1'), 'B': 2}
reduction in reduce_MUL1() ((arg1 * 2) * 2) -> (arg1 * 4)
reduction in reduce_ADD1() (arg1 + arg1) -> (arg1 * 2)
reduction in reduce_ADD1() (arg1 + arg1) -> (arg1 * 2)
reduction in reduce_ADD1() ((arg1 * 2) + (arg1 * 2)) -> ((arg1 * 2) * 2)
{'A': 2, 'X': ('EXPR_SYMBOL', 'arg1'), 'B': 2}
reduction in reduce_MUL1() ((arg1 * 2) * 2) -> (arg1 * 4)
reduction in reduce_ADD1() ((arg1 * 4) + (arg1 * 4)) -> ((arg1 * 4) * 2)
{'A': 4, 'X': ('EXPR_SYMBOL', 'arg1'), 'B': 2}
reduction in reduce_MUL1() ((arg1 * 4) * 2) -> (arg1 * 8)
going to reduce (arg1 * 8)
...
result=(arg1 * 8)
```

The dictionary has keys we supplied plus expressions matcher found. We then can use them to create new expression and return it. Numbers are just summed while forming second operand to `"*"` operation.

Now a real-world optimization technique—optimizing GCC replaced multiplication by 31 by shifting and subtraction operations:

```
mov    rax, rdi
sal    rax, 5
sub    rax, rdi
```

Without reduction functions, our decompiler will translate this into  $((arg1 \ll 5) - arg1)$ . We can replace shifting left by multiplication:

```
# X<<n -> X*(2^n)
def reduce_SHL1 (expr):
    m=match (expr, create_binary_expr ("<<", bind_expr ("X"), bind_value ("Y")))
    if m==None:
        return expr # no match

    return dbg_print_reduced_expr ("reduce_SHL1", expr, create_binary_expr ("*", m["X"], create_val_expr (1<<m["Y"])))
```

Now we getting  $((arg1 * 32) - arg1)$ . We can add another reduction function:

```
# (X*n)-X -> X*(n-1)
def reduce_SUB3 (expr):
    m=match (expr, create_binary_expr ("-",
        create_binary_expr ("*", bind_expr("X1"), bind_value ("N")),
        bind_expr("X2")))

    if m!=None and match (m["X1"], m["X2"])!=None:
        return dbg_print_reduced_expr ("reduce_SUB3", expr, create_binary_expr ("*", m["X1"], create_val_expr (m["N"]-1)))
    else:
        return expr # no match
```

Matcher will return two X's, and we must be assured that they are equal. In fact, in previous versions of this toy decompiler, I did comparison with plain "=", and it worked. But we can reuse `match()` function for the same purpose, because it will process commutative operations better. For example, if X1 is "Q+1" and X2 is "1+Q", expressions are equal, but plain "=" will not work. On the other side, `match()` function, when encounter "+" operation (or another commutative operation), and it fails with comparison, it will also try swapped operand and will try to compare again.

However, to understand it easier, for a moment, you can imagine there is "=" instead of the second `match()`.

Anyway, here is what we've got:

```
working out tests/mul31_GCC.s
going to reduce ((arg1 << 5) - arg1)
reduction in reduce_SHL1() (arg1 << 5) -> (arg1 * 32)
reduction in reduce_SUB3() ((arg1 * 32) - arg1) -> (arg1 * 31)
going to reduce (arg1 * 31)
...
result=(arg1 * 31)
```

Another optimization technique is often seen in ARM thumb code: AND-ing a value with a value like 0xFFFFF0, is implemented using shifts:

```
mov rax, rdi
shr rax, 4
shl rax, 4
```

This code is quite common in ARM thumb code, because it's a headache to encode 32-bit constants using couple of 16-bit thumb instructions, while single 16-bit instruction can shift by 4 bits left or right.

Also, the expression  $(x \gg 4) \ll 4$  can be jokingly called as "twitching operator": I've heard the "--i++" expression was called like this in Russian-speaking social networks, it was some kind of meme ("operator podergivaniya").

Anyway, these reduction functions will be used:

```
# X>>n -> X / (2^n)
...
def reduce_SHR2 (expr):
    m=match(expr, create_binary_expr(">>", bind_expr("X"), bind_value("Y")))
    if m==None or m["Y"]>=64:
        return expr # no match

    return dbg_print_reduced_expr ("reduce_SHR2", expr, create_binary_expr ("/", m["X"],
        create_val_expr (1<<m["Y"])))

...

# X<<n -> X*(2^n)
def reduce_SHL1 (expr):
    m=match (expr, create_binary_expr ("<<", bind_expr ("X"), bind_value ("Y")))
    if m==None:
        return expr # no match

    return dbg_print_reduced_expr ("reduce_SHL1", expr, create_binary_expr ("*", m["X"], create_val_expr (1<<m["Y"])))

...

# FIXME: slow
# returns True if n=2^x or popcnt(n)=1
def is_2n(n):
    return bin(n).count("1")==1
```

```
# AND operation using DIV/MUL or SHL/SHR
# (X / (2^n)) * (2^n) -> X & ((2^n)-1)
def reduce_MUL2 (expr):
    m=match(expr, create_binary_expr ("*", create_binary_expr ("/", bind_expr("X"), bind_value("N1")),
        bind_value("N2")))
    if m==None or m["N1"]!=m["N2"] or is_2n(m["N1"])==False: # short-circuit expression
        return expr # no match

    return dbg_print_reduced_expr("reduce_MUL2", expr, create_binary_expr("&", m["X"],
        create_val_expr(~(m["N1"]-1)&0xffffffffffffffff)))
```

Now the result:

```
working out tests/AND_by_shifts2.s
going to reduce ((arg1 >> 4) << 4)
reduction in reduce_SHR2() (arg1 >> 4) -> (arg1 / 16)
reduction in reduce_SHL1() ((arg1 / 16) << 4) -> ((arg1 / 16) * 16)
reduction in reduce_MUL2() ((arg1 / 16) * 16) -> (arg1 & 0xffffffffffffffff)
going to reduce (arg1 & 0xffffffffffffffff)
...
result=(arg1 & 0xffffffffffffffff)
```

### 9.4.1 Division using multiplication

Division is often replaced by multiplication for performance reasons.

From school-level arithmetics, we can remember that division by 3 can be replaced by multiplication by  $\frac{1}{3}$ . In fact, sometimes compilers do so for floating-point arithmetics, for example, FDIV instruction in x86 code can be replaced by FMUL. At least MSVC 6.0 will replace division by 3 by multiplication by  $\frac{1}{3}$  and sometimes it's hard to be sure, what operation was in original source code.

But when we operate over integer values and CPU registers, we can't use fractions. However, we can rework fraction:

$$result = \frac{x}{3} = x \cdot \frac{1}{3} = x \cdot \frac{1 \cdot MagicNumber}{3 \cdot MagicNumber}$$

Given the fact that division by  $2^n$  is very fast, we now should find that *MagicNumber*, for which the following equation will be true:  $2^n = 3 \cdot MagicNumber$ .

This code performing division by 10:

```
mov    rax, rdi
movabs rdx, 0cccccccccccccdh
mul     rdx
shr     rdx, 3
mov     rax, rdx
```

Division by  $2^{64}$  is somewhat hidden: lower 64-bit of product in RAX is not used (dropped), only higher 64-bit of product (in RDX) is used and then shifted by additional 3 bits.

RDX register is set during processing of MUL/IMUL like this:

```
def handle_unary_MUL_IMUL (registers, opl):
    opl_expr=register_or_number_in_string_to_expr (registers, opl)
    result=create_binary_expr ("*", registers["rax"], opl_expr)
    registers["rax"]=result
    registers["rdx"]=create_binary_expr (">>", result, create_val_expr(64))
```

In other words, the assembly code we have just seen multiplies by  $\frac{0cccccccccccccdh}{2^{64+3}}$ , or divides by  $\frac{2^{64+3}}{0cccccccccccccdh}$ . To find divisor we just have to divide numerator by denominator.

```
# n = magic number
# m = shifting coefficient
# return = 1 / (n / 2^m) = 2^m / n
def get_divisor (n, m):
    return (2**float(m))/float(n)

# (X*n)>>m, where m>=64 -> X/...
def reduce_div_by_MUL (expr):
    m=match (expr, create_binary_expr(">>", create_binary_expr ("*", bind_expr("X"), bind_value("N")),
        bind_value("M")))
```

```

if m==None:
    return expr # no match

divisor=get_divisor(m["N"], m["M"])
return dbg_print_reduced_expr ("reduce_div_by_MUL", expr, create_binary_expr ("/", m["X"], create_val_expr (
    int(divisor))))

```

This works, but we have a problem: this rule takes  $(arg1 * 0xffffffffcccccd) \gg 64$  expression first and finds divisor to be equal to 1.25. This is correct: result is shifted by 3 bits after (or divided by 8), and  $1.25 \cdot 8 = 10$ . But our toy decompiler doesn't support real numbers.

We can solve this problem in the following way: if divisor has fractional part, we postpone reducing, with a hope, that two subsequent right shift operations will be reduced into single one:

```

# (X*n)>>m, where m>=64 -> X/...
def reduce_div_by_MUL (expr):
    m=match (expr, create_binary_expr(">>", create_binary_expr ("*", bind_expr("X"), bind_value("N")),
        bind_value("M")))
    if m==None:
        return expr # no match

    divisor=get_divisor(m["N"], m["M"])
    if math.floor(divisor)==divisor:
        return dbg_print_reduced_expr ("reduce_div_by_MUL", expr, create_binary_expr ("/", m["X"],
            create_val_expr (int(divisor))))
    else:
        print "reduce_div_by_MUL(): postponing reduction, because divisor=", divisor
        return expr

```

That works:

```

working out tests/div_by_mult10_unsigned.s
going to reduce (((arg1 * 0xffffffffcccccd) >> 64) >> 3)
reduce_div_by_MUL(): postponing reduction, because divisor= 1.25
reduction in reduce_SHR1() (((arg1 * 0xffffffffcccccd) >> 64) >> 3) -> ((arg1 * 0xffffffffcccccd) >> 67)
going to reduce ((arg1 * 0xffffffffcccccd) >> 67)
reduction in reduce_div_by_MUL() ((arg1 * 0xffffffffcccccd) >> 67) -> (arg1 / 10)
going to reduce (arg1 / 10)
result=(arg1 / 10)

```

I don't know if this is best solution. In early version of this decompiler, it processed input expression in two passes: first pass for everything except division by multiplication, and the second pass for the latter. I don't know which way is better. Or maybe we could support real numbers in expressions?

Couple of words about better understanding division by multiplication. Many people miss "hidden" division by  $2^{32}$  or  $2^{64}$ , when lower 32-bit part (or 64-bit part) of product is not used (or just dropped). Also, there is misconception that modulo inverse is used here. This is close, but not the same thing. Extended Euclidean algorithm is usually used to find *magic coefficient*, but in fact, this algorithm is rather used to solve the equation. You can solve it using any other method. Also, needless to mention, the equation is unsolvable for some divisors, because this is diophantine equation (i.e., equation allowing result to be only integer), since we work on integer CPU registers, after all.

## 9.5 Obfuscation/deobfuscation

Despite simplicity of our decompiler, we can see how to deobfuscate (or optimize) using several simple tricks.

For example, this piece of code does nothing:

```

mov rax, rdi
xor rax, 12345678h
xor rax, 0deadbeefh
xor rax, 12345678h
xor rax, 0deadbeefh

```

We would need these rules to tame it:

```

# (X^n)^m -> X^(n^m)
def reduce_XOR4 (expr):
    m=match(expr,
        create_binary_expr("^",
            create_binary_expr ("^", bind_expr("X"), bind_value("N")),
            bind_value("M")))

    if m!=None:

```



```

        return dbg_print_reduced_expr ("reduce_XOR4", expr, create_binary_expr ("^", m["X"],
            create_val_expr (m["N"]^m["M"])))
    else:
        return expr # no match

...

# X op 0 -> X, where op is ADD, OR, XOR, SUB
def reduce_op_0 (expr):
    # try each:
    for op in ["+", "|", "^", "-"]:
        m=match(expr, create_binary_expr(op, bind_expr("X"), create_val_expr (0)))
        if m!=None:
            return dbg_print_reduced_expr ("reduce_op_0", expr, m["X"])

    # default:
    return expr # no match

```

```

working out tests/t9_obf.s
going to reduce (((arg1 ^ 0x12345678) ^ 0xdeadbeef) ^ 0x12345678) ^ 0xdeadbeef)
reduction in reduce_XOR4() ((arg1 ^ 0x12345678) ^ 0xdeadbeef) -> (arg1 ^ 0xcc99e897)
reduction in reduce_XOR4() ((arg1 ^ 0xcc99e897) ^ 0x12345678) -> (arg1 ^ 0xdeadbeef)
reduction in reduce_XOR4() ((arg1 ^ 0xdeadbeef) ^ 0xdeadbeef) -> (arg1 ^ 0x0)
going to reduce (arg1 ^ 0x0)
reduction in reduce_op_0() (arg1 ^ 0x0) -> arg1
going to reduce arg1
result=arg1

```

This piece of code can be deobfuscated (or optimized) as well:

```

; toggle last bit

    mov rax, rdi
    mov rbx, rax
    mov rcx, rbx
    mov rsi, rcx
    xor rsi, 12345678h
    xor rsi, 12345679h
    mov rax, rsi

```

```

working out tests/t7_obf.s
going to reduce ((arg1 ^ 0x12345678) ^ 0x12345679)
reduction in reduce_XOR4() ((arg1 ^ 0x12345678) ^ 0x12345679) -> (arg1 ^ 1)
going to reduce (arg1 ^ 1)
result=(arg1 ^ 1)

```

I also used *aha!*<sup>62</sup> superoptimizer to find weird piece of code which does nothing.

*Aha!* is so called superoptimizer, it tries various piece of codes in brute-force manner, in attempt to find shortest possible alternative for some mathematical operation. While sane compiler developers use superoptimizers for this task, I tried it in opposite way, to find oddest pieces of code for some simple operations, including **NOP** operation. In past, I've used it to find weird alternative to XOR operation (7.5).

So here is what *aha!* can find for **NOP**:

```

; do nothing (as found by aha)

    mov rax, rdi
    and rax, rax
    or rax, rax

```

```

# X & X -> X
def reduce_AND3 (expr):
    m=match (expr, create_binary_expr ("&", bind_expr ("X1"), bind_expr ("X2")))
    if m!=None and match (m["X1"], m["X2"])!=None:
        return dbg_print_reduced_expr("reduce_AND3", expr, m["X1"])
    else:
        return expr # no match

...

# X | X -> X

```

<sup>62</sup><http://www.hackersdelight.org/aha/aha.pdf>

```
def reduce_OR1 (expr):
    m=match (expr, create_binary_expr ("|", bind_expr ("X1"), bind_expr ("X2")))
    if m!=None and match (m["X1"], m["X2"])!=None:
        return dbg_print_reduced_expr("reduce_OR1", expr, m["X1"])
    else:
        return expr # no match
```

```
working out tests/t11_obf.s
going to reduce ((arg1 & arg1) | (arg1 & arg1))
reduction in reduce_AND3() (arg1 & arg1) -> arg1
reduction in reduce_AND3() (arg1 & arg1) -> arg1
reduction in reduce_OR1() (arg1 | arg1) -> arg1
going to reduce arg1
result=arg1
```

This is weirder:

```
; do nothing (as found by aha)

;Found a 5-operation program:
; neg r1,rx
; neg r2,rx
; neg r3,r1
; or r4,rx,2
; and r5,r4,r3
; Expr: ((x | 2) & -(x))

    mov rax, rdi
    neg rax
    neg rax
    or rdi, 2
    and rax, rdi
```

Rules added (I used "NEG" string to represent sign change and to be different from subtraction operation, which is just minus ("-")):

```
# (op(op X)) -> X, where both ops are NEG or NOT
def reduce_double_NEG_or_NOT (expr):
    # try each:
    for op in ["NEG", "~"]:
        m=match (expr, create_unary_expr (op, create_unary_expr (op, bind_expr("X"))))
        if m!=None:
            return dbg_print_reduced_expr ("reduce_double_NEG_or_NOT", expr, m["X"])

    # default:
    return expr # no match

...

# X & (X | ...) -> X
def reduce_AND2 (expr):
    m=match (expr, create_binary_expr ("&", create_binary_expr ("|", bind_expr ("X1"), bind_expr ("REST")),
        bind_expr ("X2")))
    if m!=None and match (m["X1"], m["X2"])!=None:
        return dbg_print_reduced_expr("reduce_AND2", expr, m["X1"])
    else:
        return expr # no match
```

```
going to reduce ((-(-arg1)) & (arg1 | 2))
reduction in reduce_double_NEG_or_NOT() (-(-arg1)) -> arg1
reduction in reduce_AND2() (arg1 & (arg1 | 2)) -> arg1
going to reduce arg1
result=arg1
```

I also forced *aha!* to find piece of code which adds 2 with no addition/subtraction operations allowed:

```
; arg1+2, without add/sub allowed, as found by aha:

;Found a 4-operation program:
; not r1,rx
; neg r2,r1
; not r3,r2
; neg r4,r3
; Expr: -(~(-(~(x))))
```

```

mov    rax, rdi
not    rax
neg    rax
not    rax
neg    rax

```

Rule:

```

# (- (~X)) -> X+1
def reduce_NEG_NOT (expr):
    m=match (expr, create_unary_expr ("NEG", create_unary_expr ("~", bind_expr("X"))))
    if m==None:
        return expr # no match

    return dbg_print_reduced_expr ("reduce_NEG_NOT", expr, create_binary_expr ("+", m["X"], create_val_expr(1)))

```

```

working out tests/add_by_not_neg.s
going to reduce (-(~(-(~arg1))))
reduction in reduce_NEG_NOT() (-(~arg1)) -> (arg1 + 1)
reduction in reduce_NEG_NOT() (-(~(arg1 + 1))) -> ((arg1 + 1) + 1)
reduction in reduce_ADD3() ((arg1 + 1) + 1) -> (arg1 + 2)
going to reduce (arg1 + 2)
result=(arg1 + 2)

```

This is artifact of two's complement system of signed numbers representation. Same can be done for subtraction (just swap NEG and NOT operations).

Now let's add some fake luggage to Fahrenheit-to-Celsius example:

```

; celsius = 5 * (fahr-32) / 9
; fake luggage:
mov    rbx, 12345h
mov    rax, rdi
sub    rax, 32
; fake luggage:
add    rbx, rax
imul   rax, 5
mov    rbx, 9
idiv   rbx
; fake luggage:
sub    rdx, rax

```

It's not a problem for our decompiler, because the noise is left in RDX register, and not used at all:

```

working out tests/fahr_to_celsius_obf1.s
line=[mov    rbx, 12345h]
rcx=arg4
rsi=arg2
rbx=0x12345
rdx=arg3
rdi=arg1
rax=initial_RAX

line=[mov    rax, rdi]
rcx=arg4
rsi=arg2
rbx=0x12345
rdx=arg3
rdi=arg1
rax=arg1

line=[sub    rax, 32]
rcx=arg4
rsi=arg2
rbx=0x12345
rdx=arg3
rdi=arg1
rax=(arg1 - 32)

line=[add    rbx, rax]
rcx=arg4
rsi=arg2
rbx=(0x12345 + (arg1 - 32))
rdx=arg3
rdi=arg1

```

```

rax=(arg1 - 32)

line=[imul      rax, 5]
rcx=arg4
rsi=arg2
rbx=(0x12345 + (arg1 - 32))
rdx=arg3
rdi=arg1
rax=((arg1 - 32) * 5)

line=[mov      rbx, 9]
rcx=arg4
rsi=arg2
rbx=9
rdx=arg3
rdi=arg1
rax=((arg1 - 32) * 5)

line=[idiv     rbx]
rcx=arg4
rsi=arg2
rbx=9
rdx=((arg1 - 32) * 5) % 9)
rdi=arg1
rax((((arg1 - 32) * 5) / 9)

line=[sub      rdx, rax]
rcx=arg4
rsi=arg2
rbx=9
rdx((((arg1 - 32) * 5) % 9) - (((arg1 - 32) * 5) / 9))
rdi=arg1
rax((((arg1 - 32) * 5) / 9)

going to reduce (((arg1 - 32) * 5) / 9)
result((((arg1 - 32) * 5) / 9)

```

We can try to pretend we affect the result with the noise:

```

; celsius = 5 * (fahr-32) / 9
; fake luggage:
mov     rbx, 12345h
mov     rax, rdi
sub     rax, 32
; fake luggage:
add     rbx, rax
imul    rax, 5
mov     rbx, 9
idiv    rbx
; fake luggage:
sub     rdx, rax
mov     rcx, rax
; OR result with garbage (result of fake luggage):
or      rcx, rdx
; the following instruction shouldn't affect result:
and     rax, rcx

```

...but in fact, it's all reduced by `reduce_AND2()` function we already saw (9.5):

```

working out tests/fahr_to_celsius_obf2.s
going to reduce (((arg1 - 32) * 5) / 9) & (((arg1 - 32) * 5) / 9) | (((arg1 - 32) * 5) % 9) - (((arg1 - 32) * 5) / 9)))
reduction in reduce_AND2() (((arg1 - 32) * 5) / 9) & (((arg1 - 32) * 5) / 9) | (((arg1 - 32) * 5) % 9) - (((arg1 - 32) * 5) / 9))) -> (((arg1 - 32) * 5) / 9)
going to reduce (((arg1 - 32) * 5) / 9)
result((((arg1 - 32) * 5) / 9)

```

We can see that deobfuscation is in fact the same thing as optimization used in compilers. We can try this function in GCC:

```

int f(int a)
{
    return ~(~a);
};

```

Optimizing GCC 5.4 (x86) generates this:

```
f:
    mov     eax, DWORD PTR [esp+4]
    add     eax, 1
    ret
```

GCC has its own rewriting rules, some of which are, probably, close to what we use here.

## 9.6 Tests

Despite simplicity of the decompiler, it's still error-prone. We need to be sure that original expression and reduced one are equivalent to each other.

### 9.6.1 Evaluating expressions

First of all, we would just evaluate (or *run*, or *execute*) expression with random values as arguments, and then compare results.

Evaluator do arithmetical operations when possible, recursively. When any symbol is encountered, its value (randomly generated before) is taken from a table.

```
un_ops={"NEG":operator.neg,
        "~":operator.invert}

bin_ops={">>":operator.rshift,
        "<<":(lambda x, c: x<<(c&0x3f)), # operator.lshift should be here, but it doesn't handle too big counts
        "&":operator.and_,
        "|":operator.or_,
        "^":operator.xor,
        "+":operator.add,
        "-":operator.sub,
        "*":operator.mul,
        "/":operator.div,
        "%":operator.mod}

def eval_expr(e, symbols):
    t=get_expr_type (e)
    if t=="EXPR_SYMBOL":
        return symbols[get_symbol(e)]
    elif t=="EXPR_VALUE":
        return get_val (e)
    elif t=="EXPR_OP":
        if is_unary_op (get_op (e)):
            return un_ops[get_op(e)](eval_expr(get_op1(e), symbols))
        else:
            return bin_ops[get_op(e)](eval_expr(get_op1(e), symbols), eval_expr(get_op2(e), symbols))
    else:
        raise AssertionError

def do_selftest(old, new):
    for n in range(100):
        symbols={"arg1":random.getrandbits(64),
                "arg2":random.getrandbits(64),
                "arg3":random.getrandbits(64),
                "arg4":random.getrandbits(64)}
        old_result=eval_expr (old, symbols)&0xffffffffffffffff # signed->unsigned
        new_result=eval_expr (new, symbols)&0xffffffffffffffff # signed->unsigned
        if old_result!=new_result:
            print "self-test failed"
            print "initial expression: "+expr_to_string(old)
            print "reduced expression: "+expr_to_string(new)
            print "initial expression result: ", old_result
            print "reduced expression result: ", new_result
            exit(0)
```

In fact, this is very close to what LISP *EVAL* function does, or even LISP interpreter. However, not all symbols are set. If the expression is using initial values from RAX or RBX (to which symbols "initial\_RAX" and "initial\_RBX" are assigned, decompiler will stop with exception, because no random values assigned to these registers, and these symbols are absent in *symbols* dictionary.

Using this test, I've suddenly found a bug here (despite simplicity of all these reduction rules). Well, no-one protected from eye strain. Nevertheless, the test has a serious problem: some bugs can be revealed only if one of arguments is 0, or 1, or -1. Maybe there are even more special cases exists.

Mentioned above *aha!* superoptimizer tries at least these values as arguments while testing: 1, 0, -1, 0x80000000, 0x7FFFFFFF, 0x80000001, 0x7FFFFFFE, 0x01234567, 0x89ABCDEF, -2, 2, -3, 3, -64, 64, -5, -31415.

Still, you cannot be sure.

### 9.6.2 Using Z3 SMT-solver for testing

So here we will try Z3 SMT-solver. SMT-solver can *prove* that two expressions are equivalent to each other.

For example, with the help of *aha!*, I've found another weird piece of code, which does nothing:

```
; do nothing (obfuscation)

;Found a 5-operation program:
;  neg  r1,rx
;  neg  r2,r1
;  sub  r3,r1,3
;  sub  r4,r3,r1
;  sub  r5,r4,r3
;  Expr: (((-(x) - 3) - -(x)) - (-(x) - 3))

    mov rax, rdi
    neg rax
    mov rbx, rax
    ; rbx=-x
    mov rcx, rbx
    sub rcx, 3
    ; rcx=-x-3
    mov rax, rcx
    sub rax, rbx
    ; rax=(-(x) - 3) - -(x)
    sub rax, rcx
```

Using toy decompiler, I've found that this piece is reduced to *arg1* expression:

```
working out tests/t5_obf.s
going to reduce ((((-arg1) - 3) - (-arg1)) - ((-arg1) - 3))
reduction in reduce_SUB2() ((-arg1) - 3) -> (-arg1 + 3)
reduction in reduce_SUB5() ((-arg1 + 3)) - (-arg1) -> ((-arg1 + 3)) + arg1
reduction in reduce_SUB2() ((-arg1) - 3) -> (-arg1 + 3)
reduction in reduce_ADD_SUB() (((-arg1 + 3)) + arg1) - ((-arg1 + 3))) -> arg1
going to reduce arg1
result=arg1
```

But is it correct? I've added a function which can output expression(s) to SMT-LIB-format, it's as simple as a function which converts expression to string.

And this is SMT-LIB-file for Z3:

```
(assert
  (forall ((arg1 (_ BitVec 64)) (arg2 (_ BitVec 64)) (arg3 (_ BitVec 64)) (arg4 (_ BitVec 64)))
    (=
      (bvsub (bvsub (bvsub (bvneg arg1) #x0000000000000003) (bvneg arg1)) (bvsub (bvneg arg1) #
        x0000000000000003))
      arg1
    )
  )
)
(check-sat)
```

In plain English terms, what we asking it to be sure, that *forall* four 64-bit arguments, two expressions are equivalent (second is just *arg1*).

The syntax maybe hard to understand, but in fact, this is very close to LISP, and arithmetical operations are named "bvsb", "bvadd", etc, because "bv" stands for *bit vector*.

While running, Z3 shows "sat", meaning "satisfiable". In other words, Z3 couldn't find counterexample for this expression.

In fact, I can rewrite this expression in the following form: *expr1 != expr2*, and we would ask Z3 to find at least one set of input arguments, for which expressions are not equal to each other:

```
(declare-const arg1 (_ BitVec 64))
(declare-const arg2 (_ BitVec 64))
(declare-const arg3 (_ BitVec 64))
(declare-const arg4 (_ BitVec 64))

(assert
  (not
    (=
      (bvsub (bvsub (bvsub (bvneg arg1) #x0000000000000003) (bvneg arg1)) (bvsub (bvneg arg1) #
        x0000000000000003))
      arg1
    )
  )
)
(check-sat)
```

Z3 says “unsat”, meaning, it couldn’t find any such counterexample. In other words, for all possible input arguments, results of these two expressions are always equal to each other.

Nevertheless, Z3 is not omnipotent. It fails to prove equivalence of the code which performs division by multiplication. First of all, I extended it so both results will have size of 128 bit instead of 64:

```
(declare-const x (_ BitVec 64))
(assert
  (forall ((x (_ BitVec 64)))
    (=
      ((_ zero_extend 64) (bvudiv x (_ bv17 64)))
      (bvlsr (bvmul ((_ zero_extend 64) x) #x0000000000000000f0f0f0f0f0f0f1) (_ bv68 128))
    )
  )
)
(check-sat)
(get-model)
```

(bv17 is just 64-bit number 17, etc. “bv” stands for “bit vector”, as opposed to integer value.)

Z3 works too long without any answer, and I had to interrupt it.

As Z3 developers mentioned, such expressions are hard for Z3 so far: <https://github.com/Z3Prover/z3/issues/514>.

Still, division by multiplication can be tested using previously described brute-force check.

## 9.7 My other implementations of toy decompiler

When I made attempt to write it in C++, of course, node in expression was represented using class. There is also implementation in pure C<sup>63</sup>, node is represented using structure.

Matchers in both C++ and C versions doesn’t return any dictionary, but instead, `bind_value()` functions takes pointer to a variable which will contain value after successful matching. `bind_expr()` takes pointer to a pointer, which will points to the part of expression, again, in case of success. I took this idea from LLVM.

Here are two pieces of code from LLVM source code with couple of reducing rules:

```
// (X >> A) << A -> X
Value *X;
if (match(Op0, m_Exact(m_Shr(m_Value(X), m_Specific(Op1))))
    return X;
```

( [lib/Analysis/InstructionSimplify.cpp](#) )

```
// (A | B) | C and A | (B | C) -> bswap if possible.
// (A >> B) | (C << D) and (A << B) | (B >> C) -> bswap if possible.
if (match(Op0, m_Or(m_Value(), m_Value())) ||
    match(Op1, m_Or(m_Value(), m_Value())) ||
    (match(Op0, m_LogicalShift(m_Value(), m_Value())) &&
     match(Op1, m_LogicalShift(m_Value(), m_Value())))) {
    if (Instruction *BSwap = MatchBSwap(I))
        return BSwap;
```

( [lib/Transforms/InstCombine/InstCombineAndOrXor.cpp](#) )

As you can see, my matcher tries to mimic LLVM. What I call *reduction* is called *folding* in LLVM. Both terms are popular.

<sup>63</sup>[https://github.com/dennis714/SAT\\_SMT\\_article/tree/master/toy\\_decompiler/files/C](https://github.com/dennis714/SAT_SMT_article/tree/master/toy_decompiler/files/C)

I have also a blog post about LLVM obfuscator, in which LLVM matcher is mentioned: <https://yurichev.com/blog/llvm/>.

Python version of toy decompiler uses strings in place where enumerate data type is used in C version (like *OP\_AND*, *OP\_MUL*, etc) and symbols used in Racket version<sup>64</sup> (like *'OP\_DIV*, etc). This may be seen as inefficient, nevertheless, thanks to strings interning, only address of strings are compared in Python version, not strings as a whole. So strings in Python can be seen as possible replacement for LISP symbols.

### 9.7.1 Even simpler toy decompiler

Knowledge of LISP makes you understand all these things naturally, without significant effort. But when I had no knowledge of it, but still tried to make a simple toy decompiler, I made it using usual text strings which holded expressions for each registers (and even memory).

So when MOV instruction copies value from one register to another, we just copy string. When arithmetical instruction occurred, we do string concatenation:

```
std::string registers[TOTAL];

...

// all 3 arguments are strings
switch (ins, op1, op2)
{
    ...
    case ADD:    registers[op1]="(" + registers[op1] + " + " + registers[op2] + ")";
                break;
    ...
    case MUL:    registers[op1]="(" + registers[op1] + " / " + registers[op2] + ")";
                break;
    ...
}
```

Now you'll have long expressions for each register, represented as strings. For reducing them, you can use plain simple regular expression matcher.

For example, for the rule  $(X*n)+(X*m) \rightarrow X*(n+m)$ , you can match (sub)string using the following regular expression:

$((.*)*(.))+(.*)*(.))$ <sup>65</sup>. If the string is matched, you're getting 4 groups (or substrings). You then just compare 1st and 3rd using string comparison function, then you check if the 2nd and 4th are numbers, you convert them to numbers, sum them and you make new string, consisting of 1st group and sum, like this:  $(" + X + "*" + (int(n) + int(m)) + ")$ .

It was naïve, clumsy, it was source of great embarrassment, but it worked correctly.

## 9.8 Difference between toy decompiler and commercial-grade one

Perhaps, someone, who currently reading this text, may rush into extending my source code. As an exercise, I would say, that the first step could be support of partial registers: i.e., AL, AX, EAX. This is tricky, but doable.

Another task may be support of FPU<sup>66</sup> x86 instructions (FPU stack modeling isn't a big deal).

The gap between toy decompiler and a commercial decompiler like Hex-Rays is still enormous. Several tricky problems must be solved, at least these:

- C data types: arrays, structures, pointers, etc. This problem is virtually non-existent for JVM<sup>67</sup> (Java, etc) and .NET decompilers, because type information is present in binary files.
- Basic blocks, C/C++ statements. Mike Van Emmerik in his thesis<sup>68</sup> shows how this can be tackled using SSA forms (which are also used heavily in compilers).
- Memory support, including local stack. Keep in mind pointer aliasing problem. Again, decompilers of JVM and .NET files are simpler here.

<sup>64</sup>Racket is Scheme (which is, in turn, LISP dialect) dialect. [https://github.com/dennis714/SAT\\_SMT\\_article/tree/master/toy\\_decompiler/files/Racket](https://github.com/dennis714/SAT_SMT_article/tree/master/toy_decompiler/files/Racket)

<sup>65</sup>This regular expression string hasn't been properly escaped, for the reason of easier readability and understanding.

<sup>66</sup>Floating-point unit

<sup>67</sup>Java Virtual Machine

<sup>68</sup>[https://yurichev.com/mirrors/vanEmmerik\\_ssa.pdf](https://yurichev.com/mirrors/vanEmmerik_ssa.pdf)



## 9.9 Further reading

There are several interesting open-source attempts to build decompiler. Both source code and theses are interesting study.

- *decomp* by Jim Reuter<sup>69</sup>.
- *DCC* by Cristina Cifuentes<sup>70</sup>.

It is interesting that this decompiler supports only one type (*int*). Maybe this is a reason why DCC decompiler produces source code with *.B* extension? Read more about B typeless language (C predecessor): <https://yurichev.com/blog/typeless/>.

- *Boomerang* by Mike Van Emmerik, Trent Waddington et al<sup>71</sup>.

As I've said, LISP knowledge can help to understand this all much easier. Here is well-known micro-interpreter of LISP by Peter Norvig, also written in Python: <https://web.archive.org/web/20161116133448/http://www.norvig.com/lispy.html>, <https://web.archive.org/web/20160305172301/http://norvig.com/lispy2.html>.

## 9.10 The files

Python version and tests: [https://github.com/dennis714/SAT\\_SMT\\_article/tree/master/toy\\_decompiler/files](https://github.com/dennis714/SAT_SMT_article/tree/master/toy_decompiler/files).

There are also C and Racket versions, but outdated.

Keep in mind—this decompiler is still at toy level, and it was tested only on tiny test files supplied.

# 10 Symbolic execution

## 10.1 Symbolic computation

Let's first start with symbolic computation<sup>72</sup>.

Some numbers can only be represented in binary system approximately, like  $\frac{1}{3}$  and  $\pi$ . If we calculate  $\frac{1}{3} \cdot 3$  step-by-step, we may have loss of significance. We also know that  $\sin(\frac{\pi}{2}) = 1$ , but calculating this expression in usual way, we can also have some noise in result. Arbitrary-precision arithmetic<sup>73</sup> is not a solution, because these numbers cannot be stored in memory as a binary number of finite length.

How we could tackle this problem? Humans reduce such expressions using paper and pencil without any calculations. We can mimic human behaviour programmatically if we will store expression as tree and symbols like  $\pi$  will be converted into number at the very last step(s).

This is what Wolfram Mathematica<sup>74</sup> does. Let's start it and try this:

```
In[]:= x + 2*8
Out[]:= 16 + x
```

Since Mathematica has no clue what  $x$  is, it's left as  $is$ , but  $2 \cdot 8$  can be reduced easily, both by Mathematica and by humans, so that is what has done. In some point of time in future, Mathematica's user may assign some number to  $x$  and then, Mathematica will reduce the expression even further.

Mathematica does this because it parses the expression and finds some known patterns. This is also called *term rewriting*<sup>75</sup>. In plain English language it may sounds like this: "if there is a  $+$  operator between two known numbers, replace this subexpression by a computed number which is sum of these two numbers, if possible". Just like humans do.

Mathematica also has rules like "replace  $\sin(\pi)$  by 0" and "replace  $\sin(\frac{\pi}{2})$  by 1", but as you can see,  $\pi$  must be preserved as some kind of symbol instead of a number.

<sup>69</sup> <http://www.program-transformation.org/Transform/DecompReadMe>, <http://www.program-transformation.org/Transform/DecompDecompiler>

<sup>70</sup> <http://www.program-transformation.org/Transform/DccDecompiler>, thesis: [https://yurichev.com/mirrors/DCC\\_decompilation\\_thesis.pdf](https://yurichev.com/mirrors/DCC_decompilation_thesis.pdf)

<sup>71</sup> <http://boomerang.sourceforge.net/>, <http://www.program-transformation.org/Transform/MikeVanEmmerik>, thesis: [https://yurichev.com/mirrors/vanEmmerik\\_ssa.pdf](https://yurichev.com/mirrors/vanEmmerik_ssa.pdf)

<sup>72</sup> [https://en.wikipedia.org/wiki/Symbolic\\_computation](https://en.wikipedia.org/wiki/Symbolic_computation)

<sup>73</sup> [https://en.wikipedia.org/wiki/Arbitrary-precision\\_arithmetic](https://en.wikipedia.org/wiki/Arbitrary-precision_arithmetic)

<sup>74</sup> Another well-known symbolic computation system are *Maxima* and *SymPy*

<sup>75</sup> <https://en.wikipedia.org/wiki/Rewriting>

So Mathematica left  $x$  as unknown value. This is, in fact, common mistake by Mathematica's users: a small typo in an input expression may lead to a huge irreducible expression with the typo left.

Another example: Mathematica left this deliberately while computing binary logarithm:

```
In[]:= Log[2, 36]
Out[]= Log[36]/Log[2]
```

Because it has a hope that at some point in future, this expression will become a subexpression in another expression and it will be reduced nicely at the very end. But if we really need a numerical answer, we can force Mathematica to calculate it:

```
In[]:= Log[2, 36] // N
Out[]= 5.16993
```

Sometimes unresolved values are desirable:

```
In[]:= Union[{a, b, a, c}, {d, a, e, b}, {c, a}]
Out[]= {a, b, c, d, e}
```

Characters in the expression are just unresolved symbols<sup>76</sup> with no connections to numbers or other expressions, so Mathematica left them *as is*.

Another real world example is symbolic integration<sup>77</sup>, i.e., finding formula for integral by rewriting initial expression using some predefined rules. Mathematica also does it:

```
In[]:= Integrate[1/(x^5), x]
Out[]= -(1/(4 x^4))
```

Benefits of symbolic computation are obvious: it is not prone to loss of significance<sup>78</sup> and round-off errors<sup>79</sup>, but drawbacks are also obvious: you need to store expression in (possible huge) tree and process it many times. Term rewriting is also slow. All these things are extremely clumsy in comparison to a fast FPU.

“Symbolic computation” is opposed to “numerical computation”, the last one is just processing numbers step-by-step, using calculator, CPU or FPU.

Some task can be solved better by the first method, some others – by the second one.

### 10.1.1 Rational data type

Some LISP implementations can store a number as a ratio/fraction<sup>80</sup>, i.e., placing two numbers in a cell (which, in this case, is called *atom* in LISP lingo). For example, you divide 1 by 3, and the interpreter, by understanding that  $\frac{1}{3}$  is an irreducible fraction<sup>81</sup>, creates a cell with 1 and 3 numbers. Some time after, you may multiply this cell by 6, and the multiplication function inside LISP interpreter may return much better result (2 without *noise*).

Printing function in interpreter can also print something like `1 / 3` instead of floating point number.

This is sometimes called “fractional arithmetic” [see Donald E. Knuth, *The Art of Computing Programming*, 3rd ed., (1997), 4.5.1, page 330].

This is not symbolic computation in any way, but this is slightly better than storing ratios/fractions as just floating point numbers.

Drawbacks are clearly visible: you need more memory to store ratio instead of a number; and all arithmetic functions are more complex and slower, because they must handle both numbers and ratios.

Perhaps, because of drawbacks, some programming languages offers separate (*rational*) data type, as language feature, or supported by a library<sup>82</sup>: Haskell, OCaml, Perl, Ruby, Python (*fractions*), Smalltalk, Java, Clojure, C/C++<sup>83</sup>.

---

<sup>76</sup>Symbol like in LISP

<sup>77</sup>[https://en.wikipedia.org/wiki/Symbolic\\_integration](https://en.wikipedia.org/wiki/Symbolic_integration)

<sup>78</sup>[https://en.wikipedia.org/wiki/Loss\\_of\\_significance](https://en.wikipedia.org/wiki/Loss_of_significance)

<sup>79</sup>[https://en.wikipedia.org/wiki/Round-off\\_error](https://en.wikipedia.org/wiki/Round-off_error)

<sup>80</sup>[https://en.wikipedia.org/wiki/Rational\\_data\\_type](https://en.wikipedia.org/wiki/Rational_data_type)

<sup>81</sup>[https://en.wikipedia.org/wiki/Irreducible\\_fraction](https://en.wikipedia.org/wiki/Irreducible_fraction)

<sup>82</sup>More detailed list: [https://en.wikipedia.org/wiki/Rational\\_data\\_type](https://en.wikipedia.org/wiki/Rational_data_type)

<sup>83</sup>By GNU Multiple Precision Arithmetic Library

## 10.2 Symbolic execution

### 10.2.1 Swapping two values using XOR

There is a well-known (but counterintuitive) algorithm for swapping two values in two variables using XOR operation without use of any additional memory/register:

```
X=X^Y
Y=Y^X
X=X^Y
```

How it works? It would be better to construct an expression at each step of execution.

```
#!/usr/bin/env python
class Expr:
    def __init__(self,s):
        self.s=s

    def __str__(self):
        return self.s

    def __xor__(self, other):
        return Expr("(" + self.s + "^" + other.s + ")")

def XOR_swap(X, Y):
    X=X^Y
    Y=Y^X
    X=X^Y
    return X, Y

new_X, new_Y=XOR_swap(Expr("X"), Expr("Y"))
print "new_X", new_X
print "new_Y", new_Y
```

It works, because Python is dynamically typed [PL](#), so the function doesn't care what to operate on, numerical values, or on objects of Expr() class.

Here is result:

```
new_X ((X^Y)^(Y^(X^Y)))
new_Y (Y^(X^Y))
```

You can remove double variables in your mind (since XORing by a value twice will result in nothing). At new\_X we can drop two X-es and two Y-es, and single Y will left. At new\_Y we can drop two Y-es, and single X will left.

### 10.2.2 Change endianness

What does this code do?

```
mov     eax, ecx
mov     edx, ecx
shl     edx, 16
and     eax, 0000ff00H
or      eax, edx
mov     edx, ecx
and     edx, 00ff0000H
shr     ecx, 16
or      edx, ecx
shl     eax, 8
shr     edx, 8
or      eax, edx
```

In fact, many reverse engineers play shell game a lot, keeping track of what is stored where, at each point of time.



Figure 15: Hieronymus Bosch – The Conjurer

Again, we can build equivalent function which can take both numerical variables and Expr() objects. We also extend Expr() class to support many arithmetical and boolean operations. Also, Expr() methods would take both Expr() objects on input and integer values.

```
#!/usr/bin/env python
class Expr:
    def __init__(self,s):
        self.s=s

    def convert_to_Expr_if_int(self, n):
        if isinstance(n, int):
            return Expr(str(n))
        if isinstance(n, Expr):
            return n
        raise AssertionError # unsupported type

    def __str__(self):
        return self.s

    def __xor__(self, other):
        return Expr("(" + self.s + "^" + self.convert_to_Expr_if_int(other).s + ")")

    def __and__(self, other):
        return Expr("(" + self.s + "&" + self.convert_to_Expr_if_int(other).s + ")")

    def __or__(self, other):
        return Expr("(" + self.s + "|" + self.convert_to_Expr_if_int(other).s + ")")

    def __lshift__(self, other):
        return Expr("(" + self.s + "<<" + self.convert_to_Expr_if_int(other).s + ")")

    def __rshift__(self, other):
        return Expr("(" + self.s + ">>" + self.convert_to_Expr_if_int(other).s + ")")
```

```
# change endianness
ecx=Expr("initial_ECX") # 1st argument
eax=ecx                # mov    eax, ecx
edx=ecx                # mov    edx, ecx
edx=edx<<16            # shl    edx, 16
eax=eax&0xff00         # and    eax, 0000ff00H
eax=eax|edx            # or     eax, edx
edx=ecx                # mov    edx, ecx
edx=edx&0x00ff0000     # and    edx, 00ff0000H
ecx=ecx>>16            # shr    ecx, 16
edx=edx|ecx            # or     edx, ecx
eax=eax<<8             # shl    eax, 8
edx=edx>>8             # shr    edx, 8
eax=eax|edx            # or     eax, edx

print eax
```

I run it:

```
((((initial_ECX&65280)|(initial_ECX<<16))<<8)|(((initial_ECX&16711680)|(initial_ECX>>16))>>8))
```

Now this is something more readable, however, a bit LISP-y at first sight. In fact, this is a function which change endianness in 32-bit word.

By the way, my Toy Decompiler can do this job as well, but operates on [AST](#) instead of plain strings: [9](#).

### 10.2.3 Fast Fourier transform

I've found one of the smallest possible FFT implementations on [reddit](#):

```
#!/usr/bin/env python
from cmath import exp,pi

def FFT(X):
    n = len(X)
    w = exp(-2*pi*1j/n)
    if n > 1:
        X = FFT(X[:2]) + FFT(X[1:2])
        for k in xrange(n/2):
            xk = X[k]
            X[k] = xk + w**k*X[k+n/2]
            X[k+n/2] = xk - w**k*X[k+n/2]
    return X

print FFT([1,2,3,4,5,6,7,8])
```

Just interesting, what value has each element on output?

```
#!/usr/bin/env python
from cmath import exp,pi

class Expr:
    def __init__(self,s):
        self.s=s

    def convert_to_Expr_if_int(self, n):
        if isinstance(n, int):
            return Expr(str(n))
        if isinstance(n, Expr):
            return n
        raise AssertionError # unsupported type

    def __str__(self):
        return self.s

    def __add__(self, other):
        return Expr("(" + self.s + "+" + self.convert_to_Expr_if_int(other).s + ")")

    def __sub__(self, other):
        return Expr("(" + self.s + "-" + self.convert_to_Expr_if_int(other).s + ")")

    def __mul__(self, other):
        return Expr("(" + self.s + "*" + self.convert_to_Expr_if_int(other).s + ")")

    def __pow__(self, other):
```

```

        return Expr("(" + self.s + "*" + self.convert_to_Expr_if_int(other).s + ")")

def FFT(X):
    n = len(X)
    # cast complex value to string, and then to Expr
    w = Expr(str(exp(-2*pi*1j/n)))
    if n > 1:
        X = FFT(X[:2]) + FFT(X[1:2])
        for k in xrange(n/2):
            xk = X[k]
            X[k] = xk + w**k*X[k+n/2]
            X[k+n/2] = xk - w**k*X[k+n/2]
    return X

input=[Expr("input_%d" % i) for i in range(8)]
output=FFT(input)
for i in range(len(output)):
    print i, ":", output[i]

```

FFT() function left almost intact, the only thing I added: complex value is converted into string and then Expr() object is constructed.

```

0 : (((input_0+(((−1−1.22464679915e−16j)**0)*input_4))+(((6.12323399574e−17−1j)**0)*(input_2+(((−1−1.22464679915
e−16j)**0)*input_6))))+(((0.707106781187−0.707106781187j)**0)*(input_1+(((−1−1.22464679915e−16j)**0)*
input_5))+(((6.12323399574e−17−1j)**0)*(input_3+(((−1−1.22464679915e−16j)**0)*input_7))))))
1 : (((input_0−(((−1−1.22464679915e−16j)**0)*input_4))+(((6.12323399574e−17−1j)**1)*(input_2−(((−1−1.22464679915
e−16j)**0)*input_6))))+(((0.707106781187−0.707106781187j)**1)*(input_1−(((−1−1.22464679915e−16j)**0)*
input_5))+(((6.12323399574e−17−1j)**1)*(input_3−(((−1−1.22464679915e−16j)**0)*input_7))))))
2 : (((input_0+(((−1−1.22464679915e−16j)**0)*input_4))−(((6.12323399574e−17−1j)**0)*(input_2+(((−1−1.22464679915
e−16j)**0)*input_6))))+(((0.707106781187−0.707106781187j)**2)*(input_1+(((−1−1.22464679915e−16j)**0)*
input_5))−(((6.12323399574e−17−1j)**0)*(input_3+(((−1−1.22464679915e−16j)**0)*input_7))))))
3 : (((input_0−(((−1−1.22464679915e−16j)**0)*input_4))−(((6.12323399574e−17−1j)**1)*(input_2−(((−1−1.22464679915
e−16j)**0)*input_6))))+(((0.707106781187−0.707106781187j)**3)*(input_1−(((−1−1.22464679915e−16j)**0)*
input_5))−(((6.12323399574e−17−1j)**1)*(input_3−(((−1−1.22464679915e−16j)**0)*input_7))))))
4 : (((input_0+(((−1−1.22464679915e−16j)**0)*input_4))+(((6.12323399574e−17−1j)**0)*(input_2+(((−1−1.22464679915
e−16j)**0)*input_6))))−(((0.707106781187−0.707106781187j)**0)*(input_1+(((−1−1.22464679915e−16j)**0)*
input_5))+(((6.12323399574e−17−1j)**0)*(input_3+(((−1−1.22464679915e−16j)**0)*input_7))))))
5 : (((input_0−(((−1−1.22464679915e−16j)**0)*input_4))+(((6.12323399574e−17−1j)**1)*(input_2−(((−1−1.22464679915
e−16j)**0)*input_6))))−(((0.707106781187−0.707106781187j)**1)*(input_1−(((−1−1.22464679915e−16j)**0)*
input_5))+(((6.12323399574e−17−1j)**1)*(input_3−(((−1−1.22464679915e−16j)**0)*input_7))))))
6 : (((input_0+(((−1−1.22464679915e−16j)**0)*input_4))−(((6.12323399574e−17−1j)**0)*(input_2+(((−1−1.22464679915
e−16j)**0)*input_6))))−(((0.707106781187−0.707106781187j)**2)*(input_1+(((−1−1.22464679915e−16j)**0)*
input_5))−(((6.12323399574e−17−1j)**0)*(input_3+(((−1−1.22464679915e−16j)**0)*input_7))))))
7 : (((input_0−(((−1−1.22464679915e−16j)**0)*input_4))−(((6.12323399574e−17−1j)**1)*(input_2−(((−1−1.22464679915
e−16j)**0)*input_6))))−(((0.707106781187−0.707106781187j)**3)*(input_1−(((−1−1.22464679915e−16j)**0)*
input_5))−(((6.12323399574e−17−1j)**1)*(input_3−(((−1−1.22464679915e−16j)**0)*input_7))))))

```

We can see subexpressions in form like  $x^0$  and  $x^1$ . We can eliminate them, since  $x^0 = 1$  and  $x^1 = x$ . Also, we can reduce subexpressions like  $x \cdot 1$  to just  $x$ .

```

def __mul__(self, other):
    op1=self.s
    op2=self.convert_to_Expr_if_int(other).s

    if op1=="1":
        return Expr(op2)
    if op2=="1":
        return Expr(op1)

    return Expr("(" + op1 + "*" + op2 + ")")

def __pow__(self, other):
    op2=self.convert_to_Expr_if_int(other).s
    if op2=="0":
        return Expr("1")
    if op2=="1":
        return Expr(self.s)

    return Expr("(" + self.s + "*" + op2 + ")")

```

```

0 : (((input_0+input_4)+(input_2+input_6))+((input_1+input_5)+(input_3+input_7)))
1 : (((input_0−input_4)+((6.12323399574e−17−1j)*(input_2−input_6))+((0.707106781187−0.707106781187j)*((input_1−
input_5)+((6.12323399574e−17−1j)*(input_3−input_7))))))
2 : (((input_0+input_4)−(input_2+input_6))+(((0.707106781187−0.707106781187j)**2)*((input_1+input_5)−(input_3+
input_7))))

```

```

3 : (((input_0-input_4)-((6.12323399574e-17-1j)*(input_2-input_6)))+(((0.707106781187-0.707106781187j)**3)*((
    input_1-input_5)-((6.12323399574e-17-1j)*(input_3-input_7)))))
4 : (((input_0+input_4)+(input_2+input_6))-((input_1+input_5)+(input_3+input_7)))
5 : (((input_0-input_4)+((6.12323399574e-17-1j)*(input_2-input_6)))-((0.707106781187-0.707106781187j)*((input_1-
    input_5)+((6.12323399574e-17-1j)*(input_3-input_7)))))
6 : (((input_0+input_4)-(input_2+input_6))-(((0.707106781187-0.707106781187j)**2)*((input_1+input_5)-(input_3+
    input_7))))
7 : (((input_0-input_4)-((6.12323399574e-17-1j)*(input_2-input_6)))-(((0.707106781187-0.707106781187j)**3)*((
    input_1-input_5)-((6.12323399574e-17-1j)*(input_3-input_7)))))

```

### 10.2.4 Cyclic redundancy check

I've always been wondering, which input bit affects which bit in the final CRC32 value.

From the [CRC<sup>84</sup>](http://web.archive.org/web/20161220015646/http://www.hackersdelight.org/crc.pdf) theory (good and concise introduction: <http://web.archive.org/web/20161220015646/http://www.hackersdelight.org/crc.pdf>) we know that CRC is shifting register with taps.

We will track each bit rather than byte or word, which is highly inefficient, but serves our purpose better:

```

#!/usr/bin/env python
import sys

class Expr:
    def __init__(self,s):
        self.s=s

    def convert_to_Expr_if_int(self, n):
        if isinstance(n, int):
            return Expr(str(n))
        if isinstance(n, Expr):
            return n
        raise AssertionError # unsupported type

    def __str__(self):
        return self.s

    def __xor__(self, other):
        return Expr("(" + self.s + "^" + self.convert_to_Expr_if_int(other).s + ")")

BYTES=1

def crc32(buf):
    #state=[Expr("init %d" % i) for i in range(32)]
    state=[Expr("1") for i in range(32)]
    for byte in buf:
        for n in range(8):
            bit=byte[n]
            to_taps=bit^state[31]
            state[31]=state[30]
            state[30]=state[29]
            state[29]=state[28]
            state[28]=state[27]
            state[27]=state[26]
            state[26]=state[25]^to_taps
            state[25]=state[24]
            state[24]=state[23]
            state[23]=state[22]^to_taps
            state[22]=state[21]^to_taps
            state[21]=state[20]
            state[20]=state[19]
            state[19]=state[18]
            state[18]=state[17]
            state[17]=state[16]
            state[16]=state[15]^to_taps
            state[15]=state[14]
            state[14]=state[13]
            state[13]=state[12]
            state[12]=state[11]^to_taps
            state[11]=state[10]^to_taps
            state[10]=state[9]^to_taps
            state[9]=state[8]
            state[8]=state[7]^to_taps
            state[7]=state[6]^to_taps

```

<sup>84</sup>Cyclic redundancy check



Here are expressions for each CRC32 bit for 1-byte buffer:

For larger buffer, expressions gets increasing exponentially. This is 0th bit of the final state for 4-byte buffer:

Expression for the 0th bit of the final state for 8-byte buffer has length of  $\approx 350KiB$ , which is, of course, can be reduced significantly (because this expression is basically XOR tree), but you can feel the weight of it.



Now we can process this expressions somehow to get a smaller picture on what is affecting what. Let's say, if we can find "in\_2\_3" substring in expression, this means that 3rd bit of 2nd byte of input affects this expression. But even more than that: since this is XOR tree (i.e., expression consisting only of XOR operations), if some input variable is occurring twice, it's *annihilated*, since  $x \oplus x = 0$ . More than that: if a variable occurred even number of times (2, 4, 8, etc), it's annihilated, but left if it's occurred odd number of times (1, 3, 5, etc).

```
for i in range(32):
    #print "state %d=%s" % (i, state[31-i])
    sys.stdout.write ("state %02d: " % i)
    for byte in range(BYTES):
        for bit in range(8):
            s="in_%d_%d" % (byte, bit)
            if str(state[31-i]).count(s) & 1:
                sys.stdout.write ("*")
            else:
                sys.stdout.write (" ")
    sys.stdout.write ("\n")
```

( [https://github.com/dennis714/SAT\\_SMT\\_article/blob/master/symbolic/4\\_CRC/2.py](https://github.com/dennis714/SAT_SMT_article/blob/master/symbolic/4_CRC/2.py) )

Now this how each bit of 1-byte input buffer affects each bit of the final CRC32 state:

```
state 00:  *
state 01:  *  *
state 02:  **  *
state 03:  **  *
state 04:  * **  *
state 05:  * **  *
state 06:      **
state 07:  *    **
state 08:  *    **
state 09:      *
state 10:  *
state 11:  *
state 12:  *  *
state 13:  **  *
state 14:  **  *
state 15:  **  *
state 16:  * ***
state 17:  ** ***
state 18:  *** ***
state 19:  *** ***
state 20:      ** **
state 21:  * **  *
state 22:  ** **
state 23:  ** **
state 24:  * * **  *
state 25:  ****  **
state 26:  ***** **
state 27:  * ***  *
state 28:  *    ***
state 29:  **    ***
state 30:  **    **
state 31:  *    *
```

This is 8\*8=64 bits of 8-byte input buffer:

```
state 00:  * * * *  * * *  * * * *  * * * *  * * * *  *
state 01:  * * * *  * * *  * * * *  * * * *  * * * *  *
state 02:  ** * * *  * * *  * * * *  * * * *  * * * *  *
state 03:  *** * * *  * * *  * * * *  * * * *  * * * *  *
state 04:  **** * * *  * * *  * * * *  * * * *  * * * *  *
state 05:  ***** * * *  * * *  * * * *  * * * *  * * * *  *
state 06:  ** ***  * * *  * * * *  * * * *  * * * *  * * * *  *
state 07:  * ** ***  * * *  * * * *  * * * *  * * * *  * * * *  *
state 08:  * ** ***  * * *  * * * *  * * * *  * * * *  * * * *  *
state 09:  *** ** *  *  * * * *  * * * *  * * * *  * * * *  *
state 10:  **  * ***  *  * * * *  * * * *  * * * *  * * * *  *
state 11:  **  * ***  *  * * * *  * * * *  * * * *  * * * *  *
state 12:  **  * ***  *  * * * *  * * * *  * * * *  * * * *  *
state 13:  **  * ***  *  * * * *  * * * *  * * * *  * * * *  *
state 14:  **  * ***  *  * * * *  * * * *  * * * *  * * * *  *
state 15:  **  * ***  *  * * * *  * * * *  * * * *  * * * *  *
state 16:  * ** * * * *  * * *  * * * *  * * * *  * * * *  *
state 17:  * ** * * * *  * * *  * * * *  * * * *  * * * *  *
```

```

state 18: * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * *
state 19: * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * *
state 20: * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * *
state 21: * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * *
state 22: * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * *
state 23: * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * *
state 24: * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * *
state 25: * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * *
state 26: * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * *
state 27: * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * *
state 28: * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * *
state 29: * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * *
state 30: * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * *
state 31: * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * *

```

### 10.2.5 Linear congruential generator

This is popular PRNG<sup>85</sup> from OpenWatcom CRT<sup>86</sup> library: <https://github.com/open-watcom/open-watcom-v2/blob/d468b609ba6ca61eeddad80dd2485e3256fc5261/bld/clib/math/c/rand.c>.

What expression it generates on each step?

```

#!/usr/bin/env python
class Expr:
    def __init__(self,s):
        self.s=s

    def convert_to_Expr_if_int(self, n):
        if isinstance(n, int):
            return Expr(str(n))
        if isinstance(n, Expr):
            return n
        raise AssertionError # unsupported type

    def __str__(self):
        return self.s

    def __xor__(self, other):
        return Expr("(" + self.s + "^" + self.convert_to_Expr_if_int(other).s + ")")

    def __mul__(self, other):
        return Expr("(" + self.s + "*" + self.convert_to_Expr_if_int(other).s + ")")

    def __add__(self, other):
        return Expr("(" + self.s + "+" + self.convert_to_Expr_if_int(other).s + ")")

    def __and__(self, other):
        return Expr("(" + self.s + "&" + self.convert_to_Expr_if_int(other).s + ")")

    def __rshift__(self, other):
        return Expr("(" + self.s + ">>" + self.convert_to_Expr_if_int(other).s + ")")

seed=Expr("initial_seed")

def rand():
    global seed
    seed=seed*1103515245+12345
    return (seed>>16) & 0x7fff

for i in range(10):
    print i, ":", rand()

```

```

0 : (((initial_seed*1103515245)+12345)>>16)&32767)
1 : ((((((initial_seed*1103515245)+12345)*1103515245)+12345)>>16)&32767)
2 : (((((((initial_seed*1103515245)+12345)*1103515245)+12345)*1103515245)+12345)>>16)&32767)
3 : ((((((((((initial_seed*1103515245)+12345)*1103515245)+12345)*1103515245)+12345)*1103515245)+12345)>>16)
&32767)
4 : (((((((((((initial_seed*1103515245)+12345)*1103515245)+12345)*1103515245)+12345)*1103515245)+12345)
*1103515245)+12345)>>16)&32767)
5 : ((((((((((((((initial_seed*1103515245)+12345)*1103515245)+12345)*1103515245)+12345)*1103515245)+12345)
*1103515245)+12345)*1103515245)+12345)>>16)&32767)

```

<sup>85</sup>Pseudorandom number generator

<sup>86</sup>C runtime library

[illegible]

Now if we once got several values from this PRNG, like 4583, 16304, 14440, 32315, 28670, 12568..., how would we recover the initial seed? The problem in fact is solving a system of equations:

```
((((initial_seed*1103515245)+12345)>>16)&32767)==4583
((((((initial_seed*1103515245)+12345)*1103515245)+12345)>>16)&32767)==16304
(((((((initial_seed*1103515245)+12345)*1103515245)+12345)*1103515245)+12345)>>16)&32767)==14440
((((((((initial_seed*1103515245)+12345)*1103515245)+12345)*1103515245)+12345)*1103515245)+12345)>>16)&32767)
==32315
```

As it turns out, Z3 can solve this system correctly using only two equations:

```
#!/usr/bin/env python
from z3 import *

s=Solver()

x=BitVec("x",32)

a=1103515245
c=12345
s.add((((x*a)+c)>>16)&32767==4583)
s.add((((((x*a)+c)*a)+c)>>16)&32767==16304)
#s.add((((((((x*a)+c)*a)+c)*a)+c)>>16)&32767==14440)
#s.add((((((((((((x*a)+c)*a)+c)*a)+c)*a)+c)>>16)&32767==32315))

s.check()
print s.model()
```

[x = 11223344]

(Though, it takes  $\approx 20$  seconds on my ancient Intel Atom netbook.)

### 10.2.6 Path constraint

## How to get weekday from UNIX timestamp?

```
#!/usr/bin/env python

input=...
SECS_DAY=24*60*60
dayno = input / SECS_DAY
wday = (dayno + 4) % 7
if wday==5:
    print "Thanks God, it's Friday!"
```

Let's say, we should find a way to run the block with `print()` call in it. What input value should be? First, let's build expression of *wday* variable:

```
#!/usr/bin/env python
class Expr:
    def __init__(self,s):
        self.s=s

    def convert_to_Expr_if_int(self, n):
        if isinstance(n, int):
            return Expr(str(n))
        if isinstance(n, Expr):
            return n
        raise AssertionError # unsupported type

    def __str__(self):
        return self.s
```

```

def __div__(self, other):
    return Expr("(" + self.s + "/" + self.convert_to_Expr_if_int(other).s + ")")

def __mod__(self, other):
    return Expr("(" + self.s + "%" + self.convert_to_Expr_if_int(other).s + ")")

def __add__(self, other):
    return Expr("(" + self.s + "+" + self.convert_to_Expr_if_int(other).s + ")")

input=Expr("input")
SECS_DAY=24*60*60
dayno = input / SECS_DAY
wday = (dayno + 4) % 7
print wday
if wday==5:
    print "Thanks God, it's Friday!"

```

```
((input/86400)+4)%7)
```

In order to execute the block, we should solve this equation:  $((\frac{input}{86400} + 4) \equiv 5 \pmod{7})$ .  
So far, this is easy task for Z3:

```

#!/usr/bin/env python
from z3 import *

s=Solver()

x=Int("x")

s.add(((x/86400)+4)%7==5)

s.check()
print s.model()

```

```
[x = 86438]
```

This is indeed correct UNIX timestamp for Friday:

```
% date --date='@86438'
Fri Jan  2 03:00:38 MSK 1970
```

Though the date back in year 1970, but it's still correct!

This is also called “path constraint”, i.e., what constraint must be satisfied to execute specific block? Several tools has “path” in their names, like “pathgrind”, [Symbolic PathFinder](#), CodeSurfer Path Inspector, etc.

Like the shell game, this task is also often encounters in practice. You can see that something dangerous can be executed inside some basic block and you're trying to deduce, what input values can cause execution of it. It may be buffer overflow, etc. Such input values are sometimes also called “inputs of death”.

Many crackmes are solved in this way, all you need is find a path into block which prints “key is correct” or something like that.

We can extend this tiny example:

```

input=...
SECS_DAY=24*60*60
dayno = input / SECS_DAY
wday = (dayno + 4) % 7
print wday
if wday==5:
    print "Thanks God, it's Friday!"
else:
    print "Got to wait a little"

```

Now we have two blocks: for the first we should solve this equation:  $((\frac{input}{86400} + 4) \equiv 5 \pmod{7})$ . But for the second we should solve inverted equation:  $((\frac{input}{86400} + 4) \not\equiv 5 \pmod{7})$ . By solving these equations, we will find two paths into both blocks.

KLEE (or similar tool) tries to find path to each [basic] block and produces “ideal” unit test. Hence, KLEE can find a path into the block which crashes everything, or reporting about correctness of the input key/license, etc. Surprisingly, KLEE can find backdoors in the very same manner.

KLEE is also called “KLEE Symbolic Virtual Machine” – by that its creators mean that the KLEE is [VM<sup>87</sup>](#) which executes a code symbolically rather than numerically (like usual [CPU](#)).

### 10.2.7 Division by zero

If division by zero is unwrapped by sanitizing check, and exception isn’t caught, it can crash process.

Let’s calculate simple expression  $\frac{x}{2y+4z-12}$ . We can add a warning into `__div__` method:

```
#!/usr/bin/env python
class Expr:
    def __init__(self,s):
        self.s=s

    def convert_to_Expr_if_int(self, n):
        if isinstance(n, int):
            return Expr(str(n))
        if isinstance(n, Expr):
            return n
        raise AssertionError # unsupported type

    def __str__(self):
        return self.s

    def __mul__(self, other):
        return Expr("(" + self.s + "*" + self.convert_to_Expr_if_int(other).s + ")")

    def __div__(self, other):
        op2=self.convert_to_Expr_if_int(other).s
        print "warning: division by zero if "+op2+"==0"
        return Expr("(" + self.s + "/" + op2 + ")")

    def __add__(self, other):
        return Expr("(" + self.s + "+" + self.convert_to_Expr_if_int(other).s + ")")

    def __sub__(self, other):
        return Expr("(" + self.s + "-" + self.convert_to_Expr_if_int(other).s + ")")

"""
      x
-----
2y + 4z - 12
"""

def f(x, y, z):
    return x/(y*2 + z*4 - 12)

print f(Expr("x"), Expr("y"), Expr("z"))
```

...so it will report about dangerous states and conditions:

```
warning: division by zero if ((y*2)+(z*4))-12==0
(x/((y*2)+(z*4))-12))
```

This equation is easy to solve, let’s try Wolfram Mathematica this time:

```
In[]:= FindInstance[{(y*2 + z*4) - 12 == 0}, {y, z}, Integers]
Out[]:= {{y -> 0, z -> 3}}
```

These values for  $y$  and  $z$  can also be called “inputs of death”.

### 10.2.8 Merge sort

How merge sort works? I have cypasted Python code from rosettacode.com almost intact:

```
#!/usr/bin/env python
class Expr:
    def __init__(self,s,i):
        self.s=s
        self.i=i
```

---

<sup>87</sup>Virtual Machine

```

def __str__(self):
    # return both symbolic and integer:
    return self.s+" (" + str(self.i)+")"

def __le__(self, other):
    # compare only integer parts:
    return self.i <= other.i

# cypasted from http://rosettacode.org/wiki/Sorting_algorithms/Merge_sort#Python
def merge(left, right):
    result = []
    left_idx, right_idx = 0, 0
    while left_idx < len(left) and right_idx < len(right):
        # change the direction of this comparison to change the direction of the sort
        if left[left_idx] <= right[right_idx]:
            result.append(left[left_idx])
            left_idx += 1
        else:
            result.append(right[right_idx])
            right_idx += 1

    if left_idx < len(left):
        result.extend(left[left_idx:])
    if right_idx < len(right):
        result.extend(right[right_idx:])
    return result

def tabs (t):
    return "\t"*t

def merge_sort(m, indent=0):
    print tabs(indent)+"merge_sort() begin. input:"
    for i in m:
        print tabs(indent)+str(i)

    if len(m) <= 1:
        print tabs(indent)+"merge_sort() end. returning single element"
        return m

    middle = len(m) // 2
    left = m[:middle]
    right = m[middle:]

    left = merge_sort(left, indent+1)
    right = merge_sort(right, indent+1)
    rt=list(merge(left, right))
    print tabs(indent)+"merge_sort() end. returning:"
    for i in rt:
        print tabs(indent)+str(i)
    return rt

# input buffer has both symbolic and numerical values:
input=[Expr("input1",22), Expr("input2",7), Expr("input3",2), Expr("input4",1), Expr("input5",8), Expr("input6",4)]
merge_sort(input)

```

But here is a function which compares elements. Obviously, it wouldn't work correctly without it.

So we can track both expression for each element and numerical value. Both will be printed finally. But whenever values are to be compared, only numerical parts will be used.

Result:

```

merge_sort() begin. input:
input1 (22)
input2 (7)
input3 (2)
input4 (1)
input5 (8)
input6 (4)
merge_sort() begin. input:
input1 (22)
input2 (7)
input3 (2)
merge_sort() begin. input:
input1 (22)
merge_sort() end. returning single element

```

```

merge_sort() begin. input:
input2 (7)
input3 (2)
    merge_sort() begin. input:
    input2 (7)
    merge_sort() end. returning single element
    merge_sort() begin. input:
    input3 (2)
    merge_sort() end. returning single element
merge_sort() end. returning:
input3 (2)
input2 (7)
merge_sort() end. returning:
input3 (2)
input2 (7)
input1 (22)
merge_sort() begin. input:
input4 (1)
input5 (8)
input6 (4)
    merge_sort() begin. input:
    input4 (1)
    merge_sort() end. returning single element
    merge_sort() begin. input:
    input5 (8)
    input6 (4)
        merge_sort() begin. input:
        input5 (8)
        merge_sort() end. returning single element
        merge_sort() begin. input:
        input6 (4)
        merge_sort() end. returning single element
    merge_sort() end. returning:
    input6 (4)
    input5 (8)
merge_sort() end. returning:
input4 (1)
input6 (4)
input5 (8)
merge_sort() end. returning:
input4 (1)
input3 (2)
input6 (4)
input2 (7)
input5 (8)
input1 (22)

```

### 10.2.9 Extending Expr class

This is somewhat senseless, nevertheless, it's easy task to extend my Expr class to support [AST](#) instead of plain strings. It's also possible to add folding steps (like I demonstrated in [Toy Decompiler: 9](#)). Maybe someone will want to do this as an exercise. By the way, the toy decompiler can be used as simple symbolic engine as well, just feed all the instructions to it and it will track contents of each register.

### 10.2.10 Conclusion

For the sake of demonstration, I made things as simple as possible. But reality is always harsh and inconvenient, so all this shouldn't be taken as a silver bullet.

The files used in this part: [https://github.com/dennis714/SAT\\_SMT\\_article/tree/master/symbolic](https://github.com/dennis714/SAT_SMT_article/tree/master/symbolic).

## 10.3 Further reading

James C. King — Symbolic Execution and Program Testing <sup>88</sup>

<sup>88</sup><https://yurichev.com/mirrors/king76symbolicexecution.pdf>

# 11 KLEE

## 11.1 Installation

KLEE building from source is tricky. Easiest way to use KLEE is to install docker<sup>89</sup> and then to run KLEE docker image<sup>90</sup>. The path where KLEE files residing can look like **/var/lib/docker/aufs/mnt/(lots of hexadecimal digits)/home/klee**.

## 11.2 School-level equation

Let's revisit school-level system of equations from (6.2).

We will force KLEE to find a path, where all the constraints are satisfied:

```
int main()
{
    int circle, square, triangle;

    klee_make_symbolic(&circle, sizeof circle, "circle");
    klee_make_symbolic(&square, sizeof square, "square");
    klee_make_symbolic(&triangle, sizeof triangle, "triangle");

    if (circle+circle!=10) return 0;
    if (circle*square+square!=12) return 0;
    if (circle*square-triangle*circle!=circle) return 0;

    // all constraints should be satisfied at this point
    // force KLEE to produce .err file:
    klee_assert(0);
};
```

```
% clang -emit-llvm -c -g klee_eq.c
...

% klee klee_eq.bc
KLEE: output directory is "/home/klee/klee-out-93"
KLEE: WARNING: undefined reference to function: klee_assert
KLEE: WARNING ONCE: calling external: klee_assert(0)
KLEE: ERROR: /home/klee/klee_eq.c:18: failed external call: klee_assert
KLEE: NOTE: now ignoring this error at this location

KLEE: done: total instructions = 32
KLEE: done: completed paths = 1
KLEE: done: generated tests = 1
```

Let's find out, where `klee_assert()` has been triggered:

```
% ls klee-last | grep err
test000001.external.err

% ktest-tool --write-ints klee-last/test000001.ktest
ktest file : 'klee-last/test000001.ktest'
args       : ['klee_eq.bc']
num objects: 3
object 0: name: b'circle'
object 0: size: 4
object 0: data: 5
object 1: name: b'square'
object 1: size: 4
object 1: data: 2
object 2: name: b'triangle'
object 2: size: 4
object 2: data: 1
```

This is indeed correct solution to the system of equations.

KLEE has *intrinsic* `klee_assume()` which tells KLEE to cut path if some constraint is not satisfied. So we can rewrite our example in such cleaner way:

<sup>89</sup><https://docs.docker.com/engine/installation/linux/ubuntu/linux/>

<sup>90</sup><http://klee.github.io/docker/>



```

int main()
{
    int circle, square, triangle;

    klee_make_symbolic(&circle, sizeof circle, "circle");
    klee_make_symbolic(&square, sizeof square, "square");
    klee_make_symbolic(&triangle, sizeof triangle, "triangle");

    klee_assume (circle+circle==10);
    klee_assume (circle*square+square==12);
    klee_assume (circle*square-triangle*circle==circle);

    // all constraints should be satisfied at this point
    // force KLEE to produce .err file:
    klee_assert(0);
};

```

## 11.3 Zebra puzzle

Let's revisit zebra puzzle from (7.2).

We just define all variables and add constraints:

```

int main()
{
    int Yellow, Blue, Red, Ivory, Green;
    int Norwegian, Ukrainian, Englishman, Spaniard, Japanese;
    int Water, Tea, Milk, OrangeJuice, Coffee;
    int Kools, Chesterfield, OldGold, LuckyStrike, Parliament;
    int Fox, Horse, Snails, Dog, Zebra;

    klee_make_symbolic(&Yellow, sizeof(int), "Yellow");
    klee_make_symbolic(&Blue, sizeof(int), "Blue");
    klee_make_symbolic(&Red, sizeof(int), "Red");
    klee_make_symbolic(&Ivory, sizeof(int), "Ivory");
    klee_make_symbolic(&Green, sizeof(int), "Green");

    klee_make_symbolic(&Norwegian, sizeof(int), "Norwegian");
    klee_make_symbolic(&Ukrainian, sizeof(int), "Ukrainian");
    klee_make_symbolic(&Englishman, sizeof(int), "Englishman");
    klee_make_symbolic(&Spaniard, sizeof(int), "Spaniard");
    klee_make_symbolic(&Japanese, sizeof(int), "Japanese");

    klee_make_symbolic(&Water, sizeof(int), "Water");
    klee_make_symbolic(&Tea, sizeof(int), "Tea");
    klee_make_symbolic(&Milk, sizeof(int), "Milk");
    klee_make_symbolic(&OrangeJuice, sizeof(int), "OrangeJuice");
    klee_make_symbolic(&Coffee, sizeof(int), "Coffee");

    klee_make_symbolic(&Kools, sizeof(int), "Kools");
    klee_make_symbolic(&Chesterfield, sizeof(int), "Chesterfield");
    klee_make_symbolic(&OldGold, sizeof(int), "OldGold");
    klee_make_symbolic(&LuckyStrike, sizeof(int), "LuckyStrike");
    klee_make_symbolic(&Parliament, sizeof(int), "Parliament");

    klee_make_symbolic(&Fox, sizeof(int), "Fox");
    klee_make_symbolic(&Horse, sizeof(int), "Horse");
    klee_make_symbolic(&Snails, sizeof(int), "Snails");
    klee_make_symbolic(&Dog, sizeof(int), "Dog");
    klee_make_symbolic(&Zebra, sizeof(int), "Zebra");

    // limits.
    if (Yellow<1 || Yellow>5) return 0;
    if (Blue<1 || Blue>5) return 0;
    if (Red<1 || Red>5) return 0;
    if (Ivory<1 || Ivory>5) return 0;
    if (Green<1 || Green>5) return 0;

    if (Norwegian<1 || Norwegian>5) return 0;
    if (Ukrainian<1 || Ukrainian>5) return 0;
    if (Englishman<1 || Englishman>5) return 0;
    if (Spaniard<1 || Spaniard>5) return 0;
    if (Japanese<1 || Japanese>5) return 0;

```

```

if (Water<1 || Water>5) return 0;
if (Tea<1 || Tea>5) return 0;
if (Milk<1 || Milk>5) return 0;
if (OrangeJuice<1 || OrangeJuice>5) return 0;
if (Coffee<1 || Coffee>5) return 0;

if (Kools<1 || Kools>5) return 0;
if (Chesterfield<1 || Chesterfield>5) return 0;
if (OldGold<1 || OldGold>5) return 0;
if (LuckyStrike<1 || LuckyStrike>5) return 0;
if (Parliament<1 || Parliament>5) return 0;

if (Fox<1 || Fox>5) return 0;
if (Horse<1 || Horse>5) return 0;
if (Snails<1 || Snails>5) return 0;
if (Dog<1 || Dog>5) return 0;
if (Zebra<1 || Zebra>5) return 0;

// colors are distinct for all 5 houses:
if (((1<<Yellow) | (1<<Blue) | (1<<Red) | (1<<Ivory) | (1<<Green))!=0x3E) return 0; // 111110

// all nationalities are living in different houses:
if (((1<<Norwegian) | (1<<Ukrainian) | (1<<Englishman) | (1<<Spaniard) | (1<<Japanese))!=0x3E) return 0;
// 111110

// so are beverages:
if (((1<<Water) | (1<<Tea) | (1<<Milk) | (1<<OrangeJuice) | (1<<Coffee))!=0x3E) return 0; // 111110

// so are cigarettes:
if (((1<<Kools) | (1<<Chesterfield) | (1<<OldGold) | (1<<LuckyStrike) | (1<<Parliament))!=0x3E) return
0; // 111110

// so are pets:
if (((1<<Fox) | (1<<Horse) | (1<<Snails) | (1<<Dog) | (1<<Zebra))!=0x3E) return 0; // 111110

// main constraints of the puzzle:

// 2.The Englishman lives in the red house.
if (Englishman!=Red) return 0;

// 3.The Spaniard owns the dog.
if (Spaniard!=Dog) return 0;

// 4.Coffee is drunk in the green house.
if (Coffee!=Green) return 0;

// 5.The Ukrainian drinks tea.
if (Ukrainian!=Tea) return 0;

// 6.The green house is immediately to the right of the ivory house.
if (Green!=Ivory+1) return 0;

// 7.The Old Gold smoker owns snails.
if (OldGold!=Snails) return 0;

// 8.Kools are smoked in the yellow house.
if (Kools!=Yellow) return 0;

// 9.Milk is drunk in the middle house.
if (Milk!=3) return 0; // i.e., 3rd house

// 10.The Norwegian lives in the first house.
if (Norwegian!=1) return 0;

// 11.The man who smokes Chesterfields lives in the house next to the man with the fox.
if (Chesterfield!=Fox+1 && Chesterfield!=Fox-1) return 0; // left or right

// 12.Kools are smoked in the house next to the house where the horse is kept.
if (Kools!=Horse+1 && Kools!=Horse-1) return 0; // left or right

// 13.The Lucky Strike smoker drinks orange juice.
if (LuckyStrike!=OrangeJuice) return 0;

// 14.The Japanese smokes Parliaments.
if (Japanese!=Parliament) return 0;

```

```

// 15.The Norwegian lives next to the blue house.
if (Norwegian!=Blue+1 && Norwegian!=Blue-1) return 0; // left or right

// all constraints are satisfied at this point
// force KLEE to produce .err file:
klee_assert(0);

return 0;
};

```

I force KLEE to find distinct values for colors, nationalities, cigarettes, etc, in the same way as I did for Sudoku earlier (7.3).

Let's run it:

```

% clang -emit-llvm -c -g klee_zebra1.c
...

% klee klee_zebra1.bc
KLEE: output directory is "/home/klee/klee-out-97"
KLEE: WARNING: undefined reference to function: klee_assert
KLEE: WARNING ONCE: calling external: klee_assert(0)
KLEE: ERROR: /home/klee/klee_zebra1.c:130: failed external call: klee_assert
KLEE: NOTE: now ignoring this error at this location

KLEE: done: total instructions = 761
KLEE: done: completed paths = 55
KLEE: done: generated tests = 55

```

It works for  $\approx 7$  seconds on my Intel Core i3-3110M 2.4GHz notebook. Let's find out path, where `klee_assert()` has been executed:

```

% ls klee-last | grep err
test000051.external.err

% ktest-tool --write-ints klee-last/test000051.ktest | less

ktest file : 'klee-last/test000051.ktest'
args       : ['klee_zebra1.bc']
num objects: 25
object 0: name: b'Yellow'
object 0: size: 4
object 0: data: 1
object 1: name: b'Blue'
object 1: size: 4
object 1: data: 2
object 2: name: b'Red'
object 2: size: 4
object 2: data: 3
object 3: name: b'Ivory'
object 3: size: 4
object 3: data: 4
...
object 21: name: b'Horse'
object 21: size: 4
object 21: data: 2
object 22: name: b'Snails'
object 22: size: 4
object 22: data: 3
object 23: name: b'Dog'
object 23: size: 4
object 23: data: 4
object 24: name: b'Zebra'
object 24: size: 4
object 24: data: 5

```

This is indeed correct solution.

`klee_assume()` also can be used this time:

```

int main()
{
    int Yellow, Blue, Red, Ivory, Green;
    int Norwegian, Ukrainian, Englishman, Spaniard, Japanese;

```

```

int Water, Tea, Milk, OrangeJuice, Coffee;
int Kools, Chesterfield, OldGold, LuckyStrike, Parliament;
int Fox, Horse, Snails, Dog, Zebra;

klee_make_symbolic(&Yellow, sizeof(int), "Yellow");
klee_make_symbolic(&Blue, sizeof(int), "Blue");
klee_make_symbolic(&Red, sizeof(int), "Red");
klee_make_symbolic(&Ivory, sizeof(int), "Ivory");
klee_make_symbolic(&Green, sizeof(int), "Green");

klee_make_symbolic(&Norwegian, sizeof(int), "Norwegian");
klee_make_symbolic(&Ukrainian, sizeof(int), "Ukrainian");
klee_make_symbolic(&Englishman, sizeof(int), "Englishman");
klee_make_symbolic(&Spaniard, sizeof(int), "Spaniard");
klee_make_symbolic(&Japanese, sizeof(int), "Japanese");

klee_make_symbolic(&Water, sizeof(int), "Water");
klee_make_symbolic(&Tea, sizeof(int), "Tea");
klee_make_symbolic(&Milk, sizeof(int), "Milk");
klee_make_symbolic(&OrangeJuice, sizeof(int), "OrangeJuice");
klee_make_symbolic(&Coffee, sizeof(int), "Coffee");

klee_make_symbolic(&Kools, sizeof(int), "Kools");
klee_make_symbolic(&Chesterfield, sizeof(int), "Chesterfield");
klee_make_symbolic(&OldGold, sizeof(int), "OldGold");
klee_make_symbolic(&LuckyStrike, sizeof(int), "LuckyStrike");
klee_make_symbolic(&Parliament, sizeof(int), "Parliament");

klee_make_symbolic(&Fox, sizeof(int), "Fox");
klee_make_symbolic(&Horse, sizeof(int), "Horse");
klee_make_symbolic(&Snails, sizeof(int), "Snails");
klee_make_symbolic(&Dog, sizeof(int), "Dog");
klee_make_symbolic(&Zebra, sizeof(int), "Zebra");

// limits.
klee_assume (Yellow>=1 && Yellow<=5);
klee_assume (Blue>=1 && Blue<=5);
klee_assume (Red>=1 && Red<=5);
klee_assume (Ivory>=1 && Ivory<=5);
klee_assume (Green>=1 && Green<=5);

klee_assume (Norwegian>=1 && Norwegian<=5);
klee_assume (Ukrainian>=1 && Ukrainian<=5);
klee_assume (Englishman>=1 && Englishman<=5);
klee_assume (Spaniard>=1 && Spaniard<=5);
klee_assume (Japanese>=1 && Japanese<=5);

klee_assume (Water>=1 && Water<=5);
klee_assume (Tea>=1 && Tea<=5);
klee_assume (Milk>=1 && Milk<=5);
klee_assume (OrangeJuice>=1 && OrangeJuice<=5);
klee_assume (Coffee>=1 && Coffee<=5);

klee_assume (Kools>=1 && Kools<=5);
klee_assume (Chesterfield>=1 && Chesterfield<=5);
klee_assume (OldGold>=1 && OldGold<=5);
klee_assume (LuckyStrike>=1 && LuckyStrike<=5);
klee_assume (Parliament>=1 && Parliament<=5);

klee_assume (Fox>=1 && Fox<=5);
klee_assume (Horse>=1 && Horse<=5);
klee_assume (Snails>=1 && Snails<=5);
klee_assume (Dog>=1 && Dog<=5);
klee_assume (Zebra>=1 && Zebra<=5);

// colors are distinct for all 5 houses:
klee_assume (((1<<Yellow) | (1<<Blue) | (1<<Red) | (1<<Ivory) | (1<<Green))==0x3E); // 111110

// all nationalities are living in different houses:
klee_assume (((1<<Norwegian) | (1<<Ukrainian) | (1<<Englishman) | (1<<Spaniard) | (1<<Japanese))==0x3E);
// 111110

// so are beverages:
klee_assume (((1<<Water) | (1<<Tea) | (1<<Milk) | (1<<OrangeJuice) | (1<<Coffee))==0x3E); // 111110

// so are cigarettes:

```

```

klee_assume (((1<<Kools) | (1<<Chesterfield) | (1<<OldGold) | (1<<LuckyStrike) | (1<<Parliament))==0x3E)
; // 111110

// so are pets:
klee_assume (((1<<Fox) | (1<<Horse) | (1<<Snails) | (1<<Dog) | (1<<Zebra))==0x3E); // 111110

// main constraints of the puzzle:

// 2.The Englishman lives in the red house.
klee_assume (Englishman==Red);

// 3.The Spaniard owns the dog.
klee_assume (Spaniard==Dog);

// 4.Coffee is drunk in the green house.
klee_assume (Coffee==Green);

// 5.The Ukrainian drinks tea.
klee_assume (Ukrainian==Tea);

// 6.The green house is immediately to the right of the ivory house.
klee_assume (Green==Ivory+1);

// 7.The Old Gold smoker owns snails.
klee_assume (OldGold==Snails);

// 8.Kools are smoked in the yellow house.
klee_assume (Kools==Yellow);

// 9.Milk is drunk in the middle house.
klee_assume (Milk==3); // i.e., 3rd house

// 10.The Norwegian lives in the first house.
klee_assume (Norwegian==1);

// 11.The man who smokes Chesterfields lives in the house next to the man with the fox.
klee_assume (Chesterfield==Fox+1 || Chesterfield==Fox-1); // left or right

// 12.Kools are smoked in the house next to the house where the horse is kept.
klee_assume (Kools==Horse+1 || Kools==Horse-1); // left or right

// 13.The Lucky Strike smoker drinks orange juice.
klee_assume (LuckyStrike==OrangeJuice);

// 14.The Japanese smokes Parliaments.
klee_assume (Japanese==Parliament);

// 15.The Norwegian lives next to the blue house.
klee_assume (Norwegian==Blue+1 || Norwegian==Blue-1); // left or right

// all constraints are satisfied at this point
// force KLEE to produce .err file:
klee_assert(0);
};

```

...and this version works slightly faster ( $\approx 5$  seconds), maybe because KLEE is aware of this *intrinsic* and handles it in a special way?

## 11.4 Sudoku

I've also rewritten Sudoku example (7.3) for KLEE:

```

1 #include <stdint.h>
2
3 /*
4 coordinates:
5 -----
6 00 01 02 | 03 04 05 | 06 07 08
7 10 11 12 | 13 14 15 | 16 17 18
8 20 21 22 | 23 24 25 | 26 27 28
9 -----
10 30 31 32 | 33 34 35 | 36 37 38
11 40 41 42 | 43 44 45 | 46 47 48
12 50 51 52 | 53 54 55 | 56 57 58

```

```

13 -----
14 60 61 62 | 63 64 65 | 66 67 68
15 70 71 72 | 73 74 75 | 76 77 78
16 80 81 82 | 83 84 85 | 86 87 88
17 -----
18 */
19
20 uint8_t cells[9][9];
21
22 // http://www.norvig.com/sudoku.html
23 // http://www.mirror.co.uk/news/weird-news/worlds-hardest-sudoku-can-you-242294
24 char *puzzle="..53.....8.....2..7..1.5..4....53...1..7...6..32...8..6.5....9..4....3.....97..";
25
26 int main()
27 {
28     klee_make_symbolic(cells, sizeof cells, "cells");
29
30     // process text line:
31     for (int row=0; row<9; row++)
32         for (int column=0; column<9; column++)
33         {
34             char c=puzzle[row*9 + column];
35             if (c!='.')
36             {
37                 if (cells[row][column]!=c-'0') return 0;
38             }
39             else
40             {
41                 // limit cells values to 1..9:
42                 if (cells[row][column]<1) return 0;
43                 if (cells[row][column]>9) return 0;
44             }
45         };
46
47     // for all 9 rows
48     for (int row=0; row<9; row++)
49     {
50
51         if (((1<<cells[row][0]) |
52             (1<<cells[row][1]) |
53             (1<<cells[row][2]) |
54             (1<<cells[row][3]) |
55             (1<<cells[row][4]) |
56             (1<<cells[row][5]) |
57             (1<<cells[row][6]) |
58             (1<<cells[row][7]) |
59             (1<<cells[row][8]))!=0x3FE ) return 0; // 11 1111 1110
60     };
61
62     // for all 9 columns
63     for (int c=0; c<9; c++)
64     {
65         if (((1<<cells[0][c]) |
66             (1<<cells[1][c]) |
67             (1<<cells[2][c]) |
68             (1<<cells[3][c]) |
69             (1<<cells[4][c]) |
70             (1<<cells[5][c]) |
71             (1<<cells[6][c]) |
72             (1<<cells[7][c]) |
73             (1<<cells[8][c]))!=0x3FE ) return 0; // 11 1111 1110
74     };
75
76     // enumerate all 9 squares
77     for (int r=0; r<9; r+=3)
78         for (int c=0; c<9; c+=3)
79         {
80             // add constraints for each 3*3 square:
81             if ((1<<cells[r+0][c+0]) |
82                 1<<cells[r+0][c+1] |
83                 1<<cells[r+0][c+2] |
84                 1<<cells[r+1][c+0] |
85                 1<<cells[r+1][c+1] |
86                 1<<cells[r+1][c+2] |
87                 1<<cells[r+2][c+0] |
88                 1<<cells[r+2][c+1] |

```

```

89         1<<cells[r+2][c+2])!=0x3FE ) return 0; // 11 1111 1110
90     };
91
92     // at this point, all constraints must be satisfied
93     klee_assert(0);
94 };

```

Let's run it:

```

% clang -emit-llvm -c -g klee_sudoku_or1.c
...

\ $ time klee klee_sudoku_or1.bc
KLEE: output directory is "/home/klee/klee-out-98"
KLEE: WARNING: undefined reference to function: klee_assert
KLEE: WARNING ONCE: calling external: klee_assert(0)
KLEE: ERROR: /home/klee/klee_sudoku_or1.c:93: failed external call: klee_assert
KLEE: NOTE: now ignoring this error at this location

KLEE: done: total instructions = 7512
KLEE: done: completed paths = 161
KLEE: done: generated tests = 161

real    3m44.111s
user    3m43.319s
sys     0m0.951s

```

Now this is really slower (on my Intel Core i3-3110M 2.4GHz notebook) in comparison to Z3Py solution (7.3).

But the answer is correct:

```

% ls klee-last | grep err
test000161.external.err

% ktest-tool --write-ints klee-last/test000161.ktest
ktest file : 'klee-last/test000161.ktest'
args       : ['klee_sudoku_or1.bc']
num objects: 1
object 0: name: b'cells'
object 0: size: 81
object 0: data: b'\x01\x04\x05\x03\x02\x07\x06\t\x08\x08\x03\t\x06\x05\x04\x01\x02\x07\x06\x07\x02\t\x01\x08\x05\x04\x03\x04\t\x06\x01\x08\x05\x03\x07\x02\x02\x01\x08\x04\x07\x03\t\x05\x06\x07\x05\x03\x02\t\x06\x04\x08\x01\x03\x06\x07\x05\x04\x02\x08\x01\t\t\x08\x04\x07\x06\x01\x02\x03\x05\x05\x02\x01\x08\x03\t\x07\x06\x04'

```

Character `\t` has code of 9 in C/C++, and KLEE prints byte array as a C/C++ string, so it shows some values in such way. We can just keep in mind that there is 9 at the each place where we see `\t`. The solution, while not properly formatted, correct indeed.

By the way, at lines 42 and 43 you may see how we tell to KLEE that all array elements must be within some limits. If we comment these lines out, we've got this:

```

% time klee klee_sudoku_or1.bc
KLEE: output directory is "/home/klee/klee-out-100"
KLEE: WARNING: undefined reference to function: klee_assert
KLEE: ERROR: /home/klee/klee_sudoku_or1.c:51: overshift error
KLEE: NOTE: now ignoring this error at this location
KLEE: ERROR: /home/klee/klee_sudoku_or1.c:51: overshift error
KLEE: NOTE: now ignoring this error at this location
KLEE: ERROR: /home/klee/klee_sudoku_or1.c:51: overshift error
KLEE: NOTE: now ignoring this error at this location
...

```

KLEE warns us that shift value at line 51 is too big. Indeed, KLEE may try all byte values up to 255 (0xFF), which are pointless to use there, and may be a symptom of error or bug, so KLEE warns about it.

Now let's use `klee_assume()` again:

```

#include <stdint.h>

/*
coordinates:
-----
00 01 02 | 03 04 05 | 06 07 08
10 11 12 | 13 14 15 | 16 17 18

```

```

20 21 22 | 23 24 25 | 26 27 28
-----
30 31 32 | 33 34 35 | 36 37 38
40 41 42 | 43 44 45 | 46 47 48
50 51 52 | 53 54 55 | 56 57 58
-----
60 61 62 | 63 64 65 | 66 67 68
70 71 72 | 73 74 75 | 76 77 78
80 81 82 | 83 84 85 | 86 87 88
-----
*/

uint8_t cells[9][9];

// http://www.norvig.com/sudoku.html
// http://www.mirror.co.uk/news/weird-news/worlds-hardest-sudoku-can-you-242294
char *puzzle="..53.....8.....2..7..1.5..4....53...1..7...6..32...8..6.5....9..4....3.....97..";

int main()
{
    klee_make_symbolic(cells, sizeof cells, "cells");

    // process text line:
    for (int row=0; row<9; row++)
        for (int column=0; column<9; column++)
        {
            char c=puzzle[row*9 + column];
            if (c!='.')
                klee_assume (cells[row][column]==c-'0');
            else
            {
                klee_assume (cells[row][column]>=1);
                klee_assume (cells[row][column]<=9);
            };
        };

    // for all 9 rows
    for (int row=0; row<9; row++)
    {
        klee_assume (((1<<cells[row][0]) |
                      (1<<cells[row][1]) |
                      (1<<cells[row][2]) |
                      (1<<cells[row][3]) |
                      (1<<cells[row][4]) |
                      (1<<cells[row][5]) |
                      (1<<cells[row][6]) |
                      (1<<cells[row][7]) |
                      (1<<cells[row][8]))==0x3FE ); // 11 1111 1110

    };

    // for all 9 columns
    for (int c=0; c<9; c++)
    {
        klee_assume (((1<<cells[0][c]) |
                      (1<<cells[1][c]) |
                      (1<<cells[2][c]) |
                      (1<<cells[3][c]) |
                      (1<<cells[4][c]) |
                      (1<<cells[5][c]) |
                      (1<<cells[6][c]) |
                      (1<<cells[7][c]) |
                      (1<<cells[8][c]))==0x3FE ); // 11 1111 1110

    };

    // enumerate all 9 squares
    for (int r=0; r<9; r+=3)
        for (int c=0; c<9; c+=3)
        {
            // add constraints for each 3*3 square:
            klee_assume ((1<<cells[r+0][c+0] |
                          1<<cells[r+0][c+1] |
                          1<<cells[r+0][c+2] |

```



```

        1<<cells[r+1][c+0] |
        1<<cells[r+1][c+1] |
        1<<cells[r+1][c+2] |
        1<<cells[r+2][c+0] |
        1<<cells[r+2][c+1] |
        1<<cells[r+2][c+2])==0x3FE ); // 11 1111 1110
    };

    // at this point, all constraints must be satisfied
    klee_assert(0);
};

```

```

% time klee klee_sudoku_or2.bc
KLEE: output directory is "/home/klee/klee-out-99"
KLEE: WARNING: undefined reference to function: klee_assert
KLEE: WARNING ONCE: calling external: klee_assert(0)
KLEE: ERROR: /home/klee/klee_sudoku_or2.c:93: failed external call: klee_assert
KLEE: NOTE: now ignoring this error at this location

KLEE: done: total instructions = 7119
KLEE: done: completed paths = 1
KLEE: done: generated tests = 1

real    0m35.312s
user    0m34.945s
sys     0m0.318s

```

That works much faster: perhaps KLEE indeed handle this *intrinsic* in a special way. And, as we see, the only one path has been found (one we actually interesting in it) instead of 161. It's still much slower than Z3Py solution, though.

## 11.5 Unit test: HTML/CSS color

The most popular ways to represent HTML/CSS color is by English name (e.g., “red”) and by 6-digit hexadecimal number (e.g., “#0077CC”). There is third, less popular way: if each byte in hexadecimal number has two doubling digits, it can be *abbreviated*, thus, “#0077CC” can be written just as “#07C”.

Let's write a function to convert 3 color components into name (if possible, first priority), 3-digit hexadecimal form (if possible, second priority), or as 6-digit hexadecimal form (as a last resort).

```

#include <string.h>
#include <stdio.h>
#include <stdint.h>

void HTML_color(uint8_t R, uint8_t G, uint8_t B, char* out)
{
    if (R==0xFF && G==0 && B==0)
    {
        strcpy (out, "red");
        return;
    };

    if (R==0x0 && G==0xFF && B==0)
    {
        strcpy (out, "green");
        return;
    };

    if (R==0 && G==0 && B==0xFF)
    {
        strcpy (out, "blue");
        return;
    };

    // abbreviated hexadecimal
    if (R>>4==(R&0xF) && G>>4==(G&0xF) && B>>4==(B&0xF))
    {
        sprintf (out, "#%X%X%X", R&0xF, G&0xF, B&0xF);
        return;
    };

    // last resort
    sprintf (out, "#%02X%02X%02X", R, G, B);
}

```

```

};

int main()
{
    uint8_t R, G, B;
    klee_make_symbolic (&R, sizeof R, "R");
    klee_make_symbolic (&G, sizeof R, "G");
    klee_make_symbolic (&B, sizeof R, "B");

    char tmp[16];

    HTML_color(R, G, B, tmp);
};

```

There are 5 possible paths in function, and let's see, if KLEE could find them all? It's indeed so:

```

% clang -emit-llvm -c -g color.c

% klee color.bc
KLEE: output directory is "/home/klee/klee-out-134"
KLEE: WARNING: undefined reference to function: sprintf
KLEE: WARNING: undefined reference to function: strcpy
KLEE: WARNING ONCE: calling external: strcpy(51867584, 51598960)
KLEE: ERROR: /home/klee/color.c:33: external call with symbolic argument: sprintf
KLEE: NOTE: now ignoring this error at this location
KLEE: ERROR: /home/klee/color.c:28: external call with symbolic argument: sprintf
KLEE: NOTE: now ignoring this error at this location

KLEE: done: total instructions = 479
KLEE: done: completed paths = 19
KLEE: done: generated tests = 5

```

We can ignore calls to strcpy() and sprintf(), because we are not really interesting in state of `out` variable. So there are exactly 5 paths:

```

% ls klee-last
assembly.ll  run.stats      test000003.ktest  test000005.ktest
info         test000001.ktest  test000003.pc     test000005.pc
messages.txt test000002.ktest  test000004.ktest  warnings.txt
run.istats   test000003.exec.err test000005.exec.err

```

1st set of input variables will result in "red" string:

```

% ktest-tool --write-ints klee-last/test000001.ktest
ktest file : 'klee-last/test000001.ktest'
args       : ['color.bc']
num objects: 3
object 0: name: b'R'
object 0: size: 1
object 0: data: b'\xff'
object 1: name: b'G'
object 1: size: 1
object 1: data: b'\x00'
object 2: name: b'B'
object 2: size: 1
object 2: data: b'\x00'

```

2nd set of input variables will result in "green" string:

```

% ktest-tool --write-ints klee-last/test000002.ktest
ktest file : 'klee-last/test000002.ktest'
args       : ['color.bc']
num objects: 3
object 0: name: b'R'
object 0: size: 1
object 0: data: b'\x00'
object 1: name: b'G'
object 1: size: 1
object 1: data: b'\xff'
object 2: name: b'B'
object 2: size: 1
object 2: data: b'\x00'

```

3rd set of input variables will result in "#010000" string:

```
% ktest-tool --write-ints klee-last/test000003.ktest
ktest file : 'klee-last/test000003.ktest'
args       : ['color.bc']
num objects: 3
object 0: name: b'R'
object 0: size: 1
object 0: data: b'\x01'
object 1: name: b'G'
object 1: size: 1
object 1: data: b'\x00'
object 2: name: b'B'
object 2: size: 1
object 2: data: b'\x00'
```

4th set of input variables will result in “blue” string:

```
% ktest-tool --write-ints klee-last/test000004.ktest
ktest file : 'klee-last/test000004.ktest'
args       : ['color.bc']
num objects: 3
object 0: name: b'R'
object 0: size: 1
object 0: data: b'\x00'
object 1: name: b'G'
object 1: size: 1
object 1: data: b'\x00'
object 2: name: b'B'
object 2: size: 1
object 2: data: b'\xff'
```

5th set of input variables will result in “#F01” string:

```
% ktest-tool --write-ints klee-last/test000005.ktest
ktest file : 'klee-last/test000005.ktest'
args       : ['color.bc']
num objects: 3
object 0: name: b'R'
object 0: size: 1
object 0: data: b'\xff'
object 1: name: b'G'
object 1: size: 1
object 1: data: b'\x00'
object 2: name: b'B'
object 2: size: 1
object 2: data: b'\x11'
```

These 5 sets of input variables can form a unit test for our function.

## 11.6 Unit test: strcmp() function

The standard `strcmp()` function from C library can return 0, -1 or 1, depending of comparison result.

Here is my own implementation of `strcmp()` :

```
int my_strcmp(const char *s1, const char *s2)
{
    int ret = 0;

    while (1)
    {
        ret = *(unsigned char *) s1 - *(unsigned char *) s2;
        if (ret!=0)
            break;
        if ((*s1==0) || (*s2)==0)
            break;
        s1++;
        s2++;
    };

    if (ret < 0)
    {
        return -1;
    } else if (ret > 0)
    {
        return 1;
    }
}
```

```

        return 1;
    }

    return 0;
}

int main()
{
    char input1[2];
    char input2[2];

    klee_make_symbolic(input1, sizeof input1, "input1");
    klee_make_symbolic(input2, sizeof input2, "input2");

    klee_assume((input1[0]>='a') && (input1[0]<='z'));
    klee_assume((input2[0]>='a') && (input2[0]<='z'));

    klee_assume(input1[1]==0);
    klee_assume(input2[1]==0);

    my_strcmp (input1, input2);
};

```

Let's find out, if KLEE is capable of finding all three paths? I intentionally made things simpler for KLEE by limiting input arrays to two 2 bytes or to 1 character + terminal zero byte.

```

% clang -emit-llvm -c -g strcmp.c

% klee strcmp.bc
KLEE: output directory is "/home/klee/klee-out-131"
KLEE: ERROR: /home/klee/strcmp.c:35: invalid klee_assume call (provably false)
KLEE: NOTE: now ignoring this error at this location
KLEE: ERROR: /home/klee/strcmp.c:36: invalid klee_assume call (provably false)
KLEE: NOTE: now ignoring this error at this location

KLEE: done: total instructions = 137
KLEE: done: completed paths = 5
KLEE: done: generated tests = 5

% ls klee-last
assembly.ll      run.stats        test000002.ktest  test000004.ktest
info             test000001.ktest test000002.pc     test000005.ktest
messages.txt     test000001.pc    test000002.user.err warnings.txt
run.istats       test000001.user.err test000003.ktest

```

The first two errors are about `klee_assume()`. These are input values on which `klee_assume()` calls are stuck. We can ignore them, or take a peek out of curiosity:

```

% ktest-tool --write-ints klee-last/test000001.ktest
ktest file : 'klee-last/test000001.ktest'
args       : ['strcmp.bc']
num objects: 2
object 0: name: b'input1'
object 0: size: 2
object 0: data: b'\x00\x00'
object 1: name: b'input2'
object 1: size: 2
object 1: data: b'\x00\x00'

% ktest-tool --write-ints klee-last/test000002.ktest
ktest file : 'klee-last/test000002.ktest'
args       : ['strcmp.bc']
num objects: 2
object 0: name: b'input1'
object 0: size: 2
object 0: data: b'a\xff'
object 1: name: b'input2'
object 1: size: 2
object 1: data: b'\x00\x00'

```

Three rest files are the input values for each path inside of my implementation of `strcmp()` :

```

% ktest-tool --write-ints klee-last/test000003.ktest
ktest file : 'klee-last/test000003.ktest'
args       : ['strcmp.bc']

```

```

num objects: 2
object 0: name: b'input1'
object 0: size: 2
object 0: data: b'b\x00'
object 1: name: b'input2'
object 1: size: 2
object 1: data: b'c\x00'

% ktest-tool --write-ints klee-last/test000004.ktest
ktest file : 'klee-last/test000004.ktest'
args       : ['strcmp.bc']
num objects: 2
object 0: name: b'input1'
object 0: size: 2
object 0: data: b'c\x00'
object 1: name: b'input2'
object 1: size: 2
object 1: data: b'a\x00'

% ktest-tool --write-ints klee-last/test000005.ktest
ktest file : 'klee-last/test000005.ktest'
args       : ['strcmp.bc']
num objects: 2
object 0: name: b'input1'
object 0: size: 2
object 0: data: b'a\x00'
object 1: name: b'input2'
object 1: size: 2
object 1: data: b'a\x00'

```

3rd is about first argument (“b”) is lesser than the second (“c”). 4th is opposite (“c” and “a”). 5th is when they are equal (“a” and “a”).

Using these 3 test cases, we’ve got full coverage of our implementation of `strcmp()`.

## 11.7 UNIX date/time

UNIX date/time<sup>91</sup> is a number of seconds that have elapsed since 1-Jan-1970 00:00 UTC. C/C++ `gmtime()` function is used to decode this value into human-readable date/time.

Here is a piece of code I’ve copypasted from some ancient version of Minix OS (<http://www.cise.ufl.edu/~cop4600/cgi-bin/lxr/http/source.cgi/lib/ansi/gmtime.c>) and reworked slightly:

```

1 #include <stdint.h>
2 #include <time.h>
3 #include <assert.h>
4
5 /*
6  * copypasted and reworked from
7  * http://www.cise.ufl.edu/~cop4600/cgi-bin/lxr/http/source.cgi/lib/ansi/loc_time.h
8  * http://www.cise.ufl.edu/~cop4600/cgi-bin/lxr/http/source.cgi/lib/ansi/misc.c
9  * http://www.cise.ufl.edu/~cop4600/cgi-bin/lxr/http/source.cgi/lib/ansi/gmtime.c
10 */
11
12 #define YEAR0          1900
13 #define EPOCH_YR       1970
14 #define SECS_DAY       (24L * 60L * 60L)
15 #define YEARSIZE(year) (LEAPYEAR(year) ? 366 : 365)
16
17 const int _ytab[2][12] =
18 {
19     { 31, 28, 31, 30, 31, 30, 31, 31, 30, 31, 30, 31 },
20     { 31, 29, 31, 30, 31, 30, 31, 31, 30, 31, 30, 31 }
21 };
22
23 const char *_days[] =
24 {
25     "Sunday", "Monday", "Tuesday", "Wednesday",
26     "Thursday", "Friday", "Saturday"
27 };
28
29 const char *_months[] =
30 {

```

<sup>91</sup>[https://en.wikipedia.org/wiki/Unix\\_time](https://en.wikipedia.org/wiki/Unix_time)

```

31     "January", "February", "March",
32     "April", "May", "June",
33     "July", "August", "September",
34     "October", "November", "December"
35 };
36
37 #define LEAPYEAR(year) (!((year) % 4) && (((year) % 100) || !((year) % 400)))
38
39 void decode_UNIX_time(const time_t time)
40 {
41     unsigned int dayclock, dayno;
42     int year = EPOCH_YR;
43
44     dayclock = (unsigned long)time % SECS_DAY;
45     dayno = (unsigned long)time / SECS_DAY;
46
47     int seconds = dayclock % 60;
48     int minutes = (dayclock % 3600) / 60;
49     int hour = dayclock / 3600;
50     int wday = (dayno + 4) % 7;
51     while (dayno >= YEARSIZE(year))
52     {
53         dayno -= YEARSIZE(year);
54         year++;
55     }
56
57     year = year - YEAR0;
58
59     int month = 0;
60
61     while (dayno >= _ytab[LEAPYEAR(year)][month])
62     {
63         dayno -= _ytab[LEAPYEAR(year)][month];
64         month++;
65     }
66
67     char *s;
68     switch (month)
69     {
70         case 0: s="January"; break;
71         case 1: s="February"; break;
72         case 2: s="March"; break;
73         case 3: s="April"; break;
74         case 4: s="May"; break;
75         case 5: s="June"; break;
76         case 6: s="July"; break;
77         case 7: s="August"; break;
78         case 8: s="September"; break;
79         case 9: s="October"; break;
80         case 10: s="November"; break;
81         case 11: s="December"; break;
82         default:
83             assert(0);
84     };
85
86     printf ("%04d-%s-%02d %02d:%02d:%02d\n", YEAR0+year, s, dayno+1, hour, minutes, seconds);
87     printf ("week day: %s\n", _days[wday]);
88 }
89
90 int main()
91 {
92     uint32_t time;
93
94     klee_make_symbolic(&time, sizeof time, "time");
95
96     decode_UNIX_time(time);
97
98     return 0;
99 }

```

Let's try it:

```

% clang -emit-llvm -c -g klee_time1.c
...
% klee klee_time1.bc

```

```

KLEE: output directory is "/home/klee/klee-out-107"
KLEE: WARNING: undefined reference to function: printf
KLEE: ERROR: /home/klee/klee_time1.c:86: external call with symbolic argument: printf
KLEE: NOTE: now ignoring this error at this location
KLEE: ERROR: /home/klee/klee_time1.c:83: ASSERTION FAIL: 0
KLEE: NOTE: now ignoring this error at this location

KLEE: done: total instructions = 101579
KLEE: done: completed paths = 1635
KLEE: done: generated tests = 2

```

Wow, assert() at line 83 has been triggered, why? Let's see a value of UNIX time which triggers it:

```

% ls klee-last | grep err
test000001.exec.err
test000002.assert.err

% ktest-tool --write-ints klee-last/test000002.ktest
ktest file : 'klee-last/test000002.ktest'
args       : ['klee_time1.bc']
num objects: 1
object 0: name: b'time'
object 0: size: 4
object 0: data: 978278400

```

Let's decode this value using UNIX date utility:

```

% date -u --date='@978278400'
Sun Dec 31 16:00:00 UTC 2000

```

After my investigation, I've found that `month` variable can hold incorrect value of 12 (while 11 is maximal, for December), because `LEAPYEAR()` macro should receive year number as 2000, not as 100. So I've introduced a bug during rewriting this function, and KLEE found it!

Just interesting, what would be if I'll replace `switch()` to array of strings, like it usually happens in concise C/C++ code?

```

...
const char *_months[] =
{
    "January", "February", "March",
    "April", "May", "June",
    "July", "August", "September",
    "October", "November", "December"
};

...

while (dayno >= _ytab[LEAPYEAR(year)][month])
{
    dayno -= _ytab[LEAPYEAR(year)][month];
    month++;
}

char *s=_months[month];

printf ("%04d-%s-%02d %02d:%02d:%02d\n", YEAR0+year, s, dayno+1, hour, minutes, seconds);
printf ("week day: %s\n", _days[wday]);

...

```

KLEE detects attempt to read beyond array boundaries:

```

% klee klee_time2.bc
KLEE: output directory is "/home/klee/klee-out-108"
KLEE: WARNING: undefined reference to function: printf
KLEE: ERROR: /home/klee/klee_time2.c:69: external call with symbolic argument: printf
KLEE: NOTE: now ignoring this error at this location
KLEE: ERROR: /home/klee/klee_time2.c:67: memory error: out of bound pointer
KLEE: NOTE: now ignoring this error at this location

KLEE: done: total instructions = 101716
KLEE: done: completed paths = 1635
KLEE: done: generated tests = 2

```

This is the same UNIX time value we've already seen:

```
% ls klee-last | grep err
test000001.exec.err
test000002.ptr.err

% ktest-tool --write-ints klee-last/test000002.ktest
ktest file : 'klee-last/test000002.ktest'
args       : ['klee_time2.bc']
num objects: 1
object 0: name: b'time'
object 0: size: 4
object 0: data: 978278400
```

So, if this piece of code can be triggered on remote computer, with this input value (*input of death*), it's possible to crash the process (with some luck, though).

OK, now I'm fixing a bug by moving year subtracting expression to line 43, and let's find, what UNIX time value corresponds to some fancy date like 2022-February-2?

```
1 #include <stdint.h>
2 #include <time.h>
3 #include <assert.h>
4
5 #define YEAR0          1900
6 #define EPOCH_YR      1970
7 #define SECS_DAY      (24L * 60L * 60L)
8 #define YEARSIZE(year) (LEAPYEAR(year) ? 366 : 365)
9
10 const int _ytab[2][12] =
11 {
12     { 31, 28, 31, 30, 31, 30, 31, 31, 30, 31, 30, 31 },
13     { 31, 29, 31, 30, 31, 30, 31, 31, 30, 31, 30, 31 }
14 };
15
16 #define LEAPYEAR(year) (!((year) % 4) && (((year) % 100) || !((year) % 400)))
17
18 void decode_UNIX_time(const time_t time)
19 {
20     unsigned int dayclock, dayno;
21     int year = EPOCH_YR;
22
23     dayclock = (unsigned long)time % SECS_DAY;
24     dayno = (unsigned long)time / SECS_DAY;
25
26     int seconds = dayclock % 60;
27     int minutes = (dayclock % 3600) / 60;
28     int hour = dayclock / 3600;
29     int wday = (dayno + 4) % 7;
30     while (dayno >= YEARSIZE(year))
31     {
32         dayno -= YEARSIZE(year);
33         year++;
34     }
35
36     int month = 0;
37
38     while (dayno >= _ytab[LEAPYEAR(year)][month])
39     {
40         dayno -= _ytab[LEAPYEAR(year)][month];
41         month++;
42     }
43     year = year - YEAR0;
44
45     if (YEAR0+year==2022 && month==1 && dayno+1==22)
46         klee_assert(0);
47 }
48 int main()
49 {
50     uint32_t time;
51
52     klee_make_symbolic(&time, sizeof time, "time");
53
54     decode_UNIX_time(time);
55 }
```



```

56     return 0;
57 }

```

```

% clang -emit-llvm -c -g klee_time3.c
...

% klee klee_time3.bc
KLEE: output directory is "/home/klee/klee-out-109"
KLEE: WARNING: undefined reference to function: klee_assert
KLEE: WARNING ONCE: calling external: klee_assert(0)
KLEE: ERROR: /home/klee/klee_time3.c:47: failed external call: klee_assert
KLEE: NOTE: now ignoring this error at this location

KLEE: done: total instructions = 101087
KLEE: done: completed paths = 1635
KLEE: done: generated tests = 1635

% ls klee-last | grep err
test000587.external.err

% ktest-tool --write-ints klee-last/test000587.ktest
ktest file : 'klee-last/test000587.ktest'
args       : ['klee_time3.bc']
num objects: 1
object 0: name: b'time'
object 0: size: 4
object 0: data: 1645488640

% date -u --date='@1645488640'
Tue Feb 22 00:10:40 UTC 2022

```

Success, but hours/minutes/seconds are seems random—they are random indeed, because, KLEE satisfied all constraints we’ve put, nothing else. We didn’t ask it to set hours/minutes/seconds to zeroes.

Let’s add constraints to hours/minutes/seconds as well:

```

...

if (YEAR0+year==2022 && month==1 && dayno+1==22 && hour==22 && minutes==22 && seconds==22)
    klee_assert(0);

...

```

Let’s run it and check ...

```

% ktest-tool --write-ints klee-last/test000597.ktest
ktest file : 'klee-last/test000597.ktest'
args       : ['klee_time3.bc']
num objects: 1
object 0: name: b'time'
object 0: size: 4
object 0: data: 1645568542

% date -u --date='@1645568542'
Tue Feb 22 22:22:22 UTC 2022

```

Now that is precise.

Yes, of course, C/C++ libraries has function(s) to encode human-readable date into UNIX time value, but what we’ve got here is KLEE working *antipode* of decoding function, *inverse function* in a way.

## 11.8 Inverse function for base64 decoder

It’s piece of cake for KLEE to reconstruct input base64 string given just base64 decoder code without corresponding encoder code. I’ve copy-pasted this piece of code from <http://www.opensource.apple.com/source/QuickTimeStreamingServer/QuickTimeStreamingServer-452/CommonUtilitiesLib/base64.c>.

We add constraints (lines 84, 85) so that output buffer must have byte values from 0 to 15. We also tell to KLEE that the Base64decode() function must return 16 (i.e., size of output buffer in bytes, line 82).

```

1 #include <string.h>
2 #include <stdint.h>
3 #include <stdbool.h>
4

```

```

5 // copyasted from http://www.opensource.apple.com/source/QuickTimeStreamingServer/QuickTimeStreamingServer-452/
  CommonUtilitiesLib/base64.c
6
7 static const unsigned char pr2six[256] =
8 {
9     /* ASCII table */
10    64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64,
11    64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64,
12    64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 62, 64, 64, 64, 63,
13    52, 53, 54, 55, 56, 57, 58, 59, 60, 61, 64, 64, 64, 64, 64, 64,
14    64, 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14,
15    15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 64, 64, 64, 64, 64,
16    64, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37, 38, 39, 40,
17    41, 42, 43, 44, 45, 46, 47, 48, 49, 50, 51, 64, 64, 64, 64, 64,
18    64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64,
19    64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64,
20    64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64,
21    64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64,
22    64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64,
23    64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64,
24    64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64,
25    64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64, 64
26 };
27
28 int Base64decode(char *bufplain, const char *bufcoded)
29 {
30     int nbytesdecoded;
31     register const unsigned char *bufin;
32     register unsigned char *bufout;
33     register int nprbytes;
34
35     bufin = (const unsigned char *) bufcoded;
36     while (pr2six[*bufin] <= 63);
37     nprbytes = (bufin - (const unsigned char *) bufcoded) - 1;
38     nbytesdecoded = ((nprbytes + 3) / 4) * 3;
39
40     bufout = (unsigned char *) bufplain;
41     bufin = (const unsigned char *) bufcoded;
42
43     while (nprbytes > 4) {
44         *(bufout++) =
45             (unsigned char) (pr2six[*bufin] << 2 | pr2six[bufin[1]] >> 4);
46         *(bufout++) =
47             (unsigned char) (pr2six[bufin[1]] << 4 | pr2six[bufin[2]] >> 2);
48         *(bufout++) =
49             (unsigned char) (pr2six[bufin[2]] << 6 | pr2six[bufin[3]]);
50         bufin += 4;
51         nprbytes -= 4;
52     }
53
54     /* Note: (nprbytes == 1) would be an error, so just ignore that case */
55     if (nprbytes > 1) {
56         *(bufout++) =
57             (unsigned char) (pr2six[*bufin] << 2 | pr2six[bufin[1]] >> 4);
58     }
59     if (nprbytes > 2) {
60         *(bufout++) =
61             (unsigned char) (pr2six[bufin[1]] << 4 | pr2six[bufin[2]] >> 2);
62     }
63     if (nprbytes > 3) {
64         *(bufout++) =
65             (unsigned char) (pr2six[bufin[2]] << 6 | pr2six[bufin[3]]);
66     }
67
68     *(bufout++) = '\0';
69     nbytesdecoded -= (4 - nprbytes) & 3;
70     return nbytesdecoded;
71 }
72
73 int main()
74 {
75     char input[32];
76     uint8_t output[16+1];
77
78     klee_make_symbolic(input, sizeof input, "input");
79

```

```

80     klee_assume(input[31]==0);
81
82     klee_assume (Base64decode(output, input)==16);
83
84     for (int i=0; i<16; i++)
85         klee_assume (output[i]==i);
86
87     klee_assert(0);
88
89     return 0;
90 }

```

```

% clang -emit-llvm -c -g klee_base64.c
...

% klee klee_base64.bc
KLEE: output directory is "/home/klee/klee-out-99"
KLEE: WARNING: undefined reference to function: klee_assert
KLEE: ERROR: /home/klee/klee_base64.c:99: invalid klee_assume call (provably false)
KLEE: NOTE: now ignoring this error at this location
KLEE: WARNING ONCE: calling external: klee_assert(0)
KLEE: ERROR: /home/klee/klee_base64.c:104: failed external call: klee_assert
KLEE: NOTE: now ignoring this error at this location
KLEE: ERROR: /home/klee/klee_base64.c:85: memory error: out of bound pointer
KLEE: NOTE: now ignoring this error at this location
KLEE: ERROR: /home/klee/klee_base64.c:81: memory error: out of bound pointer
KLEE: NOTE: now ignoring this error at this location
KLEE: ERROR: /home/klee/klee_base64.c:65: memory error: out of bound pointer
KLEE: NOTE: now ignoring this error at this location
...

```

We're interesting in the second error, where `klee_assert()` has been triggered:

```

% ls klee-last | grep err
test000001.user.err
test000002.external.err
test000003.ptr.err
test000004.ptr.err
test000005.ptr.err

% ktest-tool --write-ints klee-last/test000002.ktest
ktest file : 'klee-last/test000002.ktest'
args       : ['klee_base64.bc']
num objects: 1
object  0: name: b'input'
object  0: size: 32
object  0: data: b'AAECAwQFBgcICQoLDA00D4\x00\xff\xff\xff\xff\xff\xff\xff\xff\x00'

```

This is indeed a real base64 string, terminated with the zero byte, just as it's requested by C/C++ standards. The final zero byte at 31th byte (starting at zeroth byte) is our deed: so that KLEE would report lesser number of errors.

The base64 string is indeed correct:

```

% echo AAECAwQFBgcICQoLDA00D4 | base64 -d | hexdump -C
base64: invalid input
00000000 00 01 02 03 04 05 06 07 08 09 0a 0b 0c 0d 0e 0f  |.....|
00000010

```

base64 decoder Linux utility I've just run blaming for "invalid input"—it means the input string is not properly padded. Now let's pad it manually, and decoder utility will no complain anymore:

```

% echo AAECAwQFBgcICQoLDA00D4== | base64 -d | hexdump -C
00000000 00 01 02 03 04 05 06 07 08 09 0a 0b 0c 0d 0e 0f  |.....|
00000010

```

The reason our generated base64 string is not padded is because base64 decoders are usually discards padding symbols ("=") at the end. In other words, they are not require them, so is the case of our decoder. Hence, padding symbols are left unnoticed to KLEE.

So we again made *antipode* or *inverse function* of base64 decoder.

## 11.9 CRC (Cyclic redundancy check)

### 11.9.1 Buffer alteration case #1

Sometimes, you need to alter a piece of data which is *protected* by some kind of checksum or CRC, and you can't change checksum or CRC value, but can alter piece of data so that checksum will remain the same.

Let's pretend, we've got a piece of data with "Hello, world!" string at the beginning and "and goodbye" string at the end. We can alter 14 characters at the middle, but for some reason, they must be in a..z limits, but we can put any characters there. CRC64 of the whole block must be 0x12345678abcdef12.

Let's see<sup>92</sup>:

```
#include <string.h>
#include <stdint.h>

uint64_t crc64(uint64_t crc, unsigned char *buf, int len)
{
    int k;

    crc = ~crc;
    while (len--)
    {
        crc ^= *buf++;
        for (k = 0; k < 8; k++)
            crc = crc & 1 ? (crc >> 1) ^ 0x42f0e1eba9ea3693 : crc >> 1;
    }
    return crc;
}

int main()
{
#define HEAD_STR "Hello, world!.. "
#define HEAD_SIZE strlen(HEAD_STR)
#define TAIL_STR " ... and goodbye"
#define TAIL_SIZE strlen(TAIL_STR)
#define MID_SIZE 14 // work
#define BUF_SIZE HEAD_SIZE+TAIL_SIZE+MID_SIZE

    char buf[BUF_SIZE];

    klee_make_symbolic(buf, sizeof buf, "buf");

    klee_assume (memcmp (buf, HEAD_STR, HEAD_SIZE)==0);

    for (int i=0; i<MID_SIZE; i++)
        klee_assume (buf[HEAD_SIZE+i]>='a' && buf[HEAD_SIZE+i]<='z');

    klee_assume (memcmp (buf+HEAD_SIZE+MID_SIZE, TAIL_STR, TAIL_SIZE)==0);

    klee_assume (crc64 (0, buf, BUF_SIZE)==0x12345678abcdef12);

    klee_assert(0);

    return 0;
}
```

Since our code uses memcmp() standard C/C++ function, we need to add `--libc=uclibc` switch, so KLEE will use its own uClibc implementation.

```
% clang -emit-llvm -c -g klee_CRC64.c
% time klee --libc=uclibc klee_CRC64.bc
```

It takes about 1 minute (on my Intel Core i3-3110M 2.4GHz notebook) and we getting this:

```
...
real    0m52.643s
user    0m51.232s
sys     0m0.239s
...
% ls klee-last | grep err
test000001.user.err
test000002.user.err
```

<sup>92</sup>There are several slightly different CRC64 implementations, the one I use here can also be different from popular ones.

```
test000003.user.err
test000004.external.err

% ktest-tool --write-ints klee-last/test000004.ktest
ktest file : 'klee-last/test000004.ktest'
args      : ['klee_CRC64.bc']
num objects: 1
object 0: name: b'buf'
object 0: size: 46
object 0: data: b'Hello, world!.. qqlicayzceamyw ... and goodbye'
```

Maybe it's slow, but definitely faster than bruteforce. Indeed,  $\log_2 26^{14} \approx 65.8$  which is close to 64 bits. In other words, one need  $\approx 14$  latin characters to encode 64 bits. And KLEE + [SMT](#) solver needs 64 bits at some place it can alter to make final CRC64 value equal to what we defined.

I tried to reduce length of the *middle block* to 13 characters: no luck for KLEE then, it has no space enough.

### 11.9.2 Buffer alteration case #2

I went sadistic: what if the buffer must contain the CRC64 value which, after calculation of CRC64, will result in the same value? Fascinatedly, KLEE can solve this. The buffer will have the following format:

```
Hello, world! <8 bytes (64-bit value)> and goodbye <6 more bytes>
```

```
int main()
{
#define HEAD_STR "Hello, world!.. "
#define HEAD_SIZE strlen(HEAD_STR)
#define TAIL_STR " ... and goodbye"
#define TAIL_SIZE strlen(TAIL_STR)
// 8 bytes for 64-bit value:
#define MID_SIZE 8
#define BUF_SIZE HEAD_SIZE+TAIL_SIZE+MID_SIZE+6

    char buf[BUF_SIZE];

    klee_make_symbolic(buf, sizeof buf, "buf");

    klee_assume (memcmp (buf, HEAD_STR, HEAD_SIZE)==0);

    klee_assume (memcmp (buf+HEAD_SIZE+MID_SIZE, TAIL_STR, TAIL_SIZE)==0);

    uint64_t mid_value=*(uint64_t*)(buf+HEAD_SIZE);
    klee_assume (crc64 (0, buf, BUF_SIZE)==mid_value);

    klee_assert(0);

    return 0;
}
```

It works:

```
% time klee --libc=uclibc klee_CRC64.bc
...
real    5m17.081s
user    5m17.014s
sys     0m0.319s

% ls klee-last | grep err
test000001.user.err
test000002.user.err
test000003.external.err

% ktest-tool --write-ints klee-last/test000003.ktest
ktest file : 'klee-last/test000003.ktest'
args      : ['klee_CRC64.bc']
num objects: 1
object 0: name: b'buf'
object 0: size: 46
object 0: data: b'Hello, world!.. T+]\xb9A\x08\x0fq ... and goodbye\xb6\x8f\x9c\xd8\xc5\x00'
```

8 bytes between two strings is 64-bit value which equals to CRC64 of this whole block. Again, it's faster than brute-force way to find it. If to decrease last spare 6-byte buffer to 4 bytes or less, KLEE works so long so I've stopped it.

### 11.9.3 Recovering input data for given CRC32 value of it

I've always wanted to do so, but everyone knows this is impossible for input buffers larger than 4 bytes. As my experiments show, it's still possible for tiny input buffers of data, which is constrained in some way.

The CRC32 value of 6-byte "SILVER" string is known: `0xDFA3DFDD`. KLEE can find this 6-byte string, if it knows that each byte of input buffer is in A..Z limits:

```
1 #include <stdint.h>
2 #include <stdbool.h>
3
4 uint32_t crc32(uint32_t crc, unsigned char *buf, int len)
5 {
6     int k;
7
8     crc = ~crc;
9     while (len--)
10     {
11         crc ^= *buf++;
12         for (k = 0; k < 8; k++)
13             crc = crc & 1 ? (crc >> 1) ^ 0xedb88320 : crc >> 1;
14     }
15     return ~crc;
16 }
17
18 #define SIZE 6
19
20 bool find_string(char str[SIZE])
21 {
22     int i=0;
23     for (i=0; i<SIZE; i++)
24         if (str[i]<'A' || str[i]>'Z')
25             return false;
26
27     if (crc32(0, &str[0], SIZE)!=0xDFA3DFDD)
28         return false;
29
30     // OK, input str is valid
31     klee_assert(0); // force KLEE to produce .err file
32     return true;
33 };
34
35 int main()
36 {
37     uint8_t str[SIZE];
38
39     klee_make_symbolic(str, sizeof str, "str");
40
41     find_string(str);
42
43     return 0;
44 }
```

```
% clang -emit-llvm -c -g klee_SILVER.c
...

% klee klee_SILVER.bc
...

% ls klee-last | grep err
test000013.external.err

% ktest-tool --write-ints klee-last/test000013.ktest
ktest file : 'klee-last/test000013.ktest'
args       : ['klee_SILVER.bc']
num objects: 1
object 0: name: b'str'
object 0: size: 6
object 0: data: b'SILVER'
```

Still, it's no magic: if to remove condition at lines 23..25 (i.e., if to relax constraints), KLEE will produce some other string, which will be still correct for the CRC32 value given.

It works, because 6 Latin characters in A..Z limits contain  $\approx 28.2$  bits:  $\log_2 26^6 \approx 28.2$ , which is even smaller value than 32. In other words, the final CRC32 value holds enough bits to recover  $\approx 28.2$  bits of input.

The input buffer can be even bigger, if each byte of it will be in even tighter constraints (decimal digits, binary digits, etc).

#### 11.9.4 In comparison with other hashing algorithms

Things are that easy for some other hashing algorithms like *Fletcher checksum*, but not for cryptographically secure ones (like MD5, SHA1, etc), they are protected from such simple cryptoanalysis. See also: [12](#).

### 11.10 LZSS decompressor

I've googled for a very simple [LZSS](#) decompressor and landed at this page: <http://www.opensource.apple.com/source/boot/boot-132/i386/boot2/lzss.c>.

Let's pretend, we're looking at unknown compressing algorithm with no compressor available. Will it be possible to reconstruct a compressed piece of data so that decompressor would generate data we need?

Here is my first experiment:

```
// copyasted from http://www.opensource.apple.com/source/boot/boot-132/i386/boot2/lzss.c
//
#include <string.h>
#include <stdint.h>
#include <stdbool.h>

// #define N          4096 /* size of ring buffer - must be power of 2 */
// #define N          32 /* size of ring buffer - must be power of 2 */
// #define F          18 /* upper limit for match_length */
// #define THRESHOLD 2    /* encode string into position and length
//                        if match_length is greater than this */
// #define NIL        N    /* index for root of binary search trees */

int
decompress_lzss(uint8_t *dst, uint8_t *src, uint32_t srclen)
{
    /* ring buffer of size N, with extra F-1 bytes to aid string comparison */
    uint8_t *dststart = dst;
    uint8_t *srcend = src + srclen;
    int i, j, k, r, c;
    unsigned int flags;
    uint8_t text_buf[N + F - 1];

    dst = dststart;
    srcend = src + srclen;
    for (i = 0; i < N - F; i++)
        text_buf[i] = ' ';
    r = N - F;
    flags = 0;
    for ( ; ; ) {
        if (((flags >>= 1) & 0x100) == 0) {
            if (src < srcend) c = *src++; else break;
            flags = c | 0xFF00; /* uses higher byte cleverly */
        } /* to count eight */
        if (flags & 1) {
            if (src < srcend) c = *src++; else break;
            *dst++ = c;
            text_buf[r++] = c;
            r &= (N - 1);
        } else {
            if (src < srcend) i = *src++; else break;
            if (src < srcend) j = *src++; else break;
            i |= ((j & 0xF0) << 4);
            j = (j & 0x0F) + THRESHOLD;
            for (k = 0; k <= j; k++) {
                c = text_buf[(i + k) & (N - 1)];
                *dst++ = c;
                text_buf[r++] = c;
                r &= (N - 1);
            }
        }
    }
    return dst - dststart;
}
```

```

int main()
{
#define COMPRESSED_LEN 15
    uint8_t input[COMPRESSED_LEN];
    uint8_t plain[24];
    uint32_t size=COMPRESSED_LEN;

    klee_make_symbolic(input, sizeof input, "input");

    decompress_lzss(plain, input, size);

    // https://en.wikipedia.org/wiki/Buffalo_buffalo_Buffalo_buffalo_buffalo_buffalo_Buffalo_buffalo
    for (int i=0; i<23; i++)
        klee_assume (plain[i]=="Buffalo buffalo Buffalo"[i]);

    klee_assert(0);

    return 0;
}

```

What I did is changing size of ring buffer from 4096 to 32, because if bigger, KLEE consumes all [RAM](#)<sup>93</sup> it can. But I've found that KLEE can live with that small buffer. I've also decreased `COMPRESSED_LEN` gradually to check, whether KLEE would find compressed piece of data, and it did:

```

% clang -emit-llvm -c -g klee_lzss.c
...

% time klee klee_lzss.bc
KLEE: output directory is "/home/klee/klee-out-7"
KLEE: WARNING: undefined reference to function: klee_assert
KLEE: ERROR: /home/klee/klee_lzss.c:122: invalid klee_assume call (provably false)
KLEE: NOTE: now ignoring this error at this location
KLEE: ERROR: /home/klee/klee_lzss.c:47: memory error: out of bound pointer
KLEE: NOTE: now ignoring this error at this location
KLEE: ERROR: /home/klee/klee_lzss.c:37: memory error: out of bound pointer
KLEE: NOTE: now ignoring this error at this location
KLEE: WARNING ONCE: calling external: klee_assert(0)
KLEE: ERROR: /home/klee/klee_lzss.c:124: failed external call: klee_assert
KLEE: NOTE: now ignoring this error at this location

KLEE: done: total instructions = 41417919
KLEE: done: completed paths = 437820
KLEE: done: generated tests = 4

real    13m0.215s
user    11m57.517s
sys     1m2.187s

% ls klee-last | grep err
test000001.user.err
test000002.ptr.err
test000003.ptr.err
test000004.external.err

% ktest-tool --write-ints klee-last/test000004.ktest
ktest file : 'klee-last/test000004.ktest'
args      : ['klee_lzss.bc']
num objects: 1
object 0: name: b'input'
object 0: size: 15
object 0: data: b'\xffBuffalo \x01b\x0f\x03\r\x05'

```

KLEE consumed  $\approx 1GB$  of RAM and worked for  $\approx 15$  minutes (on my Intel Core i3-3110M 2.4GHz notebook), but here it is, a 15 bytes which, if decompressed by our copypasted algorithm, will result in desired text!

During my experimentation, I've found that KLEE can do even more cooler thing, to find out size of compressed piece of data:

```

int main()
{
    uint8_t input[24];
    uint8_t plain[24];
    uint32_t size;

```

<sup>93</sup>Random-access memory



```

    klee_make_symbolic(input, sizeof input, "input");
    klee_make_symbolic(&size, sizeof size, "size");

    decompress_lzss(plain, input, size);

    for (int i=0; i<23; i++)
        klee_assume (plain[i]=="Buffalo buffalo Buffalo"[i]);

    klee_assert(0);

    return 0;
}

```

...but then KLEE works much slower, consumes much more RAM and I had success only with even smaller pieces of desired text.

So how [LZSS](#) works? Without peeking into Wikipedia, we can say that: if [LZSS](#) compressor observes some data it already had, it replaces the data with a link to some place in past with size. If it observes something yet unseen, it puts data as is. This is theory. This is indeed what we've got. Desired text is three "Buffalo" words, the first and the last are equivalent, but the second is *almost* equivalent, differing with first by one character.

That's what we see:

```
'\xffBuffalo \x01b\x0f\x03\r\x05'
```

Here is some control byte (0xff), "Buffalo" word is placed as *is*, then another control byte (0x01), then we see beginning of the second word ("b") and more control bytes, perhaps, links to the beginning of the buffer. These are command to decompressor, like, in plain English, "copy data from the buffer we've already done, from that place to that place", etc.

Interesting, is it possible to meddle into this piece of compressed data? Out of whim, can we force KLEE to find a compressed data, where not just "b" character has been placed as *is*, but also the second character of the word, i.e., "bu"?

I've modified main() function by adding `klee_assume()` : now the 11th byte of input (compressed) data (right after "b" byte) must have "u". I has no luck with 15 byte of compressed data, so I increased it to 16 bytes:

```

int main()
{
#define COMPRESSED_LEN 16
    uint8_t input[COMPRESSED_LEN];
    uint8_t plain[24];
    uint32_t size=COMPRESSED_LEN;

    klee_make_symbolic(input, sizeof input, "input");

    klee_assume(input[11]=='u');

    decompress_lzss(plain, input, size);

    for (int i=0; i<23; i++)
        klee_assume (plain[i]=="Buffalo buffalo Buffalo"[i]);

    klee_assert(0);

    return 0;
}

```

...and voilà: KLEE found a compressed piece of data which satisfied our whimsical constraint:

```

% time klee klee_lzss.bc
KLEE: output directory is "/home/klee/klee-out-9"
KLEE: WARNING: undefined reference to function: klee_assert
KLEE: ERROR: /home/klee/klee_lzss.c:97: invalid klee_assume call (provably false)
KLEE: NOTE: now ignoring this error at this location
KLEE: ERROR: /home/klee/klee_lzss.c:47: memory error: out of bound pointer
KLEE: NOTE: now ignoring this error at this location
KLEE: ERROR: /home/klee/klee_lzss.c:37: memory error: out of bound pointer
KLEE: NOTE: now ignoring this error at this location
KLEE: WARNING ONCE: calling external: klee_assert(0)
KLEE: ERROR: /home/klee/klee_lzss.c:99: failed external call: klee_assert
KLEE: NOTE: now ignoring this error at this location

```

```

KLEE: done: total instructions = 36700587
KLEE: done: completed paths = 369756
KLEE: done: generated tests = 4

real    12m16.983s
user    11m17.492s
sys      0m58.358s

% ktest-tool --write-ints klee-last/test000004.ktest
ktest file : 'klee-last/test000004.ktest'
args       : ['klee_lzss.bc']
num objects: 1
object 0: name: b'input'
object 0: size: 16
object 0: data: b'\xffBuffalo \x13bu\x10\x02\r\x05'

```

So now we find a piece of compressed data where two strings are placed as is: “Buffalo” and “bu”.

```
'\xffBuffalo \x13bu\x10\x02\r\x05'
```

Both pieces of compressed data, if feeded into our cypasted function, produce “Buffalo buffalo Buffalo” text string.

Please note, I still have no access to [LZSS](#) compressor code, and I didn’t get into [LZSS](#) decompressor details yet.

Unfortunately, things are not that cool: KLEE is very slow and I had success only with small pieces of text, and also ring buffer size had to be decreased significantly (original [LZSS](#) decompressor with ring buffer of 4096 bytes cannot decompress correctly what we found).

Nevertheless, it’s very impressive, taking into account the fact that we’re not getting into internals of this specific LZSS decompressor. Once more time, we’ve created *antipode* of decompressor, or *inverse function*.

Also, as it seems, KLEE isn’t very good so far with decompression algorithms (but who’s good then?). I’ve also tried various JPEG/PNG/GIF decoders (which, of course, has decompressors), starting with simplest possible, and KLEE had stuck.

## 11.11 strtodx() from RetroBSD

Just found this function in RetroBSD: <https://github.com/RetroBSD/retrobsd/blob/master/src/libc/stdlib/strtod.c>. It converts a string into floating point number for given radix.

```

1 #include <stdio.h>
2
3 // my own version, only for radix 10:
4 int isdigitx (char c, int radix)
5 {
6     if (c>='0' && c<='9')
7         return 1;
8     return 0;
9 };
10
11 /*
12  * double strtodx (char *string, char **endPtr, int radix)
13  *     This procedure converts a floating-point number from an ASCII
14  *     decimal representation to internal double-precision format.
15  *
16  * Original sources taken from 386bsd and modified for variable radix
17  * by Serge Vakulenko, <vak@kiaes.ru>.
18  *
19  * Arguments:
20  * string
21  *     A decimal ASCII floating-point number, optionally preceded
22  *     by white space. Must have form "-I.FE-X", where I is the integer
23  *     part of the mantissa, F is the fractional part of the mantissa,
24  *     and X is the exponent. Either of the signs may be "+", "-", or
25  *     omitted. Either I or F may be omitted, or both. The decimal point
26  *     isn't necessary unless F is present. The "E" may actually be an "e",
27  *     or "E", "S", "s", "F", "f", "D", "d", "L", "l".
28  *     E and X may both be omitted (but not just one).
29  *
30  * endPtr
31  *     If non-NULL, store terminating character's address here.
32  *
33  * radix

```

```

34 *      Radix of floating point, one of 2, 8, 10, 16.
35 *
36 * The return value is the double-precision floating-point
37 * representation of the characters in string. If endPtr isn't
38 * NULL, then *endPtr is filled in with the address of the
39 * next character after the last one that was part of the
40 * floating-point number.
41 */
42 double strtodx (char *string, char **endPtr, int radix)
43 {
44     int sign = 0, expSign = 0, fracSz, fracOff, i;
45     double fraction, dblExp, *powTab;
46     register char *p;
47     register char c;
48
49     /* Exponent read from "EX" field. */
50     int exp = 0;
51
52     /* Exponent that derives from the fractional part. Under normal
53      * circumstances, it is the negative of the number of digits in F.
54      * However, if I is very long, the last digits of I get dropped
55      * (otherwise a long I with a large negative exponent could cause an
56      * unnecessary overflow on I alone). In this case, fracExp is
57      * incremented one for each dropped digit. */
58     int fracExp = 0;
59
60     /* Number of digits in mantissa. */
61     int mantSize;
62
63     /* Number of mantissa digits BEFORE decimal point. */
64     int decPt;
65
66     /* Temporarily holds location of exponent in string. */
67     char *pExp;
68
69     /* Largest possible base 10 exponent.
70      * Any exponent larger than this will already
71      * produce underflow or overflow, so there's
72      * no need to worry about additional digits. */
73     static int maxExponent = 307;
74
75     /* Table giving binary powers of 10.
76      * Entry is 10^2^i. Used to convert decimal
77      * exponents into floating-point numbers. */
78     static double powersOf10[] = {
79         1e1, 1e2, 1e4, 1e8, 1e16, 1e32, //1e64, 1e128, 1e256,
80     };
81     static double powersOf2[] = {
82         2, 4, 16, 256, 65536, 4.294967296e9, 1.8446744073709551616e19,
83         //3.4028236692093846346e38, 1.1579208923731619542e77, 1.3407807929942597099e154,
84     };
85     static double powersOf8[] = {
86         8, 64, 4096, 2.81474976710656e14, 7.9228162514264337593e28,
87         //6.2771017353866807638e57, 3.9402006196394479212e115, 1.5525180923007089351e231,
88     };
89     static double powersOf16[] = {
90         16, 256, 65536, 1.8446744073709551616e19,
91         //3.4028236692093846346e38, 1.1579208923731619542e77, 1.3407807929942597099e154,
92     };
93
94     /*
95      * Strip off leading blanks and check for a sign.
96      */
97     p = string;
98     while (*p==' ' || *p=='\t')
99         ++p;
100     if (*p == '-') {
101         sign = 1;
102         ++p;
103     } else if (*p == '+')
104         ++p;
105
106     /*
107      * Count the number of digits in the mantissa (including the decimal
108      * point), and also locate the decimal point.
109      */

```

```

110 decPt = -1;
111 for (mantSize=0; ; ++mantSize) {
112     c = *p;
113     if (!isdigitx (c, radix)) {
114         if (c != '.' || decPt >= 0)
115             break;
116         decPt = mantSize;
117     }
118     ++p;
119 }
120
121 /*
122  * Now suck up the digits in the mantissa. Use two integers to
123  * collect 9 digits each (this is faster than using floating-point).
124  * If the mantissa has more than 18 digits, ignore the extras, since
125  * they can't affect the value anyway.
126  */
127 pExp = p;
128 p -= mantSize;
129 if (decPt < 0)
130     decPt = mantSize;
131 else
132     --mantSize;          /* One of the digits was the point. */
133
134 switch (radix) {
135 default:
136 case 10: fracSz = 9; fracOff = 1000000000; powTab = powersOf10; break;
137 case 2:  fracSz = 30; fracOff = 1073741824; powTab = powersOf2; break;
138 case 8:  fracSz = 10; fracOff = 1073741824; powTab = powersOf8; break;
139 case 16: fracSz = 7;  fracOff = 268435456;  powTab = powersOf16; break;
140 }
141 if (mantSize > 2 * fracSz)
142     mantSize = 2 * fracSz;
143 fracExp = decPt - mantSize;
144 if (mantSize == 0) {
145     fraction = 0.0;
146     p = string;
147     goto done;
148 } else {
149     int frac1, frac2;
150
151     for (frac1=0; mantSize>fracSz; --mantSize) {
152         c = *p++;
153         if (c == '.')
154             c = *p++;
155         frac1 = frac1 * radix + (c - '0');
156     }
157     for (frac2=0; mantSize>0; --mantSize) {
158         c = *p++;
159         if (c == '.')
160             c = *p++;
161         frac2 = frac2 * radix + (c - '0');
162     }
163     fraction = (double) fracOff * frac1 + frac2;
164 }
165
166 /*
167  * Skim off the exponent.
168  */
169 p = pExp;
170 if (*p=='E' || *p=='e' || *p=='S' || *p=='s' || *p=='F' || *p=='f' ||
171     *p=='D' || *p=='d' || *p=='L' || *p=='l') {
172     ++p;
173     if (*p == '-') {
174         expSign = 1;
175         ++p;
176     } else if (*p == '+')
177         ++p;
178     while (isdigitx (*p, radix))
179         exp = exp * radix + (*p++ - '0');
180 }
181 if (expSign)
182     exp = fracExp - exp;
183 else
184     exp = fracExp + exp;
185

```

```

186      /*
187      * Generate a floating-point number that represents the exponent.
188      * Do this by processing the exponent one bit at a time to combine
189      * many powers of 2 of 10. Then combine the exponent with the
190      * fraction.
191      */
192      if (exp < 0) {
193          expSign = 1;
194          exp = -exp;
195      } else
196          expSign = 0;
197      if (exp > maxExponent)
198          exp = maxExponent;
199      dblExp = 1.0;
200      for (i=0; exp; exp>>=1, ++i)
201          if (exp & 01)
202              dblExp *= powTab[i];
203      if (expSign)
204          fraction /= dblExp;
205      else
206          fraction *= dblExp;
207
208  done:
209      if (endPtr)
210          *endPtr = p;
211
212      return sign ? -fraction : fraction;
213  }
214
215  #define BUFSIZE 10
216  int main()
217  {
218      char buf[BUFSIZE];
219      klee_make_symbolic (buf, sizeof buf, "buf");
220      klee_assume(buf[9]==0);
221
222      strtodx (buf, NULL, 10);
223  };

```

( [https://github.com/dennis714/SAT\\_SMT\\_article/blob/master/KLEE/strtodx.c](https://github.com/dennis714/SAT_SMT_article/blob/master/KLEE/strtodx.c) )

Interestingly, KLEE cannot handle floating-point arithmetic, but nevertheless, found something:

```

...
KLEE: ERROR: /home/klee/klee_test.c:202: memory error: out of bound pointer
...

% ktest-tool klee-last/test003483.ktest
ktest file : 'klee-last/test003483.ktest'
args       : ['klee_test.bc']
num objects: 1
object 0: name: b'buf'
object 0: size: 10
object 0: data: b'-.0E-66\x00\x00\x00'

```

As it seems, string “-.0E-66” makes out of bounds array access (read) at line 202. While further investigation, I’ve found that `powers0f10[]` array is too short: 6th element (started at 0th) has been accessed. And here we see part of array commented (line 79)! Probably someone’s mistake?

## 11.12 Unit testing: simple expression evaluator (calculator)

I has been looking for simple expression evaluator (calculator in other words) which takes expression like “2+2” on input and gives answer. I’ve found one at <http://stackoverflow.com/a/13895198>. Unfortunately, it has no bugs, so I’ve introduced one: a token buffer ( `buf[]` at line 31) is smaller than input buffer ( `input[]` at line 19).

```

1 // copied from http://stackoverflow.com/a/13895198 and reworked
2
3 // Bare bones scanner and parser for the following LL(1) grammar:
4 // expr -> term { [+ -] term } ; An expression is terms separated by add ops.
5 // term -> factor { [* /] factor } ; A term is factors separated by mul ops.

```

```

6 // factor -> unsigned_factor      ; A signed factor is a factor,
7 //      | - unsigned_factor      ; possibly with leading minus sign
8 // unsigned_factor -> ( `expr ` ) ; An unsigned factor is a parenthesized expression
9 //      | NUMBER                  ; or a number
10 //
11 // The parser returns the floating point value of the expression.
12
13 #include <string.h>
14 #include <stdio.h>
15 #include <stdlib.h>
16 #include <stdint.h>
17 #include <stdbool.h>
18
19 char input[128];
20 int input_idx=0;
21
22 char my_getchar()
23 {
24     char rt=input[input_idx];
25     input_idx++;
26     return rt;
27 };
28
29 // The token buffer. We never check for overflow! Do so in production code.
30 // it's deliberately smaller than input[] so KLEE could find buffer overflow
31 char buf[64];
32 int n = 0;
33
34 // The current character.
35 int ch;
36
37 // The look-ahead token. This is the 1 in LL(1).
38 enum { ADD_OP, MUL_OP, LEFT_PAREN, RIGHT_PAREN, NOT_OP, NUMBER, END_INPUT } look_ahead;
39
40 // Forward declarations.
41 void init(void);
42 void advance(void);
43 int expr(void);
44 void error(char *msg);
45
46 // Parse expressions, one per line.
47 int main(void)
48 {
49     // take input expression from input[]
50     //input[0]=0;
51     //strcpy (input, "2+2");
52     klee_make_symbolic(input, sizeof input, "input");
53     input[127]=0;
54
55     init();
56     while (1)
57     {
58         int val = expr();
59         printf("%d\n", val);
60
61         if (look_ahead != END_INPUT)
62             error("junk after expression");
63         advance(); // past end of input mark
64     }
65     return 0;
66 }
67
68 // Just die on any error.
69 void error(char *msg)
70 {
71     fprintf(stderr, "Error: %s. Exiting.\n", msg);
72     exit(1);
73 }
74
75 // Buffer the current character and read a new one.
76 void read()
77 {
78     buf[n++] = ch;
79     buf[n] = '\0'; // Terminate the string.
80     ch = my_getchar();
81 }

```

```

82
83 // Ignore the current character.
84 void ignore()
85 {
86     ch = my_getchar();
87 }
88
89 // Reset the token buffer.
90 void reset()
91 {
92     n = 0;
93     buf[0] = '\0';
94 }
95
96 // The scanner. A tiny deterministic finite automaton.
97 int scan()
98 {
99     reset();
100 START:
101     // first character is digit?
102     if (isdigit (ch))
103         goto DIGITS;
104
105     switch (ch)
106     {
107         case ' ': case '\t': case '\r':
108             ignore();
109             goto START;
110
111         case '-': case '+': case '^':
112             read();
113             return ADD_OP;
114
115         case '~':
116             read();
117             return NOT_OP;
118
119         case '*': case '/': case '%':
120             read();
121             return MUL_OP;
122
123         case '(':
124             read();
125             return LEFT_PAREN;
126
127         case ')':
128             read();
129             return RIGHT_PAREN;
130
131         case 0:
132         case '\n':
133             ch = ' '; // delayed ignore()
134             return END_INPUT;
135
136         default:
137             printf ("bad character: 0x%x\n", ch);
138             exit(0);
139     }
140
141 DIGITS:
142     if (isdigit (ch))
143     {
144         read();
145         goto DIGITS;
146     }
147     else
148         return NUMBER;
149 }
150
151 // To advance is just to replace the look-ahead.
152 void advance()
153 {
154     look_ahead = scan();
155 }
156
157 // Clear the token buffer and read the first look-ahead.

```

```

158 void init()
159 {
160     reset();
161     ignore(); // junk current character
162     advance();
163 }
164
165 int get_number(char *buf)
166 {
167     char *endptr;
168
169     int rt=strtoul (buf, &endptr, 10);
170
171     // is the whole buffer has been processed?
172     if (strlen(buf)!=endptr-buf)
173     {
174         fprintf (stderr, "invalid number: %s\n", buf);
175         exit(0);
176     };
177     return rt;
178 };
179
180 int unsigned_factor()
181 {
182     int rtn = 0;
183     switch (look_ahead)
184     {
185         case NUMBER:
186             rtn=get_number(buf);
187             advance();
188             break;
189
190         case LEFT_PAREN:
191             advance();
192             rtn = expr();
193             if (look_ahead != RIGHT_PAREN) error("missing ')'");
194             advance();
195             break;
196
197         default:
198             printf("unexpected token: %d\n", look_ahead);
199             exit(0);
200     }
201     return rtn;
202 }
203
204 int factor()
205 {
206     int rtn = 0;
207     // If there is a leading minus...
208     if (look_ahead == ADD_OP && buf[0] == '-')
209     {
210         advance();
211         rtn = -unsigned_factor();
212     }
213     else
214         rtn = unsigned_factor();
215
216     return rtn;
217 }
218
219 int term()
220 {
221     int rtn = factor();
222     while (look_ahead == MUL_OP)
223     {
224         switch(buf[0])
225         {
226             case '*':
227                 advance();
228                 rtn *= factor();
229                 break;
230
231             case '/':
232                 advance();
233                 rtn /= factor();

```





Maybe this is not impressive result, nevertheless, it's yet another reminder that division and remainder operations must be wrapped somehow in production code to avoid possible crash.

## 11.13 Regular expressions

I've always wanted to generate possible strings for given regular expression. This is not so hard if to dive into regular expression matcher theory and details, but can we force RE matcher to do this?

I took lightest RE engine I've found: <https://github.com/cesanta/slre>, and wrote this:

```
int main(void)
{
    char s[6];
    klee_make_symbolic(s, sizeof s, "s");
    s[5]=0;
    if (slre_match("^\\d[a-c]+(x|y|z)", s, 5, NULL, 0, 0)==5)
        klee_assert(0);
}
```

So I wanted a string consisting of digit, "a" or "b" or "c" (at least one character) and "x" or "y" or "z" (one character). The whole string must have size of 5 characters.

```
% klee --libc=uclibc slre.bc
...
KLEE: ERROR: /home/klee/slre.c:445: failed external call: klee_assert
KLEE: NOTE: now ignoring this error at this location
...

% ls klee-last | grep err
test000014.external.err

% ktest-tool --write-ints klee-last/test000014.ktest
ktest file : 'klee-last/test000014.ktest'
args       : ['slre.bc']
num objects: 1
object  0: name: b's'
object  0: size: 6
object  0: data: b'5aaax\xff'
```

This is indeed correct string. "\xff" is at the place where terminal zero byte should be, but RE engine we use ignores the last zero byte, because it has buffer length as a passed parameter. Hence, KLEE doesn't *reconstruct* final byte.

Can we get more? Now we add additional constraint:

```
int main(void)
{
    char s[6];
    klee_make_symbolic(s, sizeof s, "s");
    s[5]=0;
    if (slre_match("^\\d[a-c]+(x|y|z)", s, 5, NULL, 0, 0)==5 &&
        strcmp(s, "5aaax")!=0)
        klee_assert(0);
}
```

```
% ktest-tool --write-ints klee-last/test000014.ktest
ktest file : 'klee-last/test000014.ktest'
args       : ['slre.bc']
num objects: 1
object  0: name: b's'
object  0: size: 6
object  0: data: b'7aaax\xff'
```

Let's say, out of whim, we don't like "a" at the 2nd position (starting at 0th):

```
int main(void)
{
    char s[6];
    klee_make_symbolic(s, sizeof s, "s");
    s[5]=0;
    if (slre_match("^\\d[a-c]+(x|y|z)", s, 5, NULL, 0, 0)==5 &&
        strcmp(s, "5aaax")!=0 &&
        s[2]!='a')
        klee_assert(0);
}
```

KLEE found a way to satisfy our new constraint:

```
% ktest-tool -write-ints klee-last/test000014.ktest
ktest file : 'klee-last/test000014.ktest'
args       : ['slre.bc']
num objects: 1
object 0: name: b's'
object 0: size: 6
object 0: data: b'7abax\xff'
```

Let's also define constraint KLEE cannot satisfy:

```
int main(void)
{
    char s[6];
    klee_make_symbolic(s, sizeof s, "s");
    s[5]=0;
    if (slre_match("^\\d[a-c]+(x|y|z)", s, 5, NULL, 0, 0)==5 &&
        strcmp(s, "5aaax")!=0 &&
        s[2]!='a' &&
        s[2]!='b' &&
        s[2]!='c')
        klee_assert(0);
}
```

It cannot indeed, and KLEE finished without reporting about `klee_assert()` triggering.

## 11.14 Exercise

Here is my crackme/keygenme, which may be tricky, but easy to solve using KLEE: <http://challenges.re/74/>.

# 12 (Amateur) cryptography

## 12.1 Serious cryptography

Let's back to the method we previously used (10.2) to construct expressions using running Python function. We can try to build expression for the output of XXTEA encryption algorithm:

```
#!/usr/bin/env python

class Expr:
    def __init__(self,s):
        self.s=s

    def __str__(self):
        return self.s

    def convert_to_Expr_if_int(self, n):
        if isinstance(n, int):
            return Expr(str(n))
        if isinstance(n, Expr):
            return n
        raise AssertionError # unsupported type

    def __xor__(self, other):
        return Expr("(" + self.s + "^" + self.convert_to_Expr_if_int(other).s + ")")

    def __mul__(self, other):
        return Expr("(" + self.s + "*" + self.convert_to_Expr_if_int(other).s + ")")

    def __add__(self, other):
        return Expr("(" + self.s + "+" + self.convert_to_Expr_if_int(other).s + ")")

    def __and__(self, other):
        return Expr("(" + self.s + "&" + self.convert_to_Expr_if_int(other).s + ")")

    def __lshift__(self, other):
        return Expr("(" + self.s + "<<" + self.convert_to_Expr_if_int(other).s + ")")

    def __rshift__(self, other):
```

```

        return Expr("(" + self.s + ">>" + self.convert_to_Expr_if_int(other).s + ")")

    def __getitem__(self, d):
        return Expr("(" + self.s + "[" + d.s + "]")")

# reworked from:

# Pure Python (2.x) implementation of the XXTEA cipher
# (c) 2009. Ivan Voras <ivoras@gmail.com>
# Released under the BSD License.

def raw_xxtea(v, n, k):

    def MX():
        return ((z>>5)^(y<<2)) + ((y>>3)^(z<<4))^(sum^y) + (k[(Expr(str(p)) & 3)^e]^z)

    y = v[0]
    sum = Expr("0")
    DELTA = 0x9e3779b9
    # Encoding only
    z = v[n-1]

    # number of rounds:
    #q = 6 + 52 / n
    q=1

    while q > 0:
        q -= 1
        sum = sum + DELTA
        e = (sum >> 2) & 3
        p = 0
        while p < n - 1:
            y = v[p+1]
            z = v[p] = v[p] + MX()
            p += 1
        y = v[0]
        z = v[n-1] = v[n-1] + MX()
    return 0

v=[Expr("input1"), Expr("input2"), Expr("input3"), Expr("input4")]
k=Expr("key")

raw_xxtea(v, 4, k)

for i in range(4):
    print i, ":", v[i]
#print len(str(v[0]))+len(str(v[1]))+len(str(v[2]))+len(str(v[3]))

```

A key is chosen according to input data, and, obviously, we can't know it during symbolic execution, so we leave expression like `k[...]`.

Now results for just one round, for each of 4 outputs:

```

0 : (input1+((((input4>>5)^(input2<<2))+((input2>>3)^(input4<<4)))^(((0+2654435769)^input2)+
((key[((0&3)^((0+2654435769)>>2)&3]))^input4))))

1 : (input2+((((input1+((((input4>>5)^(input2<<2))+((input2>>3)^(input4<<4)))^(((0+2654435769)^
input2)+((key[((0&3)^((0+2654435769)>>2)&3]))^input4))))>>5)^(input3<<2))+((input3>>3)^(input1+
((((input4>>5)^(input2<<2))+((input2>>3)^(input4<<4)))^(((0+2654435769)^input2)+((key[((0&3)^
(((0+2654435769)>>2)&3]))^input4))))<<4)))^(((0+2654435769)^input3)+((key[((1&3)^((0+2654435769)>>2)&
3]))^input1+((((input4>>5)^(input2<<2))+((input2>>3)^(input4<<4)))^(((0+2654435769)^input2)+
((key[((0&3)^((0+2654435769)>>2)&3]))^input4))))))

2 : (input3+((((input2+((((input1+((((input4>>5)^(input2<<2))+((input2>>3)^(input4<<4)))^
(((0+2654435769)^input2)+((key[((0&3)^((0+2654435769)>>2)&3]))^input4))))>>5)^(input3<<2))+
((input3>>3)^(input1+((((input4>>5)^(input2<<2))+((input2>>3)^(input4<<4)))^(((0+2654435769)^
input2)+((key[((0&3)^((0+2654435769)>>2)&3]))^input4))))<<4)))^(((0+2654435769)^input3)+
((key[((1&3)^((0+2654435769)>>2)&3]))^input1+((((input4>>5)^(input2<<2))+((input2>>3)^(input4<<4)))
^(((0+2654435769)^input2)+((key[((0&3)^((0+2654435769)>>2)&3]))^input4))))>>5)^(input4<<2))+
((input4>>3)^(input2+((((input1+((((input4>>5)^(input2<<2))+((input2>>3)^(input4<<4)))^
(((0+2654435769)^input2)+((key[((0&3)^((0+2654435769)>>2)&3]))^input4))))>>5)^(input3<<2))+
((input3>>3)^(input1+
((((input4>>5)^(input2<<2))+((input2>>3)^(input4<<4)))^(((0+2654435769)^input2)+((key[((0&3)^
(((0+2654435769)>>2)&3]))^input4))))<<4)))^(((0+2654435769)^input3)+((key[((1&3)^((0+2654435769)>>2)&3]))
^input1+((((input4>>5)^(input2<<2))+((input2>>3)^(input4<<4)))^(((0+2654435769)^input2)+((key[((0&3)^
(((0+2654435769)>>2)&3]))^input4))))<<4)))^(((0+2654435769)^input4)+((key[((2&3)^((0+2654435769)>>2)&
3]))^input2+((((input1+((((input4>>5)^(input2<<2))+((input2>>3)^(input4<<4)))^(((0+2654435769)^

```



In other words, theoretically, you can build system of equation like this:  $MD5(x) = 12341234\dots$ , but expressions are so huge so it's impossible to solve them. Yes, cryptographers are fully aware of this and one of the goals of the successful cipher is to make expressions as big as possible, using reasonable time and size of algorithm.

Nevertheless, you can find numerous papers about breaking these cryptosystems with reduced number of rounds: when expression isn't *exploded* yet, sometimes it's possible. This cannot be applied in practice, but such experience has some interesting theoretical results.

### 12.1.1 Attempts to break “serious” crypto

CryptoMiniSat itself exist to support XOR operation, which is ubiquitous in cryptography.

- Bitcoin mining with SAT solver: <http://jheusser.github.io/2013/02/03/satcoin.html>, <https://github.com/msoos/sha256-sat-bitcoin>.
- Alexander Semenov, attempts to break A5/1, etc. (Russian presentation)
- Vegard Nossrum - SAT-based preimage attacks on SHA-1
- Algebraic Attacks on the Crypto-1 Stream Cipher in MiFare Classic and Oyster Cards
- Attacking Bivium Using SAT Solvers
- Extending SAT Solvers to Cryptographic Problems
- Applications of SAT Solvers to Cryptanalysis of Hash Functions
- Algebraic-Differential Cryptanalysis of DES

## 12.2 Amateur cryptography

This is what you can find in serial numbers, license keys, executable file packers, CTF<sup>94</sup>, malware, etc. Sometimes even ransomware (but rarely nowadays, in 2017).

Amateur cryptography is often can be broken using SMT solver, or even KLEE.

Amateur cryptography is usually based not on theory, but on visual complexity: if its creator getting results which are seems chaotic enough, often, one stops to improve it further. This is security not even on obscurity, but on chaotic mess. This is also sometimes called “The Fallacy of Complex Manipulation” (see RFC4086).

Devising your own cryptoalgorithm is a very tricky thing to do. This can be compared to devising your own PRNG. Even famous Donald Knuth in 1959 constructed one, and it was visually very complex, but, as it turns out in practice, it has very short cycle of length 3178. [See also: The Art of Computer Programming vol.II page 4, (1997).]

The very first problem is that making an algorithm which can generate very long expressions is tricky thing itself. Common error is to use operations like XOR and rotations/permutations, which can't help much. Even worse: some people think that XORing a value several times can be better, like:  $(x \oplus 1234) \oplus 5678$ . Obviously, these two XOR operations (or more precisely, any number of it) can be reduced to a single one. Same story about applied operations like addition and subtraction—they all also can be reduced to single one.

Real cryptoalgorithms, like IDEA, can use several operations from different groups, like XOR, addition and multiplication. Applying them all in specific order will make resulting expression irreducible.

When I prepared this part, I tried to make an example of such amateur hash function:

```
// copied from http://blog.regehr.org/archives/1063
uint32_t rotr32b (uint32_t x, uint32_t n)
{
    assert (n<32);
    if (!n) return x;
    return (x<<n) | (x>>(32-n));
}

uint32_t rotr32b (uint32_t x, uint32_t n)
{
    assert (n<32);
    if (!n) return x;
    return (x>>n) | (x<<(32-n));
}
```

<sup>94</sup>Capture the Flag

```

void megahash (uint32_t buf[4])
{
    for (int i=0; i<4; i++)
    {
        uint32_t t0=buf[0]^0x12345678^buf[1];
        uint32_t t1=buf[1]^0xabcdef01^buf[2];
        uint32_t t2=buf[2]^0x23456789^buf[3];
        uint32_t t3=buf[3]^0x0abcdef0^buf[0];

        buf[0]=rotl32b(t0, 1);
        buf[1]=rotr32b(t1, 2);
        buf[2]=rotl32b(t2, 3);
        buf[3]=rotr32b(t3, 4);
    };
};

int main()
{
    uint32_t buf[4];
    klee_make_symbolic(buf, sizeof buf);
    megahash (buf);
    if (buf[0]==0x18f71ce6          // or whatever
        && buf[1]==0xf37c2fc9
        && buf[2]==0x1cfe96fe
        && buf[3]==0x8c02c75e)
        klee_assert(0);
};

```

KLEE can break it with little effort. Functions of such complexity is common in shareware, which checks license keys, etc.

Here is how we can make its work harder by making rotations dependent of inputs, and this makes number of possible inputs much, much bigger:

```

void megahash (uint32_t buf[4])
{
    for (int i=0; i<16; i++)
    {
        uint32_t t0=buf[0]^0x12345678^buf[1];
        uint32_t t1=buf[1]^0xabcdef01^buf[2];
        uint32_t t2=buf[2]^0x23456789^buf[3];
        uint32_t t3=buf[3]^0x0abcdef0^buf[0];

        buf[0]=rotl32b(t0, t1&0x1F);
        buf[1]=rotr32b(t1, t2&0x1F);
        buf[2]=rotl32b(t2, t3&0x1F);
        buf[3]=rotr32b(t3, t0&0x1F);
    };
};

```

Addition (or [modular addition](#), as cryptographers say) can make thing even harder:

```

void megahash (uint32_t buf[4])
{
    for (int i=0; i<4; i++)
    {
        uint32_t t0=buf[0]^0x12345678^buf[1];
        uint32_t t1=buf[1]^0xabcdef01^buf[2];
        uint32_t t2=buf[2]^0x23456789^buf[3];
        uint32_t t3=buf[3]^0x0abcdef0^buf[0];

        buf[0]=rotl32b(t0, t2&0x1F)+t1;
        buf[1]=rotr32b(t1, t3&0x1F)+t2;
        buf[2]=rotl32b(t2, t1&0x1F)+t3;
        buf[3]=rotr32b(t3, t2&0x1F)+t4;
    };
};

```

As an exercise, you can try to make a block cipher which KLEE wouldn't break. This is quite sobering experience. But even if you can, this is not a panacea, there is an array of other cryptoanalytical methods to break it.

Summary: if you deal with amateur cryptography, you can always give KLEE and SMT solver a try. Even more: sometimes you have only decryption function, and if algorithm is simple enough, KLEE or SMT solver can reverse things back.

One funny thing to mention: if you try to implement amateur cryptoalgorithm in Verilog/VHDL language to run it on [FPGA<sup>95</sup>](#), maybe in brute-force way, you can find that [EDA<sup>96</sup>](#) tools can optimize many things during synthesis (this is the word they use for “compilation”) and can leave cryptoalgorithm much smaller/simpler than it was. Even if you try to implement DES algorithm *in bare metal* with a fixed key, Altera Quartus can optimize first round of it and make it smaller than others.

### 12.2.1 Bugs

Another prominent feature of amateur cryptography is bugs. Bugs here often left uncaught because output of encrypting function visually looked “good enough” or “obfuscated enough”, so a developer stopped to work on it.

This is especially feature of hash functions, because when you work on block cipher, you have to do two functions (encryption/decryption), while hashing function is single.

Weirdest ever amateur encryption algorithm I once saw, encrypted only odd bytes of input block, while even bytes left untouched, so the input plain text has been partially seen in the resulting encrypted block. It was encryption routine used in license key validation. Hard to believe someone did this on purpose. Most likely, it was just an unnoticed bug.

### 12.2.2 XOR ciphers

Simplest possible amateur cryptography is just application of XOR operation using some kind of table. Maybe even simpler. This is a real algorithm I once saw:

```
for (i=0; i<size; i++)
    buf[i]=buf[i]^(31*(i+1));
```

This is not even encryption, rather concealing or hiding.

Some other examples of simple XOR-cipher cryptanalysis are present in the “Reverse Engineering for Beginners” [97](#) book.

### 12.2.3 Other features

**Tables** There are often table(s) with pseudorandom data, which is/are used chaotically.

**Checksumming** End-users can have proclivity to changing license codes, serial numbers, etc., with a hope this could affect behaviour of software. So there is often some kind of checksum: starting at simple summing and [CRC](#). This is close to [MAC<sup>98</sup>](#) in real cryptography.

**Entropy level** Maybe (much) lower, despite the fact that data looks random.

### 12.2.4 Examples

- A popular FLEXlm license manager was based on a simple amateur cryptoalgorithm (before they switched to [ECC<sup>99</sup>](#)), which can be cracked easily.
- Pegasus Mail Password Decoder: <http://phrack.org/issues/52/3.html> - a very typical example.
- You can find a lot of blog posts about breaking CTF-level crypto using Z3, etc. Here is one of them: <http://doar-e.github.io/blog/2015/08/18/keygenning-with-klée/>.
- Another: [Automated algebraic cryptanalysis with OpenREIL and Z3](#). By the way, this solution tracks state of each register at each EIP/RIP, this is almost the same as [SSA](#), which is heavily used in compilers and worth learning.
- Many examples of amateur cryptography I’ve taken from an old Fravia website: [https://yurichev.com/mirrors/amateur\\_crypto\\_examples\\_from\\_Fravia/](https://yurichev.com/mirrors/amateur_crypto_examples_from_Fravia/).

<sup>95</sup>Field-programmable gate array

<sup>96</sup>Electronic design automation

<sup>97</sup><http://beginners.re>

<sup>98</sup>Message authentication code

<sup>99</sup>Elliptic curve cryptography



## 12.3 Case study: simple hash function

(This piece of text was initially added to my “Reverse Engineering for Beginners” book ([beginners.re](http://beginners.re)) at March 2014)<sup>100</sup>.

Here is one-way hash function, that converted a 64-bit value to another and we need to try to reverse its flow back.

### 12.3.1 Manual decompiling

Here its assembly language listing in IDA:

```
sub_401510    proc near
              ; ECX = input
              mov     rdx, 5D7E0D1F2E0F1F84h
              mov     rax, rcx           ; input
              imul    rax, rdx
              mov     rdx, 388D76AEE8CB1500h
              mov     ecx, eax
              and     ecx, 0Fh
              ror     rax, cl
              xor     rax, rdx
              mov     rdx, 0D2E9EE7E83C4285Bh
              mov     ecx, eax
              and     ecx, 0Fh
              rol     rax, cl
              lea     r8, [rax+rdx]
              mov     rdx, 8888888888888889h
              mov     rax, r8
              mul     rdx
              shr     rdx, 5
              mov     rax, rdx
              lea     rcx, [r8+rdx*4]
              shl     rax, 6
              sub     rcx, rax
              mov     rax, r8
              rol     rax, cl
              ; EAX = output
              retn
sub_401510    endp
```

The example was compiled by GCC, so the first argument is passed in ECX.

If you don't have Hex-Rays, or if you distrust to it, you can try to reverse this code manually. One method is to represent the CPU registers as local C variables and replace each instruction by a one-line equivalent expression, like:

```
uint64_t f(uint64_t input)
{
    uint64_t rax, rbx, rcx, rdx, r8;

    ecx=input;

    rdx=0x5D7E0D1F2E0F1F84;
    rax=rcx;
    rax*=rdx;
    rdx=0x388D76AEE8CB1500;
    rax=_rotr(rax, rax&0xF); // rotate right
    rax^=rdx;
    rdx=0xD2E9EE7E83C4285B;
    rax=_rotl(rax, rax&0xF); // rotate left
    r8=rax+rdx;
    rdx=0x8888888888888889;
    rax=r8;
    rax*=rdx;
    rdx=rdx>>5;
    rax=rdx;
    rcx=r8+rdx*4;
    rax=rax<<6;
    rcx=rcx-rax;
    rax=r8
    rax=_rotl (rax, rcx&0xFF); // rotate left
```

<sup>100</sup>This example was also used by Murphy Berzish in his lecture about SAT and SMT: <http://mirror.csclub.uwaterloo.ca/csclub/mtrberzi-sat-smt-slides.pdf>, <http://mirror.csclub.uwaterloo.ca/csclub/mtrberzi-sat-smt.mp4>

```
};  
    return rax;
```

If you are careful enough, this code can be compiled and will even work in the same way as the original. Then, we are going to rewrite it gradually, keeping in mind all registers usage. Attention and focus is very important here—any tiny typo may ruin all your work!  
Here is the first step:

```
uint64_t f(uint64_t input)  
{  
    uint64_t rax, rbx, rcx, rdx, r8;  
  
    ecx=input;  
  
    rdx=0x5D7E0D1F2E0F1F84;  
    rax=rcx;  
    rax*=rdx;  
    rdx=0x388D76AEE8CB1500;  
    rax=_rotr(rax, rax&0xF); // rotate right  
    rax^=rdx;  
    rdx=0xD2E9EE7E83C4285B;  
    rax=_lrotl(rax, rax&0xF); // rotate left  
    r8=rax+rdx;  
  
    rdx=0x8888888888888889;  
    rax=r8;  
    rax*=rdx;  
    // RDX here is a high part of multiplication result  
    rdx=rdx>>5;  
    // RDX here is division result!  
    rax=rdx;  
  
    rcx=r8+rdx*4;  
    rax=rax<<6;  
    rcx=rcx-rax;  
    rax=r8  
    rax=_lrotl (rax, rcx&0xFF); // rotate left  
    return rax;  
};
```

Next step:

```
uint64_t f(uint64_t input)  
{  
    uint64_t rax, rbx, rcx, rdx, r8;  
  
    ecx=input;  
  
    rdx=0x5D7E0D1F2E0F1F84;  
    rax=rcx;  
    rax*=rdx;  
    rdx=0x388D76AEE8CB1500;  
    rax=_rotr(rax, rax&0xF); // rotate right  
    rax^=rdx;  
    rdx=0xD2E9EE7E83C4285B;  
    rax=_lrotl(rax, rax&0xF); // rotate left  
    r8=rax+rdx;  
  
    rdx=0x8888888888888889;  
    rax=r8;  
    rax*=rdx;  
    // RDX here is a high part of multiplication result  
    rdx=rdx>>5;  
    // RDX here is division result!  
    rax=rdx;  
  
    rcx=(r8+rdx*4)-(rax<<6);  
    rax=r8  
    rax=_lrotl (rax, rcx&0xFF); // rotate left  
    return rax;  
};
```

We can spot the division using multiplication. Indeed, let's calculate the divider in Wolfram Mathematica:

Listing 1: Wolfram Mathematica

```
In[1]:=N[2^(64 + 5)/16^8888888888888889]
Out[1]:=60.
```

We get this:

```
uint64_t f(uint64_t input)
{
    uint64_t rax, rbx, rcx, rdx, r8;

    ecx=input;

    rdx=0x5D7E0D1F2E0F1F84;
    rax=rcx;
    rax*=rdx;
    rdx=0x388D76AEE8CB1500;
    rax=_rotr(rax, rax&0xF); // rotate right
    rax^=rdx;
    rdx=0xD2E9EE7E83C4285B;
    rax=_lrotl(rax, rax&0xF); // rotate left
    r8=rax+rdx;

    rax=rdx=r8/60;

    rcx=(r8+rax*4)-(rax*64);
    rax=r8
    rax=_lrotl (rax, rcx&0xFF); // rotate left
    return rax;
};
```

One more step:

```
uint64_t f(uint64_t input)
{
    uint64_t rax, rbx, rcx, rdx, r8;

    rax=input;
    rax*=0x5D7E0D1F2E0F1F84;
    rax=_rotr(rax, rax&0xF); // rotate right
    rax^=0x388D76AEE8CB1500;
    rax=_lrotl(rax, rax&0xF); // rotate left
    r8=rax+0xD2E9EE7E83C4285B;

    rcx=r8-(r8/60)*60;
    rax=r8
    rax=_lrotl (rax, rcx&0xFF); // rotate left
    return rax;
};
```

By simple reducing, we finally see that it's calculating the remainder, not the quotient:

```
uint64_t f(uint64_t input)
{
    uint64_t rax, rbx, rcx, rdx, r8;

    rax=input;
    rax*=0x5D7E0D1F2E0F1F84;
    rax=_rotr(rax, rax&0xF); // rotate right
    rax^=0x388D76AEE8CB1500;
    rax=_lrotl(rax, rax&0xF); // rotate left
    r8=rax+0xD2E9EE7E83C4285B;

    return _lrotl (r8, r8 % 60); // rotate left
};
```

We end up with this fancy formatted source-code:

```
#include <stdio.h>
#include <stdint.h>
#include <stdbool.h>
#include <string.h>
#include <intrin.h>

#define C1 0x5D7E0D1F2E0F1F84
#define C2 0x388D76AEE8CB1500
#define C3 0xD2E9EE7E83C4285B
```

```

uint64_t hash(uint64_t v)
{
    v*=C1;
    v=_lrotr(v, v&0xF); // rotate right
    v^=C2;
    v=_lrotl(v, v&0xF); // rotate left
    v+=C3;
    v=_lrotl(v, v % 60); // rotate left
    return v;
};

int main()
{
    printf ("%llu\n", hash(...));
};

```

Since we are not cryptoanalysts we can't find an easy way to generate the input value for some specific output value. The rotate instruction's coefficients look frightening—it's a warranty that the function is not bijective, it is rather surjective, it has collisions, or, speaking more simply, many inputs may be possible for one output.

Brute-force is not solution because values are 64-bit ones, that's beyond reality.

### 12.3.2 Now let's use the Z3

Still, without any special cryptographic knowledge, we may try to break this algorithm using Z3.

Here is the Python source code:

```

1  #!/usr/bin/env python
2
3  from z3 import *
4
5  C1=0x5D7E0D1F2E0F1F84
6  C2=0x388D76AEE8CB1500
7  C3=0xD2E9EE7E83C4285B
8
9  inp, i1, i2, i3, i4, i5, i6, outp = BitVecs('inp i1 i2 i3 i4 i5 i6 outp', 64)
10
11  s = Solver()
12  s.add(i1==inp*C1)
13  s.add(i2==RotateRight (i1, i1 & 0xF))
14  s.add(i3==i2 ^ C2)
15  s.add(i4==RotateLeft(i3, i3 & 0xF))
16  s.add(i5==i4 + C3)
17  s.add(outp==RotateLeft (i5, URem(i5, 60)))
18
19  s.add(outp==10816636949158156260)
20
21  print s.check()
22  m=s.model()
23  print m
24  print (" inp=0x%X" % m[inp].as_long())
25  print (" outp=0x%X" % m[outp].as_long())

```

This is going to be our first solver.

We see the variable definitions on line 7. These are just 64-bit variables. `i1..i6` are intermediate variables, representing the values in the registers between instruction executions.

Then we add the so-called constraints on lines 10..15. The last constraint at 17 is the most important one: we are going to try to find an input value for which our algorithm will produce 10816636949158156260.

*RotateRight*, *RotateLeft*, *URem*—are functions from the Z3 Python API, not related to Python language.

Then we run it:

```

...>python.exe 1.py
sat
[i1 = 3959740824832824396,
 i3 = 8957124831728646493,
 i5 = 10816636949158156260,
 inp = 1364123924608584563,
 outp = 10816636949158156260,
 i4 = 14065440378185297801,
 i2 = 4954926323707358301]
inp=0x12EE577B63E80B73
outp=0x961C69FF0AEFD7E4

```

“sat” mean “satisfiable”, i.e., the solver was able to find at least one solution. The solution is printed in the square brackets. The last two lines are the input/output pair in hexadecimal form. Yes, indeed, if we run our function with `0x12EE577B63E80B73` as input, the algorithm will produce the value we were looking for.

But, as we noticed before, the function we work with is not bijective, so there may be other correct input values. The Z3 is not capable of producing more than one result, but let’s hack our example slightly, by adding line 19, which implies “look for any other results than this”:

```

1  #!/usr/bin/env python
2
3  from z3 import *
4
5  C1=0x5D7E0D1F2E0F1F84
6  C2=0x388D76AEE8CB1500
7  C3=0xD2E9EE7E83C4285B
8
9  inp, i1, i2, i3, i4, i5, i6, outp = BitVecs('inp i1 i2 i3 i4 i5 i6 outp', 64)
10
11 s = Solver()
12 s.add(i1==inp*C1)
13 s.add(i2==RotateRight (i1, i1 & 0xF))
14 s.add(i3==i2 ^ C2)
15 s.add(i4==RotateLeft(i3, i3 & 0xF))
16 s.add(i5==i4 + C3)
17 s.add(outp==RotateLeft (i5, URem(i5, 60)))
18
19 s.add(outp==10816636949158156260)
20
21 s.add(inp!=0x12EE577B63E80B73)
22
23 print s.check()
24 m=s.model()
25 print m
26 print (" inp=0x%X" % m[inp].as_long())
27 print ("outp=0x%X" % m[outp].as_long())

```

Indeed, it finds another correct result:

```

...>python.exe 2.py
sat
[i1 = 3959740824832824396,
 i3 = 8957124831728646493,
 i5 = 10816636949158156260,
 inp = 10587495961463360371,
 outp = 10816636949158156260,
 i4 = 14065440378185297801,
 i2 = 4954926323707358301]
inp=0x92EE577B63E80B73
outp=0x961C69FF0AEFD7E4

```

This can be automated. Each found result can be added as a constraint and then the next result will be searched for. Here is a slightly more sophisticated example:

```

1  #!/usr/bin/env python
2
3  from z3 import *
4
5  C1=0x5D7E0D1F2E0F1F84
6  C2=0x388D76AEE8CB1500
7  C3=0xD2E9EE7E83C4285B
8
9  inp, i1, i2, i3, i4, i5, i6, outp = BitVecs('inp i1 i2 i3 i4 i5 i6 outp', 64)
10
11 s = Solver()
12 s.add(i1==inp*C1)
13 s.add(i2==RotateRight (i1, i1 & 0xF))
14 s.add(i3==i2 ^ C2)
15 s.add(i4==RotateLeft(i3, i3 & 0xF))
16 s.add(i5==i4 + C3)
17 s.add(outp==RotateLeft (i5, URem(i5, 60)))
18
19 s.add(outp==10816636949158156260)
20
21 # cypasted from http://stackoverflow.com/questions/11867611/z3py-checking-all-solutions-for-equation
22 result=[]

```

```

23 while True:
24     if s.check() == sat:
25         m = s.model()
26         print m[inp]
27         result.append(m)
28         # Create a new constraint the blocks the current model
29         block = []
30         for d in m:
31             # d is a declaration
32             if d.arity() > 0:
33                 raise Z3Exception("uninterpreted functions are not supported")
34             # create a constant from declaration
35             c=d()
36             if is_array(c) or c.sort().kind() == Z3_UNINTERPRETED_SORT:
37                 raise Z3Exception("arrays and uninterpreted sorts are not supported")
38             block.append(c != m[d])
39         s.add(Or(block))
40     else:
41         print "results total=",len(result)
42         break

```

We got:

```

1364123924608584563
1234567890
9223372038089343698
4611686019661955794
13835058056516731602
3096040143925676201
12319412180780452009
7707726162353064105
16931098199207839913
1906652839273745429
11130024876128521237
15741710894555909141
6518338857701133333
5975809943035972467
15199181979890748275
10587495961463360371
results total= 16

```

So there are 16 correct input values for `0x92EE577B63E80B73` as a result.

The second is 1234567890—it is indeed the value which was used by me originally while preparing this example.

Let's also try to research our algorithm a bit more. Acting on a sadistic whim, let's find if there are any possible input/output pairs in which the lower 32-bit parts are equal to each other?

Let's remove the *outp* constraint and add another, at line 17:

```

1  #!/usr/bin/env python
2
3  from z3 import *
4
5  C1=0x5D7E0D1F2E0F1F84
6  C2=0x388D76AEE8CB1500
7  C3=0xD2E9EE7E83C4285B
8
9  inp, i1, i2, i3, i4, i5, i6, outp = BitVecs('inp i1 i2 i3 i4 i5 i6 outp', 64)
10
11  s = Solver()
12  s.add(i1==inp*C1)
13  s.add(i2==RotateRight (i1, i1 & 0xF))
14  s.add(i3==i2 ^ C2)
15  s.add(i4==RotateLeft(i3, i3 & 0xF))
16  s.add(i5==i4 + C3)
17  s.add(outp==RotateLeft (i5, URem(i5, 60)))
18
19  s.add(outp & 0xFFFFFFFF == inp & 0xFFFFFFFF)
20
21  print s.check()
22  m=s.model()
23  print m
24  print (" inp=0x%X" % m[inp].as_long())
25  print ("outp=0x%X" % m[outp].as_long())

```

It is indeed so:

```
sat
[i1 = 14869545517796235860,
 i3 = 8388171335828825253,
 i5 = 6918262285561543945,
 inp = 1370377541658871093,
 outp = 14543180351754208565,
 i4 = 10167065714588685486,
 i2 = 5541032613289652645]
inp=0x13048F1D12C00535
outp=0xC9D3C17A12C00535
```

Let's be more sadistic and add another constraint: last 16 bits must be 0x1234 :

```
1 #!/usr/bin/env python
2
3 from z3 import *
4
5 C1=0x5D7E0D1F2E0F1F84
6 C2=0x388D76AEE8CB1500
7 C3=0xD2E9EE7E83C4285B
8
9 inp, i1, i2, i3, i4, i5, i6, outp = BitVecs('inp i1 i2 i3 i4 i5 i6 outp', 64)
10
11 s = Solver()
12 s.add(i1==inp*C1)
13 s.add(i2==RotateRight (i1, i1 & 0xF))
14 s.add(i3==i2 ^ C2)
15 s.add(i4==RotateLeft(i3, i3 & 0xF))
16 s.add(i5==i4 + C3)
17 s.add(outp==RotateLeft (i5, URem(i5, 60)))
18
19 s.add(outp & 0xFFFFFFFF == inp & 0xFFFFFFFF)
20 s.add(outp & 0xFFFF == 0x1234)
21
22 print s.check()
23 m=s.model()
24 print m
25 print (" inp=0x%X" % m[inp].as_long())
26 print ("outp=0x%X" % m[outp].as_long())
```

Oh yes, this possible as well:

```
sat
[i1 = 2834222860503985872,
 i3 = 2294680776671411152,
 i5 = 17492621421353821227,
 inp = 461881484695179828,
 outp = 419247225543463476,
 i4 = 2294680776671411152,
 i2 = 2834222860503985872]
inp=0x668EEC35F961234
outp=0x5D177215F961234
```

Z3 works very fast and it implies that the algorithm is weak, it is not cryptographic at all (like the most of the amateur cryptography).

## 13 SAT-solvers

SMT vs. SAT is like high level PL vs. assembly language. The latter can be much more efficient, but it's hard to program in it.

SAT is abbreviation of "Boolean satisfiability problem". The problem is to find such a set of variables, which, if plugged into boolean expression, will result in "true".

### 13.1 CNF form

CNF<sup>101</sup> is a *normal form*.

Any boolean expression can be converted to *normal form* and CNF is one of them. The CNF expression is a bunch of clauses (sub-expressions) consisting of terms (variables), ORs and NOTs, all of which are then

<sup>101</sup>[https://en.wikipedia.org/wiki/Conjunctive\\_normal\\_form](https://en.wikipedia.org/wiki/Conjunctive_normal_form)

glued together with AND into a full expression. There is a way to memorize it: **CNF** is “AND of ORs” (or “product of sums”) and **DNF**<sup>102</sup> is “OR of ANDs” (or “sum of products”).

Example is:  $(\neg A \vee B) \wedge (C \vee \neg D)$ .

$\vee$  stands for OR (logical disjunction<sup>103</sup>), “+” sign is also sometimes used for OR.

$\wedge$  stands for AND (logical conjunction<sup>104</sup>). It is easy to memorize:  $\wedge$  looks like “A” letter. “.” is also sometimes used for AND.

$\neg$  is negation (NOT).

## 13.2 Example: 2-bit adder

**SAT**-solver is merely a solver of huge boolean equations in CNF form. It just gives the answer, if there is a set of input values which can satisfy CNF expression, and what input values must be.

Here is a 2-bit adder for example:




Figure 16: 2-bit adder circuit

The adder in its simplest form: it has no carry-in and carry-out, and it has 3 XOR gates and one AND gate. Let's try to figure out, which sets of input values will force adder to set both two output bits? By doing quick memory calculation, we can see that there are 4 ways to do so:  $0 + 3 = 3$ ,  $1 + 2 = 3$ ,  $2 + 1 = 3$ ,  $3 + 0 = 3$ . Here is also truth table, with these rows highlighted:

<sup>102</sup>Disjunctive normal form

<sup>103</sup>[https://en.wikipedia.org/wiki/Logical\\_disjunction](https://en.wikipedia.org/wiki/Logical_disjunction)

<sup>104</sup>[https://en.wikipedia.org/wiki/Logical\\_conjunction](https://en.wikipedia.org/wiki/Logical_conjunction)



	aH	aL	bH	bL	qH	qL
$3+3 = 6 \equiv 2 \pmod{4}$	1	1	1	1	1	0
$3+2 = 5 \equiv 1 \pmod{4}$	1	1	1	0	0	1
$3+1 = 4 \equiv 0 \pmod{4}$	1	1	0	1	0	0
$3+0 = 3 \equiv 3 \pmod{4}$	1	1	0	0	1	1
$2+3 = 5 \equiv 1 \pmod{4}$	1	0	1	1	0	1
$2+2 = 4 \equiv 0 \pmod{4}$	1	0	1	0	0	0
$2+1 = 3 \equiv 3 \pmod{4}$	1	0	0	1	1	1
$2+0 = 2 \equiv 2 \pmod{4}$	1	0	0	0	1	0
$1+3 = 4 \equiv 0 \pmod{4}$	0	1	1	1	0	0
$1+2 = 3 \equiv 3 \pmod{4}$	0	1	1	0	1	1
$1+1 = 2 \equiv 2 \pmod{4}$	0	1	0	1	1	0
$1+0 = 1 \equiv 1 \pmod{4}$	0	1	0	0	0	1
$0+3 = 3 \equiv 3 \pmod{4}$	0	0	1	1	1	1
$0+2 = 2 \equiv 2 \pmod{4}$	0	0	1	0	1	0
$0+1 = 1 \equiv 1 \pmod{4}$	0	0	0	1	0	1
$0+0 = 0 \equiv 0 \pmod{4}$	0	0	0	0	0	0

Let's find, what SAT-solver can say about it?

First, we should represent our 2-bit adder as CNF expression.

Using Wolfram Mathematica, we can express 1-bit expression for both adder outputs:

```
In[]:=AdderQ0[aL_,bL_]=Xor[aL,bL]
```

```
Out[]:=aL ∨ bL
```

```
In[]:=AdderQ1[aL_,aH_,bL_,bH_]=Xor[And[aL,bL],Xor[aH,bH]]
```

```
Out[]:=aH ∨ bH ∨ (aL && bL)
```

We need such expression, where both parts will generate 1's. Let's use Wolfram Mathematica find all instances of such expression (I glued both parts with And):

```
In[]:=Boole[SatisfiabilityInstances[And[AdderQ0[aL,bL],AdderQ1[aL,aH,bL,bH]],{aL,aH,bL,bH},4]]
```

```
Out[]:={1,1,0,0},{1,0,0,1},{0,1,1,0},{0,0,1,1}
```

Yes, indeed, Mathematica says, there are 4 inputs which will lead to the result we need. So, Mathematica can also be used as SAT solver.

Nevertheless, let's proceed to CNF form. Using Mathematica again, let's convert our expression to CNF form:

```
In[]:=cnf=BooleanConvert[And[AdderQ0[aL,bL],AdderQ1[aL,aH,bL,bH]],`CNF']
```

```
Out[]:=(!aH ∥ !bH) && (aH ∥ bH) && (!aL ∥ !bL) && (aL ∥ bL)
```

Looks more complex. The reason of such verbosity is that CNF form doesn't allow XOR operations.

### 13.2.1 MiniSat

For starters, we can try MiniSat<sup>105</sup>. The standard way to encode CNF expression for MiniSat is to enumerate all OR parts at each line. Also, MiniSat doesn't support variable names, just numbers. Let's enumerate our variables: 1 will be aH, 2 - aL, 3 - bH, 4 - bL.

Here is what I've got when I converted Mathematica expression to the MiniSat input file:

```
p cnf 4 4
-1 -3 0
1 3 0
-2 -4 0
2 4 0
```

Two 4's at the first lines are number of variables and number of clauses respectively. There are 4 lines then, each for each OR clause. Minus before variable number meaning that the variable is negated. Absence of minus - not negated. Zero at the end is just terminating zero, meaning end of the clause.

In other words, each line is OR-clause with optional negations, and the task of MiniSat is to find such set of input, which can satisfy all lines in the input file.

That file I named as *adder.cnf* and now let's try MiniSat:

```
% minisat -verb=0 adder.cnf results.txt
SATISFIABLE
```

The results are in *results.txt* file:

```
SAT
-1 -2 3 4 0
```

This means, if the first two variables (aH and aL) will be *false*, and the last two variables (bH and bL) will be set to *true*, the whole CNF expression is satisfiable. Seems to be true: if bH and bL are the only inputs set to *true*, both resulting bits are also has *true* states.

Now how to get other instances? SAT-solvers, like SMT solvers, produce only one solution (or *instance*).

MiniSat uses PRNG and its initial seed can be set explicitly. I tried different values, but result is still the same. Nevertheless, CryptoMiniSat in this case was able to show all possible 4 instances, in chaotic order, though. So this is not very robust way.

Perhaps, the only known way is to negate solution clause and add it to the input expression. We've got -1 -2 3 4, now we can negate all values in it (just toggle minuses: 1 2 -3 -4) and add it to the end of the input file:

```
p cnf 4 5
-1 -3 0
1 3 0
-2 -4 0
2 4 0
1 2 -3 -4
```

Now we've got another result:

```
SAT
1 2 -3 -4 0
```

This means, aH and aL must be both *true* and bH and bL must be *false*, to satisfy the input expression. Let's negate this clause and add it again:

```
p cnf 4 6
-1 -3 0
1 3 0
-2 -4 0
2 4 0
1 2 -3 -4
-1 -2 3 4 0
```

The result is:

```
SAT
-1 2 3 -4 0
```

aH=false, aL=true, bH=true, bL=false. This is also correct, according to our truth table.

Let's add it again:

<sup>105</sup><http://minisat.se/MiniSat.html>

```
p cnf 4 7
-1 -3 0
1 3 0
-2 -4 0
2 4 0
1 2 -3 -4
-1 -2 3 4 0
1 -2 -3 4 0
```

```
SAT
1 -2 -3 4 0
```

$aH=true$ ,  $aL=false$ ,  $bH=false$ ,  $bL=true$ . This is also correct.  
This is fourth result. There are shouldn't be more. What if to add it?

```
p cnf 4 8
-1 -3 0
1 3 0
-2 -4 0
2 4 0
1 2 -3 -4
-1 -2 3 4 0
1 -2 -3 4 0
-1 2 3 -4 0
```

Now MiniSat just says “UNSATISFIABLE” without any additional information in the resulting file.  
Our example is tiny, but MiniSat can work with huge CNF expressions.

### 13.2.2 CryptoMiniSat

XOR operation is absent in CNF form, but crucial in cryptographical algorithms. Simplest possible way to represent single XOR operation in CNF form is:  $(\neg x \vee \neg y) \wedge (x \vee y)$  – not that small expression, though, many XOR operations in single expression can be optimized better.

One significant difference between MiniSat and CryptoMiniSat is that the latter supports clauses with XOR operations instead of ORs, because CryptoMiniSat has aim to analyze crypto algorithms<sup>106</sup>. XOR clauses are handled by CryptoMiniSat in a special way without translating to OR clauses.

You need just to prepend a clause with “x” in CNF file and OR clause is then treated as XOR clause by CryptoMiniSat. As of 2-bit adder, this smallest possible XOR-CNF expression can be used to find all inputs where both output adder bits are set:

$$(aH \oplus bH) \wedge (aL \oplus bL)$$

This is .cnf file for CryptoMiniSat:

```
p cnf 4 2
x1 3 0
x2 4 0
```

Now I run CryptoMiniSat with various random values to initialize its PRNG ...

```
% cryptominisat4 --verb 0 --random 0 XOR_adder.cnf
s SATISFIABLE
v 1 2 -3 -4 0
% cryptominisat4 --verb 0 --random 1 XOR_adder.cnf
s SATISFIABLE
v -1 -2 3 4 0
% cryptominisat4 --verb 0 --random 2 XOR_adder.cnf
s SATISFIABLE
v 1 -2 -3 4 0
% cryptominisat4 --verb 0 --random 3 XOR_adder.cnf
s SATISFIABLE
v 1 2 -3 -4 0
% cryptominisat4 --verb 0 --random 4 XOR_adder.cnf
s SATISFIABLE
v -1 2 3 -4 0
% cryptominisat4 --verb 0 --random 5 XOR_adder.cnf
s SATISFIABLE
v -1 2 3 -4 0
% cryptominisat4 --verb 0 --random 6 XOR_adder.cnf
s SATISFIABLE
```

<sup>106</sup><http://www.msoos.org/xor-clauses/>

```
v -1 -2 3 4 0
% cryptominisat4 --verb 0 --random 7 XOR_adder.cnf
s SATISFIABLE
v 1 -2 -3 4 0
% cryptominisat4 --verb 0 --random 8 XOR_adder.cnf
s SATISFIABLE
v 1 2 -3 -4 0
% cryptominisat4 --verb 0 --random 9 XOR_adder.cnf
s SATISFIABLE
v 1 2 -3 -4 0
```

Nevertheless, all 4 possible solutions are:

```
v -1 -2 3 4 0
v -1 2 3 -4 0
v 1 -2 -3 4 0
v 1 2 -3 -4 0
```

...the same as reported by MiniSat.

### 13.3 Picosat

At least Picosat can enumerate all possible solutions without crutches I just shown:

```
% picosat --all adder.cnf
s SATISFIABLE
v -1 -2 3 4 0
s SATISFIABLE
v -1 2 3 -4 0
s SATISFIABLE
v 1 2 -3 -4 0
s SATISFIABLE
v 1 -2 -3 4 0
s SOLUTIONS 4
```

### 13.4 Eight queens puzzle

Eight queens is a very popular puzzle and often used for measuring performance of SAT solvers. The problem is to place 8 queens on chess board so they will not attack each other. For example:

```
| | | | * | | | |
| | | | | | * | |
| | * | | | | | |
| * | | | | * | | |
| | | * | | | | |
| | | | | | | * |
```

Let's try to figure out how to solve it.

#### 13.4.1 POPCNT1

One important function we will (often) use is `POPCNT1`. This is a function which returns *True* if one single of inputs is *True* and others are *False*. It will return *False* otherwise.

In my other examples, I've used Wolfram Mathematica to generate CNF clauses for it, for example: [13.7](#). What expression Mathematica offers as `POPCNT1` function with 8 inputs?

```
(!a||!b)&&(!a||!c)&&(!a||!d)&&(!a||!e)&&(!a||!f)&&(!a||!g)&&(!a||!h)&&(a||b||c||d||e||f||g||h)&&
(!b||!c)&&(!b||!d)&&(!b||!e)&&(!b||!f)&&(!b||!g)&&(!b||!h)&&(!c||!d)&&(!c||!e)&&(!c||!f)&&(!c||!g)&&
(!c||!h)&&(!d||!e)&&(!d||!f)&&(!d||!g)&&(!d||!h)&&(!e||!f)&&(!e||!g)&&(!e||!h)&&(!f||!g)&&(!f||!h)&&(!g||!h)
```

We can clearly see that the expression consisting of all possible variable pairs (negated) plus enumeration of all variables (non-negated). In plain English terms, this means: "no pair can be equal to two *True*'s AND at least one *True* must be present among all variables".

This is how it works: if two variables will be *True*, negated they will be both *False*, and this clause will not be evaluated to *True*, which is our ultimate goal. If one of variables is *True*, both negated will produce one *True* and one *False* (fine). If both variables are *False*, both negated will produce two *True*'s (again, fine).

Here is how I can generate clauses for the function using *itertools* module from Python, which provides many important functions from combinatorics:

```
# naive/pairwise encoding
def AtMost1(self, lst):
    for pair in itertools.combinations(lst, r=2):
        self.add_clause([self.neg(pair[0]), self.neg(pair[1])])

def POPCNT1(self, lst):
    self.AtMost1(lst)
    self.OR_always(lst)
```

`AtMost1()` function enumerates all possible pairs using *itertools* function *combinations()*.

`POPCNT1()` function does the same, but also adds a final clause, which forces at least one *True* variable to be present.

What clauses will be generated for 5 variables (1..5)?

```
p cnf 5 11
-2 -5 0
-2 -4 0
-4 -5 0
-2 -3 0
-1 -4 0
-1 -5 0
-1 -2 0
-1 -3 0
-3 -4 0
-3 -5 0
1 2 3 4 5 0
```

Yes, these are all possible pairs of 1..5 numbers + all 5 numbers.

We can get all solutions using Picosat:

```
% picosat --all popcnt1.cnf
s SATISFIABLE
v -1 -2 -3 -4 5 0
s SATISFIABLE
v -1 -2 -3 4 -5 0
s SATISFIABLE
v -1 -2 3 -4 -5 0
s SATISFIABLE
v -1 2 -3 -4 -5 0
s SATISFIABLE
v 1 -2 -3 -4 -5 0
s SOLUTIONS 5
```

5 possible solutions indeed.

### 13.4.2 Eight queens

Now let's get back to eight queens.

We can assign 64 variables to  $8 \cdot 8 = 64$  cells. Cell occupied with queen will be *True*, vacant cell will be *False*.

The problem of placing non-attacking (each other) queens on chess board (of any size) can be stated in plain English like this:

- one single queen must be present at each row;
- one single queen must be present at each column;
- zero or one queen must be present at each diagonal (empty diagonals can be present in valid solution).

These rules can be translated like that:

- `POPCNT1(each row)==True`
- `POPCNT1(each column)==True`
- `AtMost1(each diagonal)==True`

Now all we need is to enumerate rows, columns and diagonals and gather all clauses:

```

#-*- coding: utf-8 -*-

#!/usr/bin/env python

import itertools, subprocess, os, frolic, Xu

SIZE=8
SKIP_SYMMETRIES=True
#SKIP_SYMMETRIES=False

def row_col_to_var(row, col):
    global first_var
    return str(row*SIZE+col+first_var)

def gen_diagonal(s, start_row, start_col, increment, limit):
    col=start_col
    tmp=[]
    for row in range(start_row, SIZE):
        tmp.append(row_col_to_var(row, col))
        col=col+increment
        if col==limit:
            break
    # ignore diagonals consisting of 1 cell:
    if len(tmp)>1:
        # we can't use POPCNT1() here, since some diagonals are empty in valid solutions.
        s.AtMost1(tmp)

def add_2D_array_as_negated_constraint(s, a):
    negated_solution=[]
    for row in range(SIZE):
        for col in range(SIZE):
            negated_solution.append(s.neg_if(a[row][col], row_col_to_var(row, col)))
    s.add_clause(negated_solution)

def main():
    global first_var

    s=Xu.Xu(False)

    _vars=s.alloc_BV(SIZE**2)
    first_var=int(_vars[0])

    # enumerate all rows:
    for row in range(SIZE):
        s.POPCNT1([row_col_to_var(row, col) for col in range(SIZE)])

    # enumerate all columns:
    # POPCNT1() could be used here as well:
    for col in range(SIZE):
        s.AtMost1([row_col_to_var(row, col) for row in range(SIZE)])

    # enumerate all diagonals:
    for row in range(SIZE):
        for col in range(SIZE):
            gen_diagonal(s, row, col, 1, SIZE) # from L to R
            gen_diagonal(s, row, col, -1, -1) # from R to L

    # find all solutions:
    sol_n=1
    while True:
        if s.solve()==False:
            print "unsat!"
            print "solutions total=", sol_n-1
            exit(0)

        # print solution:
        print "solution number", sol_n, ":"

        # get solution and make 2D array of bools:
        solution_as_2D_bool_array=[]
        for row in range(SIZE):
            solution_as_2D_bool_array.append ([s.get_var_from_solution(row_col_to_var(row, col)) for col in
                                                range(SIZE)])

        # print 2D array:

```

```

for row in range(SIZE):
    tmp=[[" ", "*"][solution_as_2D_bool_array[row][col]]+"|" for col in range(SIZE)]
    print "|"+"".join(tmp)

# add 2D array as negated constraint:
add_2D_array_as_negated_constraint(s, solution_as_2D_bool_array)

# if we skip symmetries, rotate/reflect solution and add them as negated constraints:
if SKIP_SYMMETRIES:
    for a in range(4):
        tmp=frolic.rotate_rect_array(solution_as_2D_bool_array, a)
        add_2D_array_as_negated_constraint(s, tmp)

        tmp=frolic.reflect_horizontally(frolic.rotate_rect_array(solution_as_2D_bool_array, a))
        add_2D_array_as_negated_constraint(s, tmp)

    sol_n=sol_n+1

main()

```

( [https://github.com/dennis714/SAT\\_SMT\\_article/blob/master/SAT/8queens/8queens.py](https://github.com/dennis714/SAT_SMT_article/blob/master/SAT/8queens/8queens.py) )

Perhaps, `gen_diagonal()` function is not very aesthetically appealing: it enumerates also subdiagonals of already enumerated longer diagonals. To prevent presence of duplicate clauses, *clauses* global variable is not a list, rather set, which allows only unique data to be present there.

Also, I've used `AtMost1` for each column, this will help to produce slightly lower number of clauses. Each column will have a queen anyway, this is implied from the first rule (`POPCNT1` for each row).

After running, we got CNF file with 64 variables and 736 clauses ([https://github.com/dennis714/SAT\\_SMT\\_article/blob/master/SAT/8queens/8queens.cnf](https://github.com/dennis714/SAT_SMT_article/blob/master/SAT/8queens/8queens.cnf)). Here is one solution:

```

% python 8queens.py
len(clauses)= 736
| | | | * | | | |
| | | | | | | * |
| | | | * | | | |
| | * | | | | | |
| * | | | | | | |
| | | * | | | | |
| | | | | | | * |

```

How many possible solutions are there? Picosat tells 92, which is indeed correct number of solutions (<https://oeis.org/A000170>).

Performance of Picosat is not impressive, probably because it has to output all the solutions. It took 34 seconds on my ancient Intel Atom 1.66GHz netbook to enumerate all solutions for 11 · 11 chess board (2680 solutions), which is way slower than my straight brute-force program: <https://yurichev.com/blog/8queens/>. Nevertheless, it's lightning fast (as other SAT solvers) in finding first solution.

The SAT instance is also small enough to be easily solved by my simplest possible backtracking SAT solver: [13.9](#).

### 13.4.3 Counting all solutions

We get a solution, negate it and add as a new constraint. In plain English language this sounds “find a solution, which is also can't be equal to the recently found/added”. We add them consequently and the process is slowing—just because a problem (*instance*) is growing and SAT solver experience hard times in find yet another solution.

### 13.4.4 Skipping symmetrical solutions

We can also add rotated and reflected (horizontally) solution, so to skip symmetrical solutions. By doing so, we're getting 12 solutions for 8\*8 board, 46 for 9\*9 board, etc. This is <https://oeis.org/A002562>.

## 13.5 Sudoku in SAT

One might think that we can encode each 1..9 number in binary form: 5 bits or variables would be enough. But there is even simpler way: allocate 9 bits, where only one bit will be *True*. The number 1 can be encoded

as [1, 0, 0, 0, 0, 0, 0, 0, 0], the number 3 as [0, 0, 1, 0, 0, 0, 0, 0, 0], etc. Seems uneconomical? Yes, but other operations would be simpler.

First of all, we'll reuse important `POPCNT1` function I've described earlier: [13.4.1](#).

The second important operation we need to invent is making 9 numbers unique. If each number is encoded as 9-bits vector, 9 numbers can form a matrix, like:

```
0 0 0 0 0 0 1 0 0 <- 1st number
0 0 0 0 0 1 0 0 0 <- 2nd number
0 1 0 0 0 0 0 0 0 <- ...
0 0 1 0 0 0 0 0 0 <- ...
0 0 0 0 0 0 0 0 1 <- ...
0 0 0 0 1 0 0 0 0 <- ...
0 0 0 0 0 0 0 1 0 <- ...
1 0 0 0 0 0 0 0 0 <- ...
0 0 0 1 0 0 0 0 0 <- 9th number
```

Now we will use a `POPCNT1` function to make each row in the matrix to contain only one *True* bit, that will preserve consistency in encoding, since no vector can contain more than 1 *True* bit, or no *True* bits at all. Then we will use a `POPCNT1` function again to make all columns in the matrix to have only one single *True* bit. That will make all rows in matrix unique, in other words, all 9 encoded numbers will always be unique.

After applying `POPCNT1` function  $9+9=18$  times we'll have 9 unique numbers in 1..9 range.

Using that operation we can make each row of Sudoku puzzle unique, each column unique and also each  $3 \cdot 3 = 9$  box.

```
#!/usr/bin/env python
import itertools, subprocess, os

# global variables:
clauses=[]
vector_names={}
last_var=1

BITS_PER_VECTOR=9

def read_lines_from_file (fname):
    f=open(fname)
    new_ar=[item.rstrip() for item in f.readlines()]
    f.close()
    return new_ar

def run_minisat (CNF_fname):
    child = subprocess.Popen(["minisat", CNF_fname, "results.txt"], stdout=subprocess.PIPE)
    child.wait()
    # 10 is SAT, 20 is UNSAT
    if child.returncode==20:
        os.remove ("results.txt")
        return None

    if child.returncode!=10:
        print "(minisat) unknown retcode: ", child.returncode
        exit(0)

    solution=read_lines_from_file("results.txt")[1].split(" ")
    os.remove ("results.txt")

    return solution

def write_CNF(fname, clauses, VARS_TOTAL):
    f=open(fname, "w")
    f.write ("p cnf "+str(VARS_TOTAL)+" "+str(len(clauses))+"\n")
    [f.write(" ".join(c)+" 0\n") for c in clauses]
    f.close()

def neg(v):
    return "-" + v

def add_popcnt1(vars):
    global clauses
    # enumerate all possible pairs
    # no pair can have both True's
```



```

    # so add "~var OR ~var2"
    for pair in itertools.combinations(vars, r=2):
        clauses.append([neg(pair[0]), neg(pair[1])])
    # at least one var must be present:
    clauses.append(vars)

def make_distinct_bits_in_vector(vec_name):
    global vector_names
    global last_var

    add_popcnt1([vector_names[(vec_name,i)] for i in range(BITS_PER_VECTOR)])

def make_distinct_vectors(vectors):
    # take each bit from all vectors, call add_popcnt1()
    for i in range(BITS_PER_VECTOR):
        add_popcnt1([vector_names[(vec,i)] for vec in vectors])

def cvt_vector_to_number(vec_name, solution):
    for i in range(BITS_PER_VECTOR):
        if vector_names[(vec_name,i)] in solution:
            # variable present in solution as non-negated (without a "-" prefix)
            return i+1
    raise AssertionError

def alloc_var():
    global last_var
    last_var=last_var+1
    return str(last_var-1)

def alloc_vector(l, name):
    global last_var
    global vector_names
    rt=[]
    for i in range(l):
        v=alloc_var()
        vector_names[(name,i)]=v
        rt.append(v)
    return rt

def add_constant(var,b):
    global clauses
    if b==True or b==1:
        clauses.append([var])
    else:
        clauses.append([neg(var)])

# vec is a list of True/False/0/1
def add_constant_vector(vec_name, vec):
    global vector_names
    for i in range(BITS_PER_VECTOR):
        add_constant (vector_names[(vec_name, i)], vec[i])

# 1 -> [1, 0, 0, 0, 0, 0, 0, 0, 0]
# 3 -> [0, 0, 1, 0, 0, 0, 0, 0, 0]
def number_to_vector(n):
    rt=[0]*(n-1)
    rt.append(1)
    rt=rt+[0]*(BITS_PER_VECTOR-len(rt))
    return rt

"""
coordinates we're using here:

+-----+-----+-----+
|11 12 13|14 15 16|17 18 19|
|21 22 23|24 25 26|27 28 29|
|31 32 33|34 35 36|37 38 39|
+-----+-----+-----+
|41 42 43|44 45 46|47 48 49|
|51 52 53|54 55 56|57 58 59|
|61 62 63|64 65 66|67 68 69|
+-----+-----+-----+
|71 72 73|74 75 76|77 78 79|
|81 82 83|84 85 86|87 88 89|
|91 92 93|94 95 96|97 98 99|
+-----+-----+-----+

```

```

"""
def make_vec_name(row, col):
    return "cell"+str(row)+str(col)

def puzzle_to_clauses (puzzle):
    # process text line:
    current_column=1
    current_row=1
    for i in puzzle:
        if i!='.':
            add_constant_vector(make_vec_name(current_row, current_column), number_to_vector(int(i)))
            current_column=current_column+1
        if current_column==10:
            current_column=1
            current_row=current_row+1

def print_solution(solution):
    for row in range(1,9+1):
        # print row:
        print " ".join([str(cvt_vector_to_number(make_vec_name(row, col), solution)) for col in range(1,9+1)])

def main():
    # allocate 9*9*9=729 variables:
    for row in range(1, 9+1):
        for col in range(1, 9+1):
            alloc_vector(9, make_vec_name(row, col))
            make_distinct_bits_in_vector(make_vec_name(row, col))

    # variables in each row are unique:
    for row in range(1, 9+1):
        make_distinct_vectors([make_vec_name(row, col) for col in range(1, 9+1)])

    # variables in each column are unique:
    for col in range(1, 9+1):
        make_distinct_vectors([make_vec_name(row, col) for row in range(1, 9+1)])

    # variables in each 3*3 box are unique:
    for row in range(1, 9+1, 3):
        for col in range(1, 9+1, 3):
            tmp=[]
            tmp.append(make_vec_name(row+0, col+0))
            tmp.append(make_vec_name(row+0, col+1))
            tmp.append(make_vec_name(row+0, col+2))
            tmp.append(make_vec_name(row+1, col+0))
            tmp.append(make_vec_name(row+1, col+1))
            tmp.append(make_vec_name(row+1, col+2))
            tmp.append(make_vec_name(row+2, col+0))
            tmp.append(make_vec_name(row+2, col+1))
            tmp.append(make_vec_name(row+2, col+2))
            make_distinct_vectors(tmp)

    # http://www.norvig.com/sudoku.html
    # http://www.mirror.co.uk/news/weird-news/worlds-hardest-sudoku-can-you-242294
    puzzle_to_clauses("..53....8.....2..7..1.5..4....53...1..7...6..32...8..6.5....9..4....3.....97..")

    print "len(clauses)=",len(clauses)
    write_CNF("1.cnf", clauses, last_var-1)
    solution=run_minisat("1.cnf")
    #os.remove("1.cnf")
    if solution==None:
        print "unsat!"
        exit(0)

    print_solution(solution)

main()

```

( [https://github.com/dennis714/SAT\\_SMT\\_article/blob/master/SAT/sudoku/sudoku\\_SAT.py](https://github.com/dennis714/SAT_SMT_article/blob/master/SAT/sudoku/sudoku_SAT.py) )

The `make_distinct_bits_in_vector()` function preserves consistency of encoding.

The `make_distinct_vectors()` function makes 9 numbers unique.

The `cvt_vector_to_number()` decodes vector to number.

The `number_to_vector()` encodes number to vector.

The `main()` function has all necessary calls to make rows/columns/ $3 \cdot 3$  boxes unique.

That works:

```
% python sudoku_SAT.py
len(clauses)= 12195
1 4 5 3 2 7 6 9 8
8 3 9 6 5 4 1 2 7
6 7 2 9 1 8 5 4 3
4 9 6 1 8 5 3 7 2
2 1 8 4 7 3 9 5 6
7 5 3 2 9 6 4 8 1
3 6 7 5 4 2 8 1 9
9 8 4 7 6 1 2 3 5
5 2 1 8 3 9 7 6 4
```

Same solution as earlier: [7.3](#).

Picosat tells this SAT instance has only one solution. Indeed, as they say, true Sudoku puzzle can have only one solution.

### 13.5.1 Getting rid of one POPCNT1 function call

To make 9 unique 1..9 numbers we can use `POPCNT1` function to make each row in matrix be unique and use `OR` boolean operation for all columns. That will have merely the same effect: all rows has to be unique to make each column to be evaluated to `True` if all variables in column are OR'ed. (I will do this in the next example: [13.6](#).)

That will make 3447 clauses instead of 12195, but somehow, SAT solvers works slower. No idea why.

## 13.6 Zebra puzzle as a SAT problem

Let's try to solve Zebra puzzle ([7.2](#)) in SAT.

I would define each variable as vector of 5 variables, as I did before in Sudoku solver: [13.5](#).

I also use `POPCNT1` function, but unlike previous example, I used Wolfram Mathematica to generate it in CNF form:

```
In[]:= tbl1=Table[PadLeft[IntegerDigits[i,2],5] ->If[Equal[DigitCount[i,2][[1]],1],1,0],{i,0,63}]
Out[]= {{0,0,0,0,0}->0,
{0,0,0,0,1}->1,
{0,0,0,1,0}->1,
{0,0,0,1,1}->0,
{0,0,1,0,0}->1,
{0,0,1,0,1}->0,
...
{1,1,1,1,0}->0,
{1,1,1,1,1}->0}

In[]:= BooleanConvert[BooleanFunction[tbl1,{a,b,c,d,e}], "CNF"]
Out[]= (!a||b)&&(!a||c)&&(!a||d)&&(!a||e)&&(a||b||c||d||e)&&(!b||c)&&(!b||d)&&(!b||e)&&(!c||d)&&(!c||e)
&&(!d||e)
```

Also, as I suggested before ([13.5.1](#)), I used `OR` operation as the second step.

```
def mathematica_to_CNF (s, d):
    for k in d.keys():
        s=s.replace(k, d[k])
    s=s.replace("!", "-").replace("||", " ").replace("(", "").replace(")", "")
    s=s.split("&&")
    return s

def add_popcnt1(v1, v2, v3, v4, v5):
    global clauses
    s="(!a||b)&& \
      (!a||c)&& \
      (!a||d)&& \
      (!a||e)&& \
      (!b||c)&& \
      (!b||d)&& \
      (!b||e)&& \
      (!c||d)&& \
```

```

    "(!c||!e)&&" \
    "(!d||!e)&&" \
    "(a||b||c||d||e)"

    clauses=clauses+mathematica_to_CNF(s, {"a":v1, "b":v2, "c":v3, "d":v4, "e":v5})

...

# k=tuple: ("high-level" variable name, number of bit (0..4))
# v=variable number in CNF
vars={}
vars_last=1

...

def alloc_distinct_variables(names):
    global vars
    global vars_last
    for name in names:
        for i in range(5):
            vars[(name,i)]=str(vars_last)
            vars_last=vars_last+1

        add_popcnt1(vars[(name,0)], vars[(name,1)], vars[(name,2)], vars[(name,3)], vars[(name,4)])

    # make them distinct:
    for i in range(5):
        clauses.append(vars[(names[0],i)] + " " + vars[(names[1],i)] + " " + vars[(names[2],i)] + " " + vars[(names[3],i)] + " " + vars[(names[4],i)])

...

alloc_distinct_variables(["Yellow", "Blue", "Red", "Ivory", "Green"])
alloc_distinct_variables(["Norwegian", "Ukrainian", "Englishman", "Spaniard", "Japanese"])
alloc_distinct_variables(["Water", "Tea", "Milk", "OrangeJuice", "Coffee"])
alloc_distinct_variables(["Kools", "Chesterfield", "OldGold", "LuckyStrike", "Parliament"])
alloc_distinct_variables(["Fox", "Horse", "Snails", "Dog", "Zebra"])

...

```

Now we have 5 boolean variables for each *high-level* variable, and each group of variables will always have distinct values.

Now let's reread puzzle description: "2.The Englishman lives in the red house.". That's easy. In my Z3 and KLEE examples I just wrote "Englishman==Red". Same story here: we just add a clauses showing that 5 boolean variables for "Englishman" must be equal to 5 booleans for "Red".

On a lowest CNF level, if we want to say that two variables must be equal to each other, we add two clauses:

$$(var1 \vee \neg var2) \wedge (\neg var1 \vee var2)$$

That means, both *var1* and *var2* values must be *False* or *True*, but they cannot be different.

```

def add_eq_clauses(var1, var2):
    global clauses
    clauses.append(var1 + " -" + var2)
    clauses.append("-" + var1 + " " + var2)

def add_eq (n1, n2):
    for i in range(5):
        add_eq_clauses(vars[(n1,i)], vars[(n2, i)])

...

# 2.The Englishman lives in the red house.
add_eq("Englishman","Red")

# 3.The Spaniard owns the dog.
add_eq("Spaniard","Dog")

# 4.Coffee is drunk in the green house.
add_eq("Coffee","Green")

...

```

Now the next conditions: "9.Milk is drunk in the middle house." (i.e., 3rd house), "10.The Norwegian lives in the first house." We can just assign boolean values directly:

```
# n=1..5
def add_eq_var_n (name, n):
    global clauses
    global vars
    for i in range(5):
        if i==n-1:
            clauses.append(vars[(name,i)]) # always True
        else:
            clauses.append("-"+vars[(name,i)]) # always False

...

# 9.Milk is drunk in the middle house.
add_eq_var_n("Milk",3) # i.e., 3rd house

# 10.The Norwegian lives in the first house.
add_eq_var_n("Norwegian",1)
```

For “Milk” we will have “0 0 1 0 0” value, for “Norwegian”: “1 0 0 0 0”.

What to do with this? “6.The green house is immediately to the right of the ivory house.” I can construct the following condition:

Ivory	Green
AND(1 0 0 0 0	0 1 0 0 0)
.. OR ..	
AND(0 1 0 0 0	0 0 1 0 0)
.. OR ..	
AND(0 0 1 0 0	0 0 0 1 0)
.. OR ..	
AND(0 0 0 1 0	0 0 0 0 1)

There is no “0 0 0 0 1” for “Ivory”, because it cannot be the last one. Now I can convert these conditions to CNF using Wolfram Mathematica:

```
In[ ]:= BooleanConvert[(a1&& !b1&&c1&&d1&&!e1&&a2&& b2&&c2&&d2&&!e2) ||
(!a1&& b1&&!c1&&d1&&!e1&&a2&& !b2&&c2&&d2&&!e2) ||
(!a1&& !b1&&c1&&d1&&!e1&&a2&& !b2&&c2&&d2&&!e2) ||
(!a1&& !b1&&c1&&d1&&!e1&&a2&& !b2&&c2&&d2&&e2) , "CNF"]

Out[ ]= (!a1||!b1)&&(!a1||!c1)&&(!a1||!d1)&&(a1||b1||c1||d1)&&a2&&(!b1||!b2)&&(!b1||!c1)&&
(!b1||!d1)&&(b1||b2||c1||d1)&&(!b2||!c1)&&(!b2||!c2)&&(!b2||!d1)&&(!b2||!d2)&&(!b2||!e2)&&
(b2||c1||c2||d1)&&(b2||c2||d1||d2)&&(b2||c2||d2||e2)&&(!c1||!c2)&&(!c1||!d1)&&(!c2||!d1)&&
(!c2||!d2)&&(!c2||!e2)&&(!d1||!d2)&&(!d2||!e2)&&!e1
```

And here is a piece of my Python code:

```
def add_right (n1, n2):
    global clauses
    s="(!a1||!b1)&&(!a1||!c1)&&(!a1||!d1)&&(a1||b1||c1||d1)&&a2&&(!b1||!b2)&&(!b1||!c1)&&(!b1||!d1)&&" \
    "(b1||b2||c1||d1)&&(!b2||!c1)&&(!b2||!c2)&&(!b2||!d1)&&(!b2||!d2)&&(!b2||!e2)&&(b2||c1||c2||d1)&&" \
    "(b2||c2||d1||d2)&&(b2||c2||d2||e2)&&(!c1||!c2)&&(!c1||!d1)&&(!c2||!d1)&&(!c2||!d2)&&(!c2||!e2)&&" \
    "(!d1||!d2)&&(!d2||!e2)&&!e1"

    clauses=clauses+mathematica_to_CNF(s, {
        "a1": vars[(n1,0)], "b1": vars[(n1,1)], "c1": vars[(n1,2)], "d1": vars[(n1,3)], "e1": vars[(n1,4)],
        "a2": vars[(n2,0)], "b2": vars[(n2,1)], "c2": vars[(n2,2)], "d2": vars[(n2,3)], "e2": vars[(n2,4)]})

...

# 6.The green house is immediately to the right of the ivory house.
add_right("Ivory", "Green")
```

What we will do with that? “11.The man who smokes Chesterfields lives in the house next to the man with the fox.” “12.Kools are smoked in the house next to the house where the horse is kept.”

We don’t know side, left or right, but we know that they are differ in one. Here is a clauses I would add:

Chesterfield	Fox
AND(0 0 0 0 1	0 0 0 1 0)
.. OR ..	
AND(0 0 0 1 0	0 0 0 0 1)
AND(0 0 0 1 0	0 0 1 0 0)
.. OR ..	
AND(0 0 1 0 0	0 1 0 0 0)
AND(0 0 1 0 0	0 0 0 1 0)

```

.. OR ..
AND(0 1 0 0 0      1 0 0 0 0)
AND(0 1 0 0 0      0 0 1 0 0)
.. OR ..
AND(1 0 0 0 0      0 1 0 0 0)

```

I can convert this into CNF using Mathematica again:

```

In[]:= BooleanConvert[(a1&&!b1&&!c1&&!d1&&!e1&&a2&&b2&&!c2&&!d2&&!e2) ||
(!a1&&b1&&!c1&&!d1&&!e1&&a2&&!b2&&!c2&&!d2&&!e2) ||
(!a1&&b1&&!c1&&!d1&&!e1&&a2&&!b2&&c2&&!d2&&!e2) ||
(!a1&&!b1&&c1&&!d1&&!e1&&a2&&b2&&!c2&&!d2&&!e2) ||
(!a1&&!b1&&c1&&!d1&&!e1&&a2&&!b2&&c2&&!d2&&!e2) ||
(!a1&&!b1&&c1&&d1&&!e1&&a2&&!b2&&c2&&!d2&&!e2) ||
(!a1&&!b1&&c1&&d1&&!e1&&a2&&!b2&&!c2&&!d2&&!e2) ||
(!a1&&!b1&&c1&&d1&&!e1&&a2&&!b2&&c2&&d2&&!e2) ,"CNF"]

Out[]:= (!a1||!b1)&&(!a1||!c1)&&(!a1||!d1)&&(!a1||!e1)&&(a1||b1||c1||d1||e1)&&(!a2||b1)&&(!a2||!b2)&&
(!a2||!c2)&&(!a2||!d2)&&(!a2||!e2)&&(a2||b2||c1||c2||d1||e1)&&(a2||b2||c2||d1||d2)&&(a2||b2||c2||d2||e2)&&
(!b1||!b2)&&(!b1||!c1)&&(!b1||!d1)&&(!b1||!e1)&&(b1||b2||c1||d1||e1)&&(!b2||!c2)&&(!b2||!d1)&&(!b2||!d2)&&
(!b2||!e1)&&(!b2||!e2)&&(!c1||!c2)&&(!c1||!d1)&&(!c1||!e1)&&(!c2||!d2)&&(!c2||!e1)&&(!c2||!e2)&&
(!d1||!d2)&&(!d1||!e1)&&(!d2||!e2)

```

And here is my code:

```

def add_right_or_left (n1, n2):
    global clauses
    s="(!a1||!b1)&&(!a1||!c1)&&(!a1||!d1)&&(!a1||!e1)&&(a1||b1||c1||d1||e1)&&(!a2||b1)&&" \
      "(!a2||!b2)&&(!a2||!c2)&&(!a2||!d2)&&(!a2||!e2)&&(a2||b2||c1||c2||d1||e1)&&(a2||b2||c2||d1||d2)&&" \
      "(a2||b2||c2||d2||e2)&&(!b1||!b2)&&(!b1||!c1)&&(!b1||!d1)&&(!b1||!e1)&&(b1||b2||c1||d1||e1)&&" \
      "(!b2||!c2)&&(!b2||!d1)&&(!b2||!d2)&&(!b2||!e1)&&(!b2||!e2)&&(!c1||!c2)&&(!c1||!d1)&&(!c1||!e1)&&" \
      "(!c2||!d2)&&(!c2||!e1)&&(!c2||!e2)&&(!d1||!d2)&&(!d1||!e1)&&(!d2||!e2)"

    clauses=clauses+mathematica_to_CNF(s, {
        "a1": vars[(n1,0)], "b1": vars[(n1,1)], "c1": vars[(n1,2)], "d1": vars[(n1,3)], "e1": vars[(n1,4)],
        "a2": vars[(n2,0)], "b2": vars[(n2,1)], "c2": vars[(n2,2)], "d2": vars[(n2,3)], "e2": vars[(n2,4)]})

    ...

# 11.The man who smokes Chesterfields lives in the house next to the man with the fox.
add_right_or_left("Chesterfield","Fox") # left or right

# 12.Kools are smoked in the house next to the house where the horse is kept.
add_right_or_left("Kools","Horse") # left or right

```

This is it! The full source code: [https://github.com/dennis714/SAT\\_SMT\\_article/blob/master/SAT/zebra/zebra\\_SAT.py](https://github.com/dennis714/SAT_SMT_article/blob/master/SAT/zebra/zebra_SAT.py).

Resulting CNF instance has 125 boolean variables and 511 clauses: [https://github.com/dennis714/SAT\\_SMT\\_article/blob/master/SAT/zebra/1.cnf](https://github.com/dennis714/SAT_SMT_article/blob/master/SAT/zebra/1.cnf). It is a piece of cake for any SAT solver. Even my toy-level SAT-solver (13.9) can solve it in ~1 second on my ancient Intel Atom netbook.

And of course, there is only one possible solution, what is acknowledged by Picosat.

```

% python zebra_SAT.py
Yellow 1
Blue 2
Red 3
Ivory 4
Green 5
Norwegian 1
Ukrainian 2
Englishman 3
Spaniard 4
Japanese 5
Water 1
Tea 2
Milk 3
OrangeJuice 4
Coffee 5
Kools 1

```

```
Chesterfield 2
OldGold 3
LuckyStrike 4
Parliament 5
Fox 1
Horse 2
Snails 3
Dog 4
Zebra 5
```

## 13.7 Cracking Minesweeper with SAT solver

See also about cracking it using Z3: [7.9](#).

### 13.7.1 Simple *population count* function

First of all, somehow we need to count neighbour bombs. The counting function is similar to *population count* function.

We can try to make [CNF](#) expression using Wolfram Mathematica. This will be a function, returning *True* if any of 2 bits of 8 inputs bits are *True* and others are *False*. First, we make truth table of such function:

```
In[]:= tbl2 =
Table[PadLeft[IntegerDigits[i, 2], 8] ->
If[Equal[DigitCount[i, 2][[1]], 2], 1, 0], {i, 0, 255}]

Out[]= {{0, 0, 0, 0, 0, 0, 0, 0} -> 0, {0, 0, 0, 0, 0, 0, 0, 1} -> 0,
{0, 0, 0, 0, 0, 0, 1, 0} -> 0, {0, 0, 0, 0, 0, 0, 1, 1} -> 1,
{0, 0, 0, 0, 0, 1, 0, 0} -> 0, {0, 0, 0, 0, 0, 1, 0, 1} -> 1,
{0, 0, 0, 0, 0, 1, 1, 0} -> 1, {0, 0, 0, 0, 0, 1, 1, 1} -> 0,
{0, 0, 0, 0, 1, 0, 0, 0} -> 0, {0, 0, 0, 0, 1, 0, 0, 1} -> 1,
{0, 0, 0, 0, 1, 0, 1, 0} -> 1, {0, 0, 0, 0, 1, 0, 1, 1} -> 0,
...
{1, 1, 1, 1, 1, 0, 1, 0} -> 0, {1, 1, 1, 1, 1, 0, 1, 1} -> 0,
{1, 1, 1, 1, 1, 1, 0, 0} -> 0, {1, 1, 1, 1, 1, 1, 0, 1} -> 0,
{1, 1, 1, 1, 1, 1, 1, 0} -> 0, {1, 1, 1, 1, 1, 1, 1, 1} -> 0}
```

Now we can make [CNF](#) expression using this truth table:

```
In[]:= BooleanConvert[
BooleanFunction[tbl2, {a, b, c, d, e, f, g, h}], "CNF"]

Out[]= (! a || ! b || ! c) && (! a || ! b || ! d) && (! a || !
b || ! e) && (! a || ! b || ! f) && (! a || ! b || ! g) && (!
a || ! b || ! h) && (! a || ! c || ! d) && (! a || ! c || !
e) && (! a || ! c || ! f) && (! a || ! c || ! g) && (! a || !
c || ! h) && (! a || ! d || ! e) && (! a || ! d || ! f) && (!
a || ! d || ! g) && (! a || ! d || ! h) && (! a || ! e || !
f) && (! a || ! e || ! g) && (! a || ! e || ! h) && (! a || !
f || ! g) && (! a || ! f || ! h) && (! a || ! g || ! h) && (a ||
b || c || d || e || f || g) && (a || b || c || d || e || f ||
h) && (a || b || c || d || e || g || h) && (a || b || c || d || f ||
g || h) && (a || b || c || e || f || g || h) && (a || b || d ||
e || f || g || h) && (a || c || d || e || f || g || h) &&
h) && (! b || ! c || ! d) && (! b || ! c || ! e) && (! b || !
c || ! f) && (! b || ! c || ! g) && (! b || ! c || ! h) && (!
b || ! d || ! e) && (! b || ! d || ! f) && (! b || ! d || !
g) && (! b || ! d || ! h) && (! b || ! e || ! f) && (! b || !
e || ! g) && (! b || ! e || ! h) && (! b || ! f || ! g) && (!
b || ! f || ! h) && (! b || ! g || ! h) && (b || c || d || e ||
f || g || h) && (! c || ! d || ! e) && (! c || ! d || ! f) && (! c || !
d || ! g) && (! c || ! d || ! h) && (! c || ! e || ! f) && (!
c || ! e || ! g) && (! c || ! e || ! h) && (! c || ! f || !
g) && (! c || ! f || ! h) && (! c || ! g || ! h) && (! d || !
e || ! f) && (! d || ! e || ! g) && (! d || ! e || ! h) && (!
d || ! f || ! g) && (! d || ! f || ! h) && (! d || ! g || !
h) && (! e || ! f || ! g) && (! e || ! f || ! h) && (! e || !
g || ! h) && (! f || ! g || ! h)
```

The syntax is similar to C/C++. Let's check it.

I wrote a Python function to convert Mathematica's output into [CNF](#) file which can be feeded to SAT solver:

```
#!/usr/bin/python

import subprocess

def mathematica_to_CNF (s, a):
    s=s.replace("a", a[0]).replace("b", a[1]).replace("c", a[2]).replace("d", a[3])
    s=s.replace("e", a[4]).replace("f", a[5]).replace("g", a[6]).replace("h", a[7])
    s=s.replace("!", "-").replace("||", " ").replace("(", "").replace(")", "")
    s=s.split ("&&")
    return s

def POPCNT2 (a):
    s="(!a||!b||!c)&&(!a||!b||!d)&&(!a||!b||!e)&&(!a||!b||!f)&&(!a||!b||!g)&&(!a||!b||!h)&&(!a||!c||!d)&&" \
    "(!a||!c||!e)&&(!a||!c||!f)&&(!a||!c||!g)&&(!a||!c||!h)&&(!a||!d||!e)&&(!a||!d||!f)&&(!a||!d||!g)&&" \
    "(!a||!d||!h)&&(!a||!e||!f)&&(!a||!e||!g)&&(!a||!e||!h)&&(!a||!f||!g)&&(!a||!f||!h)&&(!a||!g||!h)&&" \
    "(a||b||c||d||e||f||g)&&(a||b||c||d||e||f||h)&&(a||b||c||d||e||g||h)&&(a||b||c||d||f||g||h)&&" \
    "(a||b||c||e||f||g||h)&&(a||b||d||e||f||g||h)&&(a||c||d||e||f||g||h)&&(!b||!c||!d)&&(!b||!c||!e)&&" \
    "(!b||!c||!f)&&(!b||!c||!g)&&(!b||!c||!h)&&(!b||!d||!e)&&(!b||!d||!f)&&(!b||!d||!g)&&(!b||!d||!h)&&" \
    "(!b||!e||!f)&&(!b||!e||!g)&&(!b||!e||!h)&&(!b||!f||!g)&&(!b||!f||!h)&&(!b||!g||!h)&&(b||c||d||e||f||g||h)" \
    "&&" \
    "(!c||!d||!e)&&(!c||!d||!f)&&(!c||!d||!g)&&(!c||!d||!h)&&(!c||!e||!f)&&(!c||!e||!g)&&(!c||!e||!h)&&" \
    "(!c||!f||!g)&&(!c||!f||!h)&&(!c||!g||!h)&&(!d||!e||!f)&&(!d||!e||!g)&&(!d||!e||!h)&&(!d||!f||!g)&&" \
    "(!d||!f||!h)&&(!d||!g||!h)&&(!e||!f||!g)&&(!e||!f||!h)&&(!e||!g||!h)&&(!f||!g||!h)"
    return mathematica_to_CNF(s, a)

clauses=POPCNT2(["1","2","3","4","5","6","7","8"])

f=open("tmp.cnf", "w")
f.write ("p cnf 8 "+str(len(clauses))+ "\n")
for c in clauses:
    f.write(c+" 0\n")
f.close()
```

It replaces a/b/c/... variables to the variable names passed (1/2/3...), reworks syntax, etc. Here is a result:

```
p cnf 8 64
-1 -2 -3 0
-1 -2 -4 0
-1 -2 -5 0
-1 -2 -6 0
-1 -2 -7 0
-1 -2 -8 0
-1 -3 -4 0
-1 -3 -5 0
-1 -3 -6 0
-1 -3 -7 0
-1 -3 -8 0
-1 -4 -5 0
-1 -4 -6 0
-1 -4 -7 0
-1 -4 -8 0
-1 -5 -6 0
-1 -5 -7 0
-1 -5 -8 0
-1 -6 -7 0
-1 -6 -8 0
-1 -7 -8 0
1 2 3 4 5 6 7 0
1 2 3 4 5 6 8 0
1 2 3 4 5 7 8 0
1 2 3 4 6 7 8 0
1 2 3 5 6 7 8 0
1 2 4 5 6 7 8 0
1 3 4 5 6 7 8 0
-2 -3 -4 0
-2 -3 -5 0
-2 -3 -6 0
-2 -3 -7 0
-2 -3 -8 0
-2 -4 -5 0
-2 -4 -6 0
-2 -4 -7 0
-2 -4 -8 0
-2 -5 -6 0
-2 -5 -7 0
```



```
-2 -5 -8 0
-2 -6 -7 0
-2 -6 -8 0
-2 -7 -8 0
2 3 4 5 6 7 8 0
-3 -4 -5 0
-3 -4 -6 0
-3 -4 -7 0
-3 -4 -8 0
-3 -5 -6 0
-3 -5 -7 0
-3 -5 -8 0
-3 -6 -7 0
-3 -6 -8 0
-3 -7 -8 0
-4 -5 -6 0
-4 -5 -7 0
-4 -5 -8 0
-4 -6 -7 0
-4 -6 -8 0
-4 -7 -8 0
-5 -6 -7 0
-5 -6 -8 0
-5 -7 -8 0
-6 -7 -8 0
```

I can run it:

```
% minisat -verb=0 tst1.cnf results.txt
SATISFIABLE

% cat results.txt
SAT
1 -2 -3 -4 -5 -6 -7 8 0
```

The variable name in results lacking minus sign is *True*. Variable name with minus sign is *False*. We see there are just two variables are *True*: 1 and 8. This is indeed correct: MiniSat solver found a condition, for which our function returns *True*. Zero at the end is just a terminal symbol which means nothing.

We can ask MiniSat for another solution, by adding current solution to the input CNF file, but with all variables negated:

```
...
-5 -6 -8 0
-5 -7 -8 0
-6 -7 -8 0
-1 2 3 4 5 6 7 -8 0
```

In plain English language, this means “give me ANY solution which can satisfy all clauses, but also not equal to the last clause we’ve just added”.

MiniSat, indeed, found another solution, again, with only 2 variables equal to *True*:

```
% minisat -verb=0 tst2.cnf results.txt
SATISFIABLE

% cat results.txt
SAT
1 2 -3 -4 -5 -6 -7 -8 0
```

By the way, *population count* function for 8 neighbours (POPCNT8) in CNF form is simplest:

```
a&&b&&c&&d&&e&&f&&g&&h
```

Indeed: it’s true if all 8 input bits are *True*.

The function for 0 neighbours (POPCNT0) is also simple:

```
!a&&!b&&!c&&!d&&!e&&!f&&!g&&!h
```

It means, it will return *True*, if all input variables are *False*.

By the way, POPCNT1 function is also simple:

```
(!a||!b)&&(!a||!c)&&(!a||!d)&&(!a||!e)&&(!a||!f)&&(!a||!g)&&(!a||!h)&&(a||b||c||d||e||f||g||h)&&
(!b||!c)&&(!b||!d)&&(!b||!e)&&(!b||!f)&&(!b||!g)&&(!b||!h)&&(!c||!d)&&(!c||!e)&&(!c||!f)&&(!c||!g)&&
(!c||!h)&&(!d||!e)&&(!d||!f)&&(!d||!g)&&(!d||!h)&&(!e||!f)&&(!e||!g)&&(!e||!h)&&(!f||!g)&&(!f||!h)&&(!g||!h)
```

There is just enumeration of all possible pairs of 8 variables (a/b, a/c, a/d, etc), which implies: no two bits must be present simultaneously in each possible pair. And there is another clause: "(a||b||c||d||e||f||g||h)", which implies: at least one bit must be present among 8 variables.

And yes, you can ask Mathematica for finding [CNF](#) expressions for any other truth table.

### 13.7.2 Minesweeper

Now we can use Mathematica to generate all *population count* functions for 0..8 neighbours.

For 9 · 9 Minesweeper matrix including invisible border, there will be 11 · 11 = 121 variables, mapped to Minesweeper matrix like this:

```

1  2  3  4  5  6  7  8  9 10 11
12 13 14 15 16 17 18 19 20 21 22
23 24 25 26 27 28 29 30 31 32 33
34 35 36 37 38 39 40 41 42 43 44

...

100 101 102 103 104 105 106 107 108 109 110
111 112 113 114 115 116 117 118 119 120 121

```

Then we write a Python script which stacks all *population count* functions: each function for each known number of neighbours (digit on Minesweeper field). Each POPCNTx() function takes list of variable numbers and outputs list of clauses to be added to the final [CNF](#) file.

As of empty cells, we also add them as clauses, but with minus sign, which means, the variable must be *False*. Whenever we try to place bomb, we add its variable as clause without minus sign, this means the variable must be *True*.

Then we execute external minisat process. The only thing we need from it is exit code. If an input [CNF](#) is **UNSAT**, it returns 20:

We use here the information from the previous solving of Minesweeper: [7.9](#).

```

#!/usr/bin/python

import subprocess

WIDTH=9
HEIGHT=9
VARS_TOTAL=(WIDTH+2)*(HEIGHT+2)

known=[
"01?10001?",
"01?100011",
"011100000",
"000000000",
"111110011",
"?????1001?",
"?????3101?",
"?????211?",
"?????????" ]

def mathematica_to_CNF (s, a):
    s=s.replace("a", a[0]).replace("b", a[1]).replace("c", a[2]).replace("d", a[3])
    s=s.replace("e", a[4]).replace("f", a[5]).replace("g", a[6]).replace("h", a[7])
    s=s.replace("!", "-").replace("||", " ").replace("(", "").replace(")", "")
    s=s.split("&&")
    return s

def POPCNT0 (a):
    s="!a&&!b&&!c&&!d&&!e&&!f&&!g&&!h"
    return mathematica_to_CNF(s, a)

def POPCNT1 (a):
    s="(!a||!b)&&(!a||!c)&&(!a||!d)&&(!a||!e)&&(!a||!f)&&(!a||!g)&&(!a||!h)&&(a||b||c||d||e||f||g||h)&&" \
    "(!b||!c)&&(!b||!d)&&(!b||!e)&&(!b||!f)&&(!b||!g)&&(!b||!h)&&(!c||!d)&&(!c||!e)&&(!c||!f)&&(!c||!g)&&" \
    "(!c||!h)&&(!d||!e)&&(!d||!f)&&(!d||!g)&&(!d||!h)&&(!e||!f)&&(!e||!g)&&(!e||!h)&&(!f||!g)&&(!f||!h)&&(!g||!h)"
    return mathematica_to_CNF(s, a)

def POPCNT2 (a):
    s="(!a||!b||!c)&&(!a||!b||!d)&&(!a||!b||!e)&&(!a||!b||!f)&&(!a||!b||!g)&&(!a||!b||!h)&&(!a||!c||!d)&&" \
    "(!a||!c||!e)&&(!a||!c||!f)&&(!a||!c||!g)&&(!a||!c||!h)&&(!a||!d||!e)&&(!a||!d||!f)&&(!a||!d||!g)&&" \

```

```

"!a||d||h)&&!a||e||f)&&!a||e||g)&&!a||e||h)&&!a||f||g)&&!a||f||h)&&!a||g||h)&&" \
"(a||b||c||d||e||f||g)&&(a||b||c||d||e||f||h)&&(a||b||c||d||e||g||h)&&(a||b||c||d||f||g||h)&&" \
"(a||b||c||e||f||g||h)&&(a||b||d||e||f||g||h)&&(a||c||d||e||f||g||h)&&(!b||c||d)&&(!b||c||e)&&" \
"(!b||c||f)&&(!b||c||g)&&(!b||c||h)&&(!b||d||e)&&(!b||d||f)&&(!b||d||g)&&(!b||d||h)&&" \
"(!b||e||f)&&(!b||e||g)&&(!b||e||h)&&(!b||f||g)&&(!b||f||h)&&(!b||g||h)&&(b||c||d||e||f||g||h)
&&" \
"(!c||d||e)&&(!c||d||f)&&(!c||d||g)&&(!c||d||h)&&(!c||e||f)&&(!c||e||g)&&(!c||e||h)&&" \
"(!c||f||g)&&(!c||f||h)&&(!c||g||h)&&(!d||e||f)&&(!d||e||g)&&(!d||e||h)&&(!d||f||g)&&" \
"(!d||f||h)&&(!d||g||h)&&(!e||f||g)&&(!e||f||h)&&(!e||g||h)&&(!f||g||h)"
return mathematica_to_CNF(s, a)

```

def POPCNT3 (a):

```

s="(!a||b||c||d)&&!a||b||c||e)&&!a||b||c||f)&&!a||b||c||g)&&!a||b||c||h)&&" \
"(!a||b||d||e)&&!a||b||d||f)&&!a||b||d||g)&&!a||b||d||h)&&!a||b||e||f)&&" \
"(!a||b||e||g)&&!a||b||e||h)&&!a||b||f||g)&&!a||b||f||h)&&!a||b||g||h)&&" \
"(!a||c||d||e)&&!a||c||d||f)&&!a||c||d||g)&&!a||c||d||h)&&!a||c||e||f)&&" \
"(!a||c||e||g)&&!a||c||e||h)&&!a||c||f||g)&&!a||c||f||h)&&!a||c||g||h)&&" \
"(!a||d||e||f)&&!a||d||e||g)&&!a||d||e||h)&&!a||d||f||g)&&!a||d||f||h)&&" \
"(!a||d||f||g||h)&&!a||e||f||g)&&!a||e||f||h)&&!a||e||g||h)&&!a||f||g||h)&&" \
"(a||b||c||d||e||f)&&(a||b||c||d||e||g)&&(a||b||c||d||e||h)&&(a||b||c||d||f||g)&&(a||b||c||d||f||h)&&" \
"(a||b||c||d||g||h)&&(a||b||c||e||f||g)&&(a||b||c||e||f||h)&&(a||b||c||e||g||h)&&(a||b||c||f||g||h)&&" \
"(a||b||d||e||f||g)&&(a||b||d||e||f||h)&&(a||b||d||e||g||h)&&(a||b||d||f||g||h)&&(a||b||e||f||g||h)&&" \
"(a||c||d||e||f||g)&&(a||c||d||e||f||h)&&(a||c||d||e||g||h)&&(a||c||d||f||g||h)&&(a||c||e||f||g||h)&&" \
"(!a||d||e||f||g||h)&&(!b||c||d||e||f||g)&&(!b||c||d||e||h)&&(!b||c||d||f||g)&&(!b||c||d||f||h)&&" \
"(!b||c||e||f||g)&&(!b||c||e||h)&&(!b||c||f||g)&&(!b||c||f||h)&&(!b||c||e||f||h)&&(!b||e||f||g||h)&&" \
"(!b||d||e||f||g)&&(!b||d||e||h)&&(!b||d||f||g)&&(!b||d||f||h)&&(!b||d||e||f||h)&&(!b||e||f||g||h)&&" \
"(!b||f||g||h)&&(!b||f||h)&&(!b||f||e||g)&&(!b||f||e||h)&&(!b||f||e||f||g||h)&&(!b||g||h)&&" \
"(b||c||e||f||g||h)&&(b||d||e||f||g||h)&&(!c||d||e||f||g)&&(!c||d||e||f||h)&&(!c||d||e||g||h)&&" \
"(!c||d||f||g)&&(!c||d||f||h)&&(!c||d||g||h)&&(!c||e||f||g)&&(!c||e||f||h)&&(!c||e||f||g||h)&&" \
"(!c||e||g||h)&&(!c||f||g||h)&&(c||d||e||f||g||h)&&(!d||e||f||g)&&(!d||e||f||h)&&" \
"(!d||e||g||h)&&(!d||f||g||h)&&(!e||f||g||h)"
return mathematica_to_CNF(s, a)

```

def POPCNT4 (a):

```

s="(!a||b||c||d||e)&&!a||b||c||d||f)&&!a||b||c||d||g)&&!a||b||c||d||h)&&" \
"(!a||b||c||e||f)&&!a||b||c||e||g)&&!a||b||c||e||h)&&!a||b||c||f||g)&&" \
"(!a||b||c||f||h)&&!a||b||c||g||h)&&!a||b||d||e||f)&&!a||b||d||e||g)&&" \
"(!a||b||d||e||h)&&!a||b||d||f||g)&&!a||b||d||f||h)&&!a||b||d||g||h)&&" \
"(!a||b||e||f||g)&&!a||b||e||f||h)&&!a||b||e||g||h)&&!a||b||f||g||h)&&" \
"(!a||c||d||e||f)&&!a||c||d||e||g)&&!a||c||d||e||h)&&!a||c||d||f||g)&&" \
"(!a||c||d||f||h)&&!a||c||d||g||h)&&!a||c||e||f||g)&&!a||c||e||f||h)&&" \
"(!a||c||e||g||h)&&!a||c||f||g||h)&&!a||d||e||f||g)&&!a||d||e||f||h)&&" \
"(!a||d||e||g||h)&&!a||d||f||g||h)&&!a||e||f||g||h)&&(a||b||c||d||e)&&(a||b||c||d||f)&&" \
"(a||b||c||d||g)&&(a||b||c||d||h)&&(a||b||c||e||f)&&(a||b||c||e||g)&&(a||b||c||e||h)&&(a||b||c||f||g)&&" \
"(a||b||c||f||h)&&(a||b||c||g||h)&&(a||b||d||e||f)&&(a||b||d||e||g)&&(a||b||d||e||h)&&(a||b||d||f||g)&&" \
"(a||b||d||f||h)&&(a||b||d||g||h)&&(a||b||e||f||g)&&(a||b||e||f||h)&&(a||b||e||g||h)&&(a||b||f||g||h)&&" \
"(a||c||d||e||f)&&(a||c||d||e||g)&&(a||c||d||e||h)&&(a||c||d||f||g)&&(a||c||d||f||h)&&(a||c||d||g||h)&&" \
"(a||c||e||f||g)&&(a||c||e||f||h)&&(a||c||e||g||h)&&(a||c||f||g||h)&&(a||d||e||f||g)&&(a||d||e||f||h)&&" \
"(a||d||e||g||h)&&(a||d||f||g||h)&&(a||e||f||g||h)&&(!b||c||d||e||f)&&(!b||c||d||e||g)&&" \
"(!b||c||d||e||f||h)&&(!b||c||d||f||g)&&(!b||c||d||f||h)&&(!b||c||d||e||f||g||h)&&" \
"(!b||c||e||f||g)&&(!b||c||e||f||h)&&(!b||c||e||g||h)&&(!b||c||e||f||h)&&(!b||c||e||f||g||h)&&" \
"(!b||d||e||f||g)&&(!b||d||e||f||h)&&(!b||d||e||g||h)&&(!b||d||e||f||g||h)&&(!b||d||e||f||g||h)&&" \
"(!b||e||f||g||h)&&(b||c||d||e||f)&&(b||c||d||e||g)&&(b||c||d||e||h)&&(b||c||d||f||g)&&" \
"(b||c||d||f||h)&&(b||c||d||g||h)&&(b||c||e||f||g)&&(b||c||e||f||h)&&(b||c||e||g||h)&&" \
"(b||c||f||g||h)&&(b||d||e||f||g)&&(b||d||e||f||h)&&(b||d||e||g||h)&&(b||d||f||g||h)&&" \
"(b||e||f||g||h)&&(!c||d||e||f||g)&&(!c||d||e||f||h)&&(!c||d||e||g||h)&&(!c||d||e||f||g||h)&&" \
"(!c||d||f||g||h)&&(!c||e||f||g||h)&&(c||d||e||f||g)&&(c||d||e||f||h)&&(c||d||e||g||h)&&" \
"(c||d||f||g||h)&&(c||e||f||g||h)&&(!d||e||f||g||h)&&(d||e||f||g||h)"
return mathematica_to_CNF(s, a)

```

def POPCNT5 (a):

```

s="(!a||b||c||d||e||f)&&!a||b||c||d||e||g)&&!a||b||c||d||e||h)&&" \
"(!a||b||c||d||f||g)&&!a||b||c||d||f||h)&&!a||b||c||d||g||h)&&" \
"(!a||b||c||e||f||g)&&!a||b||c||e||f||h)&&!a||b||c||e||g||h)&&" \
"(!a||b||c||f||g||h)&&!a||b||d||e||f||g)&&!a||b||d||e||f||h)&&" \
"(!a||b||d||e||g||h)&&!a||b||d||f||g||h)&&!a||b||e||f||g||h)&&" \
"(!a||c||d||e||f||g)&&!a||c||d||e||f||h)&&!a||c||d||e||g||h)&&" \
"(!a||c||d||f||g||h)&&!a||c||e||f||g||h)&&!a||d||e||f||g||h)&&" \
"(a||b||c||d)&&(a||b||c||e)&&(a||b||c||f)&&(a||b||c||g)&&(a||b||c||h)&&(a||b||d||e)&&" \
"(a||b||d||f)&&(a||b||d||g)&&(a||b||d||h)&&(a||b||e||f)&&(a||b||e||g)&&(a||b||e||h)&&" \
"(a||b||f||g)&&(a||b||f||h)&&(a||b||g||h)&&(a||c||d||e)&&(a||c||d||f)&&(a||c||d||g)&&" \
"(a||c||d||h)&&(a||c||e||f)&&(a||c||e||g)&&(a||c||e||h)&&(a||c||f||g)&&(a||c||f||h)&&" \
"(a||c||g||h)&&(a||d||e||f)&&(a||d||e||g)&&(a||d||e||h)&&(a||d||f||g)&&(a||d||f||h)&&" \
"(a||d||g||h)&&(a||e||f||g)&&(a||e||f||h)&&(a||e||g||h)&&(a||f||g||h)&&(!b||c||d||e||f||g)&&" \
"(!b||c||d||e||f||h)&&(!b||c||d||e||g||h)&&(!b||c||d||f||g||h)&&"

```

```

    "(!b||c||e||f||g||h)&&(!b||d||e||f||g||h)&&(b||c||d||e)&&(b||c||d||f)&&" \
    "(b||c||d||g)&&(b||c||d||h)&&(b||c||e||f)&&(b||c||e||g)&&(b||c||e||h)&&(b||c||f||g)&&" \
    "(b||c||f||h)&&(b||c||g||h)&&(b||d||e||f)&&(b||d||e||g)&&(b||d||e||h)&&(b||d||f||g)&&" \
    "(b||d||f||h)&&(b||d||g||h)&&(b||e||f||g)&&(b||e||f||h)&&(b||e||g||h)&&(b||f||g||h)&&" \
    "(!c||d||e||f||g||h)&&(c||d||e||f)&&(c||d||e||g)&&(c||d||e||h)&&(c||d||f||g)&&" \
    "(c||d||f||h)&&(c||d||g||h)&&(c||e||f||g)&&(c||e||f||h)&&(c||e||g||h)&&(c||f||g||h)&&" \
    "(d||e||f||g)&&(d||e||f||h)&&(d||e||g||h)&&(d||f||g||h)&&(e||f||g||h)"
    return mathematica_to_CNF(s, a)

def POPCNT6 (a):
    s="(!a||b||c||d||e||f||g)&&(!a||b||c||d||e||f||h)&&(!a||b||c||d||e||g||h)&&" \
    "(!a||b||c||d||e||f||g||h)&&(!a||b||c||e||f||g||h)&&(!a||b||d||e||f||g||h)&&" \
    "(!a||c||d||e||f||g||h)&&(a||b||c)&&(a||b||d)&&(a||b||e)&&(a||b||f)&&(a||b||g)&&(a||b||h)&&" \
    "(a||c||d)&&(a||c||e)&&(a||c||f)&&(a||c||g)&&(a||c||h)&&(a||d||e)&&(a||d||f)&&(a||d||g)&&" \
    "(a||d||h)&&(a||e||f)&&(a||e||g)&&(a||e||h)&&(a||f||g)&&(a||f||h)&&(a||g||h)&&" \
    "(!b||c||d||e||f||g||h)&&(b||c||d)&&(b||c||e)&&(b||c||f)&&(b||c||g)&&(b||c||h)&&(b||d||e)&&" \
    "(b||d||f)&&(b||d||g)&&(b||d||h)&&(b||e||f)&&(b||e||g)&&(b||e||h)&&(b||f||g)&&(b||f||h)&&(b||g||h)&&" \
    "(c||d||e)&&(c||d||f)&&(c||d||g)&&(c||d||h)&&(c||e||f)&&(c||e||g)&&(c||e||h)&&(c||f||g)&&(c||f||h)&&" \
    "(c||g||h)&&(d||e||f)&&(d||e||g)&&(d||e||h)&&(d||f||g)&&(d||f||h)&&(d||g||h)&&" \
    "(e||f||g)&&(e||f||h)&&(e||g||h)&&(f||g||h)"
    return mathematica_to_CNF(s, a)

def POPCNT7 (a):
    s="(!a||b||c||d||e||f||g||h)&&(a||b)&&(a||c)&&(a||d)&&(a||e)&&(a||f)&&(a||g)&&(a||h)&&(b||c)&&" \
    "(b||d)&&(b||e)&&(b||f)&&(b||g)&&(b||h)&&(c||d)&&(c||e)&&(c||f)&&(c||g)&&(c||h)&&(d||e)&&(d||f)&&(d||g)&&" \
    "(d||h)&&(e||f)&&(e||g)&&(e||h)&&(f||g)&&(f||h)&&(g||h)"
    return mathematica_to_CNF(s, a)

def POPCNT8 (a):
    s="a&&b&&c&&d&&e&&f&&g&&h"
    return mathematica_to_CNF(s, a)

POPCNT_functions=[POPCNT0, POPCNT1, POPCNT2, POPCNT3, POPCNT4, POPCNT5, POPCNT6, POPCNT7, POPCNT8]

def coords_to_var (row, col):
    # we always use SAT variables as strings, anyway.
    # the 1st variables is 1, not 0
    return str(row*(WIDTH+2)+col+1)

def chk_bomb(row, col):
    clauses=[]

    # make empty border
    # all variables are negated (because they must be False)
    for c in range(WIDTH+2):
        clauses.append ("-"+coords_to_var(0,c))
        clauses.append ("-"+coords_to_var(HEIGHT+1,c))
    for r in range(HEIGHT+2):
        clauses.append ("-"+coords_to_var(r,0))
        clauses.append ("-"+coords_to_var(r,WIDTH+1))

    for r in range(1,HEIGHT+1):
        for c in range(1,WIDTH+1):
            t=known[r-1][c-1]
            if t in "012345678":
                # cell at r, c is empty (False):
                clauses.append ("-"+coords_to_var(r,c))
                # we need an empty border so the following expression would work for all possible cells:
                neighbours=[coords_to_var(r-1, c-1), coords_to_var(r-1, c), coords_to_var(r-1, c+1),
                    coords_to_var(r, c-1),
                    coords_to_var(r, c+1), coords_to_var(r+1, c-1), coords_to_var(r+1, c), coords_to_var(r
                        +1, c+1)]
                clauses=clauses+POPCNT_functions[int(t)](neighbours)

    # place a bomb
    clauses.append (coords_to_var(row,col))

    f=open("tmp.cnf", "w")
    f.write ("p cnf "+str(VARS_TOTAL)+" "+str(len(clauses))+"\n")
    for c in clauses:
        f.write(c+" 0\n")
    f.close()

    child = subprocess.Popen(["minisat", "tmp.cnf"], stdout=subprocess.PIPE)
    child.wait()

```

```
# 10 is SAT, 20 is UNSAT
if child.returncode==20:
    print "row=%d, col=%d, unsat!" % (row, col)

for r in range(1,HEIGHT+1):
    for c in range(1,WIDTH+1):
        if known[r-1][c-1]=="?":
            chk_bomb(r, c)
```

( [https://github.com/dennis714/SAT\\_SMT\\_article/blob/master/SAT/minesweeper/minesweeper\\_SAT.py](https://github.com/dennis714/SAT_SMT_article/blob/master/SAT/minesweeper/minesweeper_SAT.py) )

The output CNF file can be large, up to  $\approx 2000$  clauses, or more, here is an example: [https://github.com/dennis714/SAT\\_SMT\\_article/blob/master/SAT/minesweeper/sample.cnf](https://github.com/dennis714/SAT_SMT_article/blob/master/SAT/minesweeper/sample.cnf).

Anyway, it works just like my previous Z3Py script:

```
row=1, col=3, unsat!
row=6, col=2, unsat!
row=6, col=3, unsat!
row=7, col=4, unsat!
row=7, col=9, unsat!
row=8, col=9, unsat!
```

...but it runs way faster, even considering overhead of executing external program. Perhaps, Z3Py version could be optimized much better?

The files, including Wolfram Mathematica notebook: [https://github.com/dennis714/SAT\\_SMT\\_article/tree/master/SAT/minesweeper](https://github.com/dennis714/SAT_SMT_article/tree/master/SAT/minesweeper).

## 13.8 Conway's "Game of Life"

### 13.8.1 Reversing back state of "Game of Life"

How could we reverse back a known state of GoL? This can be solved by brute-force, but this is extremely slow and inefficient.

Let's try to use SAT solver.

First, we need to define a function which will tell, if the new cell will be created/born, preserved/stay or died. Quick refresher: cell is born if it has 3 neighbours, it stays alive if it has 2 or 3 neighbours, it dies in any other case.

This is how I can define a function reflecting state of a new cell in the next state:

```
if center==true:
    return popcnt2(neighbours) || popcnt3(neighbours)
if center==false
    return popcnt3(neighbours)
```

We can get rid of "if" construction:

```
result=(center=true && (popcnt2(neighbours) || popcnt3(neighbours))) || (center=false && popcnt3(neighbours))
```

...where "center" is state of central cell, "neighbours" are 8 neighbouring cells, popcnt2 is a function which returns True if it has exactly 2 bits on input, popcnt3 is the same, but for 3 bits (just like these were used in my "Minesweeper" example (13.7)).

Using Wolfram Mathematica, I first create all helper functions and truth table for the function, which returns *true*, if a cell must be present in the next state, or *false* if not:

```
In[1]:= popcount[n_Integer]:=IntegerDigits[n,2] // Total
In[2]:= popcount2[n_Integer]:=Equal[popcount[n],2]
In[3]:= popcount3[n_Integer]:=Equal[popcount[n],3]
In[4]:= newcell[center_Integer,neighbours_Integer]:=(center==1 && (popcount2[neighbours]|| popcount3[neighbours]))||
(center==0 && popcount3[neighbours])

In[13]:= NewCellIsTrue=Flatten[Table[Join[{center},PadLeft[IntegerDigits[neighbours,2],8]] ->
Boole[newcell[center, neighbours]],{neighbours,0,255},{center,0,1}]]

Out[13]= {{0,0,0,0,0,0,0,0,0}->0,
{1,0,0,0,0,0,0,0,0}->0,
```

```
{0,0,0,0,0,0,0,0,0,1}->0,
{1,0,0,0,0,0,0,0,0,1}->0,
{0,0,0,0,0,0,0,0,1,0}->0,
{1,0,0,0,0,0,0,0,1,0}->0,
{0,0,0,0,0,0,0,0,1,1}->0,
{1,0,0,0,0,0,0,0,1,1}->1,
...
```

Now we can create a **CNF**-expression out of truth table:

```
In[14]:= BooleanConvert[BooleanFunction[NewCellIsTrue,{center,a,b,c,d,e,f,g,h}], "CNF"]
Out[14]= (!a||b||c||d)&&(!a||b||c||e)&&(!a||b||c||f)&&(!a||b||c||g)&&(!a||b||c||h)&&
(!a||b||d||e)&&(!a||b||d||f)&&(!a||b||d||g)&&(!a||b||d||h)&&(!a||b||e||f)&&
(!a||b||e||g)&&(!a||b||e||h)&&(!a||b||f||g)&&(!a||b||f||h)&&(!a||b||g||h)&&
(!a||c||d||e)&&(!a||c||d||f)&&(!a||c||d||g)&&(!a||c||d||h)&&(!a||c||e||f)&&
(!a||c||e||g)&&(!a||c||e||h)&&(!a||c||f||g)&&(!a||c||f||h)&&
...
```

Also, we need a second function, *inverted one*, which will return *true* if the cell must be absent in the next state, or *false* otherwise:

```
In[15]:= NewCellIsFalse=Flatten[Table[Join[{center},PadLeft[IntegerDigits[neighbours,2],8]] ->
Boole[Not[newcell[center, neighbours]]],{neighbours,0,255},{center,0,1}]]
Out[15]= {{0,0,0,0,0,0,0,0,0,0}->1,
{1,0,0,0,0,0,0,0,0,0}->1,
{0,0,0,0,0,0,0,0,0,1}->1,
{1,0,0,0,0,0,0,0,0,1}->1,
{0,0,0,0,0,0,0,0,1,0}->1,
...

In[16]:= BooleanConvert[BooleanFunction[NewCellIsFalse,{center,a,b,c,d,e,f,g,h}], "CNF"]
Out[16]= (!a||b||c||d||e||f||g||h)&&(!a||b||c||d||e||f||g||h)&&(!a||b||c||d||e||f||g||h)&&
(!a||b||c||d||e||f||g||h)&&(!a||b||c||d||e||f||g||h)&&(!a||b||c||d||e||f||g||h)&&
(!a||b||center||d||e||f||g||h)&&(!a||b||c||d||e||f||g||h)&&(!a||b||c||d||e||f||g||h)&&
(!a||b||c||d||e||f||g||h)&&(!a||b||c||d||e||f||g||h)&&(!a||b||c||d||e||f||g||h)&&
(!a||b||c||d||e||f||g||h)&&(!a||b||c||d||e||f||g||h)&&(!a||b||c||d||e||f||g||h)&&
...
```

Using the very same way as in my “Minesweeper” example, I can convert **CNF** expression to list of clauses:

```
def mathematica_to_CNF (s, center, a):
    s=s.replace("center", center)
    s=s.replace("a", a[0]).replace("b", a[1]).replace("c", a[2]).replace("d", a[3])
    s=s.replace("e", a[4]).replace("f", a[5]).replace("g", a[6]).replace("h", a[7])
    s=s.replace("!", "-").replace("||", " ").replace("(", "").replace(")", "")
    s=s.split("&&")
    return s
```

And again, as in “Minesweeper”, there is an invisible border, to make processing simpler. **SAT** variables are also numbered as in previous example:

```
1 2 3 4 5 6 7 8 9 10 11
12 13 14 15 16 17 18 19 20 21 22
23 24 25 26 27 28 29 30 31 32 33
34 35 36 37 38 39 40 41 42 43 44
...
100 101 102 103 104 105 106 107 108 109 110
111 112 113 114 115 116 117 118 119 120 121
```

Also, there is a visible border, always fixed to *False*, to make things simpler.

Now the working source code. Whenever we encounter “\*” in `final_state[]`, we add clauses generated by `cell_is_true()` function, or `cell_is_false()` if otherwise. When we get a solution, it is negated and added to the list of clauses, so when minisat is executed next time, it will skip solution which was already printed.

```
...
def cell_is_false (center, a):
```



```

W=len(final_state[0]) # WIDTH

print "HEIGHT=", H, "WIDTH=", W

VARS_TOTAL=W*H+1
VAR_FALSE=str(VARS_TOTAL)

def try_again (clauses):
    # rules for the main part of grid
    for r in range(H):
        for c in range(W):
            if final_state[r][c]=="*":
                clauses=clauses+cell_is_true(coords_to_var(r, c, H, W), get_neighbours(r, c, H, W))
            else:
                clauses=clauses+cell_is_false(coords_to_var(r, c, H, W), get_neighbours(r, c, H, W))

    # cells behind visible grid must always be false:
    for c in range(-1, W+1):
        for r in [-1,H]:
            clauses=clauses+cell_is_false(coords_to_var(r, c, H, W), get_neighbours(r, c, H, W))
    for c in [-1,W]:
        for r in range(-1, H+1):
            clauses=clauses+cell_is_false(coords_to_var(r, c, H, W), get_neighbours(r, c, H, W))

    write_CNF("tmp.cnf", clauses, VARS_TOTAL)

    print "%d clauses" % len(clauses)

    solution=run_minisat ("tmp.cnf")
    os.remove("tmp.cnf")
    if solution==None:
        print "unsat!"
        exit(0)

    grid=SAT_solution_to_grid(solution, H, W)

    print_grid(grid)
    write_RLE(grid)

    return grid

clauses=[]
# always false:
clauses.append ("-"+VAR_FALSE)

while True:
    solution=try_again(clauses)
    clauses.append(negate_clause(grid_to_clause(solution, H, W)))
    print ""

```

( [https://github.com/dennis714/SAT\\_SMT\\_article/blob/master/SAT/GoL/reverse1.py](https://github.com/dennis714/SAT_SMT_article/blob/master/SAT/GoL/reverse1.py) )  
 Here is the result:

```

HEIGHT= 3 WIDTH= 3
2525 clauses
.*.
*.*
.*.
1.rle written

2526 clauses
.**
*..
*.*
2.rle written

2527 clauses
**
..*
*.*
3.rle written

2528 clauses
*.*
*..
.**

```



```
4.rle written
```

```
2529 clauses
```

```
*.*  
.*  
**.
```

```
5.rle written
```

```
2530 clauses
```

```
*.*  
.*  
*.*
```

```
6.rle written
```

```
2531 clauses
```

```
unsat!
```

The first result is the same as initial state. Indeed: this is “still life”, i.e., state which will never change, and it is correct solution. The last solution is also valid.

Now the problem: 2nd, 3rd, 4th and 5th solutions are equivalent to each other, they just mirrored or rotated. In fact, this is reflectional<sup>107</sup> (like in mirror) and rotational<sup>108</sup> symmetries. We can solve this easily: we will take each solution, reflect and rotate it and add them negated to the list of clauses, so minisat will skip them during its work:

```
...
```

```
while True:
```

```
    solution=try_again(clauses)
```

```
    clauses.append(negate_clause(grid_to_clause(solution, H, W)))
```

```
    clauses.append(negate_clause(grid_to_clause(reflect_vertically(solution), H, W)))
```

```
    clauses.append(negate_clause(grid_to_clause(reflect_horizontally(solution), H, W)))
```

```
    # is this square?
```

```
    if W==H:
```

```
        clauses.append(negate_clause(grid_to_clause(rotate_square_array(solution,1), H, W)))
```

```
        clauses.append(negate_clause(grid_to_clause(rotate_square_array(solution,2), H, W)))
```

```
        clauses.append(negate_clause(grid_to_clause(rotate_square_array(solution,3), H, W)))
```

```
    print ""
```

```
...
```

( [https://github.com/dennis714/SAT\\_SMT\\_article/blob/master/SAT/GoL/reverse2.py](https://github.com/dennis714/SAT_SMT_article/blob/master/SAT/GoL/reverse2.py) )

Functions `reflect_vertically()`, `reflect_horizontally` and `rotate_squarearray()` are simple array manipulation routines.

Now we get just 3 solutions:

```
HEIGHT= 3 WIDTH= 3
```

```
2525 clauses
```

```
*.*  
*.*  
.*.
```

```
1.rle written
```

```
2531 clauses
```

```
.*  
*.*  
*.*
```

```
2.rle written
```

```
2537 clauses
```

```
*.*  
.*  
*.*
```

```
3.rle written
```

```
2543 clauses
```

```
unsat!
```

This one has only one single ancestor:

<sup>107</sup>[https://en.wikipedia.org/wiki/Reflection\\_symmetry](https://en.wikipedia.org/wiki/Reflection_symmetry)

<sup>108</sup>[https://en.wikipedia.org/wiki/Rotational\\_symmetry](https://en.wikipedia.org/wiki/Rotational_symmetry)

```

final_state=[
" * ",
" * ",
" * "]
_PRE_END

_PRE_BEGIN
HEIGHT= 3 WIDTH= 3
2503 clauses
...
***
...
1.rle written

2509 clauses
unsat!

```

This is oscillator, of course.  
How many states can lead to such picture?

```

final_state=[
" * ",
"  ",
" ** ",
" * ",
" * ",
" *** "]

```

28, these are few:

```

HEIGHT= 6 WIDTH= 5
5217 clauses
.*.*.
.*.*.
.**.*
.*.*.
.*.*.
.*.*.
.**.*
1.rle written

5220 clauses
.*.*.
.*.*.
.**.*
.*.*.
*.*.*
.**.*
2.rle written

5223 clauses
.*.*
.**.*
.**.*
*.*.*
.**.*
3.rle written

5226 clauses
.*.*
.**.*
.**.*
*.*.*
.*.*
.**.*
4.rle written
...

```

Now the biggest, "space invader":

```

final_state=[
"          ",
"   *      *   ",
"   *      *   "]

```

HEIGHT= 10 WIDTH= 13

16472 clauses

```
2.rle written
```

16475 clauses

```
3.rle written
```

I don't know how many possible states can lead to "space invader", perhaps, too many. Had to stop it. And it slows down during execution, because number of clauses is increasing (because of negating solutions addition).

### 13.8.2 Finding “still lives”

“Still life” in terms of GoL is a state which doesn’t change at all.

```
In[17]:= stilllife=Flatten[Table[Join[{center},PadLeft[IntegerDigits[neighbours,2],8]]->
```

```
Boole[Boole[newcell[center,neighbours]]==center],{neighbours,0,255},{center,0,1}]]
```

```
Out[17]= {{0,0,0,0,0,0,0,0,0}->1,
```

$$\{1, 0, 0, 0, 0, 0, 0, 0, 0\} \rightarrow 0,$$
$$\{0,0,0,0,0,0,0,0,1\} \rightarrow 1,$$
$$\{1, 0, 0, 0, 0, 0, 0, 0, 1\} \rightarrow 0,$$

• • •

<sup>109</sup><http://golly.sourceforge.net/>



```

    return mathematica_to_CNF(s, center, a)

def try_again (clauses):
    # rules for the main part of grid
    for r in range(H):
        for c in range(W):
            clauses=clauses+stillife(coords_to_var(r, c, H, W), get_neighbours(r, c, H, W))

    # cells behind visible grid must always be false:
    for c in range(-1, W+1):
        for r in [-1,H]:
            clauses=clauses+cell_is_false(coords_to_var(r, c, H, W), get_neighbours(r, c, H, W))
    for c in [-1,W]:
        for r in range(-1, H+1):
            clauses=clauses+cell_is_false(coords_to_var(r, c, H, W), get_neighbours(r, c, H, W))

    write_CNF("tmp.cnf", clauses, VARS_TOTAL)

    print "%d clauses" % len(clauses)

    solution=run_minisat ("tmp.cnf")
    os.remove("tmp.cnf")
    if solution==None:
        print "unsat!"
        exit(0)

    grid=SAT_solution_to_grid(solution, H, W)
    print_grid(grid)
    write_RLE(grid)

    return grid

clauses=[]
# always false:
clauses.append ("-"+VAR_FALSE)

while True:
    solution=try_again(clauses)
    clauses.append(negate_clause(grid_to_clause(solution, H, W)))
    clauses.append(negate_clause(grid_to_clause(reflect_vertically(solution), H, W)))
    clauses.append(negate_clause(grid_to_clause(reflect_horizontally(solution), H, W)))
    # is this square?
    if W==H:
        clauses.append(negate_clause(grid_to_clause(rotate_square_array(solution,1), H, W)))
        clauses.append(negate_clause(grid_to_clause(rotate_square_array(solution,2), H, W)))
        clauses.append(negate_clause(grid_to_clause(rotate_square_array(solution,3), H, W)))
    print ""

```

( [https://github.com/dennis714/SAT\\_SMT\\_article/blob/master/SAT/GoL/stillife1.py](https://github.com/dennis714/SAT_SMT_article/blob/master/SAT/GoL/stillife1.py) )

What we've got for 2 · 2?

```

1881 clauses
..
..
1.rle written

1887 clauses
**
**
2.rle written

1893 clauses
unsat!

```

Both solutions are correct: empty square will progress into empty square (no cells are born). 2 · 2 box is also known “still life”.

What about 3 · 3 square?

```

2887 clauses
...
...
...
1.rle written

2893 clauses

```

```

.**
.**
...
2.rle written

2899 clauses
.**
*.*
**
3.rle written

2905 clauses
.*
*.*
**
4.rle written

2911 clauses
.*
*.*
.*
5.rle written

2917 clauses
unsat!

```

Here is a problem: we see familiar  $2 \cdot 2$  box, but shifted. This is indeed correct solution, but we don't interested in it, because it has been already seen.

What we can do is add another condition. We can force minisat to find solutions with no empty rows and columns. This is easy. These are SAT variables for  $5 \cdot 5$  square:

```

1  2  3  4  5
6  7  8  9 10
11 12 13 14 15
16 17 18 19 20
21 22 23 24 25

```

Each clause is “OR” clause, so all we have to do is to add 5 clauses:

```

1 OR 2 OR 3 OR 4 OR 5
6 OR 7 OR 8 OR 9 OR 10
...

```

That means that each row must have at least one *True* value somewhere. We can also do this for each column as well.

```

...

# each row must contain at least one cell!
for r in range(H):
    clauses.append(" ".join([coords_to_var(r, c, H, W) for c in range(W)]))

# each column must contain at least one cell!
for c in range(W):
    clauses.append(" ".join([coords_to_var(r, c, H, W) for r in range(H)]))

...

```

( [https://github.com/dennis714/SAT\\_SMT\\_article/blob/master/SAT/GoL/stilllife2.py](https://github.com/dennis714/SAT_SMT_article/blob/master/SAT/GoL/stilllife2.py) )

Now we can see that  $3 \cdot 3$  square has 3 possible “still lives”:

```

2893 clauses
.*
*.*
**
1.rle written

2899 clauses
.*
*.*
.*
2.rle written

```

```

2905 clauses
.**
*.*
**
3.rle written

2911 clauses
unsat!

```

4 · 4 has 7:

```

4169 clauses
..**
...*
***.
*...
1.rle written

4175 clauses
..**
..*.
*.*.
**..
2.rle written

4181 clauses
..**
.*.*
*.*.
**..
3.rle written

4187 clauses
..*.
.*.*
*.*.
**..
4.rle written

4193 clauses
.**.
*..*
*.*.
.*..
5.rle written

4199 clauses
..*.
.*.*
*.*.
.*..
6.rle written

4205 clauses
.**.
*..*
*..*
.**.
7.rle written

4211 clauses
unsat!

```

When I try large squares, like  $20 \cdot 20$ , funny things happen. First of all, minisat finds solutions not very pleasing aesthetically, but still correct, like:

```

61033 clauses
....**.**.**.**.*
**.*.*.*.*.*.*
*
*.....
*
**.....
*
*.....
*
**.....
*
*.....

```

```

.*
**
*
.*
**
*
.*
.*
.*
.*
***
*
1.rle written
...

```

Indeed: all rows and columns has at least one *True* value.  
Then minisat begins to add smaller “still lives” into the whole picture:

```

61285 clauses
**          **          **          **
**      *      *      *      *      *
          **      *      *
. . . . .
**          *          *          **
      **          *          *
. . . . .
      *      *          *          *
          *      *          *          *
**      *      *          **          **
**      *      *          *          *
*      *          *          *          *
*          *          *          *          *
          *          *          *          *
      *          *          *          *
. . . . .
*          **          *          *          *
**      *          *          *          *
*      *          *          *          *
**      *          *          *          *
***      ***      *      *          *
          *          ***      **      **
**      *          *          *          *
*      **      **          *          *
. . . . .
**      *          *          *          *
**      *          *          *          *
**      *          *          *          *
**      *          *          *          *
43.rle written

```

In other words, result is a square consisting of smaller “still lives”. It then altering these parts slightly, shifting back and forth. Is it cheating? Anyway, it does it in a strict accordance to rules we defined.

But we want *denser* picture. We can add a rule: in all 5-cell chunks there must be at least one *True* cell. To achieve this, we just split the whole square by 5-cell chunks and add clause for each:

```
...
# make result denser:
lst=[]
for r in range(H):
    for c in range(W):
        lst.append(coords_to_var(r, c, H, W))
# divide them all by chunks and add to clauses:
CHUNK_LEN=5
for c in list_partition(lst,len(lst)/CHUNK_LEN):
    tmp=" ".join(c)
    clauses.append(tmp)
...
```

( [https://github.com/dennis714/SAT\\_SMT\\_article/blob/master/SAT/GoL/stilllife.py](https://github.com/dennis714/SAT_SMT_article/blob/master/SAT/GoL/stilllife.py) )

This is indeed denser:

61113 clauses



61119 clauses

2.rle written

```
2.rle written
```

• • •

Let's try more dense, one mandatory *true* cell per each 4-cell chunk:

61133 clauses
---------------

```
1.rle written
```

61139 clauses

...and even more: one cell per each 3-cell chunk:

[illegible]

This is most dense. Unfortunately, it's impossible to construct "still life" with one mandatory *true* cell per each 2-cell chunk.

### 13.8.3 The source code

Source code and Wolfram Mathematica notebook: [https://github.com/dennis714/SAT\\_SMT\\_article/tree/master/SAT/GoL](https://github.com/dennis714/SAT_SMT_article/tree/master/SAT/GoL).

## 13.9 Simplest SAT solver in ~120 lines

This is simplest possible backtracking SAT solver written in Python (not a [DPLL<sup>110</sup>](#) one). It uses the same backtracking algorithm you can find in many simple Sudoku and 8 queens solvers. It works significantly slower, but due to its extreme simplicity, it can also count solutions. For example, it can count all solutions of 8 queens problem ([13.4](#)).

Also, there are 70 solutions for POPCNT4 function <sup>111</sup> (the function is true if any 4 of its input 8 variables are true):

```
SAT
-1 -2 -3 -4 5 6 7 8 0
SAT
-1 -2 -3 4 -5 6 7 8 0
SAT
-1 -2 -3 4 5 -6 7 8 0
SAT
-1 -2 -3 4 5 6 -7 8 0
...
SAT
1 2 3 -4 -5 6 -7 -8 0
SAT
1 2 3 -4 5 -6 -7 -8 0
SAT
1 2 3 4 -5 -6 -7 -8 0
UNSAT
solutions= 70
```

It was also tested on my SAT-based Minesweeper cracker ([13.7](#)), and finishes in reasonable time (though, slower than MiniSat by a factor of ~10).

On bigger CNF instances, it gets stuck, though.

The source code:

```
#!/usr/bin/env python

count_solutions=True
#count_solutions=False

import sys

def read_text_file (fname):
    with open(fname) as f:
        content = f.readlines()
    return [x.strip() for x in content]

def read_DIMACS (fname):
    content=read_text_file(fname)

    header=content[0].split(" ")

    assert header[0]=="p" and header[1]=="cnf"
    variables_total, clauses_total = int(header[2]), int(header[3])

    # array idx=number (of line) of clause
    # val=list of terms
    # term can be negative signed integer
    clauses=[]
    for c in content[1:]:
        clause=[]
        for var_s in c.split(" "):
            var=int(var_s)
            if var!=0:
                clause.append(var)
        clauses.append(clause)

    # this is variables index.
    # for each variable, it has list of clauses, where this variable is used.
    # key=variable
    # val=list of numbers of clause
    variables_idx={}
    for i in range(len(clauses)):
```

<sup>110</sup>Davis-Putnam-Logemann-Loveland

<sup>111</sup>[https://github.com/dennis714/SAT\\_SMT\\_article/blob/master/SAT/backtrack/POPCNT4.cnf](https://github.com/dennis714/SAT_SMT_article/blob/master/SAT/backtrack/POPCNT4.cnf)

```

        for term in clauses[i]:
            variables_idx.setdefault(abs(term), []).append(i)

    return clauses, variables_idx

# clause=list of terms. signed integer. -x means negated.
# values=list of values: from 0th: [F,F,F,F,T,F,T,...]
def eval_clause (terms, values):
    try:
        # we search for at least one True
        for t in terms:
            # variable is non-negated:
            if t>0 and values[t-1]==True:
                return True
            # variable is negated:
            if t<0 and values[(-t)-1]==False:
                return True
        # all terms enumerated at this point
        return False
    except IndexError:
        # values[] has not enough values
        # None means "maybe"
        return None

def chk_vals(clauses, variables_idx, vals):
    # check only clauses which affected by the last (new/changed) value, ignore the rest
    # because since we already got here, all other values are correct, so no need to recheck them
    idx_of_last_var=len(vals)
    # variable can be absent in index, because no clause uses it:
    if idx_of_last_var not in variables_idx:
        return True
    # enumerate clauses which has this variable:
    for clause_n in variables_idx[idx_of_last_var]:
        clause=clauses[clause_n]
        # if any clause evaluated to False, stop checking, new value is incorrect:
        if eval_clause (clause, vals)==False:
            return False
    # all clauses evaluated to True or None ("maybe")
    return True

def print_vals(vals):
    # enumerate all vals[]
    # prepend "-" if vals[i] is False (i.e., negated).
    print "".join(["-", ""][vals[i]] + str(i+1) + " " for i in range(len(vals))])+"0"

clauses, variables_idx = read_DIMACS(sys.argv[1])

solutions=0

def backtrack(vals):
    global solutions

    if len(vals)==len(variables_idx):
        # we reached end - all values are correct
        print "SAT"
        print_vals(vals)
        if count_solutions:
            solutions=solutions+1
            # go back, if we need more solutions:
            return
        else:
            exit(10) # as in MiniSat
        return

    for next in [False, True]:
        # add new value:
        new_vals=vals+[next]
        if chk_vals(clauses, variables_idx, new_vals):
            # new value is correct, try add another one:
            backtrack(new_vals)
        else:
            # new value (False) is not correct, now try True (variable flip):
            continue

# try to find all values:
backtrack([])

```

```
print "UNSAT"
if count_solutions:
    print "solutions=", solutions
exit(20) # as in MiniSat
```

As you can see, all it does is enumerate all possible solutions, but prunes search tree as early as possible. This is backtracking.

The files: [https://github.com/dennis714/SAT\\_SMT\\_article/tree/master/SAT/backtrack](https://github.com/dennis714/SAT_SMT_article/tree/master/SAT/backtrack).

Some comments: [https://www.reddit.com/r/compsci/comments/6jn3th/simplest\\_sat\\_solver\\_in\\_120\\_lines/](https://www.reddit.com/r/compsci/comments/6jn3th/simplest_sat_solver_in_120_lines/).

## 13.10 Integer factorization using SAT solver

See also: integer factorization using Z3 SMT solver (7.18).

We are going to simulate electronic circuit of binary multiplier in SAT and then ask solver, what multiplier's inputs must be so the output will be a desired number? If this situation is impossible, the desired number is prime.

First we should build multiplier out of adders.

### 13.10.1 Binary adder in SAT

Simple binary adder usually consists of full-adders and one half-adder. These are basic elements of adders.




Figure 17: Half-adder

( The image has been taken from [Wikipedia](#). )

Full-adder:





Figure 18: Full-adder

( The image has been taken from [Wikipedia](#). )


Here is a 4-bit ripple-carry adder:

X3	Y3		X2	Y2		X1	Y1		X0	Y0
v	v		v	v		v	v		v	v



It can be used for most tasks.

Here is a 4-bit ripple-carry adder with carry-in:



What carries are? 4-bit adder can sum up two numbers up to 0b1111 (15). 15+15=30 and this is 0b11110, i.e., 5 bits. Lowest 4 bits is a sum. 5th most significant bit is not a part of sum, but is a carry bit.

If you sum two numbers on x86 CPU, CF flag is a carry bit connected to **ALU**<sup>112</sup>. It is set if a resulting sum is bigger than it can be fit into result.

Now you can also need carry-in. Again, x86 CPU has ADC instruction, it takes CF flag state. It can be said, CF flag is connected to adder's carry-in input. Hence, combining two ADD and ADC instructions you can sum up 128 bits on 64-bit CPU.

By the way, this is a good explanation of "carry-ripple". The very first full-adder's result is depending on the carry-out of the previous full-adder. Hence, adders cannot work in parallel. This is a problem of simplest possible adder, other adders can solve this.

To represent full-adders in CNF form, we can use Wolfram Mathematica. I've taken truth table for full-adder from [Wikipedia](#):

Inputs			Outputs	
A	B	Cin	Cout	Sum
0	0	0	0	0
0	0	1	0	1
0	1	0	0	1
0	1	1	1	0
1	0	0	0	1
1	0	1	1	0
1	1	0	1	0
1	1	1	1	1

In Mathematica, I'm setting "->1" if row is correct and "->0" if not correct.

```
In[59]:= FaTbl = {{0, 0, 0, 0, 0} -> 1, {0, 0, 0, 0, 1} ->
0, {0, 0, 0, 1, 0} -> 0, {0, 0, 0, 1, 1} -> 0, {0, 0, 1, 0, 0} ->
0, {0, 0, 1, 0, 1} -> 1, {0, 0, 1, 1, 0} -> 0, {0, 0, 1, 1, 1} ->
0, {0, 1, 0, 0, 0} -> 0, {0, 1, 0, 0, 1} -> 1, {0, 1, 0, 1, 0} ->
0, {0, 1, 0, 1, 1} -> 0, {0, 1, 1, 0, 0} -> 0, {0, 1, 1, 0, 1} ->
0, {0, 1, 1, 1, 0} -> 1, {0, 1, 1, 1, 1} -> 0, {1, 0, 0, 0, 0} ->
0, {1, 0, 0, 0, 1} -> 1, {1, 0, 0, 1, 0} -> 0, {1, 0, 0, 1, 1} ->
0, {1, 0, 1, 0, 0} -> 0, {1, 0, 1, 0, 1} -> 0, {1, 0, 1, 1, 0} ->
0, {1, 0, 1, 1, 1} -> 0, {1, 1, 0, 0, 0} -> 0, {1, 1, 0, 0, 1} ->
0, {1, 1, 0, 1, 0} -> 1, {1, 1, 0, 1, 1} -> 0, {1, 1, 1, 0, 0} ->
0, {1, 1, 1, 0, 1} -> 0, {1, 1, 1, 1, 0} -> 0, {1, 1, 1, 1, 1} -> 1}
```

...

```
In[60]:= BooleanConvert[
BooleanFunction[FaTbl, {a, b, cin, cout, s}], "CNF"]
```

```
Out[60]= (! a || ! b || ! cin || s) && (! a || ! b ||
cout) && (! a || ! cin || cout) && (! a || cout || s) && (a || b ||
cin || ! s) && (a || b || ! cout) && (a ||
cin || ! cout) && (a || ! cout || ! s) && (! b || ! cin ||
cout) && (! b || cout || s) && (b ||
cin || ! cout) && (b || ! cout || ! s) && (! cin || cout ||
s) && (cin || ! cout || ! s)
```

These clauses can be used as full-adder.  
Here is it:

```
# full-adder, as found by Mathematica using truth table:
def FA (self, a,b,cin):
    s=self.create_var()
    cout=self.create_var()

    self.add_clause([self.neg(a), self.neg(b), self.neg(cin), s])
    self.add_clause([self.neg(a), self.neg(b), cout])
    self.add_clause([self.neg(a), self.neg(cin), cout])
    self.add_clause([self.neg(a), cout, s])
    self.add_clause([a, b, cin, self.neg(s)])
    self.add_clause([a, b, self.neg(cout)])
    self.add_clause([a, cin, self.neg(cout)])
    self.add_clause([a, self.neg(cout), self.neg(s)])
    self.add_clause([self.neg(b), self.neg(cin), cout])
    self.add_clause([self.neg(b), cout, s])
    self.add_clause([b, cin, self.neg(cout)])
    self.add_clause([b, self.neg(cout), self.neg(s)])
    self.add_clause([self.neg(cin), cout, s])
    self.add_clause([cin, self.neg(cout), self.neg(s)])

    return s, cout
```

And the adder:

```
# bit order: [MSB..LSB]
# n-bit adder:
def adder(self, X,Y):
    assert len(X)==len(Y)
    # first full-adder could be half-adder
    # start with lowest bits:
    inputs=frolic.rvr(list(zip(X,Y)))
    carry=self.const_false
    sums=[]
    for pair in inputs:
        # "carry" variable is replaced at each iteration.
        # so it is used in the each FA() call from the previous FA() call.
        s, carry = self.FA(pair[0], pair[1], carry)
        sums.append(s)
    return frolic.rvr(sums), carry
```

### 13.10.2 Binary multiplier in SAT

Remember school-level long division? This multiplier works in a same way, but for binary digits.  
Here is example of multiplying 0b1101 (X) by 0b0111 (Y):

```

      LSB
      |
      v
    1101 <- X
    -----
LSB 0|    0000
    1|    1101
    1|    1101
    1|    1101
    ^
    |
    Y
```

If bit from Y is zero, a row is zero. If bit from Y is non-zero, a row is equal to X, but shifted each time. Then you just sum up all rows (which are called "partial products".)

This is 4-bit binary multiplier. It takes 4-bit inputs and produces 8-bit output:

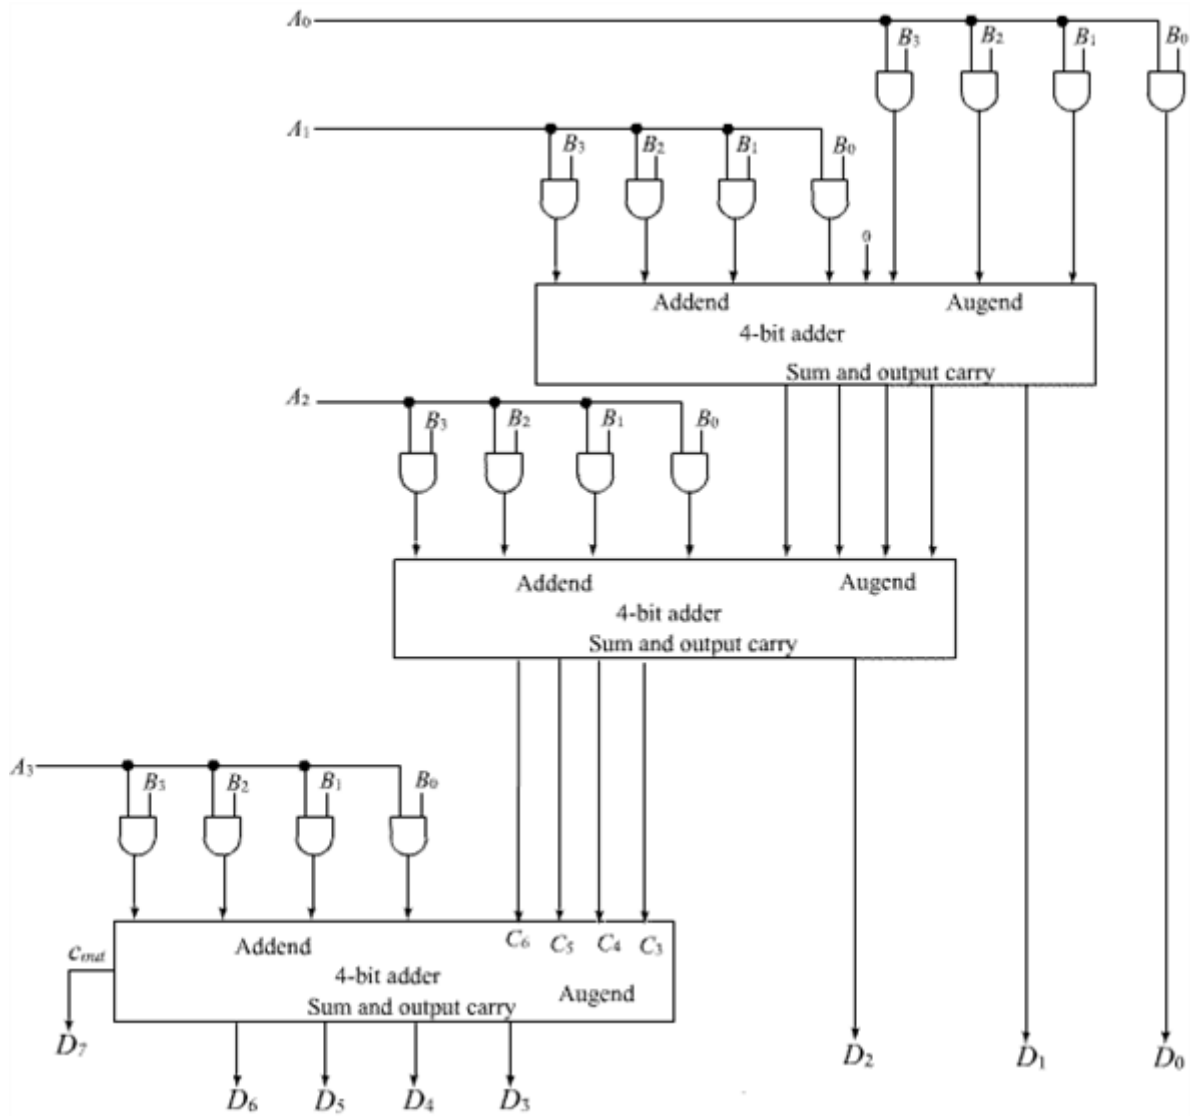


Figure 19: 4-bit binary multiplier

( The image has been taken from <http://www.chegg.com/homework-help/binary-multiplier-multiplies-two> )

I would build separate block, "multiply by one bit" as a latch for each partial product:

```
def AND_Tseitin(self, v1, v2, out):
    self.add_clause([self.neg(v1), self.neg(v2), out])
    self.add_clause([v1, self.neg(out)])
    self.add_clause([v2, self.neg(out)])

def AND(self, v1, v2):
    out = self.create_var()
    self.AND_Tseitin(v1, v2, out)
    return out

...

# bit is 0 or 1.
# i.e., if it's 0, output is 0 (all bits)
# if it's 1, output=input
def mult_by_bit(self, X, bit):
    return [self.AND(i, bit) for i in X]

# bit order: [MSB..LSB]
# build multiplier using adders and mult_by_bit blocks:
def multiplier(self, X, Y):
    assert len(X) == len(Y)
    out = []
```



```

#initial:
prev=[self.const_false]*len(X)
# first adder can be skipped, but I left thing "as is" to make it simpler
for Y_bit in frolic.rvr(Y):
    s, carry = self.adder(self.mult_by_bit(X, Y_bit), prev)
    out.append(s[-1])
    prev=[carry] + s[:-1]

return prev + frolic.rvr(out)

```

AND gate is constructed here using Tseitin transformations. This is quite popular way of encoding gates in CNF form, by adding additional variable: [https://en.wikipedia.org/wiki/Tseitin\\_transformation](https://en.wikipedia.org/wiki/Tseitin_transformation). In fact, full-adder can be constructed without Mathematica, using logic gates, and encoded by Tseitin transformation.

### 13.10.3 Glueing all together

```

#!/usr/bin/env python3

import itertools, subprocess, os, math, random
from operator import mul
import frolic, Xu

def factor(n):
    print ("factoring %d" % n)

    # size of inputs.
    # in other words, how many bits we have to allocate to store 'n'?
    input_bits=int(math.ceil(math.log(n,2)))
    print ("input_bits=%d" % input_bits)

    s=Xu.Xu(False)

    factor1,factor2=s.alloc_BV(input_bits),s.alloc_BV(input_bits)
    product=s.multiplier(factor1,factor2)

    # at least one bit in each input must be set, except lowest.
    # hence we restrict inputs to be greater than 1
    s.fix(s.OR(factor1[:-1]), True)
    s.fix(s.OR(factor2[:-1]), True)

    # output has a size twice as bigger as each input:
    s.fix_BV(product, Xu.n_to_BV(n,input_bits*2))

    if s.solve()==False:
        print ("%d is prime (unsat)" % n)
        return [n]

    # get inputs of multiplier:
    factor1_n=Xu.BV_to_number(s.get_BV_from_solution(factor1))
    factor2_n=Xu.BV_to_number(s.get_BV_from_solution(factor2))

    print ("factors of %d are %d and %d" % (n, factor1_n, factor2_n))
    # factor factors recursively:
    rt=sorted(factor (factor1_n) + factor (factor2_n))
    assert reduce(mul, rt, 1)==n
    return rt

# infinite test:
def test():
    while True:
        print (factor (random.randrange(1000000000000)))

#test()

print (factor(1234567890))

```

I just connect our number to output of multiplier and ask SAT solver to find inputs. If it's UNSAT, this is prime number. Then we factor factors recursively.

Also, we want block input factors of 1, because obviously, we do not interesting in the fact that  $n*1=n$ . I'm using wide OR gates for this.

Output:

```
% python factor_SAT.py
factoring 1234567890
input_bits=31
factors of 1234567890 are 2 and 617283945
factoring 2
input_bits=1
2 is prime (unsat)
factoring 617283945
input_bits=30
factors of 617283945 are 3 and 205761315
factoring 3
input_bits=2
3 is prime (unsat)
factoring 205761315
input_bits=28
factors of 205761315 are 3 and 68587105
factoring 3
input_bits=2
3 is prime (unsat)
factoring 68587105
input_bits=27
factors of 68587105 are 5 and 13717421
factoring 5
input_bits=3
5 is prime (unsat)
factoring 13717421
input_bits=24
factors of 13717421 are 3607 and 3803
factoring 3607
input_bits=12
3607 is prime (unsat)
factoring 3803
input_bits=12
3803 is prime (unsat)
[2, 3, 3, 5, 3607, 3803]
```

So,  $1234567890 = 2 \cdot 3 \cdot 3 \cdot 5 \cdot 3607 \cdot 3803$ .

It works way faster than by Z3 solution, but still slow. It can factor numbers up to maybe  $\sim 2^{40}$ , while Wolfram Mathematica can factor  $\sim 2^{80}$  easily.

The full source code: [https://github.com/dennis714/yurichev.com/blob/master/blog/factor\\_SAT/factor\\_SAT.py](https://github.com/dennis714/yurichev.com/blob/master/blog/factor_SAT/factor_SAT.py).

#### 13.10.4 Division using multiplier

Hard to believe, but why we couldn't define one of factors and ask SAT solver to find another factor? Then it will divide numbers! But, unfortunately, this is somewhat impractical, since it will work only if remainder is zero:

```
#!/usr/bin/env python3

import itertools, subprocess, os, math, random
from operator import mul
import frolic, Xu

def div(dividend, divisor):

    # size of inputs.
    # in other words, how many bits we have to allocate to store 'n'?
    input_bits=int(math.ceil(math.log(dividend,2)))
    print ("input_bits=%d" % input_bits)

    s=Xu.Xu(False)

    factor1,factor2=s.alloc_BV(input_bits),s.alloc_BV(input_bits)
    product=s.multiplier(factor1,factor2)

    # connect divisor to one of multiplier's input:
    s.fix_BV(factor1, Xu.n_to_BV(divisor,input_bits))
    # output has a size twice as bigger as each input.
    # connect dividend to multiplier's output:
    s.fix_BV(product, Xu.n_to_BV(dividend,input_bits*2))
```

```

if s.solve()==False:
    print ("remainder!=0 (unsat)")
    return None

# get 2nd input of multiplier, which is quotient:
return Xu.BV_to_number(s.get_BV_from_solution(factor2))

print (div (12345678901234567890123456789*12345, 12345))

```

The full source code: [https://github.com/dennis714/yurichev.com/blob/master/blog/factor\\_SAT/div.py](https://github.com/dennis714/yurichev.com/blob/master/blog/factor_SAT/div.py).

It works very fast, but still, slower than conventional ways.

### 13.10.5 Breaking RSA!

It's not a problem to build multiplier with 4096 bit inputs and 8192 output, but it will not work in practice. Still, you can break toy-level demonstrational RSA problems with key less than  $2^{40}$  or something like that (or larger, using Wolfram Mathematica).

### 13.10.6 Further reading

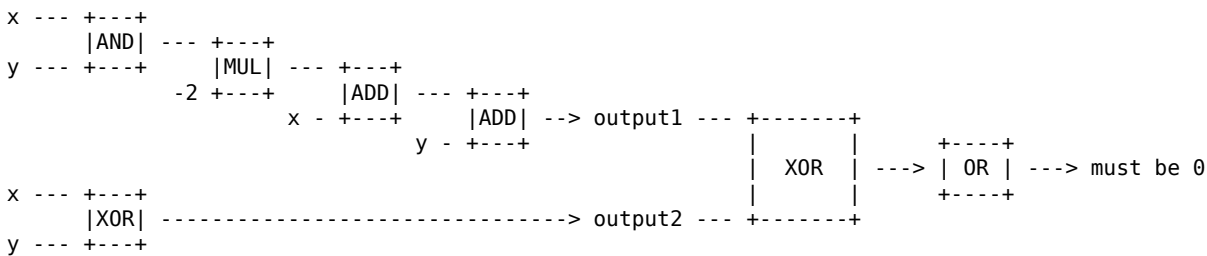
[1](#), [2](#), [3](#).

## 13.11 Proving bizarre XOR alternative using SAT solver

I once wrote about quite bizarre XOR alternative I've found using aha! superoptimizer: [7.5](#).

Now let's try to prove it using SAT.

We would build an electric circuit for  $x \oplus y = -2 * (x \& y) + (x + y)$  like that:



So it has two parts: generic XOR block and a block which must be equivalent to XOR. Then we compare its outputs using XOR and OR. If outputs of these parts are always equal to each other for all possible x and y, output of the whole block must be 0.

I do otherwise, I'm trying to find such an input pair, for which output will be 1:

```

def chk1():
    input_bits=8

    s=Xu.Xu(False)

    x,y=s.alloc_BV(input_bits),s.alloc_BV(input_bits)
    step1=s.BV_AND(x,y)
    minus_2=[s.const_true]*(input_bits-1)+[s.const_false]
    product=s.multiplier(step1,minus_2)[input_bits:]
    result1=s.adder(s.adder(product, x)[0], y)[0]

    result2=s.BV_XOR(x,y)

    s.fix(s.OR(s.BV_XOR(result1, result2)), True)

    if s.solve()==False:
        print ("unsat")
        return

    print ("sat")
    print ("x=%x" % Xu.BV_to_number(s.get_BV_from_solution(x)))
    print ("y=%x" % Xu.BV_to_number(s.get_BV_from_solution(y)))
    print ("step1=%x" % Xu.BV_to_number(s.get_BV_from_solution(step1)))
    print ("product=%x" % Xu.BV_to_number(s.get_BV_from_solution(product)))

```

```
print ("result1=%x" % Xu.BV_to_number(s.get_BV_from_solution(result1)))
print ("result2=%x" % Xu.BV_to_number(s.get_BV_from_solution(result2)))
print ("minus_2=%x" % Xu.BV_to_number(s.get_BV_from_solution(minus_2)))
```

The full source code: [https://github.com/dennis714/yurichev.com/blob/master/blog/XOR\\_SAT/XOR\\_SAT.py](https://github.com/dennis714/yurichev.com/blob/master/blog/XOR_SAT/XOR_SAT.py).

SAT solver returns "unsat", meaning, it could find such a pair. In other words, it couldn't find a counterexample. So the circuit always outputs 0, for all possible inputs, meaning, outputs of two parts are always the same.

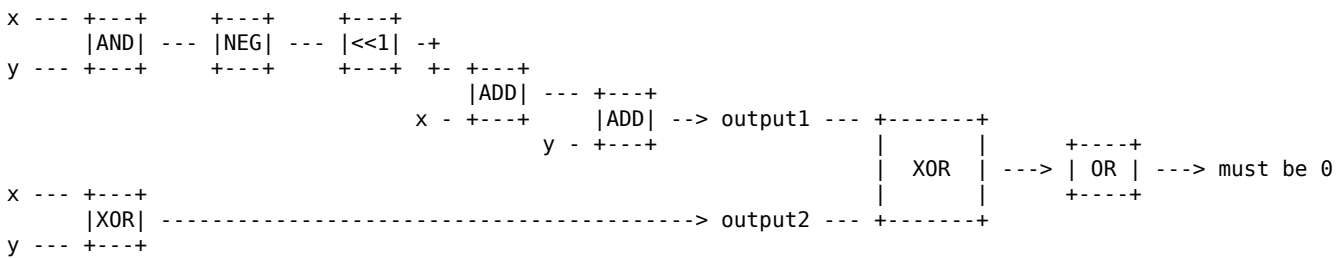
Modify the circuit, and the program will find such a state, and print it.

That circuit also called "miter". According to Google translate, one of meaning of this word is:

a joint made between two pieces of wood or other material at an angle of 90°, such that the line of junction bisects this angle.

It's also slow, because multiplier block is used: so we use small 8-bit x's and y's.

But the whole thing can be rewritten:  $x \oplus y = x + y - (x \& y) \ll 1$ . And subtraction is addition, but with one negated operand. So,  $x \oplus y = -(x \& y) \ll 1 + (x + y)$  or  $x \oplus y = (x \& y) * 2 - (x + y)$ .



**NEG** is negation block, in two's complement system. It just inverts all bits and adds 1:

```
def NEG(self, x):
    # invert all bits
    tmp=self.BV_NOT(x)
    # add 1
    one=self.alloc_BV(len(tmp))
    self.fix_BV(one,n_to_BV(1, len(tmp)))
    return self.adder(tmp, one)[0]
```

Shift by one bit does nothing except rewiring.

That works way faster, and can prove correctness for 64-bit x's and y's, or for even bigger input values:

```
def chk2():
    input_bits=64

    s=Xu.Xu(False)

    x,y=s.alloc_BV(input_bits),s.alloc_BV(input_bits)
    step1=s.BV_AND(x,y)
    step2=s.shift_left_1(s.NEG(step1))

    result1=s.adder(s.adder(step2, x)[0], y)[0]

    result2=s.BV_XOR(x,y)

    s.fix(s.OR(s.BV_XOR(result1, result2)), True)

    if s.solve()==False:
        print ("unsat")
        return

    print ("sat")
    print ("x=%x" % Xu.BV_to_number(s.get_BV_from_solution(x)))
    print ("y=%x" % Xu.BV_to_number(s.get_BV_from_solution(y)))
    print ("step1=%x" % Xu.BV_to_number(s.get_BV_from_solution(step1)))
    print ("step2=%x" % Xu.BV_to_number(s.get_BV_from_solution(step2)))
    print ("result1=%x" % Xu.BV_to_number(s.get_BV_from_solution(result1)))
    print ("result2=%x" % Xu.BV_to_number(s.get_BV_from_solution(result2)))
```

The source code: [https://github.com/dennis714/yurichev.com/blob/master/blog/XOR\\_SAT/XOR\\_SAT.py](https://github.com/dennis714/yurichev.com/blob/master/blog/XOR_SAT/XOR_SAT.py).

## 14 MaxSAT

MaxSAT problem is a problem where as many clauses should be satisfied, as possible, but maybe not all.

(Usual) clauses which *must* be satisfied, called *hard clauses*. Clauses which *should* be satisfied, called *soft clauses*.

MaxSAT solver tries to satisfy all *hard clauses* and as much *soft clauses*, as possible.

\*.wcnf files are used, the format is almost the same as in DIMACS files, like:

```
p wcnf 207 796 208
208 1 0
208 2 0
208 3 0
208 4 0

...

1 -152 0
1 -153 0
1 -154 0
1 155 0
1 -156 0
1 -157 0
```

Each clause is written as in DIMACS file, but the first number is weight. MaxSAT solver tries to maximize clauses with bigger weights first.

If the weight has *top weight*, the clause is *hard clause* and must always be satisfied. *Top weight* is set in header. In our case, it's 208.

Some well-known MaxSAT solvers are Open-WBO<sup>113</sup>, etc.

### 14.1 Gray code in MaxSAT

This is remake of gray code generator for Z3 (7.17).

Here is also *ch[]* table, but we add soft clauses for it here. The goal is to make as many *False*'s in *ch[]* table, as possible.

```
#!/usr/bin/env python3

import subprocess, os, itertools
import frolic, Xu

BITS=5

# how many times a run of bits for each bit can be changed (max).
# it can be 4 for 4-bit Gray code or 8 for 5-bit code.
# 12 for 6-bit code (maybe even less)

ROWS=2**BITS
MASK=ROWS-1 # 0x1f for 5 bits, 0xf for 4 bits, etc

def do_all():
    s=Xu.Xu(maxsat=True)

    code=[s.alloc_BV(BITS) for r in range(ROWS)]
    ch=[s.alloc_BV(BITS) for r in range(ROWS)]

    # each rows must be different from a previous one and a next one by 1 bit:
    for i in range(ROWS):
        # get bits of the current row:
        lst1=[code[i][bit] for bit in range(BITS)]
        # get bits of the next row.
        # important: if the current row is the last one, (last+1)&MASK==0, so we overlap here:
        lst2=[code[(i+1)&MASK][bit] for bit in range(BITS)]
        s.hamming1(lst1, lst2)

    # no row must be equal to any another row:
    for i in range(ROWS):
        for j in range(ROWS):
            if i==j:
                continue
```

<sup>113</sup><http://sat.inesc-id.pt/open-wbo>

```

        lst1=[code[i][bit] for bit in range(BITS)]
        lst2=[code[j][bit] for bit in range(BITS)]
        s.fix_BV_NEQ(lst1, lst2)

# 1 in ch[] table means that run of 1's has been changed to run of 0's, or back.
# "run" change detected using simple XOR:
for i in range(ROWS):
    for bit in range(BITS):
        # row overlapping works here as well.
        # we add here "soft" constraint with weight=1:
        s.fix_soft(s.EQ(ch[i][bit], s.XOR(code[i][bit],code[(i+1)&MASK][bit])), False, weight=1)

if s.solve()==False:
    print ("unsat")
    exit(0)

print ("code table:")

for i in range(ROWS):
    tmp=""
    for bit in range(BITS):
        t=s.get_var_from_solution(code[i][BITS-1-bit])
        if t:
            tmp=tmp+"*"
        else:
            tmp=tmp+" "
    print (tmp)

# get statistics:
stat={}

for i in range(ROWS):
    for bit in range(BITS):
        x=s.get_var_from_solution(ch[i][BITS-1-bit])
        if x==0:
            # increment if bit is present in dict, set 1 if not present
            stat[bit]=stat.get(bit, 0)+1

print ("stat (bit number: number of changes): ")
print (stat)

do_all()

```

So it does, for 5-bit Gray code:

code table:

```

****
*****
* ***
    ***
    **
    ***
    **
    *
    **
**  *
*** *
***
****
* **
*  *
*
*  *
* * *
* *
*
*  *
*
*
**
* **
** **
** *
*  *
* **

```

```

*  *
** *
stat (bit number: number of changes):
{0: 6, 1: 4, 2: 6, 3: 6, 4: 10}

```

## 15 Acronyms used

<b>CNF</b> Conjunctive normal form .....	4
<b>DNF</b> Disjunctive normal form .....	151
<b>DSL</b> Domain-specific language .....	5
<b>CPRNG</b> Cryptographically Secure Pseudorandom Number Generator .....	25
<b>SMT</b> Satisfiability modulo theories .....	1
<b>SAT</b> Boolean satisfiability problem .....	1
<b>LCG</b> Linear congruential generator .....	1
<b>PL</b> Programming Language .....	5
<b>OOP</b> Object-oriented programming .....	68
<b>SSA</b> Static single assignment form .....	60
<b>CPU</b> Central processing unit .....	62
<b>FPU</b> Floating-point unit .....	87
<b>PRNG</b> Pseudorandom number generator .....	97
<b>CRT</b> C runtime library .....	97
<b>CRC</b> Cyclic redundancy check .....	94
<b>AST</b> Abstract syntax tree .....	69
<b>AKA</b> Also Known As .....	1
<b>CTF</b> Capture the Flag .....	141
<b>ISA</b> Instruction Set Architecture .....	62
<b>CSP</b> Constraint satisfaction problem .....	11

<b>CS</b> Computer science .....	4
<b>DAG</b> Directed acyclic graph.....	35
<b>NOP</b> No Operation .....	66
<b>JVM</b> Java Virtual Machine.....	87
<b>VM</b> Virtual Machine.....	100
<b>LZSS</b> Lempel-Ziv-Storer-Szymanski .....	42
<b>RAM</b> Random-access memory .....	127
<b>FPGA</b> Field-programmable gate array .....	143
<b>EDA</b> Electronic design automation.....	143
<b>MAC</b> Message authentication code .....	143
<b>ECC</b> Elliptic curve cryptography .....	143
<b>API</b> Application programming interface .....	5
<b>NSA</b> National Security Agency .....	25
<b>DPLL</b> Davis-Putnam-Logemann-Loveland .....	186